

MATHEMATICS FOR SOFTWARE ENGINEERING

Mathematics for Software Engineering

Authors: Richard Brooks & Eduard Fekete

Date: May, 2025

Version: 2.0

Contents

Chapter 1 Basic Arithmetic and Functions	1
1.1 Factorisation and the Order of Operations	1
1.2 Fractions	3
1.3 Exponents, Radicals and Surds	6
1.4 Using Formulae and Substitution	9
1.5 Rearranging Formulae	10
1.6 Functions	12
1.7 Graphical Identification of Function Types	17
1.8 Logarithms	22
Chapter 2 Number Systems	30
2.1 Binary Numbers	30
2.2 Octal and Hexadecimal	40
2.3 Converting Between Systems	42
Chapter 3 Set Theory	47
3.1 What is a Set?	47
3.2 Important Sets: The Number Systems	48
3.3 Relationships Between Sets	50
3.4 Properties of Sets	52
3.5 Operations on Sets	53
3.6 Cartesian Products and Tuples	56
3.7 Proving Set Equalities	57
3.8 Computer Representation of Sets	60
Chapter 4 Combinatorics and Probability Theory	61
4.1 Sample Space and Events	61
4.2 Counting Principles	64
4.3 Basic Probability	69
4.4 Probability of Joint Events and Set Operations	71
Chapter 5 Conditional Probability and Bayes' Theorem	75
5.1 Conditional Probability	75
5.2 Multiplication and Total Probability Rules	80
5.3 Independence	83
5.4 Bayes' Theorem	85
Chapter 6 Descriptive Statistics	89
6.1 Introduction to Descriptive Statistics	89
6.2 Describing Data Sets	91
6.3 Summarizing Data	97

6.4 Understanding Data Distributions	103
Chapter 7 Linear Equations in Linear Algebra	109
7.1 Systems of Linear Equations	109
7.2 Elementary Row Operations and Echelon Forms	115
Chapter 8 Vectors and Matrices	120
8.1 Vectors in \mathbb{R}^2 and \mathbb{R}^n	120
8.2 Matrix Equations	125
8.3 Solution Sets of Linear Systems	129
8.4 Linear Independence	133
Chapter 9 Matrix Algebra	136
9.1 Matrix Arithmetic	136
9.2 Matrix Transpose	142
9.3 Invertible Matrices	144
9.4 The Invertible Matrix Theorem	150
Chapter 10 Differential Calculus	152
10.1 Limits and Continuity: The Foundation of Calculus	153
10.2 The Derivative: Measuring the Rate of Change	155
10.3 Rules of Differentiation	156
10.4 Applications of Derivatives	164
Bibliography	171
Appendix A Important Concepts	172

Chapter 1 Basic Arithmetic and Functions

In the study of mathematics, a solid understanding of basic arithmetic operations and the concept of functions forms the foundation for more advanced topics. Arithmetic involves the manipulation of numbers through fundamental operations such as addition, subtraction, multiplication, and division. These operations are not only essential in everyday calculations but also serve as the building blocks for more complex mathematical procedures.

Functions, on the other hand, represent a crucial concept in mathematics, serving as a bridge between arithmetic and higher-level mathematical analysis. A function can be thought of as a special relationship between two sets, where each input (from the domain) is associated with exactly one output (from the co-domain). Understanding functions and their properties allows us to model and solve real-world problems with greater precision and flexibility.

In this section, we will explore the basic arithmetic rules and introduce the concept of functions, including their definitions, notations, and key properties. We will also discuss how these concepts are applied in various contexts, setting the stage for more advanced mathematical discussions.

1.1 Factorisation and the Order of Operations

Recall that when computing a product such as

$$a(b + c) = ab + ac$$

we distribute the factor a across each term inside the parentheses. We call this **the distributive property**. However, it is often advantageous or necessary to perform the reverse operation, known as factorisation. Factorisation involves expressing the sum $ab + ac$ in its factored form $a(b + c)$.

Mathematically, the expression $a(b + c)$ is considered more simplified or "better" than $ab + ac$. To understand why, we must distinguish between **terms** and **factors**. Terms are separated by addition or subtraction, while factors are separated by multiplication or division. For instance, the expression

$$8 - 5 + 3$$

consists of three terms (8, -5, and 3), whereas the expression

$$2 \times 5 \times 3$$

consists of three factors (2, 5, and 3). Consider the expression

$$5 + 2 + 7 \times 8$$

which consists of three terms (5, 2, and 7×8), and note that one of the terms itself contains two factors (7 and 8). Similarly, the expression

$$5(2 + 4)(-9)$$

consists of three factors (5, $2 + 4$, and -9), with one factor, $2 + 4$, containing two terms.

The advantage of expressing mathematical expressions solely in terms of factors lies in the ability to simplify them more effectively. This concept will be elaborated throughout this chapter, especially when dealing with fractions and equations. Consider the following examples of factorisation:

$$8 - 5 + 3$$

consists of three terms (8, -5 and 3), while the expression

$$2 \times 5 \times 3$$

consists of three factors (2, 5 and 3). The expression

$$5 + 2 + 7 \times 8$$

consists of three terms (5, 2 and 7×8) and one of the terms consists of two factors (7×8), while the expression

$$5(2 + 4)(-9)$$

consists of three factors (5, $(2 + 4)$, (-9)) and one of the factors consists of two terms ($(2 + 4)$).

It can be advantageous to express terms solely in factors because this allows us to simplify the expressions. This will become clearer as we progress through the lesson, and it is particularly important when working with fractions and equations. Here are some more examples of factorisation:

Example 1.1 Factorisation

$$ab - ac = a(b - c)$$

$$-ab - ac = a(-b - c)$$

$$-ab - ac = -a(b + c)$$

$$ab - ac + aa - aa = a(b - c + a - a)$$

$$abc - ab + aba = ab(c - 1 + a)$$

$$ab - abc - a = a(b - bc - 1) = b(1 - c) - 1$$

When evaluating mathematical expressions, it is crucial to follow a specific order of operations to ensure accurate results. The correct sequence for performing these operations is as follows:

- 1. Brackets (Parentheses):** First, perform all operations inside brackets or parentheses.
- 2. Exponents and Radicals:** Next, evaluate exponents (powers) and radicals (roots).
- 3. Multiplication and Division:** Then, perform multiplication and division from left to right as they appear.
- 4. Addition and Subtraction:** Finally, execute addition and subtraction from left to right as they appear.

Let's consider examples for each operation to illustrate the order of operations:

Example 1.2 Brackets

Evaluate the expression: $(2 + 3) \times 4$

$$(2 + 3) \times 4 = 5 \times 4 = 20$$

Example 1.3 Exponents and Radicals

Evaluate the expression: $3^2 + \sqrt{16}$

$$3^2 + \sqrt{16} = 9 + 4 = 13$$

Example 1.4 Multiplication and Division

Evaluate the expression: $6 \div 2 \times 3$. According to the standard order of operations, we perform division and multiplication from left to right:

$$6 \div 2 \times 3 = (6 \div 2) \times 3 = 3 \times 3 = 9$$

Example 1.5 Addition and Subtraction

Evaluate the expression: $8 - 3 + 2$

$$8 - 3 + 2 = 5 + 2 = 7$$

1.2 Fractions

By definition, a fraction always consists of (at least) two factors. The first factor we will call the **numerator** and is the “top part” of the fraction. The bottom part we will call the **denominator**. Perhaps the most important rule when working with fractions is that two fractions can only be added or subtracted if they have identical denominators. Also, the denominator must never be equal to 0.

Example 1.6

$$\frac{1}{x^2 - 2}$$

Here we must make sure that $x^2 - 2 \neq 0$ which means that we may only use values different from $\pm\sqrt{2}$.

Proposition 1.1 (Rules for Calculations with Fractions)

For $a, b, c, m \in \mathbb{R}$, with $a, b, c, m \neq 0$ where required, the following identities hold:

$$(1) \quad \frac{a}{b} \times m = \frac{am}{b}$$

$$(2) \quad \frac{a}{b} \div m = \frac{a}{bm}$$

$$(3) \quad m \div \frac{a}{b} = \frac{mb}{a}$$

$$(4) \quad \frac{a}{b} \times \frac{c}{a} = \frac{c}{b}$$

$$(5) \quad \frac{a}{b} \div \frac{c}{a} = \frac{a^2}{bc}$$

$$(6) \quad \frac{a}{b} = \frac{ac}{bc}$$

$$(7) \quad \frac{a}{b} + \frac{c}{a} = \frac{a^2 + bc}{ab}$$



To extend or reduce a fraction, we must multiply or divide by the same numbers in the denominator and numerator:

Example 1.7 Extending or reducing fractions

$$\frac{72}{144} = \frac{72 \div 12}{144 \div 12} = \frac{6}{12} = \frac{6 \div 6}{12 \div 6} = \frac{1}{2}$$

$$\frac{2}{3} = \frac{2 \times 6}{3 \times 6} = \frac{12}{18}$$

$$\frac{2x}{2xx} = \frac{2}{2x}$$

$$\frac{3a}{6a + 3b} = \frac{3a}{3(2a + b)} = \frac{a}{2a + b}$$

To factorise an expression, all terms must be divided or multiplied uniformly. This implies that it is not possible to simplify the following expression any further, even though it might be tempting:

$$\frac{a}{2a + b} \neq \frac{1}{2 + b}$$

For proper factorisation of $\frac{a}{2a+b}$, you must divide a into all terms in the denominator:

$$\frac{a}{2a + b} = \frac{1}{2 + \frac{b}{a}}$$

As illustrated above, the multiplication of fractions is straightforward: you multiply the numerators and denominators with each other, respectively.

Example 1.8 Multiplication of fractions

Consider the fractions $\frac{a}{b}$ and $\frac{c}{d}$. Their product is:

$$\frac{a}{b} \times \frac{c}{d} = \frac{a \times c}{b \times d}$$

For instance, if $a = 2$, $b = 3$, $c = 4$, and $d = 5$, then:

$$\frac{2}{3} \times \frac{4}{5} = \frac{2 \times 4}{3 \times 5} = \frac{8}{15}$$

Example 1.9 Dividing fractions

$$\frac{6}{7} \div \frac{4}{21} = \frac{6}{7} \times \frac{21}{4} = \frac{126}{28} = \frac{63}{14} = \frac{9}{2}$$

$$\frac{1}{2} \div 3 = \frac{1}{2} \div \frac{3}{1} = \frac{1}{2} \times \frac{1}{3} = \frac{1}{6}$$

$$\frac{2}{3} \div \frac{8}{9} = \frac{2}{3} \times \frac{9}{8} = \frac{18}{24} = \frac{3}{4}$$

$$\frac{8}{9} \div 16 = \frac{8}{9} \times \frac{1}{16} = \frac{8}{144} = \frac{4}{72} = \frac{1}{18}$$

Adding and subtracting fractions seems to cause more problems than multiplication and division. The key is to find a common denominator between the fractions and then remember the above-mentioned rule about extending fractions.

Example 1.10 Adding and subtracting fractions

$$\frac{1}{5} + \frac{2}{5} = \frac{1+2}{5} = \frac{3}{5}$$

$$\frac{1}{4} + \frac{2}{3} = \frac{1 \times 3}{4 \times 3} + \frac{2 \times 4}{3 \times 4} = \frac{3}{12} + \frac{8}{12} = \frac{3+8}{12} = \frac{11}{12}$$

$$\frac{7}{12} - \frac{5}{8} = \frac{7 \times 8}{12 \times 8} - \frac{5 \times 12}{8 \times 12} = \frac{56}{96} - \frac{60}{96} = \frac{56-60}{96} = \frac{-4}{96} = \frac{-1}{24}$$

Note: Never do

$$\frac{a}{b} + \frac{c}{d} = \frac{a+b}{b+d}$$

Example 1.11

$$\frac{x}{x-1} \times \frac{2}{x(x+4)} = \frac{2x}{(x-1)x(x+4)} = \frac{2}{(x-1)(x+4)}$$

$$\frac{2}{x-1} \div \frac{x}{x-1} = \frac{2}{x-1} \times \frac{x-1}{x} = \frac{2(x-1)}{(x-1)x} = \frac{2}{x}$$

$$\frac{x+1}{x^2+2} + \frac{x-6}{x^2+2} = \frac{x+1+x-6}{x^2+2} = \frac{2x-5}{x^2+2}$$

Remember, when a fraction is preceded by a minus sign, all signs in the numerator must be changed accordingly. This is similar to how you would change all signs within parentheses when they are preceded by a minus sign.

Consider the expression:

$$-\frac{a-b}{c}$$

To correctly handle the negative sign, change all the signs in the numerator:

$$-\frac{a-b}{c} = \frac{-a+b}{c}$$

For example, if $a = 5$ and $b = 3$, then:

$$-\frac{5-3}{c} = \frac{-5+3}{c} = \frac{-2}{c}$$

Example 1.12

$$\frac{x+1}{x^2+2} - \frac{x-6}{x^2+2} = \frac{x+1-x+6}{x^2+2} = \frac{7}{x^2+2}$$

1.3 Exponents, Radicals and Surds

An **exponent** is a shortcut for repeated multiplication of the same number:

Example 1.13 Exponentiation

$$4 \times 4 \times 4 \times 4 \times 4 = 4^5$$

$$x \times x \times x \times x \times x = x^5$$

Radicals, or **roots**, represent the inverse operation of applying exponents. A radical is any number expressed with the radical symbol $\sqrt{}$. Specifically, applying a radical can reverse the effect of an exponent, and vice versa. For instance, squaring 2 yields 4, and taking the square root of 4 returns 2. Similarly, squaring 3 results in 9, and the square root of 9 brings us back to 3.

Example 1.14 Taking the root

$$\sqrt{a} \times \sqrt{a} = (\sqrt{a})^2 = a$$

$$\sqrt{a} = b \implies (\sqrt{a})^2 = b^2 \iff a = b^2$$

A **surd** is a type of radical that is both real and irrational, examples include $\sqrt{2}$, $\sqrt{3}$, $\sqrt{5}$, and $\sqrt{6}$.

Numbers can be raised to powers other than 2, such as cubing (raising to the third power), or even raising to the fourth power, the 100th power, and so forth. Correspondingly, you can take the cube root of a number, the fourth root, the 100th root, and so on. To indicate a root other than a square root, the same radical symbol is used, but with a number called the index inserted into the radical sign, typically positioned within the "check mark" part.

Example 1.15 Index and argument

$$4^3 = 64 \iff \sqrt[3]{64} = 4$$

In this example, the "3" inside the radical sign is the **index** of the radical. The "64" is referred to as the argument of the **radical**, also known as the **radicand**. Since square roots are the most common type of radicals, the index is usually omitted for square roots. Although " $\sqrt[2]{2}$ " would be technically correct, it is rarely used in practice.

Proposition 1.2 (Rules for calculations involving radicals)

- (1) $\sqrt[n]{xy} = \sqrt[n]{x} \times \sqrt[n]{y}$ where $x, y \geq 0$
- (2) $\sqrt[n]{\frac{x}{y}} = \frac{\sqrt[n]{x}}{\sqrt[n]{y}}$ where $x \geq 0$ and $y > 0$
- (3) $\sqrt{x^2} = |x|$ where $x \in \mathbb{R}$
- (4) $(\sqrt[n]{x})^n = x$ If $x < 0$ and $n \in \mathbb{N}$, then $\sqrt[n]{x}$ is not defined
- (5) $\sqrt[n]{-x} = -\sqrt[n]{x}$ where $x \geq 0$ and $n \in \mathbb{N}$ is odd



Raising a number to a **power**, also known as **exponentiation**, is a fundamental mathematical operation that involves multiplying a number by itself a certain number of times as we saw above. The **base** is the number being multiplied, and the **exponent** indicates how many times the base is used as a factor. For example, a^n means that the base a is multiplied by itself n times. Exponentiation is a powerful tool in mathematics, with a few essential rules that govern its application.

Proposition 1.3 (Properties of Integer Exponents)

Let $n, m \in \mathbb{Z}$. Then the following hold (with $x, y \in \mathbb{R}$ and nonzero where stated):

- (1) $x^n \cdot x^m = x^{n+m}$,
- (2) $\frac{x^n}{x^m} = x^{n-m}$ with $x \neq 0$,
- (3) $x^n \cdot y^n = (xy)^n$,
- (4) $\frac{x^n}{y^n} = \left(\frac{x}{y}\right)^n$ with $y \neq 0$,
- (5) $(x^n)^m = x^{nm}$,
- (6) $x^1 = x$.



Some of these rules allow the concept of powers to be extended so that the exponent may be any integer. If you set $n = m$ in rule (2), you get:

$$\frac{x^n}{x^n} = x^{n-n} = x^0$$

But since $\frac{x^n}{x^n} = 1$, we obtain

$$x^0 = 1$$

Thus, the concept of exponentiation is extended to include $n \in \mathbb{N} \cup \{0\}$. If you now set $n = 0$ again in rule (2), you get:

$$\frac{x^0}{x^m} = x^{0-m} = x^{-m}$$

But according to the previous calculation, $x^0 = 1$. Therefore, you obtain

$$\frac{1}{x^m} = x^{-m}$$

As demonstrated in **Proposition 1.3**, these rules provide a complete framework for manipulating expressions with any integer exponents. As a consequence, the following additional rules can be derived:

Proposition 1.4 (More Properties of Integer Exponents)

Let $n, m \in \mathbb{Z}$. Then the following hold (with $x, y \in \mathbb{R}$ and nonzero where stated):

$$(7) \quad x^0 = 1 \quad x \neq 0$$

$$(8) \quad \frac{1}{x^m} = x^{-m} \quad x \neq 0$$



Let us illustrate these rules with a couple of examples.

Example 1.16 Reduce the following expression

$$(3xy^6)^3 = 3^3 \cdot x^3 \cdot (y^6)^3 = 27x^3y^{18}$$

Example 1.17 Reduce the following expression

$$\frac{a^{-4}b^3}{a^7b^{-5}} = \frac{a^{-4}}{a^7} \cdot \frac{b^3}{b^{-5}} = a^{-4-7} \cdot b^{3-(-5)} = a^{-11} \cdot b^8 = \frac{b^8}{a^{11}}$$

For positive numbers x , the concept of exponentiation can be further extended to apply when the exponent is a rational number. Any rational number $r \in \mathbb{Q}$ can be written as $r = \frac{m}{n}$, where $m \in \mathbb{Z}$ and $n \in \mathbb{N}$. For $x > 0$, we now define

$$y = x^r = x^{\frac{m}{n}}$$

From this, you obtain (using, among other things, rule (5)):

$$y^n = \left(x^{\frac{m}{n}}\right)^n = x^{\frac{m}{n} \cdot n} = x^m$$

Finally, by using the concept of radicals, we obtain

$$y^n = x^m \iff y = \sqrt[n]{x^m}$$

Note that because $x > 0$, it follows that $y > 0$ as well. We are now ready to state ***the extended concept of exponentiation***.

Definition 1.1 (The Extended Concept of Exponentiation)

Let $m \in \mathbb{Z}$ and let $n \in \mathbb{N}$ such that $\frac{m}{n} \in \mathbb{Q}$. Then the following applies:

$$x^{\frac{m}{n}} = \sqrt[n]{x^m} \quad x > 0$$

And more specifically, the following holds true

$$x^{\frac{1}{n}} = \sqrt[n]{x} \quad x > 0$$



The denominator of a rational exponent corresponds to the index of the radical, while the numerator remains as the exponent of the base. Conversely, the index of a radical can be transformed into the denominator of an exponent in an equivalent exponential expression. This property allows us to convert any radical expression into an exponential form, providing a powerful tool for simplification.

Example 1.18

$$\sqrt[5]{x^3} = x^{\frac{3}{5}} \quad \text{vs.} \quad \sqrt[3]{x^5} = x^{\frac{5}{3}}$$

$$\frac{1}{\sqrt[7]{x^3}} = x^{-\frac{3}{7}}$$

$$\frac{1}{\sqrt[3]{x^2}} = (x^2)^{-\frac{2}{3}}$$

This property can also be reversed: any rational exponent can be rewritten as a radical expression by using the denominator as the radical's index. The ability to interchange between exponential and radical forms enables us to evaluate expressions that were previously difficult to handle by converting them into radicals.

Example 1.19

$$27^{-\frac{4}{3}} = \frac{1}{\sqrt[3]{27^4}} = \frac{1}{(\sqrt[3]{27})^4} = \frac{1}{3^4} = \frac{1}{81}$$

One of the greatest advantages of converting a radical expression into an exponential form is that it allows us to apply all the properties of exponents to simplify the expression. The following examples illustrate how various properties can be utilised to simplify expressions with rational exponents.

Example 1.20

$$a^{\frac{2}{3}}b^{\frac{1}{2}}a^{\frac{1}{6}}b^{\frac{1}{5}} = a^{\frac{2}{3} + \frac{1}{6}}b^{\frac{1}{2} + \frac{1}{5}} = a^{\frac{4}{6} + \frac{1}{6}}b^{\frac{5}{10} + \frac{2}{10}} = a^{\frac{5}{6}}b^{\frac{7}{10}}$$

$$(x^{\frac{1}{3}}x^{\frac{2}{5}})^{\frac{3}{4}} = x^{\frac{1}{3} \times \frac{3}{4}}x^{\frac{2}{5} \times \frac{3}{4}} = x^{\frac{3}{12}}x^{\frac{6}{20}} = x^{\frac{1}{4}}x^{\frac{3}{10}} = x^{\frac{11}{20}}$$

$$\frac{x^{\frac{4}{2}}x^{\frac{4}{6}}x^{\frac{1}{2}}x^{\frac{5}{6}}}{x^{\frac{7}{2}}x^0} = 2x^{\frac{4}{2} + \frac{1}{2}}x^{\frac{4}{6} + \frac{5}{6}}x^{\frac{7}{2}} = 2x^{\frac{5}{2}}x^{\frac{9}{6}}x^{\frac{7}{2}} = 2x^{-1}x^{\frac{3}{2}} = 2x^{\frac{1}{2}}$$

$$(25x^{\frac{1}{3}}x^{\frac{2}{5}})^{-\frac{1}{2}} = (25x^{\frac{5}{15}}x^{\frac{4}{10}})^{-\frac{1}{2}} = (25x^{-\frac{7}{15}}x^{\frac{19}{10}})^{-\frac{1}{2}} = \left(\frac{9}{25x^{-\frac{7}{15}}x^{\frac{19}{10}}}\right)^{\frac{1}{2}} = \frac{9}{2} \cdot 25x^{-\frac{7}{30}}x^{\frac{19}{20}} = \frac{3x^{\frac{7}{30}}}{5x^{\frac{19}{20}}}$$

It is important to remember that when simplifying expressions with rational exponents, we are applying the same exponent rules that are used for integer exponents. The only difference is that we must also adhere to the rules for fractions.

1.4 Using Formulae and Substitution

In the study of engineering, physical quantities are often related to each other through formulas. These formulas consist of variables and constants that represent the physical quantities in question. To evaluate a formula, one must substitute numerical values for the variables.

For example, Ohm's law provides a formula that relates the voltage, v , across a resistor with a resistance value R , to the current i flowing through it. The formula is given by

$$v = iR$$

This formula allows us to calculate the voltage v if the values for i and R are known. For instance, if $i = 13 \text{ A}$ and $R = 5 \Omega$, then

$$v = iR = (13)(5) = 65$$

Thus, the voltage is 65 V.

This example highlights the importance of paying close attention to the units of any physical quantities involved. A formula is only valid if a consistent set of units is used.

Example 1.21 Inserting into formulae

The kinetic energy K of an object with mass M moving at speed v can be calculated using the formula:

$$K = \frac{1}{2}Mv^2$$

Calculate the kinetic energy of an object with a mass of 5 kg moving at a speed of 2 m s⁻¹.

Solution:

$$K = \frac{1}{2}Mv^2 = \frac{1}{2}(5)(2^2) = 10$$



In the SI system, the unit of energy is the joule, so the kinetic energy of the object is 10 joules.

Example 1.22 Inserting into formulae

The area A of a circle with radius r can be calculated using the formula $A = \pi r^2$.

Alternatively, if the diameter d of the circle is known, the equivalent formula can be used:

$$A = \frac{\pi d^2}{4}$$

Calculate the area of a circle with a diameter of 0.1 m. The value of π is pre-programmed in your calculator.

Solution:

$$A = \frac{\pi(0.1)^2}{4} = 0.00785 \text{ m}^2$$



Example 1.23 Inserting into formulae

The volume V of a circular cylinder is equal to its cross-sectional area A multiplied by its length h .

Calculate the volume of a cylinder with a diameter of 0.1 m and a length of 0.3 m.

Solution:

$$V = Ah = \frac{\pi(0.1)^2}{4} \times 0.3 = 0.00236$$

The volume is 0.00236 m³.



1.5 Rearranging Formulae

In the formula for the area of a circle, $A = \pi r^2$, the variable A is referred to as the subject of the formula. A variable is considered the subject if it appears by itself on one side of the equation, usually on the left-hand side, and nowhere else in the formula. If we are asked to transpose the formula for r , or solve for r , we must rearrange the equation so that r becomes the subject. When transposing a formula, any operation performed on one side must also be applied to the other side. There are five key rules to follow during this process.

Rules for rearranging formulae

The following operations can be performed on both sides of the formula:

- Add the same quantity to both sides
- Subtract the same quantity from both sides
- Multiply both sides by the same quantity - remember to multiply all terms
- Divide both sides by the same quantity - remember to divide all terms
- Apply a function to both sides, such as squaring or finding the reciprocal

Example 1.24 Transpose the formula $p = 5t - 17$ to make t the subject.

Solution: To isolate t on the left-hand side, proceed in steps using the five rules. First, add 17 to both sides of the equation $p = 5t - 17$:

$$p + 17 = 5t - 17 + 17$$

Simplifying, we get:

$$p + 17 = 5t$$

Next, divide both sides by 5 to isolate t :

$$\frac{p + 17}{5} = t$$

Thus, the formula for t is:

$$t = \frac{p + 17}{5}$$



Example 1.25 Transpose the formula $\sqrt{2q} = p$ to solve for q .

Solution: First, square both sides to eliminate the square root around $2q$. Note that $(\sqrt{2q})^2 = 2q$. This gives:

$$2q = p^2$$

Next, divide both sides by 2 to solve for q :

$$q = \frac{p^2}{2}$$



Problem 1.1 Transpose the formula $v = \sqrt{t^2 + w}$ to solve for w . To isolate w , follow these steps:

- First, square both sides to eliminate the square root around $t^2 + w$:

$$v^2 = t^2 + w$$

- Next, subtract t^2 from both sides to isolate w :

$$v^2 - t^2 = w$$

- Finally, write down the formula for w :

$$w = v^2 - t^2$$

Example 1.26 Transpose the formula $x = \frac{1}{y}$ to solve for y .

Solution: To isolate y , notice that y appears in the denominator. Multiplying both sides by y removes the fraction:

$$yx = y \times \frac{1}{y}$$

This simplifies to:

$$yx = 1$$

Finally, divide both sides by x to solve for y :

$$y = \frac{1}{x}$$

Alternatively, you can simply invert both sides directly to obtain:

$$y = \frac{1}{x}$$



Example 1.27 Make R the subject of the formula:

$$\frac{2}{R} = \frac{3}{x+y}$$

Solution: Since R appears in a fraction, invert both sides:

$$\frac{R}{2} = \frac{x+y}{3}$$

Multiplying both sides by 2 yields: $R = \frac{2(x+y)}{3}$



Example 1.28 Make R the subject of the formula:

$$\frac{1}{R} = \frac{1}{R_1} + \frac{1}{R_2}$$

Solution: The two terms on the right-hand side can be combined:

$$\frac{1}{R_1} + \frac{1}{R_2} = \frac{R_2 + R_1}{R_1 R_2}$$

The formula then becomes:

$$\frac{1}{R} = \frac{R_2 + R_1}{R_1 R_2}$$

Finally, inverting both sides gives:

$$R = \frac{R_1 R_2}{R_2 + R_1}$$



1.6 Functions

In mathematics, a function assigns each element of one set to a specific element of another set (which may be the same set). For example, consider a Mathematics for Software Engineering class where each student is assigned a grade from the set $\{12, 10, 7, 4, 02\}$. Suppose the grades are as follows:

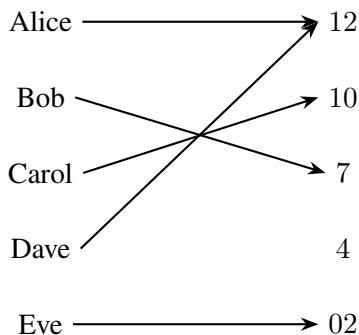


Figure 1.1: Example of a function mapping names to numbers.

This assignment of grades, illustrated in [Figure 1.1](#), exemplifies a function.

Functions play a crucial role in mathematics and computer science. They define discrete structures such as sequences and strings and are used to analyse the time complexity of algorithms. Many computer programs are designed to compute values of functions. Recursive functions, defined in terms of themselves, are especially significant in computer science. This section provides an overview of the fundamental concepts of functions needed in the mathematics for software engineering.

Definition 1.2

Let A and B be nonempty sets. A function f from A to B is an assignment of exactly one element of B to each element of A . We write $f(a) = b$ if b is the unique element of B assigned by the function f to the element a of A . If f is a function from A to B , we write $f : A \rightarrow B$.



Remark: Functions are sometimes also called mappings or transformations.

Functions can be specified in various ways. Sometimes, we explicitly state the assignments, as shown in [Figure 1.1](#). Often, a formula such as $f(x) = x + 1$ is used to define a function. In other cases, a computer program may specify the function.

Definition 1.3

If f is a function from A to B , we say that A is the **domain** of f and B is the **co-domain** of f . If $f(a) = b$, we say that b is the **image** of a and a is a **preimage** of b . The **range**, or image, of f is the set of all images of elements of A . Also, if f is a function from A to B , we say that f **maps** A to B .



When defining a function, we specify its domain, co-domain, and the mapping of elements from the domain to the co-domain. Two functions are equal if they have the same domain, the same co-domain, and map each element of their domain to the same element in the co-domain.

It's important to note that altering the domain or co-domain results in a different function. Similarly, changing the mapping of elements also produces a different function.

The following examples illustrate various functions. In each example, we describe the domain, co-domain, range, and the assignment of values to the elements of the domain.

Example 1.29 What are the domain, co-domain, and range of the function that assigns grades to students described in the first paragraph of the introduction of this section?

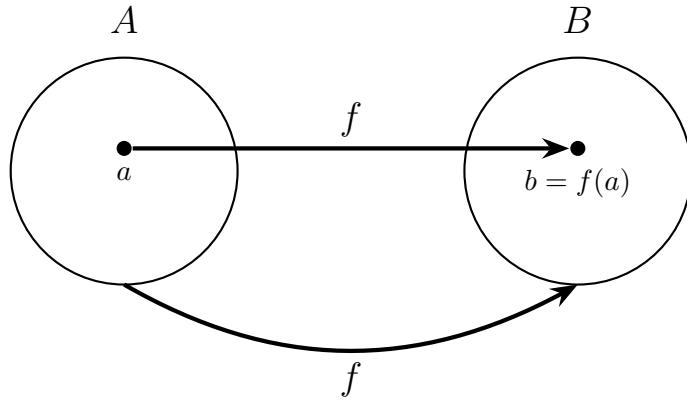


Figure 1.2: A function f mapping an element a from set A to an element $b = f(a)$ in set B .

Solution: Let G be the function that assigns a grade to a student in our Software engineering mathematics class. Note that $G(\text{Alice}) = 12$, for instance. The domain of G is the set {Alice, Bob, Carol, David, Eve}, and the co-domain is the set {12, 10, 7, 4, 02}. The range of G is the set {12, 10, 7, 02}, because each grade except 4 is assigned to some student. ◀

Example 1.30 Let f be the function that assigns the last two bits of a bit string of length 2 or greater to that string. For example, $f(11010) = 10$. Then, the domain of f is the set of all bit strings of length 2 or greater, and both the co-domain and range are the set {00, 01, 10, 11}.

Example 1.31 Let $f : \mathbb{Z} \rightarrow \mathbb{Z}$ assign the square of an integer to this integer. Then, $f(x) = x^2$, where the domain of f is the set of all integers, the co-domain of f is the set of all integers, and the range of f is the set of all integers that are perfect squares, namely, $\{0, 1, 4, 9, \dots\}$.

One-to-One and Onto Functions

In mathematics, functions are a fundamental concept used to describe the relationship between two sets. However, not all functions behave the same way. To understand these differences, we introduce the concepts of one-to-one (injective) and onto (surjective) functions.

Some functions never assign the same value to two different domain elements. These functions are said to be **one-to-one**.

Definition 1.4 (One-to-One functions (Injective))

A function $f : A \rightarrow B$ is called **one-to-one** (or **injective**) if different elements in A map to different elements in B . In other words, if $f(a_1) = f(a_2)$, then $a_1 = a_2$. This property ensures that no two distinct elements in A are mapped to the same element in B .



Graphically, a function is one-to-one if no horizontal line intersects the graph of the function at more than one point.

Definition 1.5 (Onto Functions (Surjective))

A function $f : A \rightarrow B$ is called **onto** (or **surjective**) if every element in B is the image of at least one element in A . In other words, for every $b \in B$, there exists at least one $a \in A$ such that $f(a) = b$. This

property ensures that the function “covers” the entire set B .



Inverse Functions

Now, consider a function $f : A \rightarrow B$ that is both one-to-one and onto. Because f is onto, every element of B is the image of some element in A . Furthermore, because f is one-to-one, every element of B is the image of a unique element of A . This unique correspondence allows us to define a new function from B to A that “reverses” the mapping given by f .

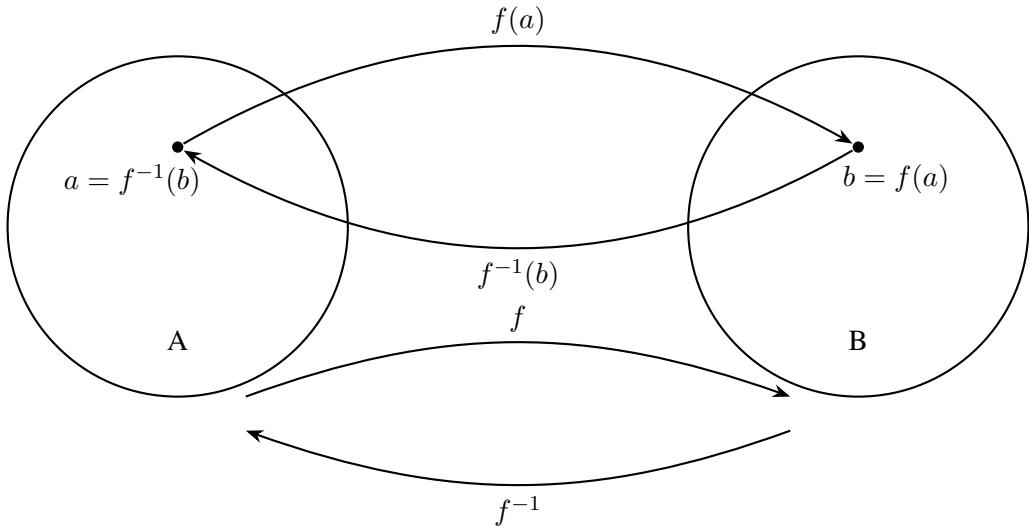


Figure 1.3: The function f^{-1} is the inverse of function f .

This new function is called the **inverse function** of f , denoted by $f^{-1} : B \rightarrow A$. The inverse function f^{-1} satisfies the following properties:

$$f(f^{-1}(b)) = b \quad \text{for every } b \in B.$$

$$f^{-1}(f(a)) = a \quad \text{for every } a \in A.$$

We can summarise these considerations in the following definition

Definition 1.6 (Inverse Functions)

Let f be a one-to-one correspondence from the set A to the set B . The inverse function of f is the function that assigns to an element b belonging to B the unique element a in A such that $f(a) = b$. The inverse function of f is denoted by f^{-1} . Hence, $f^{-1}(b) = a$ when $f(a) = b$.



These properties show that f^{-1} effectively undoes the work of f , mapping each element of B back to the corresponding element in A .

Example 1.32

Let $f : \mathbb{R} \rightarrow \mathbb{R}$ be defined by $f(x) = 2x + 3$. We can check that f is both one-to-one and onto:

- **One-to-One:** If $f(x_1) = f(x_2)$, then $2x_1 + 3 = 2x_2 + 3$. Subtracting 3 from both sides gives $2x_1 = 2x_2$, and dividing by 2 yields $x_1 = x_2$. Thus, f is one-to-one.

- **Onto:** Given any $y \in \mathbb{R}$, we can solve $y = 2x + 3$ for x to find $x = \frac{y-3}{2}$. Since this x exists for every y , f is onto.

Since f is both one-to-one and onto, it has an inverse function f^{-1} defined by

$$f^{-1}(y) = \frac{y-3}{2}.$$

Composite Functions

In mathematics, functions can be combined to form new functions. One important way of combining functions is through the composition of functions. The composite of two functions is essentially applying one function to the results of another.

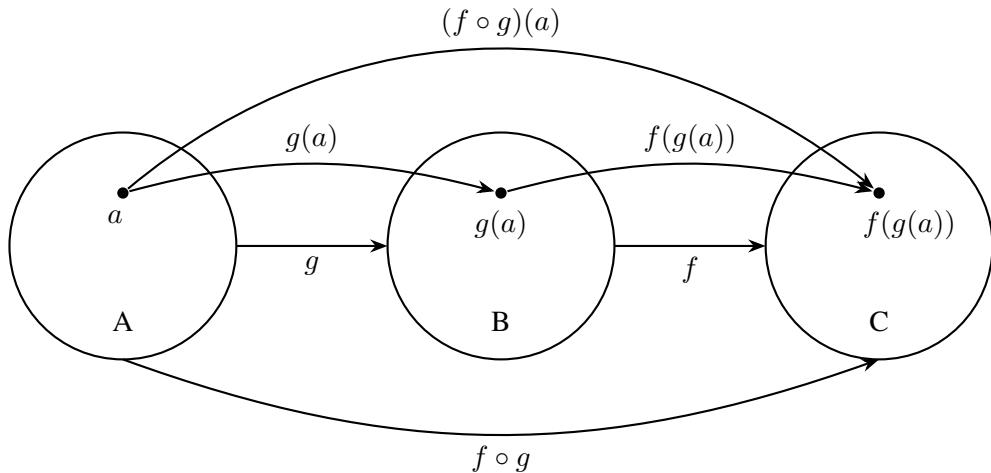


Figure 1.4: The composition of f and g .

Definition 1.7

Let $f : B \rightarrow C$ and $g : A \rightarrow B$ be two functions. The **composite function** of f and g , denoted by $f \circ g$, is a function from A to C defined by

$$(f \circ g)(x) = f(g(x)),$$

for every $x \in A$.



In other words, the composite function $f \circ g$ means that you first apply the function g to the input x , and then apply the function f to the result of $g(x)$. In **Figure 1.4** the composition of functions is shown.

Example 1.33

Consider the functions $f(x) = 2x + 3$ and $g(x) = x^2$. The composite function $f \circ g$ is given by:

$$(f \circ g)(x) = f(g(x)) = f(x^2) = 2x^2 + 3.$$

Here, the function $g(x)$ squares the input x , and then the function $f(x)$ multiplies the result by 2 and adds 3.

Now, let's reverse the composition and compute $g \circ f$:

$$(g \circ f)(x) = g(f(x)) = g(2x + 3) = (2x + 3)^2.$$

Notice that $f \circ g$ and $g \circ f$ are generally different functions, illustrating that the composition of functions is not commutative.

Example 1.34 Let g be the function from the set $\{a, b, c\}$ to itself such that $g(a) = b$, $g(b) = c$, and $g(c) = a$. Let f be the function from the set $\{a, b, c\}$ to the set $\{1, 2, 3\}$ such that $f(a) = 3$, $f(b) = 2$, and $f(c) = 1$. What is the composition of f and g , and what is the composition of g and f ?

Solution: The composition $f \circ g$ is defined by $(f \circ g)(a) = f(g(a)) = f(b) = 2$, $(f \circ g)(b) = f(g(b)) = f(c) = 1$, and $(f \circ g)(c) = f(g(c)) = f(a) = 3$.

Note that $g \circ f$ is not defined, because the range of f is not a subset of the domain of g . ◀

Example 1.35 label: Let f and g be the functions from the set of integers to the set of integers defined by $f(x) = 2x + 3$ and $g(x) = 3x + 2$. What is the composition of f and g ? What is the composition of g and f ?

Solution: Both the compositions $f \circ g$ and $g \circ f$ are defined. Moreover,

$$(f \circ g)(x) = f(g(x)) = f(3x + 2) = 2(3x + 2) + 3 = 6x + 7$$

and

$$(g \circ f)(x) = g(f(x)) = g(2x + 3) = 3(2x + 3) + 2 = 6x + 11.$$
◀

Remark: Even though $f \circ g$ and $g \circ f$ are defined for the functions f and g in example 35, $f \circ g$ and $g \circ f$ are not equal. In other words, the commutative law does not hold for the composition of functions.

When the composition of a function and its inverse is formed, in either order, an identity function is obtained. To see this, suppose that f is a one-to-one correspondence from the set A to the set B . Then the inverse function f^{-1} exists and is a one-to-one correspondence from B to A . The inverse function reverses the correspondence of the original function, so $f^{-1}(b) = a$ when $f(a) = b$, and $f(a) = b$ when $f^{-1}(b) = a$.

Hence,

$$(f^{-1} \circ f)(a) = f^{-1}(f(a)) = f^{-1}(b) = a,$$

and

$$(f \circ f^{-1})(b) = f(f^{-1}(b)) = f(a) = b.$$

Consequently, $f^{-1} \circ f = I_A$ and $f \circ f^{-1} = I_B$, where I_A and I_B are the identity functions on the sets A and B , respectively. That is, $(f^{-1})^{-1} = f$.

Example 1.36 If $f : \mathbb{R} \rightarrow \mathbb{R}$ is defined as $f(x) = 2x + 3$, then $f^{-1}(x) = \frac{x-3}{2}$. The composition $f \circ f^{-1}$ would be the identity function $I_{\mathbb{R}}$ on the real numbers, meaning $f(f^{-1}(x)) = x$ for all $x \in \mathbb{R}$.

1.7 Graphical Identification of Function Types

Understanding the behaviour of different types of functions is fundamental in mathematics. Functions can be classified based on their graphical patterns, which provide valuable insights into their characteristics. In this section, we will explore various types of functions, including linear, quadratic, exponential, and more. By examining their graphs, we can identify key features such as intercepts, slopes, curvature, and asymptotic behaviour, enabling us to distinguish between these different types of functions effectively.

Linear Functions

The general equation for a linear function is given by

$$y = ax + b \quad (\text{often written as } y = mx + b),$$

where a (or m) represents the slope and b is the y -intercept. The domain of this function is all real numbers. This equation is in slope-intercept form because a (or m) gives the slope and b gives the y -intercept. If $a = 0$, the function simplifies to $y = b$, which is a constant function.

The parent function for a linear equation is

$$y = x.$$

The transformed function can be written in the point-slope form as

$$y = y_1 + a(x - x_1),$$

where the graph contains the point (x_1, y_1) and has slope a . In this form:

- a is the vertical dilation (slope),
- y_1 represents the vertical translation,
- x_1 represents the horizontal translation.

This point-slope form can also be written as

$$y - y_1 = a(x - x_1),$$

where the coordinates of the fixed point (x_1, y_1) appear with a negative sign. The form $y = y_1 + a(x - x_1)$ expresses y explicitly in terms of x , making it easier to enter into a graphing calculator.

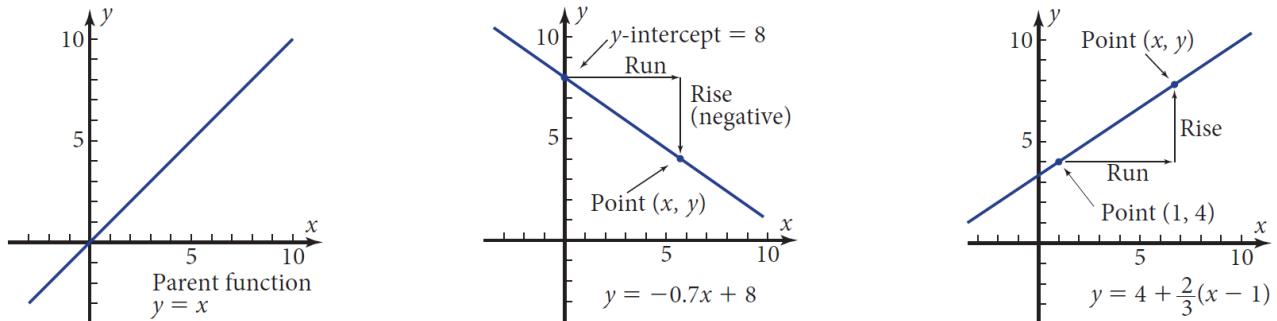


Figure 1.5: Linear functions

The graph of a linear function is a straight line. The parent function $y = x$ is shown on the left in Figure 1.5, the slope-intercept form in the middle, and the point-slope form on the right.

For the slope-intercept form: "Start at b on the y -axis, move x units horizontally, and rise ax units vertically."
For the point-slope form: "Start at (x_1, y_1) , move $(x - x_1)$ units horizontally, and rise $a(x - x_1)$ units vertically."

Quadratic Functions

The general equation for a quadratic function is given by

$$y = ax^2 + bx + c,$$

where $a \neq 0$, and a , b , and c are constants. The domain of this function is all real numbers.

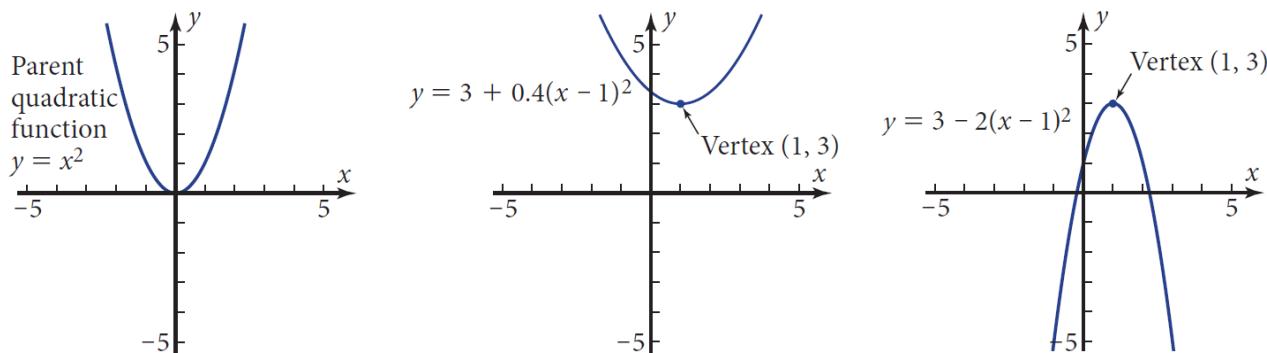


Figure 1.6: Quadratic functions

The parent function for a quadratic equation is

$$y = x^2,$$

where the vertex of the parabola is at the origin $(0, 0)$.

The transformed function can be written in vertex form as

$$y = k + a(x - h)^2,$$

where the vertex of the parabola is located at (h, k) . In this form:

- k represents the vertical translation,
- h represents the horizontal translation,
- a represents the vertical dilation.

Vertex form can also be written as

$$y - k = a(x - h)^2,$$

but expressing y explicitly in terms of x makes the equation easier to enter into a graphing calculator.

The graph of a quadratic function is a parabola (from the Greek word for "along the path of a ball"). The parabola is concave up if $a > 0$ and concave down if $a < 0$. This behaviour is illustrated in [Figure 1.6](#).

Power Functions

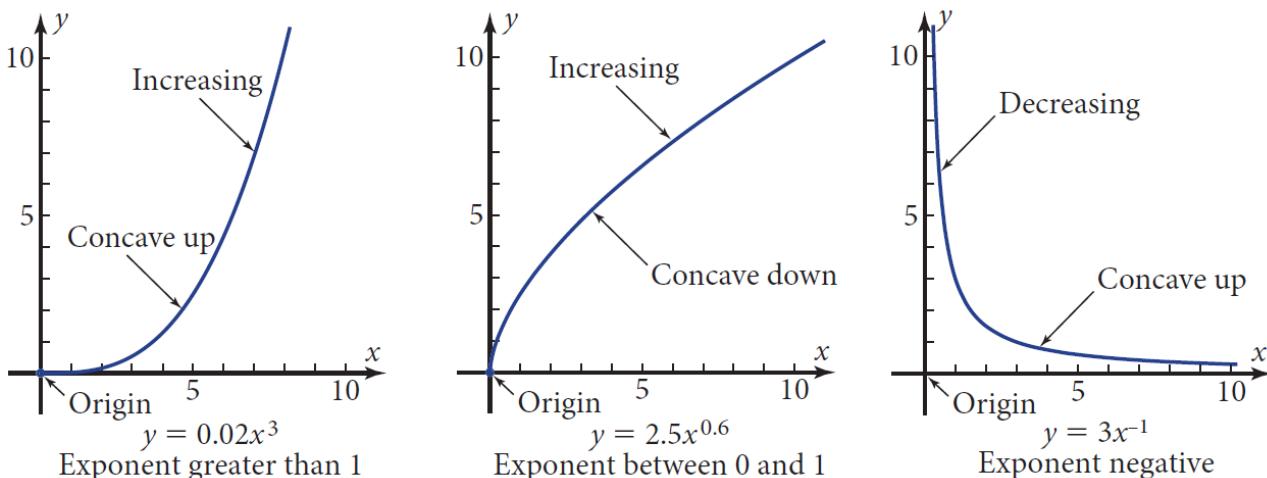
The general equation for a power function is given by

$$y = ax^b,$$

where a and b are nonzero constants. The domain of the function depends on the value of b :

- If $b > 0$, the domain is all real numbers.
- If $b < 0$, the domain excludes $x = 0$ to avoid division by zero.
- If b is not an integer, the domain usually excludes negative numbers to avoid taking roots of negative numbers.

In most applications, the domain is restricted to non-negative numbers.

**Figure 1.7:** Power functions

The parent function for a power function is

$$y = x^b.$$

For the general power function $y = ax^b$:

- If $b > 0$, then y varies directly with the b th power of x , meaning y is directly proportional to the b th power of x .
- If $b < 0$, then y varies inversely with the b th power of x , meaning y is inversely proportional to the b th power of x .

The dilation factor a serves as the proportionality constant.

The translated form of a power function is

$$y = d + a(x - c)^b,$$

where c and d are the horizontal and vertical translations, respectively. This can be compared with the translated forms of linear and quadratic functions:

$$\begin{aligned} y &= y_1 + a(x - x_1) && \text{(linear function),} \\ y &= k + a(x - h)^2 && \text{(quadratic function).} \end{aligned}$$

Unless otherwise stated, "power function" will imply the untranslated form, $y = ax^b$.

Figure 1.7 shows the graphs of power functions for different values of b . In all cases, $a > 0$. The shape and concavity of the graph depend on the value of b :

- If $b > 0$, the graph contains the origin.
- If $b < 0$, the graph has the axes as asymptotes.
- The function is increasing if $b > 0$ and decreasing if $b < 0$.
- The graph is concave up if $b > 1$ or $b < 0$, and concave down if $0 < b < 1$.

The concavity of the graph describes the rate at which y increases. For $b > 0$, concave up indicates that y is increasing at an increasing rate, while concave down indicates that y is increasing at a decreasing rate.

Exponential Functions

The general equation for an exponential function is given by

$$y = ab^x,$$

where a and b are constants, $a \neq 0$, $b > 0$, and $b \neq 1$. The domain of this function is all real numbers.

The parent function for an exponential equation is

$$y = b^x,$$

where the asymptote is the x -axis.

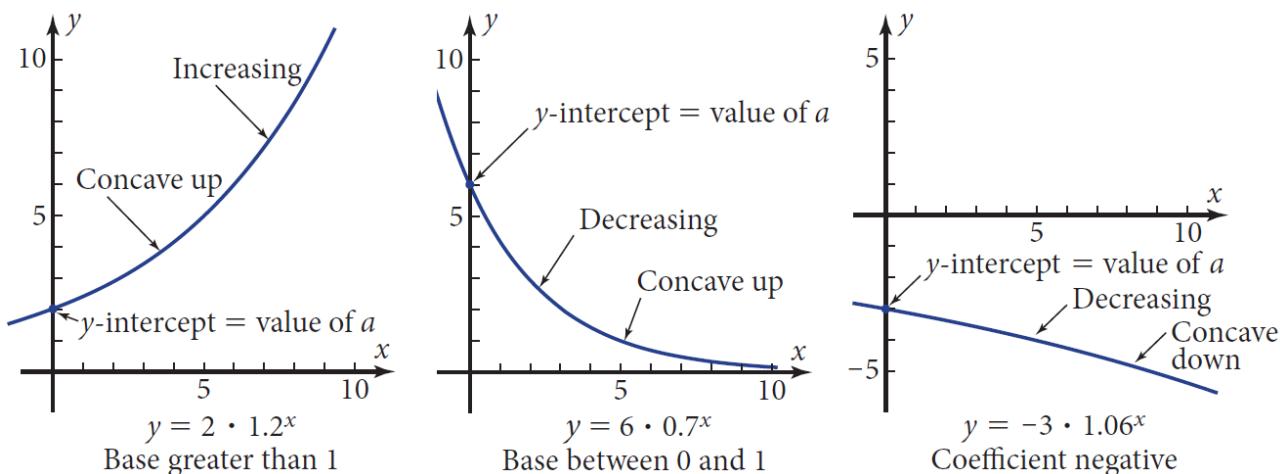


Figure 1.8: Exponential functions

In the equation $y = ab^x$, we say that " y varies exponentially with x ." This means that y changes by a constant factor b for each unit increase in x .

The translated form of the exponential function is

$$y = ab^x + c,$$

where the asymptote is the line $y = c$. Unless otherwise stated, "exponential function" will refer to the untranslated form $y = ab^x$.

Figure 1.8 illustrates exponential functions for different values of a and b . The key properties of the graph are as follows:

- The constant a is the y -intercept of the graph.
- The function is increasing if $b > 1$ and decreasing if $0 < b < 1$, provided $a > 0$.
- If $a < 0$, the function's behavior is reversed: it is decreasing if $b > 1$ and increasing if $0 < b < 1$.
- The graph is concave up if $a > 0$ and concave down if $a < 0$.

Mathematicians often use one of two particular constants as the base for an exponential function: either 10, which is the base of the decimal system, or the naturally occurring number e , which approximately equals 2.71828. These bases are significant in various mathematical applications.

Definition 1.8 (Special Exponential Functions)

$$y = a \cdot 10^{bx} \quad \text{base-10 exponential function}$$

$$y = a \cdot e^{bx} \quad \text{natural (base-}e\text{) exponential function,}$$

where a and b are constants and the domain is all real numbers.



To generalise the exponential function, the variable in the exponent is often multiplied by a constant. The (untranslated) general forms of these exponential functions are given below:

$$y = a \cdot 10^{bx} \quad \text{and} \quad y = a \cdot e^{bx}$$

These functions can be further generalised by incorporating translations in both the x - and y -directions. The translated forms are:

$$y = a \cdot 10^{b(x-c)} + d \quad \text{and} \quad y = a \cdot e^{b(x-c)} + d$$

The base- e exponential function, in particular, has a significant advantage when studying calculus, as the rate of change of e^x is equal to e^x itself.

1.8 Logarithms

Any positive number can be written as a power of 10. For instance,

$$3 = 10^{0.477\dots}$$

$$5 = 10^{0.6989\dots}$$

$$15 = 10^{1.1760\dots}$$

The exponents $0.4771\dots$, $0.6989\dots$, and $1.1760\dots$ are called the base-10 logarithms of 3, 5, and 15, respectively:

$$\log 3 = 0.4771\dots$$

$$\log 5 = 0.6989\dots$$

$$\log 15 = 1.1760\dots$$

To better understand the meaning of logarithms, press LOG 3 on your calculator. You will get:

$$\log 3 = 0.477121254\dots$$

Then, without rounding, raise 10 to this power. You will obtain:

$$10^{0.477121254\dots} = 3$$

The powers of 10 have the normal properties of exponentiation. For instance,

$$\begin{aligned} 15 &= (3)(5) = \left(10^{0.4771\dots}\right) \left(10^{0.6989\dots}\right) \\ &= 10^{0.4771\dots+0.6599\dots} \\ &= 10^{1.1760\dots} \end{aligned}$$

This means $10^{0.4771\dots+0.6599\dots} = 10^{1.1760\dots}$. Here, you add the exponents while keeping the same base. You can verify with your calculator that $10^{1.1760\dots}$ indeed equals 15.

From this example, you can infer that logarithms have the same properties as exponents. This is expected because logarithms *are* exponents. For instance,

$$\log(3 \cdot 5) = \log 3 + \log 5 \quad \text{The logarithm of a product equals the sum of the logarithms of the factors.}$$

From the values given earlier, you can also show that:

$$\log \frac{15}{3} = \log 15 - \log 3 \quad \text{The logarithm of a quotient.}$$

This property is reasonable because you divide powers of equal bases by subtracting the exponents:

$$\frac{15}{3} = \frac{10^{1.1760...}}{10^{0.477...}} = 10^{1.1760... - 0.4771...} = 10^{0.6989...} = 5$$

Since a power can be written as a product, you can find the logarithm of a power as follows:

$$\begin{aligned} \log 34 &= \log(3 \cdot 3 \cdot 3 \cdot 3) = \log 3 + \log 3 + \log 3 + \log 3 \\ &= 4 \log 3 \quad \text{Combine like terms.} \end{aligned}$$

The logarithm of a power equals the exponent of that power times the logarithm of the base. To verify this result, observe that $3^4 = 81$. Press $4 \times \text{LOG } 3$ on your calculator, and you'll find it equals $1.9084\dots$.

Definition 1.9 (Base-10 Logarithms)

$$\log x = y \iff 10^y = x$$

Verbally: $\log x$ is the exponent in the power of 10 that gives x



The term logarithm comes from the Greek words *logos*, meaning "ratio," and *arithmos*, meaning "number." Before the invention of calculators, base-10 logarithms were calculated approximately using infinite series and recorded in tables. Products involving many factors, such as

$$(357)(4.367)(22.4)(3.142)$$

could be calculated by adding their logarithms (exponents) rather than tediously multiplying several pairs of numbers. This method was invented by Englishman Henry Briggs (1561–1630) and Scotsman John Napier (1550–1616). The name logarithm, thus, reflects this "logical way to do arithmetic".

Properties of base-10 logarithms

- Log of a Product:

$$\log xy = \log x + \log y$$

Verbally: The log of a product equals the sum of the logs of the factors.

- Log of a Quotient:

$$\log \frac{x}{y} = \log x - \log y$$

Verbally: The log of a quotient equals the log of the numerator minus the log of the denominator.

- Log of a Power:

$$\log x^y = y \log x$$

Verbally: The log of a power equals the exponent times the log of the base.

Example 1.37 Find x if $\log_{10} 10^{3.721} = x$

Solution: By definition, the logarithm is the exponent of 10. So $x = 3.721$.

Example 1.38 Find x if $0.258 = 10^x$

Solution: By definition, x , the exponent of 10, is the logarithm of 0.258.

$$x = \log_{10} 0.258 = -0.5883\dots$$

The most important thing to remember about logarithms is this

A logarithm is an exponent.

Logarithms with Any Base: The Change-of-Base Property

If $x = 10^y$, then y is the base-10 logarithm of x . Similarly, if $x = 2^y$, then y is the base-2 logarithm of x . The only difference between these logarithms is the number that serves as the base. To distinguish among logarithms with different bases, the base is written as a subscript after the abbreviation "log." For instance:

$$\begin{aligned} 3 &= \log_2 8 \Leftrightarrow 2^3 = 8, \\ 4 &= \log_3 81 \Leftrightarrow 3^4 = 81, \\ 2 &= \log_{10} 100 \Leftrightarrow 10^2 = 100. \end{aligned}$$

The symbol $\log_2 8$ is pronounced "log to the base 2 of 8." The symbol $\log_{10} 100$ is, of course, equivalent to $\log 100$, as defined in the previous section. Note that in all cases, a logarithm represents an exponent.

Definition 1.10 (Logarithm with Any Base)

Algebraically:

$$\log_b x = y \text{ if and only if } b^y = x, \quad \text{where } b > 0, b \neq 1, \text{ and } x > 0$$

Verbally:

$$\log_b x = y \text{ means that } y \text{ is the exponent of } b \text{ that gives } x \text{ as the answer.}$$



The way you pronounce the symbol for logarithm gives you a way to remember the definition. The next two examples show you how to do this.

Example 1.39 Write $\log_5 c = a$ in exponential form.

Solution:

Think this:

- " $\log_5 \dots$ " is read as "log base 5 . . .," meaning 5 is the base.
- A logarithm is an exponent. Since the log equals a , a must be the exponent.
- The "answer" obtained from 5^a is the argument of the logarithm, denoted as c .

Write only this:

$$5^a = c$$

Example 1.40 Write $z^4 = m$ in logarithmic form.

Solution: $\log_z m = 4$

Two bases of logarithms are used frequently enough to have their own key on most calculators. One is the base-10 logarithm, also known as the common logarithm, as discussed in the previous section. The other is the base- e logarithm, known as the natural logarithm, where $e = 2.71828\dots$, a naturally occurring number (like π) that will be advantageous in your future mathematical studies.

The symbol $\ln x$ (pronounced "el en of x ") is used for natural logarithms, and is defined as:

$$\ln x = \log_e x$$

Definition 1.11 (Common Logarithm and Natural Logarithm)

Common: The symbol $\log x$ means $\log_{10} x$.

Natural: The symbol $\ln x$ means $\log_e x$, where e is a constant equal to $2.71828182845\dots$



Example 1.41 Find $\log_5 17$. Check your answer by an appropriate numerical method.

Solution: Let $x = \log_5 17$.

$$5^x = 17$$

$$\log_{10} 5^x = \log_{10} 17$$

$$x \log_{10} 5 = \log_{10} 17$$

$$x = \frac{\log_{10} 17}{\log_{10} 5} = 1.7603\dots$$

$$\log_5 17 = 1.7603\dots$$

$$5^{1.7603\dots} = 17$$

In this example, note that the base-5 logarithm of a number is directly proportional to the base-10 logarithm of that number. The conclusion of the example can be expressed as follows:

$$\log_5 17 = \frac{1}{\log_{10} 5} \cdot \log_{10} 17 = 1.4306\dots \log_{10} 17$$

To find the base-5 logarithm of any number, simply multiply its base-10 logarithm by 1.4306... (that is, divide by $\log_{10} 5$).

This proportional relationship is known as the change-of-base property. From the results of Example 3, you can write:

$$\log_5 17 = \frac{\log_{10} 17}{\log_{10} 5}$$

Notice that the logarithm with the desired base is isolated on the left side of the equation, while the two logarithms on the right side share the same base—typically one that is available on your calculator. The box below illustrates this property for bases a and b with argument x :

The Change-of-Base Property of Logarithms

$$\log_a x = \frac{\log_b x}{\log_b a} \quad \text{or} \quad \log_a x = \frac{1}{\log_b a} (\log_b x)$$

Example 1.42 Find $\ln 29$ using the change-of-base property with base-10 logarithms. Check your answer directly by pressing $\ln 29$ on your calculator.

Solution:

$$\ln 29 = \frac{\log 29}{\log e} = \frac{1.4623\dots}{0.4342\dots} = 3.3672\dots$$

Directly: $\ln 29 = 3.3672\dots$,

which agrees with the answer we got using the change-of-base property. 

Properties of Logarithms

The Logarithm of a Power:

$$\log_b x^y = y \log_b x$$

Verbally: The logarithm of a power equals the product of the exponent and the logarithm of the base. The Logarithm of a Product:

$$\log_b(xy) = \log_b x + \log_b y$$

Verbally: The logarithm of a product equals the sum of the logarithms of the factors. The

Logarithm of a Quotient:

$$\log_b \frac{x}{y} = \log_b x - \log_b y$$

Verbally: The logarithm of a quotient equals the logarithm of the numerator minus the logarithm of the denominator.

Solving Exponential and Logarithmic Equations

Logarithms provide a way to solve an equation with a variable in the exponent or to solve an equation that already contains logarithms. We will demonstrate this through the next few examples.

Example 1.43 Solve the exponential equation $7^{3x} = 983$ algebraically, using logarithms.

Solution:

$$7^{3x} = 983$$

$$\log 7^{3x} = \log 983$$

Take the base-10 logarithm of both sides.

$$3x \log 7 = \log 983$$

Apply the logarithm power property.

$$x = \frac{\log 983}{3 \log 7}$$

Divide both sides by the coefficient of x.

$$x = 1.1803\dots$$



Example 1.44 Solve the equation

$$\log_2(x - 1) + \log_2(x - 3) = 3$$

Solution:

$$\begin{aligned}
 \log_2(x-1) + \log_2(x-3) &= 3 \\
 \log_2[(x-1)(x-3)] &= 3 && \text{Apply the logarithm of a product property.} \\
 2^3 &= (x-1)(x-3) && \text{Use the definition of logarithm.} \\
 8 &= x^2 - 4x + 3 && \text{Expand the product.} \\
 x^2 - 4x - 5 &= 0 && \text{Reduce one side to zero. Use the symmetric} \\
 &&& \text{property of equality.} \\
 (x-5)(x+1) &= 0 && \text{Solve by factoring.} \\
 x = 5 \quad \text{or} \quad x &= -1
 \end{aligned}$$

We need to be cautious here because the solutions in the final step are the solutions of the quadratic equation, and we must make sure they are also solutions of the original logarithmic equation. Check by substituting the solutions into the original equation.

If $x = 5$, then

$$\begin{aligned}
 \log_2(5-1) + \log_2(5-3) \\
 = \log_2 4 + \log_2 2 \\
 = 2 + 1 = 3
 \end{aligned}$$

If $x = -1$, then

$$\begin{aligned}
 \log_2(-1-1) + \log_2(-1-3) \\
 = \log_2(-2) + \log_2(-4) \\
 \text{which is undefined.}
 \end{aligned}$$



Example 1.45 Solve the equation

$$e^{2x} - 3e^x + 2 = 0$$

Solution:

$$\begin{aligned}
 e^{2x} - 3e^x + 2 &= 0 \\
 (e^x)^2 - 3e^x + 2 &= 0
 \end{aligned}$$

We realise that this is a quadratic equation in the variable e^x . Using the quadratic formula, you get

$$\begin{aligned}
 e^x &= \frac{+3 \pm \sqrt{9-4(2)}}{2} = \frac{3 \pm 1}{2} \\
 e^x &= 2 \text{ or } e^x = 1
 \end{aligned}$$

You now have to solve these two equations.

$$\begin{aligned}
 e^x &= 2 & e^x &= 1 \\
 x = \ln 2 &= 0.6931\dots & x &= 0
 \end{aligned}$$

Check:

$$\begin{aligned}
 e^{2\ln 2} - 3e^{\ln 2} + 2 &= (e^0)^2 - 3e^0 + 2 \\
 &= (e^{\ln 2})^2 - 3e^{\ln 2} + 2 & &= 1^2 - 3(1) + 2 = 0 \\
 &= 2^2 - 3(2) + 2 = 0
 \end{aligned}$$

Both solutions are correct.



Example 1.46 Solve the logarithmic equation $\ln(x + 3) + \ln(x + 5) = 0$

Solution:

$$\begin{aligned}\ln(x + 3) + \ln(x + 5) &= 0 \\ \ln[(x + 3)(x + 5)] &= 0 \\ (x + 3)(x + 5) &= e^0 = 1 \\ x^2 + 8x + 15 &= 1 \\ x^2 + 8x + 14 &= 0 \\ x = -2.5857\dots \quad \text{or} \quad x &= -5.4142\dots\end{aligned}$$

Check:

$$\begin{aligned}x = -2.5857\dots : \\ \ln(-2.5857\dots + 3) + \ln(-2.5857\dots + 5) \\ = \ln(0.4142\dots) + \ln(2.4142\dots) \\ = -0.8813\dots + 0.8813\dots = 0\end{aligned}$$

which is ok.

$$\begin{aligned}x = -5.4142\dots : \\ \ln(-5.4142\dots + 3) + \ln(-5.4142\dots + 5) \\ = \ln(-2.4142\dots) + \ln(-0.4142\dots)\end{aligned}$$

which is undefined.

The only valid solution is $x = -2.5857\dots$



Chapter 2 Number Systems

We are so accustomed to working within the decimal system that we often forget it is a relatively recent invention and was once considered revolutionary. It is time to carefully examine how we represent numbers. Typically, we use the decimal system, where a number like 3459 is shorthand for $3 \times 1000 + 4 \times 100 + 5 \times 10 + 9$. The position of each digit is crucial, as it allows us to distinguish between values like 30 and 3. The decimal system is a **positional numeral system**, meaning it has designated positions for units, tens, hundreds, and so forth. Each digit's position implies the multiplier (a power of ten) that should be used with that digit, and each position has a value ten times that of the position to its right.

Notice that we can save space by writing 1000 as 10^3 , where the exponent 3 indicates the number of zeros. Thus, $100000 = 10^5$. If the exponent is negative, it represents a fraction, e.g., $10^{-3} = \frac{1}{1000}$. Perhaps the most ingenious aspect of the positional system was the addition of the decimal point, which allows us to include decimal fractions. For example, the number 123.456 is equivalent to:

$$1 \times 100 + 2 \times 10 + 3 \times 1 + 4 \times \frac{1}{10} + 5 \times \frac{1}{100} + 6 \times \frac{1}{1000}.$$

This can be visualised as:

Multiplier:	...	10^2	10^1	10^0	.	10^{-1}	10^{-2}	10^{-3}	...
Digits:	...	1	2	3	.	4	5	6	...

↑
Decimal Point:

However, there is no inherent reason why we must use powers of 10, or base 10. The Babylonians, for instance, used base 60, and base 12 was very common in medieval Europe. Today, the most widely used numeral systems are summarised in [Table 2.1](#)

Numeral system	Symbols	Base	Additional information
Decimal	0-9	10	-
Binary	0, 1	2	-
Hexadecimal	0-9, A-F	16	$A \equiv 10, B \equiv 11, C \equiv 12, D \equiv 13, E \equiv 14, F \equiv 15$
Octal	0-7	8	-

Table 2.1: Summary of Common Numeral Systems

We begin by focusing on binary which will also receive the most detailed attention in this chapter.

2.1 Binary Numbers

In the binary scale, we express numbers in powers of 2 rather than the 10s of the decimal scale. For some numbers, this is easy. Recall $2^0 = 1$,

As in decimal, we write this with the position of the digit representing the power, the first place after the decimal being the 2^0 position, the next the 2^1 , and so on. To convert a decimal number to binary, we can use the `mod` operator.

Decimal number	In powers of 2	Power of 2				Binary number
		3	2	1	0	
8	$= 2^3$	1	0	0	0	1000
7	$= 2^2 + 2^1 + 2^0$	0	1	1	1	111
6	$= 2^2 + 2^1$	0	1	1	0	110
5	$= 2^2 + 2^0$	0	1	0	1	101
4	$= 2^2$	0	1	0	0	100
3	$= 2^1 + 2^0$	0	0	1	1	11
2	$= 2^1$	0	0	1	0	10
1	$= 2^0$	0	0	0	1	1

Table 2.2: Decimal Numbers in Binary Representation

As an example, consider 88 in decimal or 88_{10} . We would like to write it as a binary number. We take the number and successively divide mod 2. See below:

Step Number n	x_n	$x_n/2$	$x_n \text{ mod } 2$
0	88	44	0
1	44	22	0
2	22	11	0
3	11	5	1
4	5	2	1
5	2	1	0
6	1	0	1

Table 2.3: Conversion of Decimal 88 to Binary

Writing the last column in reverse, that is from the bottom up, we have 1011000, which is the binary form of 88, i.e., $88_{10} = 1011000_2$.

Binary decimals are less common but quite possible. Thus, 101.1011 is just $2^2 + 2^0 + 2^{-1} + 2^{-3} + 2^{-4}$, which is, after some calculation, 5.6875. We have seen how to turn the integer part of a decimal number into a binary number, and we can do the same with a decimal fraction. Consider 0.6875. As before, we draw up a table:

Step Number n	x_n	$x_n \times 2$	$\lfloor x_n \times 2 \rfloor$
0	0.6875	1.375	1
1	0.375	0.75	0
2	0.75	1.5	1
3	0.5	1	1

Table 2.4: Conversion of Decimal Fraction 0.6875 to Binary

Giving, reading down, $0.6875_{10} = 1011_2$.

Binary Expansion

The process outlined in the previous section is called **binary expansion** and refers to the representation of a number in the binary (base-2) numeral system. Every decimal number can be expressed as a sum of powers of 2, where each power corresponds to a binary digit (bit) in the number's binary form.

Let's reconsider the decimal number 88. To find its binary expansion, we identify the largest power of 2 less than or equal to 88 and continue subtracting powers of 2 until we reach 0.

First, we note that $2^6 = 64$ is the largest power of 2 less than 88:

$$88 = 64 + 24$$

Next, we find that $2^4 = 16$ is the largest power of 2 less than 24:

$$24 = 16 + 8$$

Finally, $2^3 = 8$ exactly matches the remainder:

$$8 = 8 + 0$$

Thus, we have:

$$88 = 2^6 + 2^4 + 2^3$$

In binary, each of these powers of 2 is represented by a '1' in the corresponding place value, with '0' in place values where no power of 2 contributes:

$$88_{10} = 1011000_2$$

To summarise:

- $2^6 = 64$ corresponds to the leftmost '1' in the binary expansion.
- $2^4 = 16$ corresponds to the next '1'.
- $2^3 = 8$ corresponds to the next '1'.
- The remaining digits are '0' because $2^5, 2^2, 2^1$, and 2^0 do not contribute to the value 88.

Thus, the binary expansion of 88 is 1011000_2 . This method of representing numbers is fundamental in computer science and digital electronics, where binary representation is the standard for data storage and processing.

Binary Operations

Binary operations are basic arithmetic operations performed on binary numbers. These operations are essential in computing and digital systems, as they form the foundation for how computers process and manipulate data.

Binary addition, subtraction, and multiplication are similar to their decimal counterparts but follow simpler rules due to the binary system's limited digits. For example, binary addition follows these rules:

$$0 + 0 = 0, \quad 0 + 1 = 1, \quad 1 + 0 = 1, \quad 1 + 1 = 10$$

In this case, $1 + 1$ results in 10_2 , which means 0 with a carry of 1 to the next higher bit. Binary subtraction and multiplication follow similar straightforward rules that are easy to implement in digital systems.

The XOR (exclusive OR) operation is another important binary operation. XOR produces a 1 if the two bits being compared are different and a 0 if they are the same:

$$0 \oplus 0 = 0, \quad 0 \oplus 1 = 1, \quad 1 \oplus 0 = 1, \quad 1 \oplus 1 = 0$$

In binary addition, the XOR operation is used to add two bits without considering any carry from a previous bit. This is because XOR effectively performs addition modulo 2, which aligns perfectly with how binary addition works. For example:

Bit 1	Bit 2	XOR (Sum)	AND (Carry)
0	0	0	0
0	1	1	0
1	0	1	0
1	1	0	1

In the case of $1 + 1$, XOR gives a sum of 0 and an AND operation (which detects the carry) gives a carry of 1, resulting in the binary number 10.

$$\begin{aligned} 0 + 0 &= 0 \\ 0 + 1 &= 1 \\ 1 + 1 &= 10 \quad \text{so we carry 1 and leave a zero} \\ 1 + 1 + 1 &= 1 + (1 + 1) = 1 + 10 = 11. \end{aligned}$$

We can write this in very much the same way as for a decimal addition:

	1	1	0	1	0	1	
+	1	0	1	1	1	0	
	1	1	0	0	1	1	Sum
↑				↑			

The right-hand arrow shows where we carry a 1. The left-hand arrow shows where we have $1 + 1 + 1$ so we carry a 1 and have a 1 left over.

As we will see below, we will often need to handle multiple carries. There are two ways to handle this which resemble the methods we know from the decimal system. We will explain using an example.

Method 1: Column-wise Binary Addition with Multiple Carries

Consider

$$\begin{array}{r} 1 \ 1 \ 1 \ 1 \ 1 \\ + \ 1 \ 1 \ 1 \ 0 \ 1 \\ + \ 1 \ 1 \ 1 \ 0 \ 1 \\ + \ 1 \ 1 \ 1 \ 1 \ 1 \end{array}$$

Step 1: Add the Rightmost Column

Start by adding the rightmost bits:

$$1 + 1 + 1 + 1 = 100_2 \quad (\text{which is binary for } 4)$$

Reading the result from right to left (i.e. from *least significant bit* (LSB) to the *most significant bit*(MSB))

- write down the 0
- carry the 0 to the next column
- carry the 1 to the third column

You end up with

$$\begin{array}{r}
 & 1 & 0 \\
 & 1 & 1 & 1 & 1 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & \hline
 & 1 & 1 & 1 & 1 & 1 \\
 & & & & 0
 \end{array}$$

Step 2: Add the Second Column from the Right

Next, add the second column:

$$0 + 1 + 0 + 0 + 1 = 10_2 \quad (\text{which is binary for } 2)$$

Reading the result from LSB to MSB:

- write down the 0
- carry the 1 to the next column

You end up with

$$\begin{array}{r}
 & 1 \\
 & 1 \\
 & 1 & 1 & 1 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & \hline
 & 1 & 1 & 1 & 1 & 1 \\
 & & & 0 & 0
 \end{array}$$

Step 3: Add the Third Column from the Right

Now, add the third column:

$$1 + 1 + 1 + 1 + 1 + 1 = 110_2 \quad (\text{which is binary for } 6)$$

Reading the result from LSB to MSB

- write down the 0
- carry the 1 to the fourth column

- carry the 1 to the fifth column

You end up with

Our sum so far:

$$\begin{array}{r}
 & 1 & 1 \\
 1 & 1 & 1 & 1 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & \hline
 & 1 & 1 & 1 & 1 & 1 \\
 & 0 & 0 & 0
 \end{array}$$

Step 4: Add the Fourth Column from the Right

Move to the fourth column:

$$1 + 1 + 1 + 1 + 1 = 101_2 \quad (\text{which is binary for } 5)$$

Reading the result from LSB to MSB

- write down the 1
- carry the 0 to the fifth column
- carry the 1 to the sixth column

You end up with

$$\begin{array}{r}
 & & 0 \\
 & 1 & 1 \\
 & 1 & 1 & 1 & 1 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & \hline
 & 1 & 1 & 1 & 1 & 1 \\
 & 1 & 0 & 0 & 0
 \end{array}$$

Step 5: Add the Leftmost Column

Add the leftmost column:

$$1 + 1 + 1 + 1 + 1 = 101_2 \quad (\text{which is binary for } 5)$$

Reading the result from LSB to MSB

- write down the 1
- carry the 0 to the sixth column
- carry the 1 to the seventh column

This results in

$$\begin{array}{r}
 & & & 0 \\
 & & 1 & 1 \\
 + & & 1 & 1 & 1 & 0 & 1 \\
 + & & 1 & 1 & 1 & 0 & 1 \\
 + & \hline & 1 & 1 & 1 & 1 & 1 \\
 & & 1 & 1 & 0 & 0 & 0
 \end{array}$$

Step 6: Add the Remaining Carries

Finally, add the remaining carries:

$$\begin{array}{r}
 & & 1 & 1 & 1 & 1 & 1 \\
 + & & 1 & 1 & 1 & 0 & 1 \\
 + & & 1 & 1 & 1 & 0 & 1 \\
 + & \hline & 1 & 1 & 1 & 1 & 1 \\
 & & 1 & 1 & 1 & 1 & 0 & 0 & 0
 \end{array}$$

The following example demonstrates the entire process by using different colors to distinguish each column and the corresponding carries they produce. Note that the last two digits in the sum are colored black, as they do not result from any specific column but are instead generated solely from the carries.

$$\begin{array}{r}
 & & 1 & 0 \\
 & & 1 & 0 & 1 \\
 & & 1 & 1 & 1 & 0 \\
 & & 1 & 1 & 1 & 1 & 1 \\
 + & & 1 & 1 & 1 & 0 & 1 \\
 + & & 1 & 1 & 1 & 0 & 1 \\
 + & \hline & 1 & 1 & 1 & 1 & 1 \\
 & & 1 & 1 & 1 & 1 & 0 & 0 & 1
 \end{array}$$

Method 2: Direct Summation and Simplification

We will illustrate the second method using the same example. In the previous case, we carried the actual binary number to the next columns. In this method, we write down 0 if the sum is even and 1 if the sum is odd. Every time a sum is a multiple of 2, we carry a 1 to the next columns, and then continue this process for each column, including the carries in the calculation of that column.

Step 1: Add the Rightmost Column

Add bits in column 1 (from counting from MSB):

$$1 + 1 + 1 + 1 = 100_2 \quad (\text{which is binary for } 4)$$

Reading the result from LSB to MSB

- write down the 0
- carry a 1 for the first multiple of 2

- carry a 1 for the second multiple of 2

This results in

$$\begin{array}{r}
 & & 1 \\
 & & 1 \\
 1 & 1 & 1 & 1 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & \hline
 & 1 & 1 & 1 & 1 & 1 \\
 & & 0
 \end{array}$$

Step 2: Add the Second Column from the Right

Add bits in column 2 (from counting from MSB):

$$1 + 1 + 1 + 1 = 100_2 \quad (\text{which is binary for } 4)$$

Reading the result from LSB to MSB

- write down the 0
- carry a 1 for the first multiple of 2
- carry a 1 for the second multiple of 2

This results in

$$\begin{array}{r}
 & & 1 \\
 & & 1 \\
 1 & 1 & 1 & 1 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & \hline
 & 1 & 1 & 1 & 1 & 1 \\
 & & 0 & 0
 \end{array}$$

Step 3: Add the Third Column from the Right

Add bits in column 3 (from counting from MSB):

$$1 + 1 + 1 + 1 + 1 + 1 = 110_2 \quad (\text{which is binary for } 6)$$

Reading the result from LSB to MSB

- write down the 0
- carry a 1 for the first multiple of 2
- carry a 1 for the second multiple of 2
- carry a 1 for the third multiple of 2

This results in

$$\begin{array}{r}
 & & 1 \\
 & & 1 \\
 & & 1 \\
 1 & 1 & 1 & 1 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & \hline
 & 1 & 1 & 1 & 1 & 1 \\
 & 0 & 0 & 0
 \end{array}$$

Step 4: Add the Fourth Column from the Right

Add bits in column 4 (from counting from MSB):

$$1 + 1 + 1 + 1 + 1 + 1 = 111_2 \quad (\text{which is binary for } 7)$$

Reading the result from LSB to MSB

- write down the 1
- carry a 1 for the first multiple of 2
- carry a 1 for the second multiple of 2
- carry a 1 for the third multiple of 2

This results in

$$\begin{array}{r}
 & & 1 \\
 & & 1 \\
 & & 1 \\
 1 & 1 & 1 & 1 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & \hline
 & 1 & 1 & 1 & 1 & 1 \\
 & 1 & 0 & 0 & 0 & 0
 \end{array}$$

Step 5: Add the Leftmost Column

Add bits in leftmost column):

$$1 + 1 + 1 + 1 + 1 + 1 = 111_2 \quad (\text{which is binary for } 7)$$

Reading the result from LSB to MSB

- write down the result in binary, i.e. 1 1 1

This results in

$$\begin{array}{r}
 & 1 & 1 & 1 & 1 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & 1 & 1 & 1 & 0 & 1 \\
 + & \hline
 & 1 & 1 & 1 & 1 & 1 \\
 & 1 & 1 & 1 & 1 & 0 & 0 & 0
 \end{array}$$

which corresponds to the result we obtained above. While not demonstrated explicitly here, subtraction works in a similar fashion.

By using one of these methods for handling multiple carries, allow us to also multiply two binary numbers.

Multiplication in Binary

Multiplication in binary is technically easier than multiplication in decimal. In binary operations, we work exclusively with two digits: 0 and 1. This means that both the the multiplier¹ and multiplicand consist of 0's and 1's (and so does the multiplicand). The process of finding the binary product is analogous to traditional multiplication in the decimal system. The four five steps involved in multiplying binary digits are:

$$0 \times 0 = 0$$

$$0 \times 1 = 0$$

$$1 \times 0 = 0$$

$$1 \times 1 = 1$$

$$1 \times 10_2 = 10_2 \quad (\text{multiplying by base } 10_2 \text{ adds a 0 to the end})$$

The last step means that $101_2 \times 10_2 = 1010_2$ which is analogous to the decimal case: $143_{10} \times 10_{10} = 1430_{10}$.

We will illustrate the process by supplying a couple of examples.

Example 2.1

$$\begin{array}{r}
 (1) & 1 & 0 & 0 \\
 \times & \underline{(2)} & 1 & 1 \\
 (3) & 1 & 0 & 0 \\
 + & \underline{(4)} & 1 & 0 & 0 & 0 \\
 (5) & 1 & 1 & 0 & 0
 \end{array}$$

Here are the steps:

- Multiply the multiplicand (line 1) by the LSB of the multiplier (line 2), which in this case is 1.
- Record this result in line 3.
- Append a 0 to line 4 to account for the shift to the next power of 2 in the multiplier.

¹The "multiplicand" is the number that has to be multiplied, and the "multiplier" is the number by which it is multiplied.

- Multiply the multiplicand (line 1) by the next bit of the multiplier (line 2), which is also 1 in this case.
- Add this result to line 4, after the 0 you appended earlier.
- Finally, sum the values in lines 3 and 4, as outlined in the previous section, to obtain the final result in line 5.

We offer two additional examples:

Example 2.2

$$\begin{array}{r}
 & 1 & 0 & 1 \\
 \times & \hline
 & 1 & 0 & 1 & 1 \\
 & \hline
 & 1 & 0 & 1 \\
 + & 1 & 0 & 1 & 0 \\
 + & 0 & 0 & 0 & 0 & 0 \\
 + & \hline
 & 1 & 0 & 1 & 0 & 0 & 0 \\
 & \hline
 & 1 & 1 & 0 & 1 & 1 & 1
 \end{array}$$

Example 2.3

$$\begin{array}{r}
 & 1 & 0 & 0 & 1 & 1 & 1 & 0 \\
 \times & \hline
 & 1 & 0 & 1 \\
 & \hline
 & 1 & 0 & 0 & 1 & 1 & 1 & 0 \\
 + & 1 & 0 & 0 & 1 & 1 & 1 & 0 & 0 \\
 + & \hline
 & 1 & 1 & 0 & 0 & 0 & 1 & 1 & 0
 \end{array}$$

In Example 2.3 notice that we omitted the row of zeroes that the second value of the multiplier would have produced, and notice even further that we added two 0's before restating the multiplicand in the sum.

Binary multiplication, like binary addition, is a core operation in computer arithmetic. By breaking the process down into manageable steps—multiplying individual bits and then summing the results—it becomes clear how similar it is to the multiplication methods we use in the decimal system. The main difference is the simplicity and efficiency of working within the binary system, where only the digits 0 and 1 are involved.

We notice how binary multiplication builds on binary addition. Each step, involving shifts and sums, essentially consists of repeated additions adjusted by powers of two. A strong grasp of binary addition naturally leads to a better understanding of binary multiplication and its applications.

2.2 Octal and Hexadecimal

Octal is a base-8 numbering system that uses the digits 0 through 7. It is closely related to binary, which is a base-2 system. The connection between the two lies in how easily binary numbers can be converted to octal and vice versa. Each octal digit corresponds to exactly three binary digits (bits), making conversions straightforward. For example, the binary number ‘110‘ converts directly to the octal digit ‘6‘. Because of this close relationship,

octal is often used as a shorthand for binary in computing, particularly in contexts where grouping binary digits in sets of three simplifies reading and interpreting binary data.

$$12_8 = 1 \cdot 8^1 + 2 \cdot 8^0 = 10_{10}$$

$$3021_8 = 3 \cdot 8^3 + 0 \cdot 8^2 + 2 \cdot 8^1 + 1 \cdot 8^0 = 1553_{10}$$

Since 8 is 2^3 , we can express it in binary:

$$3 \rightarrow 011$$

$$0 \rightarrow 000$$

$$2 \rightarrow 010$$

$$1 \rightarrow 001$$

$$\text{Thus, } 3021_8 = 011000010001_2 = 11000010001_2$$

We obtain the final result by removing leading zeros.

Hexadecimal is a base-16 numbering system that uses sixteen distinct symbols: the digits 0-9 and the letters A-F, where A represents 10, B represents 11, and so on up to F, which represents 15. Hexadecimal is closely related to binary because each hexadecimal digit corresponds exactly to four binary digits (bits). This direct relationship makes it easy to convert between the two systems. For example, the hexadecimal digit ‘A’ translates to the binary sequence ‘1010’. Due to this efficiency in grouping, hexadecimal is often used in computing as a more compact and readable way to represent binary data, particularly in areas like memory addresses and colour codes in web design.

$$123_{16} = 1 \cdot 16^2 + 2 \cdot 16^1 + 3 \cdot 16^0 = 256 + 32 + 3 = 291_{10}$$

$$A2E_{16} = 10 \cdot 16^2 + 2 \cdot 16^1 + 14 \cdot 16^0 = 2560 + 32 + 14 = 2606_{10}$$

Since 16 is 2^4 , we can, for instance, express $A2E_{16}$ in binary:

$$A \rightarrow 1010$$

$$2 \rightarrow 0010$$

$$E \rightarrow 1110$$

$$\text{Thus, } A2E_{16} = 101000101110_2$$

Similarly we get $5EB52_{16}$ as

$$5 \rightarrow 0101$$

$$E \rightarrow 1110$$

$$B \rightarrow 1011$$

$$5 \rightarrow 0101$$

$$2 \rightarrow 0010$$

$$\text{Thus, } 5EB52_{16} = 01011110101101010010_2$$

Again, notice that we removed the leading 0's from 5_{16} when writing the result.

2.3 Converting Between Systems

Understanding how to convert numbers between binary, decimal, octal, and hexadecimal systems is essential in computer science and digital electronics. Each system is a different base, and each has its own applications. Here's a step-by-step guide to help you convert numbers from one system to another.

Decimal to Binary Conversion

To convert a decimal number to binary:

1. Divide the decimal number by 2.
2. Record the remainder (it will be 0 or 1).
3. Divide the quotient by 2 and record the remainder.
4. Repeat until the quotient is 0.
5. The binary number is the sequence of remainders read from bottom to top.

Example 2.4 Convert 23_{10} to binary.

Solution:

$$\begin{aligned} 23 \div 2 &= 11 \quad \text{remainder } 1 \\ 11 \div 2 &= 5 \quad \text{remainder } 1 \\ 5 \div 2 &= 2 \quad \text{remainder } 1 \\ 2 \div 2 &= 1 \quad \text{remainder } 0 \\ 1 \div 2 &= 0 \quad \text{remainder } 1 \end{aligned}$$

Thus, $23_{10} = 10111_2$.

Decimal to Octal Conversion

To convert a decimal number to octal:

1. Divide the decimal number by 8.
2. Record the remainder.
3. Divide the quotient by 8 and record the remainder.
4. Repeat until the quotient is 0.
5. The octal number is the sequence of remainders read from bottom to top.

Example 2.5 Convert 78_{10} to octal.

Solution:

$$\begin{aligned} 78 \div 8 &= 9 \quad \text{remainder } 6 \\ 9 \div 8 &= 1 \quad \text{remainder } 1 \\ 1 \div 8 &= 0 \quad \text{remainder } 1 \end{aligned}$$

Thus, $78_{10} = 116_8$.

Decimal to Hexadecimal Conversion

To convert a decimal number to hexadecimal:

1. Divide the decimal number by 16.
2. Record the remainder (use A, B, C, D, E, F for remainders 10, 11, 12, 13, 14, 15 respectively).
3. Divide the quotient by 16 and record the remainder.
4. Repeat until the quotient is 0.
5. The hexadecimal number is the sequence of remainders read from bottom to top.

Example 2.6 Convert 255_{10} to hexadecimal.

Solution:

$$255 \div 16 = 15 \text{ remainder } 15 \quad (\text{F})$$

$$15 \div 16 = 0 \text{ remainder } 15 \quad (\text{F})$$

Thus, $255_{10} = FF_{16}$. 

Binary to Decimal Conversion

To convert a binary number to decimal:

1. Multiply each bit by 2 raised to the power of its position, starting from 0 on the right.
2. Sum all the products.

Notice that this amounts to the method outlined above about binary expansion.

Example 2.7 Convert 1101_2 to decimal.

Solution:

$$1 \cdot 2^3 + 1 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 = 8 + 4 + 0 + 1 = 13_{10}$$

Binary to Octal Conversion

To convert a binary number to octal:

1. Group the binary digits into sets of three, starting from the right. Add leading zeros if necessary.
2. Convert each group of three binary digits to its octal equivalent.

Example 2.8 Convert 110110_2 to octal.

Solution:

$$110 \rightarrow 6$$

$$110 \rightarrow 6$$

Thus, $110110_2 = 66_8$. 

Binary to Hexadecimal Conversion

To convert a binary number to hexadecimal:

1. Group the binary digits into sets of four, starting from the right. Add leading zeros if necessary.
2. Convert each group of four binary digits to its hexadecimal equivalent.

Example 2.9 Convert 10110101_2 to hexadecimal.

Solution:

$$1011 \rightarrow B$$

$$0101 \rightarrow 5$$

Thus, $10110101_2 = B5_{16}$. 

Octal to Binary Conversion

To convert an octal number to binary: Convert each octal digit to its 3-bit binary equivalent!

Example 2.10 Convert 57_8 to binary.

Solution:

$$5 \rightarrow 101$$

$$7 \rightarrow 111$$

Thus, $57_8 = 101111_2$. 

Octal to Decimal Conversion

To convert an octal number to decimal:

1. Multiply each digit by 8 raised to the power of its position, starting from 0 on the right.
2. Sum all the products.

Example 2.11 Convert 157_8 to decimal.

Solution:

$$1 \cdot 8^2 + 5 \cdot 8^1 + 7 \cdot 8^0 = 64 + 40 + 7 = 111_{10}$$

Octal to Hexadecimal Conversion

To convert an octal number to hexadecimal:

1. First, convert the octal number to binary.
2. Then, convert the binary number to hexadecimal by grouping the binary digits in sets of four.

Example 2.12 Convert 157_8 to hexadecimal.

Solution:

$$1 \rightarrow 001$$

$$5 \rightarrow 101$$

$$7 \rightarrow 111$$

Thus, $157_8 = 00110111_2 = 6F_{16}$. 

Hexadecimal to Binary Conversion

To convert a hexadecimal number to binary: Convert each hexadecimal digit to its 4-bit binary equivalent.

Example 2.13 Convert $2B_{16}$ to binary.

Solution:

$$2 \rightarrow 0010$$

$$B \rightarrow 1011$$

Thus, $2B_{16} = 00101011_2$. 

Hexadecimal to Decimal Conversion

To convert a hexadecimal number to decimal:

1. Multiply each digit by 16 raised to the power of its position, starting from 0 on the right.
2. Sum all the products.

Example 2.14 Convert $2B_{16}$ to decimal.

Solution:

$$2 \cdot 16^1 + 11 \cdot 16^0 = 32 + 11 = 43_{10}$$

Hexadecimal to Octal Conversion

To convert a hexadecimal number to octal:

1. First, convert the hexadecimal number to binary.
2. Then, convert the binary number to octal by grouping the binary digits in sets of three.

Example 2.15 Convert $2B_{16}$ to octal.

Solution:

$$2 \rightarrow 0010$$

$$B \rightarrow 1011$$

Thus, $2B_{16} = 00101011_2 = 53_8$. 

Final Thoughts on Conversion

The concept of expansion plays a central role in these conversions. Whether you are expanding a decimal number into its binary, octal, or hexadecimal form, or converting a binary number into its octal or hexadecimal equivalent, you are more or less expressing the number in terms of powers of the base. The expansion method is essentially the same for each system as it boils down to dividing by the highest power of the base recursively:

$$\begin{aligned} 7562_{10} &= 1 \cdot 16^3 + 3466 = 1 \cdot 16^3 + 13 \cdot 16^2 + 138 \\ &= 1 \cdot 16^3 + 13 \cdot 16^2 + 8 \cdot 16^1 + 10 \cdot 16^0 = 1D8A \end{aligned}$$

By understanding these expansions and the relationships between these number systems, you can efficiently switch between them, allowing you to represent and manipulate data in the most suitable format for any given situation.

Chapter 3 Set Theory

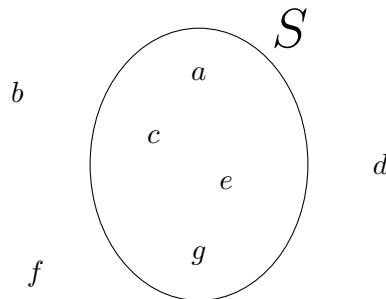
In software engineering, we are constantly working with collections of things: users in a system, records in a database, or nodes in a network. Set theory provides the formal mathematical language to describe and manipulate these collections with precision and clarity. It is the bedrock upon which many core computer science concepts—from database query languages like SQL to data structures and algorithmic logic—are built.

This chapter introduces the fundamental principles of set theory. We will begin with the simple, intuitive idea of a set and explore the formal notation used to define them. We will then cover the essential relationships between sets, such as subsets, and the core operations used to combine them, including unions, intersections, and complements. These operations directly correspond to the logical operators (OR, AND, NOT) that govern the flow of your code. Finally, we will explore ordered collections called tuples and introduce a simple method for proving set equalities.

3.1 What is a Set?

A set is an unordered collection of distinct objects. The objects within a set are called its **elements** or **members**. We can think of a set as a simple container where items are grouped together, and the order in which we list them does not matter. For example, the set of primary colors can be written as {red, yellow, blue} or equally as {blue, red, yellow}.

Sets are a cornerstone of modern mathematics and a fundamental concept in computer science. They form the logical basis for everything from database query languages and data structures to the specification of programming language types.



The elements of set S are a, c, e , and g .

Figure 3.1: A set S containing four elements. The objects b, d , and f are not elements of S .

Specifying a Set

There are two primary ways to describe a set: by explicitly listing its members or by defining a property that its members must satisfy.

Listing Notation (Roster Method)

The most direct way to define a set is by listing all its elements between curly braces, $\{\}$. This is known as the **roster method**.

For example:

- The set of the first five letters of the alphabet is $A = \{a, b, c, d, e\}$.
- The set of the first three positive integers is $C = \{1, 2, 3\}$.
- A set can contain different types of elements: $D = \{\text{Alice}, 42, \pi\}$.

When using the roster method, there are two fundamental rules:

1. **Order does not matter.** A set is defined only by the elements it contains, not by the sequence in which they are listed. For example, $\{1, 2, 3\}$ is the exact same set as $\{3, 1, 2\}$.
2. **Each element must be unique.** An element is either in a set or it is not. Listing an element more than once is redundant and does not change the set. For instance, the set $\{a, a, b, c, c\}$ is simply $\{a, b, c\}$.

Set-Builder Notation

When listing every element is impractical or impossible (for example, with infinite sets), we use **set-builder notation**. This method defines a set by stating a property or rule that its elements must satisfy. The notation uses a vertical bar ‘|’ or a colon ‘:’, which is read as “such that.”

The general structure is:

$\{\text{variable} \mid \text{a property the variable must satisfy}\}$

For example:

- $A = \{l \mid l \text{ is a vowel in the English alphabet}\}$

This is read as: “ A is the set of all elements l such that l is a vowel in the English alphabet.” This is another way of writing $A = \{a, e, i, o, u\}$.

- $C = \{n \mid n \in \mathbb{Z} \text{ and } 0 < n < 4\}$

This is read as: “ C is the set of all numbers n such that n is an integer and n is greater than 0 and less than 4.” This defines the set $\{1, 2, 3\}$.

- $E = \{x \mid x \text{ is an even integer}\}$

This defines the infinite set of all even integers: $\{\dots, -4, -2, 0, 2, 4, \dots\}$.

Set-builder notation is extremely powerful because it allows us to define large, complex, or even infinite sets with a short and precise description.

3.2 Important Sets: The Number Systems

Numbers are the foundation of mathematics and computation. The different categories of numbers we use every day, from counting items to measuring continuous values, can be formally defined as sets. Understanding these fundamental sets is crucial for any software engineer, as they underpin data types, arithmetic logic, and numerical algorithms.

The Main Number Sets

The following sets are some of the most important in mathematics and are used throughout science and engineering.

Natural Numbers (\mathbb{N}) The set of positive integers used for counting: $\{1, 2, 3, \dots\}$. Sometimes, this set is defined to include 0. Due to this ambiguity, it is often clearer to specify *positive integers* or *non-negative integers*.

Integers (\mathbb{Z}) The set of all positive and negative whole numbers, including zero: $\{\dots, -2, -1, 0, 1, 2, \dots\}$.

Rational Numbers (\mathbb{Q}) The set of all numbers that can be expressed as a fraction $\frac{p}{q}$, where p and q are integers and $q \neq 0$. This includes all integers and terminating or repeating decimals. Examples: $\frac{1}{2}, -5, 0.25$.

Real Numbers (\mathbb{R}) The set of all numbers on the number line. It includes both rational numbers and irrational numbers (like π or $\sqrt{2}$), which cannot be expressed as simple fractions.

Complex Numbers (\mathbb{C}) The set of all numbers that can be expressed in the form $a + bi$, where a and b are real numbers and i is the imaginary unit, satisfying $i^2 = -1$.

Visualizing the relationships between these sets is key to understanding them. A Venn diagram shows the hierarchy of how these sets are nested within one another, while a number line illustrates how they cover or populate the continuum of values.

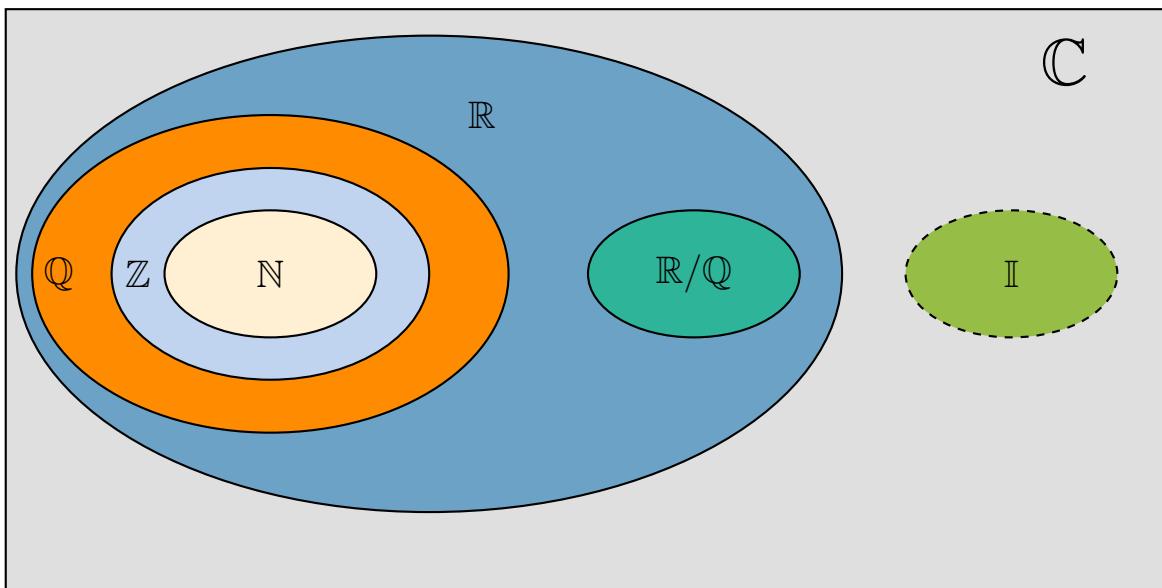


Figure 3.2: A Venn diagram illustrating the hierarchical relationship between the major number sets.

Remark: The Venn diagram in Figure 3.2 shows that the set of real numbers (\mathbb{R}) is composed exclusively of rational (\mathbb{Q}) and irrational numbers. There is no real number that is not one or the other. This is why the set of irrational numbers is formally denoted as $\mathbb{R} \setminus \mathbb{Q}$, which means "the set of all real numbers, excluding the rational numbers."

Interval Notation

In many applications, we need to refer to a continuous range of real numbers. **Interval notation** is a convenient shorthand for describing such subsets of \mathbb{R} .

An interval is defined by its two endpoints. We use square brackets '[']' to indicate that an endpoint is included

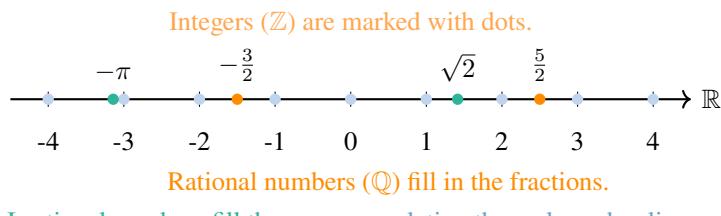


Figure 3.3: The real number line, populated by integers, rationals, and irrationals.

in the set, and parentheses ‘(‘ ‘)’ to indicate that it is excluded.

Closed Interval: Includes both endpoints.

$$[a, b] = \{x \in \mathbb{R} \mid a \leq x \leq b\}$$

Open Interval: Excludes both endpoints.

$$(a, b) = \{x \in \mathbb{R} \mid a < x < b\}$$

Half-Open Intervals: Includes one endpoint but not the other.

$$[a, b) = \{x \in \mathbb{R} \mid a \leq x < b\} \quad \text{and} \quad (a, b] = \{x \in \mathbb{R} \mid a < x \leq b\}$$

Example 3.1 Interval Notation

- $[-2, 3]$ is the set of all real numbers from -2 to 3 , including -2 and 3 .
- $(-2, 3)$ is the set of all real numbers between -2 and 3 .
- $[0, 100)$ represents all numbers from 0 up to (but not including) 100 .

Intervals can also be unbounded, extending towards positive or negative infinity (∞). Since infinity is not a number, it is always excluded with a parenthesis.

$$(a, \infty) = \{x \in \mathbb{R} \mid x > a\} \quad \text{and} \quad (-\infty, b] = \{x \in \mathbb{R} \mid x \leq b\}$$

Example 3.2 Unbounded Interval

- $(3, \infty)$ represents all real numbers greater than 3 .
- $(-\infty, 5)$ represents all real numbers less than 5 .
- $(-\infty, 0]$ represents the set of all non-positive real numbers.

3.3 Relationships Between Sets

Understanding a set is not just about its elements, but also how it relates to other sets. This section defines the fundamental relationships that allow us to compare and classify sets.

Subsets and Proper Subsets

One of the most basic relationships is that of inclusion, where one set is contained within another.

Definition 3.1 (Subset)

A set A is a **subset** of a set B if every element of A is also an element of B . We write this as $A \subseteq B$.



For example, if $A = \{1, 2\}$ and $B = \{1, 2, 3\}$, then $A \subseteq B$ because every element in A is also in B . By this definition, every set is a subset of itself (i.e., $A \subseteq A$).

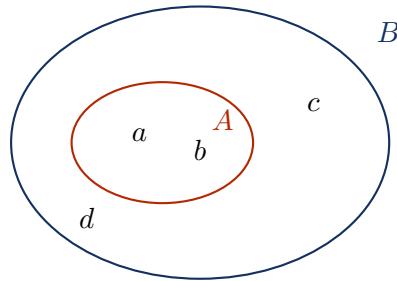


Figure 3.4: Set $A = \{a, b\}$ is a subset of set $B = \{a, b, c, d\}$, denoted $A \subseteq B$.

Sometimes we want to specify that a set is a subset of another but is not equal to it.

Definition 3.2 (Proper Subset)

A set A is a **proper subset** of a set B if $A \subseteq B$ and $A \neq B$. This means that B must contain at least one element that is not in A . We write this as $A \subset B$.



Using the previous example, since B contains the element 3 which is not in A , we can say that A is a proper subset of B , or $A \subset B$.

The Universal Set and the Empty Set

Two special sets act as the boundaries for set theory: the set containing everything and the set containing nothing.

Definition 3.3 (Universal Set)

The **universal set**, denoted by U , is the set of all possible elements under consideration in a given context. All other sets in that context are considered subsets of the universal set.



The universal set is represented in Venn diagrams by a rectangle that encloses all other sets. For example, if we are discussing integers, the universal set would be $U = \mathbb{Z}$. If we were discussing students at a university, U would be the set of all enrolled students.

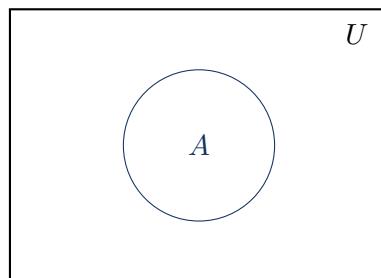


Figure 3.5: The universal set U contains all elements and sets under consideration.

Definition 3.4 (Empty Set)

The **empty set** (or **null set**) is the unique set containing no elements. It is denoted by \emptyset or by $\{\}$.



The empty set has a crucial property:

The empty set is a subset of every set.

This is because there are no elements in \emptyset that are not in any other set A . Therefore, for any set A , it is always true that $\emptyset \subseteq A$.

Disjoint Sets

Sometimes, sets have no relationship at all because they are entirely separate.

Definition 3.5 (Disjoint Sets)

Two sets, A and B , are **disjoint** if they have no elements in common. In other words, their intersection is the empty set: $A \cap B = \emptyset$.



For example, consider the set of even integers, $E = \{\dots, -2, 0, 2, \dots\}$, and the set of odd integers, $O = \{\dots, -3, -1, 1, 3, \dots\}$. These two sets are disjoint.

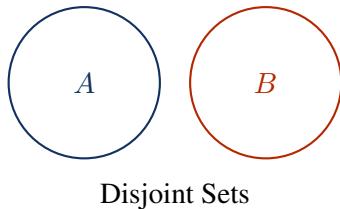


Figure 3.6: Sets A and B are disjoint because they do not overlap.

3.4 Properties of Sets

Beyond the relationships between sets, we can also describe their intrinsic properties. The two most fundamental properties are a set's size (its cardinality) and the collection of all its possible subsets (its power set).

Cardinality

The most basic property of a finite set is its size.

Definition 3.6 (Cardinality)

The **cardinality** of a finite set A , denoted $|A|$, is the number of distinct elements in the set.



For example:

- If $A = \{a, b, c, d\}$, then $|A| = 4$.
- If $B = \{n \in \mathbb{Z} \mid 0 < n < 5\}$, then $B = \{1, 2, 3, 4\}$ and $|B| = 4$.
- For the empty set, $|\emptyset| = 0$.

The concept of cardinality can be extended to infinite sets, but that is a more advanced topic beyond the scope of this chapter. For our purposes, cardinality is a simple count of the elements.

Power Sets

One of the most powerful constructs in set theory is the idea of creating a set that contains all possible subsets of another set.

Definition 3.7 (Power Set)

The **power set** of a set A , denoted $\mathcal{P}(A)$, is the set of all subsets of A . The elements of the power set are themselves sets.



The power set always includes the empty set (\emptyset) and the set A itself.

Example 3.3 Finding the Power Set of $A = \{1, 2, 3\}$

To find $\mathcal{P}(A)$, we list all possible subsets of A , grouped by their cardinality:

- Subsets of size 0: $\{\emptyset\}$
- Subsets of size 1: $\{\{1\}, \{2\}, \{3\}\}$
- Subsets of size 2: $\{\{1, 2\}, \{1, 3\}, \{2, 3\}\}$
- Subsets of size 3: $\{\{1, 2, 3\}\}$

Combining all these subsets into a single set gives us the power set:

$$\mathcal{P}(A) = \{\emptyset, \{1\}, \{2\}, \{3\}, \{1, 2\}, \{1, 3\}, \{2, 3\}, \{1, 2, 3\}\}$$

Remark: For a finite set A with cardinality $|A| = n$, the cardinality of its power set is $|\mathcal{P}(A)| = 2^n$.

This is a crucial formula for computer science. The reason is that for each of the n elements in set A , we can make a binary choice: either we **include** it in a subset or we **exclude** it. With two choices for each of the n elements, there are $2 \times 2 \times \dots \times 2$ (n times), or 2^n , total possible combinations, which corresponds to the total number of possible subsets.

3.5 Operations on Sets

Just as we can perform arithmetic operations on numbers, we can perform operations on sets to create new sets. These operations form the foundation of set algebra. For a software engineer, the most powerful insight is that set operations are a direct parallel to the logical operations of **Boolean algebra**. Every rule you learn for sets has an equivalent rule in logic and digital circuit design.

The Duality of Sets and Boolean Algebra

The relationship between set theory and Boolean algebra is so direct that they are considered "dually isomorphic." This means they are structurally identical. Understanding one is understanding the other.

The key translations are:

Set Theory		Boolean Algebra / Logic
Union (\cup)	\iff	Boolean Sum (+) or OR
Intersection (\cap)	\iff	Boolean Product (\cdot) or AND
Complement (A^c)	\iff	Complementation (\bar{x}) or NOT
The Universal Set (U)	\iff	True (1)
The Empty Set (\emptyset)	\iff	False (0)

We will now explore the primary set operations, keeping this duality in mind.

Intersection

The intersection of two sets contains only the elements that are common to both sets. It corresponds to the logical AND operator.

Definition 3.8 (Intersection)

The **intersection** of sets A and B , denoted $A \cap B$, is the set containing all elements that are in both A and B .

$$A \cap B = \{x \mid x \in A \text{ and } x \in B\}$$

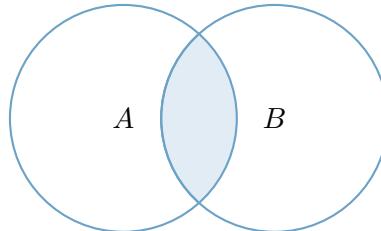


Figure 3.7: The shaded region represents the intersection $A \cap B$.

Union

The union of two sets contains all the elements that appear in either set (or both). It corresponds to the logical OR operator.

Definition 3.9 (Union)

The **union** of sets A and B , denoted $A \cup B$, is the set containing all elements that are in A , or in B , or in both.

$$A \cup B = \{x \mid x \in A \text{ or } x \in B\}$$

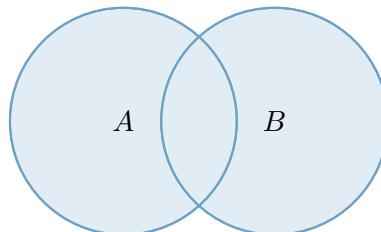


Figure 3.8: The shaded region represents the union $A \cup B$.

Set Difference

The difference between two sets contains the elements that are in the first set but *not* in the second set.

Definition 3.10 (Set Difference)

The **difference** of set A and set B , denoted $A \setminus B$, is the set containing all elements that are in A but not in B .

$$A \setminus B = \{x \mid x \in A \text{ and } x \notin B\}$$

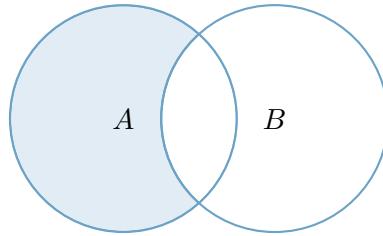


Figure 3.9: The shaded region represents the difference $A \setminus B$.

Symmetric Difference

The symmetric difference contains all elements that are in one set or the other, but not in both. It corresponds to the logical XOR operator.

Definition 3.11 (Symmetric Difference)

The **symmetric difference** of sets A and B , denoted $A \oplus B$, is the set of elements which are in either of the sets, but not in their intersection.

$$A \oplus B = (A \cup B) \setminus (A \cap B)$$

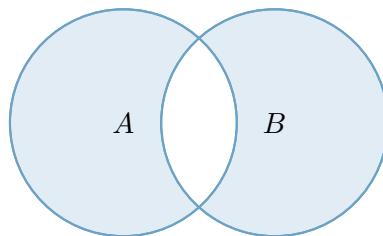


Figure 3.10: The shaded region represents the symmetric difference $A \oplus B$.

Complement

The complement of a set contains all the elements in the universal set that are *not* in the set itself. It corresponds to the logical NOT operator.

Definition 3.12 (Complement)

The **complement** of a set A , denoted A^c , is the set of all elements in the universal set U that are not in A .

$$A^c = U \setminus A$$

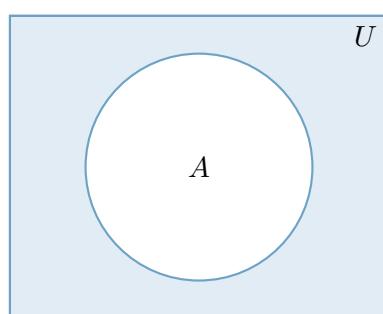


Figure 3.11: The shaded region represents the complement A^c .

3.6 Cartesian Products and Tuples

While sets are unordered collections, we often need to work with ordered collections in computer science, such as coordinates, database records, or structured data. The Cartesian product is the set operation that allows us to create these ordered structures.

Cartesian Product

The Cartesian product creates a new set from two or more existing sets, consisting of all possible ordered combinations of their elements.

Definition 3.13 (Cartesian Product)

The **Cartesian product** of sets A and B , denoted $A \times B$, is the set of all possible ordered pairs (a, b) , where the first element a is from A and the second element b is from B .

$$A \times B = \{(a, b) \mid a \in A \text{ and } b \in B\}$$



The name comes from the Cartesian coordinate system, where any point on a 2D plane can be represented by an ordered pair (x, y) from the Cartesian product of the real numbers, $\mathbb{R} \times \mathbb{R}$.

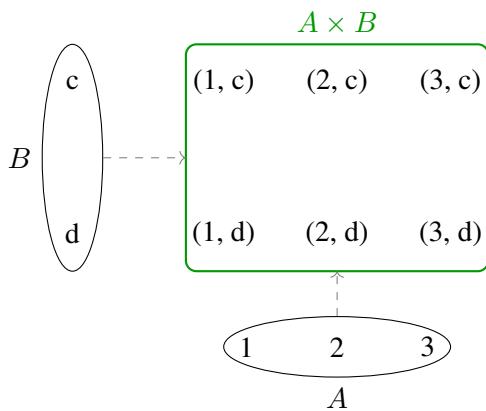


Figure 3.12: The Cartesian product of $A = \{1, 2, 3\}$ and $B = \{c, d\}$ results in a set of 6 ordered pairs.

Remark: The order of the sets in a Cartesian product matters. The set $A \times B$ is generally not equal to the set $B \times A$. For example, the pair $(1, c)$ is in $A \times B$, but the pair $(c, 1)$ would be in $B \times A$.

Tuples

The elements of a Cartesian product are called **tuples**. If the product involves n sets, its elements are called **n-tuples**.

- A **2-tuple**, such as (a, b) , is more commonly known as an **ordered pair**.
- A **3-tuple** has the form (a, b, c) .
- An **n-tuple** has the form (a_1, a_2, \dots, a_n) .

The defining characteristic of a tuple is that **order matters**. This makes tuples fundamentally different from sets.

- For sets: $\{1, 2, 3\} = \{3, 2, 1\}$
- For tuples: $(1, 2, 3) \neq (3, 2, 1)$

This property makes tuples ideal for representing data where the position of an element carries meaning, such as a database record '(UserID, Name, Email)'.

3.7 Proving Set Equalities

In software development, simplifying complex conditional logic is crucial for writing efficient and readable code. Similarly, in set theory, we often need to prove that two different expressions describe the exact same set. There are two primary methods for this: using membership tables and applying set identities.

Method 1: Membership Tables

A membership table is a tool used to prove that two set expressions are equal by checking every possible combination of an element's membership in the constituent sets.

This method is the set-theory equivalent of using a **truth table** in Boolean algebra. Instead of checking for TRUE or FALSE, we check if an element is a member (represented by a 1) or not a member (represented by a 0) of a set. If the columns for two different set expressions are identical in all rows, the expressions are proven to be equal.

Example 3.4 Showing that $A \cap B = B \setminus (B \setminus A)$

To prove this equality, we construct a membership table for all combinations of membership in sets A and B .

A	B	$A \cap B$	$B \setminus A$	$B \setminus (B \setminus A)$
1	1	1	0	1
1	0	0	0	0
0	1	0	1	0
0	0	0	0	0

Solution: (Example 3.4) Since the column for $A \cap B$ is identical to the column for $B \setminus (B \setminus A)$ for all possible membership combinations, the two expressions are equal. ◀

Method 2: Set Identities

While membership tables are effective, they can become very large as the number of sets increases. A more algebraic approach is to use **set identities** to simplify one expression until it matches the other. These identities are fundamental laws that govern how set operations behave.

Table 3.1 shows the most important set identities. Notice the striking similarity to the laws of Boolean algebra — they are structurally identical.

Table 3.1: Fundamental Set Identities and their Boolean Algebra Counterparts.

Identity Name	Set Identity	Boolean Identity
Identity Laws	$A \cup \emptyset = A$	$x + 0 = x$
	$A \cap U = A$	$x \cdot 1 = x$
Domination Laws	$A \cup U = U$	$x + 1 = 1$
	$A \cap \emptyset = \emptyset$	$x \cdot 0 = 0$
Idempotent Laws	$A \cup A = A$	$x + x = x$
	$A \cap A = A$	$x \cdot x = x$
Complement Laws	$A \cup A^c = U$	$x + \bar{x} = 1$
	$A \cap A^c = \emptyset$	$x \cdot \bar{x} = 0$
Commutative Laws	$A \cup B = B \cup A$	$x + y = y + x$
	$A \cap B = B \cap A$	$x \cdot y = y \cdot x$
Associative Laws	$(A \cup B) \cup C = A \cup (B \cup C)$	$(x + y) + z = x + (y + z)$
	$(A \cap B) \cap C = A \cap (B \cap C)$	$(x \cdot y) \cdot z = x \cdot (y \cdot z)$
Distributive Laws	$A \cap (B \cup C) = (A \cap B) \cup (A \cap C)$	$x \cdot (y + z) = (x \cdot y) + (x \cdot z)$
	$A \cup (B \cap C) = (A \cup B) \cap (A \cup C)$	$x + (y \cdot z) = (x + y) \cdot (x + z)$
De Morgan's Laws	$(A \cap B)^c = A^c \cup B^c$	$\overline{x \cdot y} = \bar{x} + \bar{y}$
	$(A \cup B)^c = A^c \cap B^c$	$\overline{x + y} = \bar{x} \cdot \bar{y}$
Absorption Laws	$A \cup (A \cap B) = A$	$x + (x \cdot y) = x$
	$A \cap (A \cup B) = A$	$x \cdot (x + y) = x$

Example 3.5 Proving $A \cup (B \cap A^c) = A \cup B$ using identities

$$\begin{aligned} A \cup (B \cap A^c) &= (A \cup B) \cap (A \cup A^c) && \text{by Distributive Law} \\ &= (A \cup B) \cap U && \text{by Complement Law} \\ &= A \cup B && \text{by Identity Law} \end{aligned}$$

Example 3.6 Determine whether $(A \cap B) \cup (A \cap B^c) = A$

Solution: Let's simplify the left side using set identities:

$$\begin{aligned} (A \cap B) \cup (A \cap B^c) &= A \cap (B \cup B^c) && \text{by Distributive Law} \\ &= A \cap U && \text{by Complement Law} \\ &= A && \text{by Identity Law} \end{aligned}$$

Since the left side simplifies to A , we have $(A \cap B) \cup (A \cap B^c) = A$. **The expressions are equal.**



Example 3.7 Determine whether $(A \cup B) \cap (A^c \cup B^c) = A \cap B$

Solution: Let's simplify the left side:

$$\begin{aligned} (A \cup B) \cap (A^c \cup B^c) &= (A \cup B) \cap (A \cap B)^c && \text{by De Morgan's Law} \\ &= (A \cup B) \cap (A \cap B)^c \end{aligned}$$

This represents all elements that are in A or B but not in both A and B simultaneously. This is the symmetric difference $A \oplus B$, not $A \cap B$.

For a concrete counterexample, let $A = \{1, 2\}$ and $B = \{2, 3\}$:

- $(A \cup B) \cap (A^c \cup B^c) = \{1, 2, 3\} \cap \{3, 1\} = \{1, 3\}$
- $A \cap B = \{2\}$

Since $\{1, 3\} \neq \{2\}$, **the expressions are not equal.**



Example 3.8 Determine whether $A \cup (B \cap C) = (A \cup B) \cap (A \cup C)$

Solution: This is actually the Distributive Law for union over intersection. Let's verify:

$$(A \cup B) \cap (A \cup C) = A \cup (B \cap C) \quad \text{by Distributive Law}$$

So we have $A \cup (B \cap C) = (A \cup B) \cap (A \cup C)$. **The expressions are equal.**



Example 3.9 Determine whether $(A \cap B^c) \cup (A^c \cap B) = (A \cup B) \cap (A \cap B)^c$

Solution: The left side is the symmetric difference $A \oplus B$. Let's check the right side:

$$\begin{aligned} (A \cup B) \cap (A \cap B)^c &= (A \cup B) \cap (A^c \cup B^c) && \text{by De Morgan's Law} \\ &= (A \cap A^c) \cup (A \cap B^c) \cup (B \cap A^c) \cup (B \cap B^c) && \text{by Distributive Law} \\ &= \emptyset \cup (A \cap B^c) \cup (A^c \cap B) \cup \emptyset && \text{by Complement Law} \\ &= (A \cap B^c) \cup (A^c \cap B) && \text{by Identity Law} \end{aligned}$$

Since both sides equal $(A \cap B^c) \cup (A^c \cap B)$, **the expressions are equal.**



3.8 Computer Representation of Sets

While set theory provides the abstract language for collections, computer science requires concrete and efficient ways to implement these ideas. For finite universal sets, one of the most elegant and performant methods is to represent sets using **bit strings**. This technique requires two conditions:

1. The universal set U must be finite.
2. The elements of U must have a fixed, agreed-upon order.

Let $U = \{a_1, a_2, \dots, a_n\}$. Any subset $A \subseteq U$ can be represented by a bit string of length n , where the i -th bit is 1 if $a_i \in A$, and 0 if $a_i \notin A$.

Example 3.10 Representing Sets as Bit Strings

Let the universal set be $U = \{1, 2, 3, 4, 5, 6, 7, 8\}$.

- The set $A = \{1, 3, 4, 8\}$ is represented by the bit string 10110001.
- The set $B = \{2, 3, 8\}$ is represented by the bit string 01100001.
- The set of all even numbers, $\{2, 4, 6, 8\}$, is 01010101.
- The empty set, \emptyset , is 00000000.

The true power of this representation is that set operations map directly to extremely fast, low-level bitwise operations that processors can execute in a single cycle.

Set Operations as Bitwise Operations

Let the bit strings for sets A and B be s_A and s_B .

- **Union** ($A \cup B$) corresponds to a bitwise **OR** operation.
- **Intersection** ($A \cap B$) corresponds to a bitwise **AND** operation.
- **Complement** (A^c) corresponds to a bitwise **NOT** (one's complement) operation.
- **Symmetric Difference** ($A \oplus B$) corresponds to a bitwise **XOR** operation.

Example 3.11 Performing Set Operations with Bit Strings

Using $U = \{1, 2, 3, 4, 5, 6, 7, 8\}$, let $A = \{1, 3, 4, 8\}$ and $B = \{2, 3, 8\}$.

	Set Representation	Bit String Representation
Set A	$\{1, 3, 4, 8\}$	10110001
Set B	$\{2, 3, 8\}$	01100001
$A \cup B$	$\{1, 2, 3, 4, 8\}$	10110001 OR 01100001 = 11110001
$A \cap B$	$\{3, 8\}$	10110001 AND 01100001 = 00100001
A^c	$\{2, 5, 6, 7\}$	NOT 10110001 = 01001110

This bit string representation is fundamental in many areas of computing, including file permissions in operating systems, database indexing, network protocols, and graphics programming, as it provides a way to manage and query collections with maximum efficiency.

Chapter 4 Combinatorics and Probability Theory

Imagine you are tasked with forming teams of 3 for a semester project in a class of 45 students. Initially, the order in which you choose the team members does not matter, so you are just concerned with combinations. The number of ways to form a team of 3 from 45 students comes out to 14,190 possibilities!

The following semester introduces the students to Scrum project management, where each team must have three specific roles: Scrum Master, Product Owner, and Development Team. This small change, specifying roles, suddenly transforms the problem from a simple *combination* into a *permutation*. Now, the number of possible ways to assign these roles leaps to 85,140!

Frustrated by the sheer number of options, the 45 students throw a party to relax. Being well-mannered, they decide that everyone should shake hands with every other person exactly once. After a few minutes, they calculate the total number of handshakes — 990. The students are once again surprised by how something as simple as shaking hands can add up so quickly.

As the night progresses, one student proposes a fun game — a random drawing for five door prizes, each unique. With 45 students in attendance and only five prizes available, the chance of winning nothing becomes a concerning 89 per cent. The students quickly realize that the odds are not in their favor.

Not ready to give up on their luck, a smaller group decides to flip a coin 10 times, with the hope of landing exactly five heads to win the game. However, when they learn that the probability of this happening is only about 25 per cent, their spirits dampen further.

The students conclude that rather than relying on chance, it is time to dive deeper into understanding combinatorics and probability theory. Armed with this knowledge, they can better predict outcomes and avoid future disappointments at both parties and project planning.

4.1 Sample Space and Events

A **random experiment** is one that can lead to different outcomes, even when repeated under the same conditions. This randomness is a fundamental aspect of many engineering tasks.

Think of it this way: Let us say you are testing the speed of a website under different conditions. Sometimes it loads quickly, and other times it is slower. Even if you are using the same code and server, things like network traffic or server load make the results vary each time.

Definition 4.1 (Random Experiment)

A random experiment is one that can give different results, even if you do everything the same each time. 

Or, imagine you are measuring the signal strength in a wireless device. You might get slightly different readings each time because of things like interference or small changes in the environment.

This randomness shows up all over the place, from software performance tests to electrical engineering experiments. It is important to expect it and include it in your thinking. Otherwise, you might make decisions based

on incomplete or misleading data. When you account for random variation, you can make smarter predictions and designs.

To model and analyze a random experiment, it is crucial to understand the set of possible outcomes that can occur. In probability theory, this set is called the **sample space**, denoted by S . A sample space can be either **discrete** (consisting of a finite or countably infinite set of outcomes) or **continuous** (containing an interval of real numbers). The exact definition of a sample space often depends on the objectives of the analysis.

An **outcome** is a single possible result of the random experiment, and an **event** is any subset of the sample space, which may consist of one or more outcomes. Below are some examples to illustrate these concepts:

Example 4.1 Network Latency

Consider an experiment where you measure the latency of data packets in a network. The sample space can be defined based on the type of measurements:

- If latency is measured as a positive real number, the sample space is continuous:

$$S = \{x \mid x > 0\}.$$

- If it is known that latency ranges between 10 and 100 milliseconds, the sample space can be refined to:

$$S = \{x \mid 10 \leq x \leq 100\}.$$

- If the objective is to categorize latency as low, medium, or high, the sample space becomes discrete:

$$S = \{\text{low, medium, high}\}.$$

- For a simple evaluation of whether the latency meets a standard threshold, the sample space can be reduced to:

$$S = \{\text{pass, fail}\}.$$

Each outcome in these sample spaces represents a single possible latency measurement, and events can be defined as sets of outcomes, such as "latency is high."

Understanding the nature of sample spaces, outcomes, and events is fundamental in probability, as it allows us to define and work with probabilities of complex scenarios in various engineering contexts. Let us summarise these key concepts:

Definition 4.2 (Sample Spaces, Outcomes, and Events)

Sample Space: The set of all possible outcomes of a random experiment is called the sample space, denoted by S . Outcomes can be discrete or continuous, depending on the nature of the experiment.

Outcome: A single possible result of a random experiment.

Event: Any subset of the sample space, which may consist of one or more outcomes.



Example 4.2 Software Release Testing

Imagine a software testing process where each test case can either pass or fail. The sample space for a single test case is discrete and can be represented as:

$$S = \{\text{pass, fail}\}.$$

If you run three test cases, the combined sample space for all possible outcomes is:

$$\begin{aligned} S = \{ & (\text{pass}, \text{pass}, \text{pass}), (\text{pass}, \text{pass}, \text{fail}), (\text{pass}, \text{fail}, \text{pass}), \\ & (\text{pass}, \text{fail}, \text{fail}), (\text{fail}, \text{pass}, \text{pass}), (\text{fail}, \text{pass}, \text{fail}), \\ & (\text{fail}, \text{fail}, \text{pass}), (\text{fail}, \text{fail}, \text{fail}) \}. \end{aligned}$$

This sample space includes all sequences of outcomes for the three tests.

- The total number of possible outcomes is $2^3 = 8$.
- An event could be defined as "at least one test fails," which would include outcomes like (fail, pass, pass), (pass, fail, fail), and others where at least one test fails.

Example 4.3 Component Quality in Manufacturing

A company manufactures electronic components, and each component is tested for compliance with quality standards. The test can return one of three outcomes: `pass`, `marginal`, or `fail`. The sample space is:

$$S = \{\text{pass}, \text{marginal}, \text{fail}\}.$$

Event: Suppose we are interested in the event that a component does not pass the quality test. This event is a set of outcomes:

$$E = \{\text{marginal}, \text{fail}\}.$$

This means that all the operations we have defined on sets translate directly into operations on events:

- $A \cup B$: the event that at least one of A or B occurs.
- $A \cap B$: the event that both A and B occur together.
- \bar{A} or A^c : the event that A does not occur.
- $A \setminus B$ or $A - B$: the event that A occurs but B does not.
- $A \Delta B$ or $A \oplus B$: the event that either A or B occurs, but not both (symmetric difference).

Thus, probability theory builds directly on set theory: probability assigns a numerical measure to these subsets of the sample space. In the next chapter, we will see how this measure is defined and used to reason about uncertainty.

We can summarise the considerations of this section in the following table:

Operation	Boolean Algebra	Logic	Set Theory
NOT	\bar{x}	$\neg x$	A^c or A'
OR	$+$	\vee	\cup
AND	\cdot	\wedge	\cap
NAND	$\overline{x \cdot y}$	$\neg(x \wedge y)$	$(A \cap B)^c$ or $\overline{A \cap B}$
NOR	$\overline{x + y}$	$\neg(x \vee y)$	$(A \cup B)^c$ or $\overline{A \cup B}$
XOR (Symmetric Difference)	$x \oplus y$	$(x \wedge \neg y) \vee (\neg x \wedge y)$	$A \Delta B$ or $A \oplus B$
Difference	$x \cdot \bar{y}$	$x \wedge \neg y$	$A - B$ or $A \setminus B$

Table 4.1: Comparison of Operators in Boolean Algebra, Logic, and Set Theory

In this book, we will use the notation A^c to denote the complement of the set A , which includes all elements not in A .

4.2 Counting Principles

In many problems across mathematics, computer science, and engineering, determining the number of ways certain events can occur is important. Whether you're arranging elements, selecting groups, or navigating through complex scenarios, counting techniques provide the foundational tools to solve these problems. These techniques go beyond simple arithmetic and allow us to tackle questions like:

- How many ways can we arrange a set of objects?
- In how many different paths can a process unfold?
- What is the probability of a specific event occurring given multiple possibilities?

Counting techniques, such as permutations, combinations, and the multiplication rule, help us quantify these possibilities systematically.

Multiplication Rule

We start out by discussing the most basic counting principle: the **multiplication rule**:

Theorem 4.1 (Multiplication Rule)

Let an operation be described as a sequence of k steps. Assume the following conditions:

- There are n_1 ways to complete step 1.
- There are n_2 ways to complete step 2 for each way of completing step 1.
- There are n_3 ways to complete step 3 for each way of completing step 2, and so on.

Then, the total number of ways to complete the entire operation is given by:

$$n_1 \times n_2 \times \cdots \times n_k.$$



Example 4.4 Suppose you are choosing a meal at a restaurant. You have the following options:

- 3 choices for the main course.
- 4 choices for the side dish.
- 2 choices for the drink.

Using the multiplication rule, the total number of ways to choose a meal is:

$$3 \times 4 \times 2 = 24.$$

Therefore, there are 24 different meal combinations available.

Example 4.5 Automobile Options

An automobile manufacturer provides vehicles equipped with selected options. Each vehicle is ordered

- With or without an automatic transmission
- With or without a sunroof
- With one of three choices of a stereo system

- With one of four exterior colors

If the sample space consists of the set of all possible vehicle types, what is the number of outcomes in the sample space?

Solution: Using the multiplication rule, we can calculate the total number of possible vehicle types by multiplying the number of choices for each option:

- 2 choices for the transmission (with or without automatic transmission)
- 2 choices for the sunroof (with or without sunroof)
- 3 choices for the stereo system
- 4 choices for the exterior color

Therefore, the total number of possible vehicle types is:

$$2 \times 2 \times 3 \times 4 = 48$$

So, there are 48 different possible vehicle types in the sample space. 

Replacement and Order in Counting

Next we turn to an important distinction between with and without replacement in counting principles:

Definition 4.3 (Counting with and without Replacement)

When counting the number of ways to select objects from a set, two common scenarios are:

- With Replacement:** An object can be selected more than once.
- Without Replacement:** Once an object is selected, it cannot be chosen again.



Example 4.6 Suppose you have a bag containing 5 different colored balls. You draw 2 balls:

- With Replacement:** The first ball is placed back in the bag before drawing the second. There are $5 \times 5 = 25$ possible outcomes.
- Without Replacement:** The first ball is not placed back, so the number of outcomes is $5 \times 4 = 20$.

Example 4.7 Three people are drawing cards one after another from a standard deck of 52 cards. The goal is to find the Ace of Spades. Let's examine the two scenarios: with replacement and without replacement.

Without Replacement

In this scenario, each card drawn is not put back into the deck, reducing the total number of cards available after each draw.

- First Draw: The first person has a $\frac{1}{52}$ chance of drawing the Ace of Spades.
- Second Draw: If the first person does not draw the Ace of Spades, there are now 51 cards left, and the second person has a $\frac{1}{51}$ chance of drawing the Ace of Spades.
- Third Draw: If the Ace of Spades has not been drawn by the first two people, the third person has a $\frac{1}{50}$ chance of drawing it.

The probabilities change with each draw because the total number of cards decreases, and previously drawn cards are not available.

With Replacement

In this scenario, each card drawn is returned to the deck and reshuffled before the next person draws. This keeps the total number of cards constant.

- First Draw: The first person has a $\frac{1}{52}$ chance of drawing the Ace of Spades.
- Second Draw: Since the card is replaced and shuffled back into the deck, the second person also has a $\frac{1}{52}$ chance of drawing the Ace of Spades.
- Third Draw: Similarly, the third person has a $\frac{1}{52}$ chance of drawing the Ace of Spades.

The probabilities remain the same for each draw because the deck is reset to its original state after each draw.

Example 4.8 Imagine you have a group of 5 students: Alice, Bob, Charlie, David, and Eve. You need to select 2 of them for different scenarios, illustrating when the order of selection matters and when it does not.

- **Order Matters:** Selecting Alice as Captain and Bob as Assistant Captain is a different outcome than selecting Bob as Captain and Alice as Assistant Captain.

Now, imagine you are simply selecting 2 students to form a study group with no specific roles assigned. Here, the order does not matter.

- **Order does not matter:** Choosing Alice and Bob is considered the same outcome as choosing Bob and Alice; there is no distinction between the two orders since there are no assigned roles.

Example 4.8 illustrates the distinction between *permutations* and *combinations*, two fundamental counting principles that are widely used in probability theory and combinatorics.

Definition 4.4 (Permutation and Combination)

- **Permutation (order matters):** Different sequences are counted as distinct outcomes, leading to a higher count.
- **Combination (order does not matter):** Sequences are treated as identical, resulting in a lower count.



We will first discuss permutations, which are used when the order of selection matters.

Permutations

Consider a set of elements, such as $S = \{a, b, c\}$. A permutation of the elements is an ordered sequence of the elements. For example, abc , acb , bac , bca , cab , and cba are all the permutations of the elements of S .

Proposition 4.1 (Permutations of n Distinct Objects)

The number of ordered arrangements (permutations) of n distinct objects is

$$n! = n \cdot (n - 1) \cdots \cdot 2 \cdot 1.$$



This outcome is a direct application of the multiplication rule. To form a permutation, you start by choosing an element for the first position from the total of n elements. Next, you choose an element for the second position from the remaining $n - 1$ elements, then for the third position from the remaining $n - 2$ elements, and continue this way until all positions are filled. Such arrangements are often called linear permutations.

Example 4.9 It is said that any shuffling of a deck of card has only happened once in history. This is because the number of ways to shuffle a deck of 52 cards is $52!$, which is an astronomically large number.

$$52! \approx 8.07 \times 10^{67}$$

Example 4.10 Suppose you have 5 different books on a shelf. You want to rearrange them in a different order. The number of ways to rearrange the books is

$$5! = 5 \times 4 \times 3 \times 2 \times 1 = 120$$

There are cases where we are only interested in arranging a subset of elements from a larger set. The formula for counting these arrangements also derives from the multiplication rule.

Theorem 4.2 (Permutations of Subsets)

For integers $n \geq r \geq 0$, the number of ordered selections of r distinct objects from n distinct objects is

$$P_r^n = P(n, r) = n \cdot (n - 1) \cdots (n - r + 1) = \frac{n!}{(n - r)!}.$$



Example 4.11 Suppose you have 5 different books on a shelf, and you want to rearrange 3 of them in a different order. The number of ways to rearrange the 3 books is

$$P_3^5 = 5 \times 4 \times 3 = 60$$

Example 4.12 There are 10 entries in a contest. Only three will win, 1st, 2nd, or 3rd prize. What are the possible results?

Solution: The number of ways to award the prizes is the number of permutations of 3 objects selected from 10, which is

$$P_3^{10} = \frac{10!}{(10 - 3)!} = \frac{10!}{7!} = 10 \times 9 \times 8 = 720$$

Therefore, there are 720 possible outcomes for awarding the prizes. 

Combinations

When the order of selection does not matter, we use the concept of combinations. Combinations are used when we are interested in selecting a subset of elements from a larger set without regard to the order in which they are selected. Let us start out with a couple of examples to illustrate the concept of combinations.

Example 4.13 Suppose you have a group of 5 students: Alice, Bob, Charlie, David, and Eve. You need to select 2 of them to form a study group. The order in which you select the students does not matter. The possible combinations are:

- Alice and Bob
- Alice and Charlie
- Alice and David
- Alice and Eve
- Bob and Charlie

- Bob and David
- Bob and Eve
- Charlie and David
- Charlie and Eve
- David and Eve

The order of the students in the study group does not matter, so the combinations are considered identical.

Example 4.14 Maria has three tickets for a concert. She'd like to use one of the tickets herself. She could then offer the other two tickets to any of four friends (Ann, Beth, Chris, Dave). How many ways can 2 people be selected from 4 to go to a concert?

Example 4.15 A circuit board has four different locations in which a component can be placed. If three identical components are to be placed on the board, how many different designs are possible?

Solution: Since you can only place one component in each slot, placing a component in any slot immediately restricts the choices for the next component.

1. Fill slots 1, 2, and 3.
2. Fill slots 1, 2, and 4.
3. Fill slots 1, 3, and 4.
4. Fill slots 2, 3, and 4.



These examples illustrate the concept of combinations, where the order of selection does not matter. The formula for combinations is derived from the permutation formula by dividing out the number of ways to arrange the r elements.

Theorem 4.3 (Combinations)

The number of combinations of r elements selected from a set of n different elements is given by

$$C_r^n = \binom{n}{r} = \frac{n!}{r!(n-r)!}$$



This is also sometimes referred to as the **binomial coefficient**, denoted by $\binom{n}{r}$, which is read as "n choose r". It is called the binomial coefficient because it appears in the binomial theorem, which expands the powers of a binomial expression:

Theorem 4.4 (Binomial Theorem)

In algebra, the binomial coefficient is used to expand powers of binomials. According to the binomial theorem

$$(a + b)^n = \sum_{k=0}^n \binom{n}{k} a^k b^{n-k}$$



The theorem states that the expansion of the binomial expression $(a + b)^n$ is the sum of the terms $\binom{n}{k} a^k b^{n-k}$ for $k = 0, 1, 2, \dots, n$. The binomial coefficient $\binom{n}{k}$ gives the number of ways to choose k elements from a set

of n elements. We will place no more emphasis on the binomial theorem here, but it is a fundamental concept in algebra and combinatorics, and is widely used in probability theory.

We will conclude our discussion of counting principles with principles of counting with replacement.

Proposition 4.2 (With replacement)

For selections from n types:

- **Ordered with replacement:** n^r outcomes.
- **Unordered with replacement:** $\binom{n+r-1}{r}$ outcomes.



4.3 Basic Probability

Probability quantifies the likelihood or chance that an outcome of a random experiment will occur. For instance, when you hear, “The chance of rain today is 30%,” it expresses our belief about the likelihood of rain. Probabilities are numbers assigned to outcomes, ranging from 0 to 1 (or equivalently, from 0% to 100%). A probability of 0 means the outcome will not happen, while a probability of 1 means it will happen for sure.

Probabilities can be interpreted in different ways:

- **Objective (or Classical) Probability:** Often referred to as classical probability, this approach is used when outcomes are equally likely, such as in rolling a fair die or flipping a coin. Probabilities are assigned based on the assumption that each outcome has an equal chance of occurring. For example, when rolling a fair six-sided die, the probability of rolling a 3 is $\frac{1}{6}$ because there are 6 equally likely outcomes (1, 2, 3, 4, 5, 6), and only one of them is a 3. The probability is the same for all observers.
- **Relative Frequency (Empirical Probability):** Empirical probability is based on observations from experiments rather than theoretical calculations. For example, if a software tester runs a stress test on a server 100 times, and it crashes 7 times, the empirical probability of a crash is $\frac{7}{100} = 0.07$. This approach relies on actual data rather than assumptions or intuition.
- **Subjective Probability:** This reflects our personal belief or degree of confidence in an outcome. Different people might assign different probabilities to the same event based on their knowledge or perspective. You and your friends discuss Denmark’s chances of winning the World Cup. Based on recent performance and team strength, you estimate a 10% chance. However, a more optimistic friend assigns a 20% chance, while another gives only 5%, considering stronger competitors. This illustrates subjective probability, where each person’s estimate varies based on personal beliefs and biases rather than objective data.

When assigning probabilities, it’s essential that the sum of all probabilities in an experiment equals 1, ensuring consistency with the relative frequency interpretation.

We start by establishing the Axioms of Probability, which lay the foundation for how probabilities are assigned to events. These axioms define the basic properties that every probability measure must satisfy.

Axiom 4.1 (Axioms of Probability)

- **Axiom 1:** For any event A , $0 \leq P(A) \leq 1$.
- **Axiom 2:** Probability of the sample space S is $P(S) = 1$.
- **Axiom 3:** If A_1, A_2, A_3, \dots are disjoint events, then $P(A_1 \cup A_2 \cup A_3 \dots) = P(A_1) + P(A_2) + \dots$

$$P(A_3) + \dots$$



The property that $0 \leq P(A) \leq 1$ is equivalent to the requirement that a relative frequency must be between 0 and 1. The property that $P(S) = 1$ is a consequence of the fact that an outcome from the sample space occurs on every trial of an experiment. Consequently, the relative frequency of S is 1. Property 3 implies that if the events A_1 and A_2 have no outcomes in common, the relative frequency of outcomes in $A_1 \cup A_2$ is the sum of the relative frequencies of the outcomes in A_1 and A_2 .

In the next sections we will see more about the probability of events and how to calculate them.

Probability of an Event

The probability of an event is a measure of the likelihood that the event will occur. It is denoted by $P(A)$, where A is the event. The probability of an event ranges from 0 to 1, where 0 indicates that the event will not occur, and 1 indicates that the event will occur for sure.

Definition 4.5 (Probability of an Event)

The probability of an event A , denoted by $P(A)$, is the likelihood that event A will occur. It is defined as the ratio of the number of favorable outcomes to the total number of outcomes in the sample space.

$$P(A) = \frac{\text{Number of favorable outcomes}}{\text{Total number of outcomes}}$$



Example 4.16 Suppose you are testing a software module with 10 different test cases. Out of these, 3 test cases are known to fail due to a bug. If you randomly select one test case to run, what is the probability that the selected test case will fail?

Solution: Here, the event A is "the test case fails."

- Number of favorable outcomes (failing test cases) = 3
- Total number of outcomes (total test cases) = 10

Using the formula:

$$P(A) = \frac{\text{Number of favorable outcomes}}{\text{Total number of outcomes}} = \frac{3}{10} = 0.3$$

Therefore, the probability that a randomly selected test case will fail is 0.3 , indicating that there is a 30% chance of failure.



Example 4.17 Imagine a software development environment where you have 50 files, consisting of 20 Python scripts, 15 Java files, and 15 configuration files. If you randomly select one file to edit, what is the probability that the file is a Python script?

Solution: Here, the event A is "the selected file is a Python script."

- Number of favorable outcomes (Python scripts) = 20

- Total number of outcomes (total files) = 50

Using the formula:

$$P(A) = \frac{\text{Number of favorable outcomes}}{\text{Total number of outcomes}} = \frac{20}{50} = 0.4$$

Thus, the probability of selecting a Python script is 0.4 , meaning there is a 40% chance of choosing a Python file from the set.



4.4 Probability of Joint Events and Set Operations

Joint events are formed by applying basic set operations to individual events. Commonly, we encounter unions of events, such as $A \cup B$; intersections of events, such as $A \cap B$; and complements of events, such as A^c . These combined events are often of particular interest, and their probabilities can frequently be derived from the probabilities of the individual events that compose them. Understanding these set operations is essential for accurately calculating the probability of joint events. In this section, we will explore how unions of events and other set operations can be used to determine the probabilities of more complex events.

When dealing with events, the intersection represents AND while the union represents OR. The probability of the intersection of events A and B , denoted as $P(A \cap B)$, can also be expressed as $P(A, B)$ or $P(AB)$.

From the axioms of probability, we can derive the following rules of probabilities:

Theorem 4.5 (Rules of Probability)

- **Complement Rule:** The probability of the complement of event A is

$$P(A^c) = 1 - P(A)$$

- **Empty Set Rule:** The probability of the empty set is 0, i.e.,

$$P(\emptyset) = 0$$

- **Addition Rule:** For any two events A and B , the probability of the union of events A and B is given by

$$P(A \cup B) = P(A) + P(B) - P(A \cap B)$$

- **Difference Rule:** The probability of the difference between events A and B is given by

$$P(A - B) = P(A) - P(A \cap B)$$

- **Subset Rule:** If A is a subset of B ($A \subset B$), then

$$P(A) \leq P(B)$$



We can obtain the Complement Rule by noting:

$$\begin{aligned} 1 &= P(S) && \text{(axiom 2)} \\ &= P(A \cup A^c) && \text{(definition of complement)} \\ &= P(A) + P(A^c) && \text{(since } A \text{ and } A^c \text{ are disjoint)} \end{aligned}$$

Since $\emptyset = S^c$, we can apply part the Complement Rule to deduce that $P(\emptyset) = 1 - P(S) = 0$. This is intuitive

because, by definition, an event occurs when the outcome of the random experiment is part of that event. However, since the empty set contains no elements, no outcome of the experiment can ever belong to it, making its probability zero.

The Difference Rule can be obtained by showing that $P(A) = P(A \cap B) + P(A - B)$. Note that the two sets $A \cap B$ and $A - B$ are disjoint and their union is A . Thus, by the third axiom of probability

$$\begin{aligned} P(A) &= P((A \cap B) \cup (A - B)) && (\text{since } A = (A \cap B) \cup (A - B)) \\ &= P(A \cap B) + P(A - B) && (\text{since } A \cap B \text{ and } A - B \text{ are disjoint}) \end{aligned}$$

The Addition Rule we obtain by noting that A and $B - A$ are disjoint sets and their union is $A \cup B$. Thus,

$$\begin{aligned} P(A \cup B) &= P(A \cup (B - A)) && (\text{since } A \cup B = A \cup (B - A)) \\ &= P(A) + P(B - A) && (\text{since } A \text{ and } B - A \text{ are disjoint}) \\ &= P(A) + P(B) - P(A \cap B) && (\text{by part the Difference Rule}) \end{aligned}$$

And finally the Subset Rule is a direct consequence of the fact that if $A \subset B$, then B can be written as the union of A and $B - A$. Since A and $B - A$ are disjoint, we have $P(B) = P(A) + P(B - A) \geq P(A)$.

We conclude this section with a few examples illustrating the application of these rules to calculate probabilities of joint events.

Example 4.18 A company has bid on two large construction projects. The company president believes that the probability of winning the first contract is 0.6, the probability of winning the second contract is 0.4, and the probability of winning both contracts is 0.2.

- (a) What is the probability that the company wins at least one contract?
- (b) What is the probability that the company wins the first contract but not the second contract?
- (c) What is the probability that the company wins neither contract?
- (d) What is the probability that the company wins exactly one contract?

Solution: Let A be the event that the company wins the first contract, and B be the event that the company wins the second contract. Given:

- $P(A) = 0.6$
- $P(B) = 0.4$
- $P(A \cap B) = 0.2$

- (a) The probability that the company wins at least one contract is the probability of the union of events A and B . Using the Addition Rule:

$$\begin{aligned} P(A \cup B) &= P(A) + P(B) - P(A \cap B) \\ &= 0.6 + 0.4 - 0.2 \\ &= 0.8 \end{aligned}$$

Therefore, the probability that the company wins at least one contract is 0.8.

- (b) The probability that the company wins the first contract but not the second contract is the probability of the difference between events A and B . Using the Difference Rule:

$$\begin{aligned}
 P(A - B) &= P(A) - P(A \cap B) \\
 &= 0.6 - 0.2 \\
 &= 0.4
 \end{aligned}$$

Therefore, the probability that the company wins the first contract but not the second contract is 0.4.

- (c) The probability that the company wins neither contract is the probability of the complement of the union of events A and B . Using the Complement Rule:

$$\begin{aligned}
 P((A \cup B)^c) &= 1 - P(A \cup B) \\
 &= 1 - 0.8 \\
 &= 0.2
 \end{aligned}$$

Therefore, the probability that the company wins neither contract is 0.2.

- (d) The probability that the company wins exactly one contract is the probability of the difference between the union of events A and B and the intersection of events A and B . Using the Difference Rule:

$$\begin{aligned}
 P((A \cup B) - (A \cap B)) &= P(A \cup B) - P(A \cap B) \\
 &= 0.8 - 0.2 \\
 &= 0.6
 \end{aligned}$$

So, the probability that the company wins exactly one contract is 0.6.



Example 4.19

- There is a 60 percent chance that it will rain today.
 - There is a 50 percent chance that it will rain tomorrow.
 - There is a 30 percent chance that it does not rain either day.
- (a) The probability that it will rain today or tomorrow
 (b) The probability that it will rain today and tomorrow.
 (c) The probability that it will rain today but not tomorrow.
 (d) The probability that it either will rain today or tomorrow, but not both.

Solution:

- (a) Let A be the event that it rains today, and B be the event that it rains tomorrow.

Given:

$$\begin{aligned}
 P(A \cup B) &= 1 - P((A \cup B)^c) && \text{by the Complement Rule} \\
 &= 1 - P(A^c \cap B^c) && \text{by De Morgan's Law} \\
 &= 1 - 0.3 \\
 &= 0.7
 \end{aligned}$$

Therefore, the probability that it will rain today or tomorrow is 0.7.

(b) The probability that it will rain today and tomorrow: this is $P(A \cap B)$. To find this we note that

$$\begin{aligned} P(A \cap B) &= P(A) + P(B) - P(A \cup B) \\ &= 0.6 + 0.5 - 0.7 \\ &= 0.4 \end{aligned}$$

(c) The probability that it will rain today but not tomorrow: this is $P(A \cap B^c)$.

$$\begin{aligned} P(A \cap B^c) &= P(A - B) \\ &= P(A) - P(A \cap B) \\ &= 0.6 - 0.4 \\ &= 0.2 \end{aligned}$$

(d) The probability that it either will rain today or tomorrow but not both: this is $P(A - B) + P(B - A)$.

We have already found $P(A - B) = .2$. Similarly, we can find $P(B - A)$:

$$\begin{aligned} P(B - A) &= P(B) - P(B \cap A) \\ &= 0.5 - 0.4 \\ &= 0.1 \end{aligned}$$

Thus,

$$\begin{aligned} P(A - B) + P(B - A) &= 0.2 + 0.1 \\ &= 0.3 \end{aligned}$$



Chapter 5 Conditional Probability and Bayes' Theorem

In this section, we introduce more advanced concepts of probability. We begin with **conditional probability**, which assesses the likelihood of an event occurring *given* that another event has already taken place. Building on this, we introduce **the multiplication rule**, a key principle for determining the probability of multiple events happening in sequence. Next, we explore **the law of total probability**, which allows us to break down and calculate probabilities across different scenarios or partitions of the sample space. To further distinguish how events interact, we examine **dependent and independent events**, clarifying how the occurrence of one event influences or does not influence another.

Finally, we look into **Bayes' Theorem**, a powerful tool for updating probabilities in light of new evidence.

5.1 Conditional Probability

Consider the following scenario: in a particular population, 5% of individuals have a specific medical condition. Therefore, the probability that a randomly selected person has the condition is 5%:

$$P(D) = 0.05,$$

where D represents the event that a person has the disease. This probability $P(D)$ is known as the *prior probability*, as it reflects our initial belief about the likelihood of the event before any additional information is obtained.

Now, imagine selecting a random person and being informed that they have tested positive for the disease. With this additional information, how should we update the probability that the person actually has the disease? In other words, what is the probability that a person has the disease given that they tested positive? Let T denote the event that a person tests positive. This conditional probability is expressed as:

$$P(D | T),$$

which represents the probability of D occurring given that T has occurred. This updated probability $P(D | T)$ is known as the *posterior probability*, as it reflects our revised belief about the likelihood of the event after taking the new evidence into account. Intuitively, it is reasonable to expect that $P(D | T)$ is greater than the prior probability $P(D)$. However, what is the exact value of $P(D | T)$? Before introducing a general formula, let us consider a simple example.

Example 5.1 I roll a fair die. Let A be the event that the outcome is a prime number, i.e., $A = \{2, 3, 5\}$. Also, let B be the event that the outcome is greater than or equal to 3, i.e., $B = \{3, 4, 5\}$.

- What is the probability of A , $P(A)$?
- What is the probability of A given B , $P(A | B)$?

Solution:

- The probability of A is the number of outcomes in A divided by the total number of outcomes. Since there are 3 prime numbers and 6 possible outcomes, we have:

$$P(A) = \frac{3}{6} = \frac{1}{2}$$

- (b) The probability of A given B is the number of outcomes in the intersection of A and B divided by the number of outcomes in B . Since the outcomes in the intersection of A and B are $\{3, 5\}$ and the outcomes in B are $\{3, 4, 5\}$, we have:

$$P(A | B) = \frac{|A \cap B|}{|B|} = \frac{|\{3, 5\}|}{|\{3, 4, 5\}|} = \frac{2}{3}$$



Having understood the basic example, we can now generalize the approach to derive a more universal formula for conditional probability. Starting from the specific case, we can manipulate the expression by dividing both the numerator and the denominator by the total number of possible outcomes, $|S|$, as shown below:

$$P(A | B) = \frac{|A \cap B|}{|B|} = \frac{\frac{|A \cap B|}{|S|}}{\frac{|B|}{|S|}} = \frac{P(A \cap B)}{P(B)}$$

Explanation:

- **Numerator:** $|A \cap B|$ represents the number of outcomes where both events A and B occur. Dividing by $|S|$ converts this count into a probability, $P(A \cap B)$.
- **Denominator:** $|B|$ is the number of outcomes where event B occurs. Similarly, dividing by $|S|$ yields the probability of event B , denoted as $P(B)$.

Thus, the conditional probability $P(A | B)$ can be expressed as the ratio of the joint probability of A and B to the probability of B .

While the above derivation assumes a finite sample space with equally likely outcomes, the resulting formula is remarkably general. It holds true regardless of whether the sample space is finite or infinite, and whether the outcomes are equally likely or not. This universality makes the formula a cornerstone of probability theory.

Definition 5.1 (Conditional Probability)

For any two events A and B with $P(B) > 0$, the conditional probability of A given B is defined as:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}, \quad P(B) > 0$$



Here is the intuition behind the formula. When we know that event B has occurred, we effectively eliminate all outcomes that are not part of B . This reduction transforms our original sample space into the subset B .

Within this new, restricted sample space B , the only way for event A to occur is if the outcome lies in the intersection of A and B , denoted by $A \cap B$. To determine the conditional probability $P(A | B)$, we calculate the ratio of the probability of both A and B occurring to the probability of B occurring alone. Mathematically, this is expressed as:

$$P(A | B) = \frac{P(A \cap B)}{P(B)}$$

This formulation ensures that the probabilities within the new sample space B are normalized, meaning the total probability sums to 1. For instance, consider the conditional probability of B given B :

$$P(B | B) = \frac{P(B \cap B)}{P(B)} = \frac{P(B)}{P(B)} = 1$$

This result intuitively confirms that if B has occurred, the probability of B occurring is certain, i.e., 1.

Note that the conditional probability $P(A | B)$ is undefined when $P(B) = 0$. This scenario implies that event B never occurs. Since conditional probability relies on the occurrence of B , if B has a probability of zero, there are no outcomes in the sample space to condition upon. Therefore, discussing the probability of A given B becomes meaningless in this context.

We can also formulate the axioms of probability in terms of conditional probability. The axioms of probability are as follows:

Axiom 5.1 (Axioms of Conditional Probability)

- **Axiom 1:** For any event A , $0 \leq P(A | B) \leq 1$.
- **Axiom 2:** Conditional probability of B given B is 1, i.e., $P(B | B) = 1$.
- **Axiom 3:** If A_1, A_2, A_3, \dots are disjoint events, then

$$P(A_1 \cup A_2 \cup A_3 \dots | B) = P(A_1 | B) + P(A_2 | B) + P(A_3 | B) + \dots$$



And we are also able to derive other rules of conditional probability from these axioms:

Theorem 5.1 (Rules of Conditional Probability)

- **Complement Rule:** The probability of the complement of event A given C is

$$P(A^c | C) = 1 - P(A | C)$$
- **Empty Set Rule:** The probability of the empty set given some event C is 0, i.e.,

$$P(\emptyset | C) = 0$$
- **Addition Rule:** For any two events A and B , the probability of the union of events A and B given another event C is

$$P(A \cup B | C) = P(A | C) + P(B | C) - P(A \cap B | C)$$

- **Difference Rule:** The probability of the difference between events A and B given another event C is

$$P(A - B | C) = P(A | C) - P(A \cap B | C)$$

- **Subset Rule:** If $A \subset B$, then

$$P(A | C) \leq P(B | C)$$



The following example illustrates the application of conditional probability and introduces the concept of **contingency tables** as well as **false positives** and **false negatives**.

Example 5.2 A researcher aims to assess the effectiveness of a diagnostic test designed to detect renal disease in patients with high blood pressure. To achieve this, she conducts the test on a sample of 137 patients, categorized as follows:

- **67 patients** with a confirmed diagnosis of renal disease.
- **70 patients** who are known to be healthy (i.e., do not have renal disease).

The diagnostic test produces one of two possible outcomes for each patient:

- **Positive:** Indicates that the patient has renal disease.
- **Negative:** Indicates that the patient does not have renal disease.

The findings are summarised in the following contingency table:

Test Results				
Truth	Positive	Negative	Total	
Renal Disease	44	23	67	
Healthy	10	60	70	
Total	54	83	137	

In this experiment:

- **True Positives:** Patients who have renal disease and tested positive.
- **False Negatives:** Patients who have renal disease but tested negative.
- **False Positives:** Healthy patients who tested positive.
- **True Negatives:** Healthy patients who tested negative.

Determine the following probabilities:

- (a) The Probability of having renal disease, $P(D)$.
- (b) The Probability of a positive test, $P(T^+)$.
- (c) The Probability of a negative test, $P(T^-)$.
- (d) If a person has renal disease, what is the probability that they test positive for the disease?
- (e) Determine the probability that a patient has renal disease given a positive test result, $P(D | T^+)$.
- (f) Determine the probability that a patient does not have renal disease given a negative test result, $P(\text{Healthy} | T^-)$.
- (g) Assess the overall accuracy of the diagnostic test.

Solution: Using the contingency table, we can calculate the following probabilities:

- (a) **Probability of Renal Disease, $P(D)$:**

$$\begin{aligned} P(D) &= \frac{\text{Number of patients with renal disease}}{\text{Total number of patients}} \\ &= \frac{67}{137} \approx 0.4883 \text{ or } 48.83\% \end{aligned}$$

- (b) **Probability of a Positive Test, $P(T^+)$:**

$$\begin{aligned} P(T^+) &= \frac{\text{Number of positive tests}}{\text{Total number of patients}} \\ &= \frac{54}{137} \approx 0.3942 \text{ or } 39.42\% \end{aligned}$$

- (c) **Probability of a Negative Test, $P(T^-)$:**

$$\begin{aligned} P(T^-) &= \frac{\text{Number of negative tests}}{\text{Total number of patients}} \\ &= \frac{83}{137} \approx 0.6058 \text{ or } 60.58\% \end{aligned}$$

- (d) **Probability of a Positive Test Given Renal Disease, $P(T^+ | D)$:**

$$P(T^+ | D) = \frac{44}{67} \approx 0.6567$$

- (e) **Probability of Renal Disease Given a Positive Test, $P(D | T^+)$:**

$$\begin{aligned} P(D | T^+) &= \frac{P(D \cap T^+)}{P(T^+)} \\ &= \frac{44}{54} = \frac{22}{27} \approx 0.8148 \end{aligned}$$

(f) **Probability of Being Healthy Given a Negative Test, $P(\text{Healthy} | T^-)$:**

$$\begin{aligned} P(\text{Healthy} | T^-) &= \frac{P(\text{Healthy} \cap T^-)}{P(T^-)} \\ &= \frac{60}{83} \approx 0.7229 \end{aligned}$$

(g) **Overall Accuracy of the Diagnostic Test:** The overall accuracy of the diagnostic test is the proportion of correct diagnoses, i.e., the sum of true positives and true negatives divided by the total number of patients:

$$\text{Overall Accuracy} = \frac{44 + 60}{137} = \frac{104}{137} \approx 0.7591$$

The results of the analysis indicate that the diagnostic test has a high probability of correctly identifying patients with renal disease (81.48%). However, the test is less effective at identifying healthy patients, with a probability of 72.29%. The overall accuracy of the test is 75.91%, reflecting the proportion of correct diagnoses across all patients. 

Here's a classic probability puzzle known as the Two-Child Problem. This scenario has appeared in various forms in the literature, each with subtle twists that lead to different results. Before looking at the calculations, take a moment to make your own predictions — you might be surprised by the outcomes!

Example 5.3 Consider a family with two children, and we are interested in the possible biological gender combinations of the children. The sample space for this situation is:

$$S = \{(G, G), (G, B), (B, G), (B, B)\}$$

where G represents a biological girl (hence 'girl') and B represents a biological boy (hence 'boy'). For simplicity, we assume that all four outcomes are equally likely.

- (a) What is the probability that both children are girls given that the first child is a girl?
- (b) We ask the father: "Do you have at least one daughter?" He responds "Yes!" Given this extra information, what is the probability that both children are girls? In other words, what is the probability that both children are girls given that we know at least one of them is a girl?

Solution: Let A be the event that both children are girls, i.e., $A = \{(G, G)\}$. Let B be the event that the first child is a girl, i.e., $B = \{(G, G), (G, B)\}$. Finally, let C be the event that at least one of the children is a girl, i.e., $C = \{(G, G), (G, B), (B, G)\}$. Since the outcomes are equally likely, we can write

$$P(A) = \frac{1}{4}$$

$$P(B) = \frac{2}{4} = \frac{1}{2}$$

$$P(C) = \frac{3}{4}$$

- (a) What is the probability that both children are girls given that the first child is a girl? This is $P(A | B)$, thus we can write

$$\begin{aligned}
 P(A | B) &= \frac{P(A \cap B)}{P(B)} \\
 &= \frac{P(A)}{P(B)} \quad (\text{since } A \subset B) \\
 &= \frac{\frac{1}{4}}{\frac{1}{2}} = \frac{1}{2}
 \end{aligned}$$

- (b) What is the probability that both children are girls given that we know at least one of them is a girl? This is $P(A | C)$, thus we can write

$$\begin{aligned}
 P(A | C) &= \frac{P(A \cap C)}{P(C)} \\
 &= \frac{P(A)}{P(C)} \quad (\text{since } A \subset C) \\
 &= \frac{\frac{1}{4}}{\frac{3}{4}} = \frac{1}{3}
 \end{aligned}$$



Remark: Many people might intuitively guess that both $P(A | B)$ and $P(A | C)$ would be 50%. However, while $P(A | B) = 50\%$, $P(A | C)$ is only 33%. This illustrates how probability can be counterintuitive. The key is to recognize that event B is a subset of event C . Specifically, B excludes the outcome (B, G) , which is included in C . As a result, C has more outcomes not in A than B , leading to a smaller $P(A | C)$ compared to $P(A | B)$.

5.2 Multiplication and Total Probability Rules

The multiplication rule is a fundamental principle in probability theory that allows us to calculate the probability of multiple events occurring in sequence. This rule is particularly useful when the events are dependent, meaning that the occurrence of one event influences the probability of the subsequent event.

Theorem 5.2 (Multiplication Rule)

For any two events A and B , the probability of both events occurring is given by:

$$P(A \cap B) = P(A)P(B | A) = P(B)P(A | B)$$



Example 5.4 Suppose a software development team is testing a new feature. There are two stages of testing: Unit Testing (A) and Integration Testing (B). The probability that a bug is detected in Unit Testing is $P(A) = 0.3$. If a bug is detected during Unit Testing, the probability that it will also be detected in Integration Testing is $P(B | A) = 0.7$. What is the probability that a bug is detected in both stages of testing?

Solution: Using the Multiplication Rule:

$$P(A \cap B) = P(A) \cdot P(B | A) = 0.3 \times 0.7 = 0.21$$

Therefore, the probability that a bug is detected in both Unit Testing and Integration Testing is 0.21.



Example 5.5 A company has implemented a two-layer security system to protect against network breaches: **Firewall (A)** and **Intrusion Detection System (IDS) (B)**. The probability that a breach attempt is detected by the Firewall is $P(A) = 0.4$. If the breach passes through the Firewall, the probability that it is detected by

the IDS is $P(B | A^c) = 0.6$. The probability that the breach is detected by both the Firewall and the IDS is $P(B | A) = 0.8$.

- What is the probability that a breach is detected by at least one of the security layers?
- What is the probability that a breach is detected by both security layers?

Solution:

- Probability of Detection by at Least One Layer:**

To find the probability that the breach is detected by at least one layer, we calculate the probability of a breach passing through both layers undetected and subtract it from 1.

$$P(\text{Undetected}) = P(A^c) \cdot P(B^c | A^c) = (1 - 0.4) \cdot (1 - 0.6) = 0.6 \cdot 0.4 = 0.24$$

Therefore, the probability of detecting the breach with at least one layer is:

$$P(\text{Detected}) = 1 - P(\text{Undetected}) = 1 - 0.24 = 0.76$$

- Probability of Detection by Both Layers:**

Using the Multiplication Rule:

$$P(A \cap B) = P(A) \cdot P(B | A) = 0.4 \cdot 0.8 = 0.32$$

The probability of detecting the breach with at least one layer is 0.76, while the probability of detecting the breach with both the Firewall and the IDS is 0.32. This discrepancy highlights the importance of considering the dependencies between events when calculating probabilities.



Remark: It might seem intuitive to some to calculate the joint probability of detection using $P(A \cap B) = P(A) \cdot P(B) = 0.4 \cdot 0.6 = 0.24$. However, this approach ignores the fact that the probability of detection by the IDS depends on the outcome of the Firewall. This is why we use the conditional probability $P(B | A)$, which properly accounts for the dependency between the two layers, yielding the correct result of 0.32.

In some scenarios, the probability of an event depends on various conditions. By knowing the conditional probabilities under these different scenarios, we can determine the overall probability of the event. For instance, consider semiconductor manufacturing, where we define:

- A as the event that a chip is highly contaminated
- B as the event that a product using the chip fails

The probability of failure for a non-contaminated chip is $P(B | A^c) = 0.005$. On the other hand, if the chip is subjected to high levels of contamination, the probability of failure is $P(B | A) = 0.10$. Suppose that in a given production run, 20% of the chips are highly contaminated ($P(A) = 0.20$). What is the probability that a product using one of these chips fails?

The probability of failure depends on whether the chip was exposed to high contamination or not. For any event A , it can be decomposed into two mutually exclusive parts: one that intersects with B and another that intersects with the complement B^c . Mathematically, we can express this as:

$$A = (B \cap A) \cup (B^c \cap A)$$

This decomposition is visualized in the Venn diagram in Figure 5.1. Since B and B^c are mutually exclusive, their intersections with A are also mutually exclusive. Thus, using the rule for the probability of the union of

mutually exclusive events and the multiplication rule, we can derive the total probability as follows:

$$P(A) = P(A \cap B) + P(A \cap B^c)$$

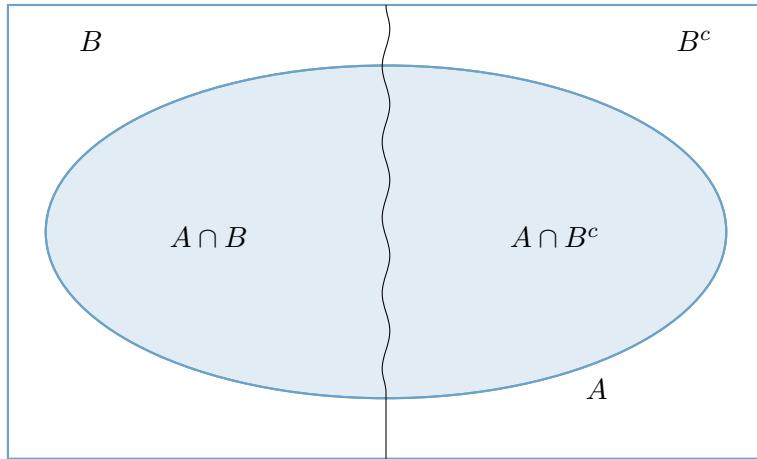


Figure 5.1: $P(A) = P(A \cap B) + P(A \cap B^c)$

We summarise this in the following theorem:

Theorem 5.3 (Law of Total Probability)

For any events A and B such that B and B^c form a partition of the sample space S , the total probability of event A is given by:

$$\begin{aligned} P(A) &= P(A \cap B) + P(A \cap B^c) \\ &= P(B)P(A | B) + P(B^c)P(A | B^c) \end{aligned}$$

For any event A and any partition of the sample space B_1, B_2, \dots, B_n such that $B_i \cap B_j = \emptyset$ for all $i \neq j$ and $\bigcup_{i=1}^n B_i = S$, the total probability of event A is given by:

$$P(A) = \sum_{i=1}^n P(A | B_i)P(B_i)$$



Let us consider some examples.

Example 5.6 A company produces two types of products: **Product A** and **Product B**. The probability that a product is defective is $P(D | A) = 0.05$ for Product A and $P(D | B) = 0.10$ for Product B. The company manufactures 60% of its products as Product A and 40% as Product B. What is the probability that a randomly selected product is defective?

Solution: Using the Law of Total Probability:

$$P(D) = P(D | A)P(A) + P(D | B)P(B)$$

Substituting the given values:

$$P(D) = 0.05 \times 0.60 + 0.10 \times 0.40 = 0.03 + 0.04 = 0.07$$

Therefore, the probability that a randomly selected product is defective is 0.07. ◀

A graphical display of partitioning an event B among a collection of mutually exclusive and exhaustive events is shown in **Figure 5.2**. The event A is partitioned into five mutually exclusive events B_1, B_2, B_3, B_4, B_5 , which

together form the sample space S . The total probability of event A is calculated as the sum of the conditional probabilities of A given each partition B_i multiplied by the probability of each partition $P(B_i)$.

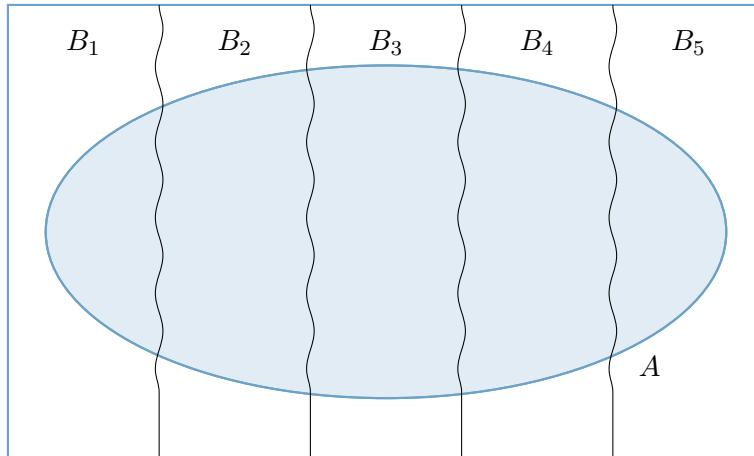


Figure 5.2: $P(A) = \sum_{i=1}^5 P(A | B_i)P(B_i)$

Example 5.7 A university offers three types of courses: **Online**, **Hybrid**, and **In-Person**. The probability that a student fails a course is $P(F | \text{Online}) = 0.10$ for Online courses, $P(F | \text{Hybrid}) = 0.05$ for Hybrid courses, and $P(F | \text{In-Person}) = 0.02$ for In-Person courses. The university offers 50% of its courses as Online, 30% as Hybrid, and 20% as In-Person. What is the probability that a randomly selected student fails a course?

Solution: Using the Law of Total Probability:

$$P(F) = P(F | \text{Online})P(\text{Online}) + P(F | \text{Hybrid})P(\text{Hybrid}) + P(F | \text{In-Person})P(\text{In-Person})$$

Substituting the given values:

$$P(F) = 0.10 \times 0.50 + 0.05 \times 0.30 + 0.02 \times 0.20 = 0.05 + 0.015 + 0.004 = 0.069$$

Therefore, the probability that a randomly selected student fails a course is 0.069. ◀

5.3 Independence

Let A represent the event that it rains tomorrow, with $P(A) = \frac{1}{3}$. Additionally, suppose I toss a fair coin, and let B be the event that it lands heads up, so $P(B) = \frac{1}{2}$.

Now, consider the probability $P(A | B)$. What do you think it would be? You might intuitively guess that $P(A | B) = P(A) = \frac{1}{3}$, and you would be correct! The coin toss outcome has no influence on the weather forecast. This means that whether B occurs or not, the probability of A remains unchanged. This scenario illustrates the concept of independent events: two events are independent if the occurrence of one does not provide any information about the other.

Let's now formalize the definition of independence.

Definition 5.2 (Independence)

Two events are considered independent if the occurrence of one event does not affect the probability of the other event. In other words, the probability of one event does not depend on the occurrence of the

other event. Two events A and B are independent if:

$$\begin{aligned} P(A | B) &= P(A) \\ P(B | A) &= P(B) \\ P(A \cap B) &= P(A)P(B) \end{aligned}$$



In summary, independence can be understood in two equivalent ways: "Independence means that the probability of the intersection of two events can be found by simply multiplying their individual probabilities," or, alternatively, "Independence means that the conditional probability of one event given the other is the same as the original, unconditioned probability."

Lemma 5.1

If A and B are independent then

- A and B^c are independent,
- A^c and B are independent,
- A^c and B^c are independent.



When dealing with the probability of the union of multiple independent events, A_1, A_2, \dots, A_n , it is often easier to find the probability of their intersection than their union. In these situations, De Morgan's Law can be quite useful:

$$A_1 \cup A_2 \cup \dots \cup A_n = (A_1^c \cap A_2^c \cap \dots \cap A_n^c)^c$$

Using this relationship, we can express the probability of the union as:

$$\begin{aligned} P(A_1 \cup A_2 \cup \dots \cup A_n) &= 1 - P(A_1^c \cap A_2^c \cap \dots \cap A_n^c) \\ &= 1 - P(A_1^c)P(A_2^c)\dots P(A_n^c) \quad (\text{since the } A_i\text{'s are independent}) \\ &= 1 - (1 - P(A_1))(1 - P(A_2))\dots(1 - P(A_n)). \end{aligned}$$

Theorem 5.4 (Independence and DeMorgan's Law)

If A_1, A_2, \dots, A_n are independent then

$$P(A_1 \cup A_2 \cup \dots \cup A_n) = 1 - (1 - P(A_1))(1 - P(A_2))\dots(1 - P(A_n))$$



Warning! A common misconception is to confuse independence with disjointness. However, these are fundamentally different concepts. Two events, A and B , are disjoint if the occurrence of one prevents the occurrence of the other, i.e., $A \cap B = \emptyset$. In this case, knowing that A has occurred gives us complete information about B namely, that B cannot occur. This dependence means that disjoint events cannot be independent.

We can extend the concept of independence to multiple events. A set of events A_1, A_2, \dots, A_n are considered independent if the occurrence of any subset of these events does not provide any information about the occurrence of the other events. Mathematically, this is expressed as:

Concept	Description	Key Formulas
Disjoint	Events A and B cannot occur at the same time	$A \cap B = \emptyset$ $P(A \cup B) = P(A) + P(B)$
Independent	Occurrence of B gives no information about A	$P(A B) = P(A)$ $P(B A) = P(B)$ $P(A \cap B) = P(A) \cdot P(B)$

Table 5.1: Comparison of Disjoint and Independent Events.**Definition 5.3 (Independence of Multiple Events)**

A set of events A_1, A_2, \dots, A_k are considered independent if the joint probability is equal to the product of the individual probabilities:

$$P(A_1 \cap A_2 \cap \dots \cap A_k) = P(A_1) \cdot P(A_2) \cdots P(A_k)$$

**Example 5.8** The Gambler's Fallacy

The **Gambler's Fallacy** is a common cognitive bias that arises when individuals believe that the outcome of a random event is influenced by previous outcomes. This fallacy is often observed in gambling scenarios, where individuals incorrectly assume that the probability of an event occurring is affected by past events.

A man tosses a fair coin eight times and observes whether the toss yields a head (H) or a tail (T) on each toss. Which of the following sequences of coin tosses is the man more likely to get a head (H) on his next toss? This one:

TTTTTTTT

or this one:

HHTHTTHH

The answer is neither as illustrated here:

$$\begin{aligned} P(H_9 | T_1 T_2 \dots T_8) &= \frac{P(H_9 \cap T_1 \cap T_2 \cap \dots \cap T_8)}{P(T_1 \cap T_2 \cap \dots \cap T_8)} = \frac{\left(\frac{1}{2}\right)^9}{\left(\frac{1}{2}\right)^8} \\ &= \frac{1}{2} \end{aligned}$$

Try to avoid falling into the trap of the Gambler's Fallacy. For example, if a fair coin lands on tails eight times in a row, it's easy to think that a head is "due" on the next toss. However, each toss is independent, and the probability remains the same — there's still a 50% chance of heads or tails, regardless of past results.

5.4 Bayes' Theorem

We are now ready to introduce one of the most powerful tools in conditional probability: Bayes' rule (or theorem). This rule is particularly useful when we know $P(A | B)$ but want to find $P(B | A)$. Starting with

the definition of conditional probability, we have:

$$P(A | B) \cdot P(B) = P(A \cap B) = P(B | A) \cdot P(A)$$

By dividing both sides by $P(A)$, we arrive at:

$$P(B | A) = \frac{P(A | B) \cdot P(B)}{P(A)}$$

This formula is famously known as Bayes' rule. In many cases, to calculate $P(A)$ in Bayes' rule, we need the law of total probability. Therefore, Bayes' rule is often presented in the form:

$$P(B | A) = \frac{P(A | B) \cdot P(B)}{P(A | B) \cdot P(B) + P(A | B^c) \cdot P(B^c)}$$

or more generally

$$P(B_j | A) = \frac{P(A | B_j) \cdot P(B_j)}{\sum_i P(A | B_i) \cdot P(B_i)}$$

where B_1, B_2, \dots, B_n form a partition of the sample space.

Theorem 5.5 (Bayes' Theorem)

For any two events A and B , where $P(A) \neq 0$, we have

$$P(B | A) = \frac{P(A | B) \cdot P(B)}{P(A)} = \frac{P(A | B) \cdot P(B)}{P(A | B) \cdot P(B) + P(A | B^c) \cdot P(B^c)}$$

If B_1, B_2, B_3, \dots form a partition of the sample space S , and A is any event with $P(A) \neq 0$, we have

$$P(B_j | A) = \frac{P(A | B_j) \cdot P(B_j)}{\sum_i P(A | B_i) \cdot P(B_i)}$$



We start with an infamous example to highlight the importance of Bayes' Theorem.

Example 5.9 False Positive Paradox

Imagine a rare disease that affects about 1 in every 10,000 people. There is a test available to detect this disease, and while the test is highly accurate, it is not perfect. Specifically:

- The probability that the test shows a positive result (indicating the disease) when the person does not have the disease is 2%.
- The probability that the test shows a negative result (indicating no disease) when the person does have the disease is 1%.

Now, suppose a randomly selected person takes the test, and the result comes back positive. What is the probability that this person actually has the disease?

Solution: Let D be the event that the person has the disease, let T^+ be the event that the test result is positive, and let T^- be the event that it is negative. We are given:

$$P(D) = \frac{1}{10,000}$$

$$P(T^+ | D^c) = 0.02 \quad (\text{False Positive Rate})$$

$$P(T^- | D) = 0.01 \quad (\text{False Negative Rate})$$

We want to compute $P(D | T^+)$. First, we need $P(T^+ | D)$, which is the probability of a true positive. Since a person with the disease will either test positive or negative, we have:

$$P(T^+ | D) = 1 - P(T^- | D) = 1 - 0.01 = 0.99$$

Now, using Bayes' rule:

$$\begin{aligned} P(D | T^+) &= \frac{P(T^+ | D)P(D)}{P(T^+ | D)P(D) + P(T^+ | D^c)P(D^c)} \\ &= \frac{0.99 \times 0.0001}{0.99 \times 0.0001 + 0.02 \times (1 - 0.0001)} \\ &= \frac{0.000099}{0.000099 + 0.02 \times 0.9999} \\ &= \frac{0.000099}{0.000099 + 0.019998} \\ &\approx 0.0049 \end{aligned}$$

This result means there is less than half a percent chance that the person actually has the disease. Despite the positive test result, the low prevalence of the disease and the test's false positive rate contribute to this counterintuitive outcome. 

Example 5.10 Bayesian networks are commonly used on the websites of high-technology manufacturers to help customers quickly diagnose issues with their products. For instance, a printer manufacturer uses data from test results to identify potential causes of printer failures. Printer failures are primarily associated with three types of problems: hardware, software, and other issues (like connectors). The probabilities of these problems are as follows:

- Probability of a hardware issue: $P(H) = 0.1$
- Probability of a software issue: $P(S) = 0.6$
- Probability of another type of issue: $P(O) = 0.3$

The likelihood of a printer failing, given each type of problem, is:

- Probability of failure given a hardware issue: $P(F | H) = 0.9$
- Probability of failure given a software issue: $P(F | S) = 0.2$
- Probability of failure given another issue: $P(F | O) = 0.5$

Given that a customer experiences a printer failure and uses the manufacturer's website to diagnose the issue, what is the most likely cause of the problem?

Solution: To determine the most likely cause of the printer failure, we need to calculate the probability of each type of problem given that a failure has occurred. This involves using Bayes' theorem to compute the posterior probabilities.

Let F denote the event of a printer failure. We want to find:

$$P(H | F), \quad P(S | F), \quad P(O | F)$$

Step 1: Calculate the Total Probability of Failure

First, we use the law of total probability to find $P(F)$:

$$P(F) = P(F | H) \cdot P(H) + P(F | S) \cdot P(S) + P(F | O) \cdot P(O)$$

Substituting the given values:

$$P(F) = (0.9 \times 0.1) + (0.2 \times 0.6) + (0.5 \times 0.3)$$

$$P(F) = 0.09 + 0.12 + 0.15 = 0.36$$

Step 2: Apply Bayes' Theorem to Find the Posterior Probabilities

Now, apply Bayes' theorem for each problem type:

1. Probability of Hardware Problem Given Failure:

$$P(H | F) = \frac{P(F | H) \cdot P(H)}{P(F)} = \frac{0.9 \times 0.1}{0.36} = \frac{0.09}{0.36} = 0.25$$

2. Probability of Software Problem Given Failure:

$$P(S | F) = \frac{P(F | S) \cdot P(S)}{P(F)} = \frac{0.2 \times 0.6}{0.36} = \frac{0.12}{0.36} = 0.3333$$

3. Probability of Other Problems Given Failure:

$$P(O | F) = \frac{P(F | O) \cdot P(O)}{P(F)} = \frac{0.5 \times 0.3}{0.36} = \frac{0.15}{0.36} = 0.4167$$

Step 3: Interpret the Results

- Hardware Problem: 25% chance
- Software Problem: Approximately 33.33% chance
- Other Problem: Approximately 41.67% chance

Conclusion: Given a printer failure, the most likely cause is an Other Problem, such as connectors, with a probability of approximately 41.67%.



Chapter 6 Descriptive Statistics

In software engineering, we constantly deal with data: performance metrics, user behavior patterns, system response times, and code complexity measures. To make sense of this information and extract meaningful insights, we need tools to summarize, organize, and visualize data. Descriptive statistics provides these essential tools, allowing us to understand the characteristics of our data without making inferences about larger populations.

Consider a software development team analyzing the performance of their web application. They collect response times for 1000 user requests and find values ranging from 50 milliseconds to 3.2 seconds. Without proper analysis, this raw data is overwhelming and uninformative. However, by applying descriptive statistics, they can determine that the average response time is 450 milliseconds, that 75% of requests complete within 600 milliseconds, and that there are a few unusually slow requests that might indicate performance issues.

This chapter introduces the fundamental concepts of descriptive statistics, focusing on measures that help us understand the central tendencies, variability, and distribution patterns in our data. We will explore how these techniques apply specifically to software engineering contexts, from analyzing algorithm performance to understanding user behavior patterns.

6.1 Introduction to Descriptive Statistics

Sometimes statistical work begins with existing data, such as precipitation records, unemployment rates, or GDP figures, which we then summarize and analyze. In other situations, data must be generated through an experiment or study. For example, to compare two teaching methods in an introductory programming course, an instructor might randomly divide students into two groups, apply a different method to each, and then compare test scores. Random assignment is essential: it ensures that differences between groups are not due to preexisting factors such as aptitude, but rather to the teaching method itself.

After the data are collected, they are summarized and visualized—for example, by reporting the average score for each group. This process is the essence of **descriptive statistics**: organising, presenting, and describing data in a meaningful way.

Population versus Sample

In statistical analysis, we distinguish between two key concepts:

Definition 6.1 (Population)

A **population** is the complete set of all possible observations or measurements of interest in a particular study. In software engineering contexts, this might include all possible execution times of an algorithm, all user sessions on a website, or all lines of code in a project.



Definition 6.2 (Sample)

A **sample** is a subset of the population that we actually observe or measure. Due to practical constraints, we often work with samples rather than entire populations.



In the teaching-method example, the students in the classroom form only a *sample* of the larger *population* of potential learners. Descriptive statistics helps us understand the sample itself, while inferential methods later allow us to use that sample to draw conclusions about the broader population. Inferential statistics lies beyond the scope of this book. We will focus on descriptive statistics in this chapter.

Example 6.1 Software Performance Analysis

Consider a web application serving millions of users daily. The **population** would be the response times for all possible user requests. However, due to computational and storage limitations, we might only collect response times for a **sample** of 10,000 requests per day. This sample should be representative of the population to draw meaningful conclusions.

Types of Data

Data can be classified into different types, each requiring different statistical approaches:

Definition 6.3 (Qualitative Data)

Qualitative data (also called categorical data) consists of non-numerical information that can be categorized. Examples include programming languages used in a project, user satisfaction ratings (satisfied/neutral/dissatisfied), or bug severity levels (critical/high/medium/low).



The essential difference between **qualitative** and **quantitative** data lies in what they represent and how they are analyzed. Qualitative data refers to categories or attributes that cannot be expressed meaningfully as numbers. Such data are used to classify or label elements, for example, programming language or bug severity. Quantitative data, on the other hand, consists of numerical values that represent counts or measurements, such as response times or number of users. Because these values are numerical, they permit mathematical operations and statistical analysis. Recognizing the distinction between qualitative and quantitative data is fundamental, as it guides the choice of statistical methods and the types of visualizations that are appropriate for analysis.

Definition 6.4 (Quantitative Data)

Quantitative data consists of numerical measurements that can be ordered and subjected to mathematical operations. This can be further divided into:

- **Discrete:** Countable values (number of bugs, lines of code, user sessions)
- **Continuous:** Measurable values that can take any value within a range (response times, memory usage, CPU utilization)



Example 6.2

- **Qualitative:** Programming language (Python, Java, C++), deployment environment (development, staging, production)
- **Quantitative Discrete:** Number of commits per day, lines of code, number of test cases
- **Quantitative Continuous:** Response time in milliseconds, memory usage in MB, CPU utilization percentage

Outliers and Extreme Data

In any dataset, most values tend to cluster around a central region, but occasionally, we encounter values that are much higher or lower than the rest. These are called **outliers** or **extreme values**. Outliers can arise for many reasons: measurement errors, unusual but valid events, or natural variability in the data.

For example, if most web requests complete in under 400 milliseconds, but one request takes 1200 milliseconds, that 1200ms value is an outlier. Outliers can have a strong influence on summary statistics like the mean, making the data appear more variable or shifting the average away from where most values lie.

It's important to look for outliers when analyzing data, as they can signal interesting phenomena (such as a rare performance bottleneck), data entry mistakes, or the need for further investigation. While there are formal methods to detect outliers, which we will discuss later, often a simple plot or a scan of the sorted data is enough to spot values that "stand out" from the rest.

In practice, understanding the context is key: sometimes outliers are errors to be corrected or removed, but other times they are the most important part of the story.

6.2 Describing Data Sets

When presenting numerical results, clarity and brevity are essential. Tables and graphs are particularly effective, as they allow the reader to quickly grasp the main characteristics of a dataset. Visual and tabular summaries often highlight important aspects such as the overall range, how tightly the values are concentrated, and whether the data appear symmetric or skewed.

Frequency Tables and Graphs

The first step in data analysis is often to organize raw data into a frequency distribution, which shows how often each value (or range of values) occurs.

Definition 6.5 (Frequency Distribution)

A **frequency distribution** is a table that shows the frequency (count) of each value or class of values in a dataset. It can be presented as:

- **Absolute frequency:** The actual count of occurrences
- **Relative frequency:** The proportion of total observations
- **Cumulative frequency:** The running total of frequencies



Example 6.3 A dataset with a relatively small number of distinct values can be conveniently presented in a frequency table. For instance, consider a dataset consisting of the starting monthly salaries (to the nearest thousand Danish Kroner) of 42 recently graduated students with B.S. degrees in software engineering.

The raw data is as follows:

37, 37, 37, 37, 38, 39, 39, 39, 40, 40, 40, 40, 40, 41, 41, 41, 41, 41, 41, 41, 41, 41, 42, 42, 42, 42, 42, 42, 42, 42, 42, 44, 44, 44, 44, 44, 46, 46, 47, 47, 47, 50

From this data, we can construct a frequency table. **Table 6.1** tells us, among other things, that the lowest starting salary of 37,000 DKK was received by four of the graduates, whereas the highest salary of 50,000 DKK

was received by a single student. The most common starting salary was 42,000 DKK, received by 10 of the students.

Salary ($\times 1000$ DKK)	Frequency
37	4
38	1
39	3
40	5
41	8
42	10
44	5
46	2
47	3
50	1

Table 6.1: Frequency Distribution of Monthly Starting Salaries

Data from a frequency table can be graphically represented by a line graph that plots the distinct data values on the horizontal axis and indicates their frequencies by the heights of vertical lines. A line graph of the data presented in [Table 6.1](#) is shown in [Figure 6.1](#).

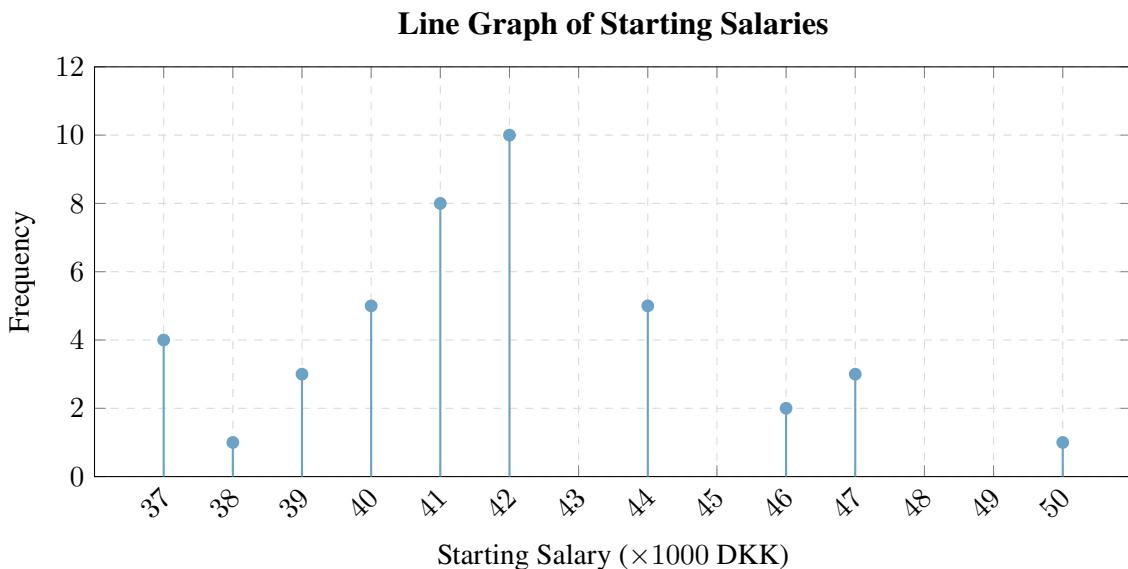
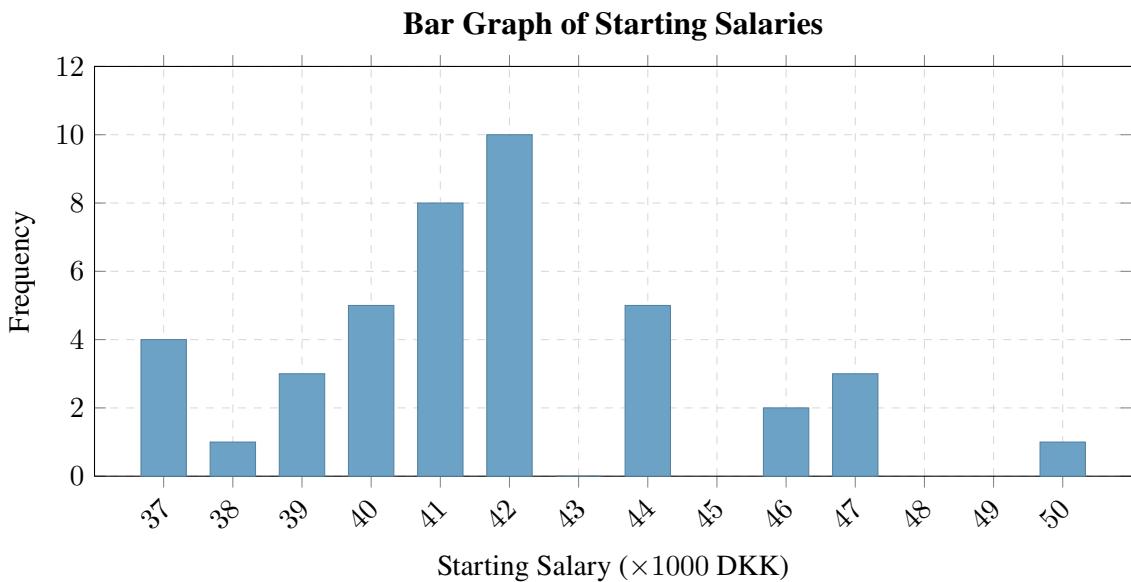
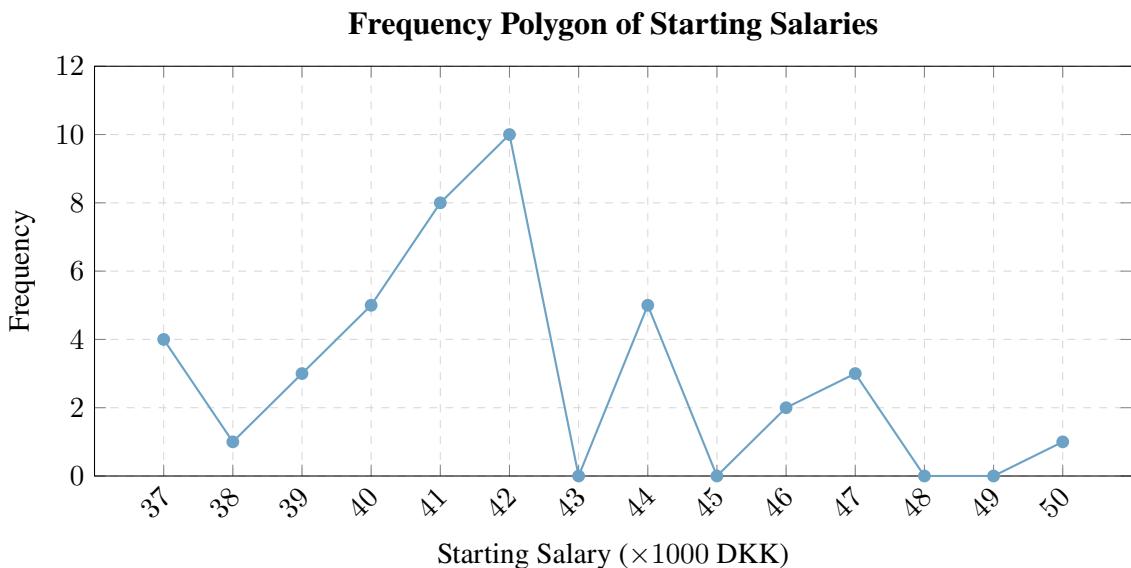


Figure 6.1: A line graph showing the frequency of different starting salaries.

When the lines in a line graph are given added thickness, the graph is called a bar graph. [Figure 6.2](#) shows a bar graph for the same salary data.

**Figure 6.2:** A bar graph of the salary data.

Another type of graph used to represent a frequency table is the frequency polygon, which plots the frequencies of the different data values on the vertical axis, and then connects the plotted points with straight lines. **Figure 6.3** shows a frequency polygon for the same salary data.

**Figure 6.3:** A frequency polygon of the salary data, connecting the points from the frequency table.

Sometimes we prefer the relative frequency over the absolute frequency. This is particularly useful when dealing with large datasets and when you want to compare values from datasets of different sizes. An extension of the relative frequency is the cumulative relative frequency, which is the running total of the relative frequencies. **Table 6.2** shows the frequency, relative frequency, and cumulative relative frequency for the salary data.

Salary	Frequency	Relative Freq.	Cumulative Rel. Freq.
37	4	$4/42 \approx 0.095$	$4/42 \approx 0.095$
38	1	$1/42 \approx 0.024$	$5/42 \approx 0.119$
39	3	$3/42 \approx 0.071$	$8/42 \approx 0.190$
40	5	$5/42 \approx 0.119$	$13/42 \approx 0.310$
41	8	$8/42 \approx 0.190$	$21/42 = 0.500$
42	10	$10/42 \approx 0.238$	$31/42 \approx 0.738$
44	5	$5/42 \approx 0.119$	$36/42 \approx 0.857$
46	2	$2/42 \approx 0.048$	$38/42 \approx 0.905$
47	3	$3/42 \approx 0.071$	$41/42 \approx 0.976$
50	1	$1/42 \approx 0.024$	$42/42 = 1.000$

Table 6.2: Frequency, Relative Frequency, and Cumulative Relative Frequency of Monthly Starting Salaries.

Grouped Data and Histograms

As we saw in [section 6.2](#), line and bar graphs are useful for showing frequencies of data values. When a dataset contains many distinct values, such as the lamp lifetimes in [Table 6.3](#), creating a frequency table for each individual value becomes impractical. A more effective approach is to group the data into a set of **class intervals**, which are also commonly referred to as **bins**.

Item Lifetimes (in Hours)										
1067	919	1196	785	1126	936	918	1156	920	948	
855	1092	1162	1170	929	950	905	972	1035	1045	
1157	1195	1195	1340	1122	938	970	1237	956	1102	
1022	978	832	1009	1157	1151	1009	765	958	902	
923	1333	811	1217	1085	896	958	1311	1037	702	
521	933	928	1153	946	858	1071	1069	830	1063	
930	807	954	1063	1002	909	1077	1021	1062	1157	
999	932	1035	944	1049	940	1122	1115	833	1320	
901	1324	818	1250	1203	1078	890	1303	1011	1102	
996	780	900	1106	704	621	854	1178	1138	951	
1187	1067	1118	1037	958	760	1101	949	992	966	
824	653	980	935	878	934	910	1058	730	980	
844	814	1103	1000	788	1143	935	1069	1170	1067	
1037	1151	863	990	1035	1112	931	970	932	904	
1026	1147	883	867	990	1258	1192	922	1150	1091	
1039	1083	1040	1289	699	1083	880	1029	658	912	
1023	984	856	924	801	1122	1292	1116	880	1173	
1134	932	938	1078	1180	1106	1184	954	824	529	
998	996	1133	765	775	1105	1081	1171	705	1425	
610	916	1001	895	709	860	1110	1149	972	1002	

Table 6.3: Life in Hours of 200 Incandescent Lamps.

The choice of how many bins to use involves a trade-off. If there are too few bins, we risk losing important

information by grouping too many distinct values together. If there are too many bins, the frequencies within each bin may become too small to reveal any clear pattern in the data's distribution.

While the optimal number of bins is a subjective choice that depends on the dataset, a selection of 5 to 10 bins is typical. It is also customary, though not required, to use bins of equal width, as this often makes interpretation easier.

The endpoints of a bin are called its **class boundaries**. We will adopt the common **left-end inclusion convention**, which specifies that a bin contains its left-end boundary but not its right-end one. For example, the bin 700–800 would contain any value x such that $700 \leq x < 800$.

Following this approach, the raw data from [Table 6.3](#) is summarized in [Table 6.4](#) using bins of length 100, starting at 500.

Class Interval	Frequency	Relative Freq.	Cumulative Rel. Freq.
500–600	2	0.010	0.010
600–700	5	0.025	0.035
700–800	12	0.060	0.095
800–900	25	0.125	0.220
900–1000	58	0.290	0.510
1000–1100	41	0.205	0.715
1100–1200	43	0.215	0.930
1200–1300	7	0.035	0.965
1300–1400	6	0.030	0.995
1400–1500	1	0.005	1.000

Table 6.4: Frequency Distribution for the Lifetimes of 200 Incandescent Lamps.

A bar graph that plots the frequencies of data grouped into bins is called a **histogram**. A key feature of a histogram is that the bars are placed directly adjacent to one another, reflecting the continuous nature of the bins. The vertical axis can represent either the absolute frequency (the count of values in each interval) or the relative frequency. The former is called a **frequency histogram**, while the latter is a **relative frequency histogram**. [Figure 6.4](#) shows a frequency histogram for the lamp lifetime data.

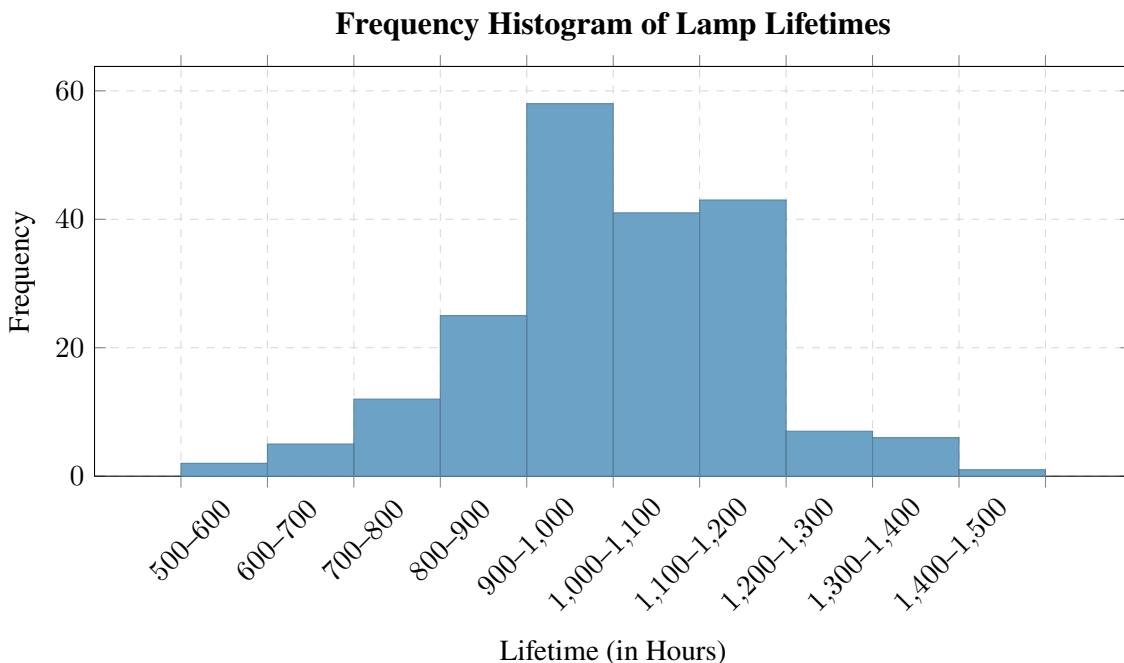


Figure 6.4: A frequency histogram for the lamp lifetime data. The adjacent bars show the distribution of data across continuous intervals.

Finally, we are sometimes interested in plotting the cumulative frequency or, more commonly, the cumulative relative frequency. Such a graph, often called an **ogive**, shows the number or proportion of data points that fall at or below a certain value.

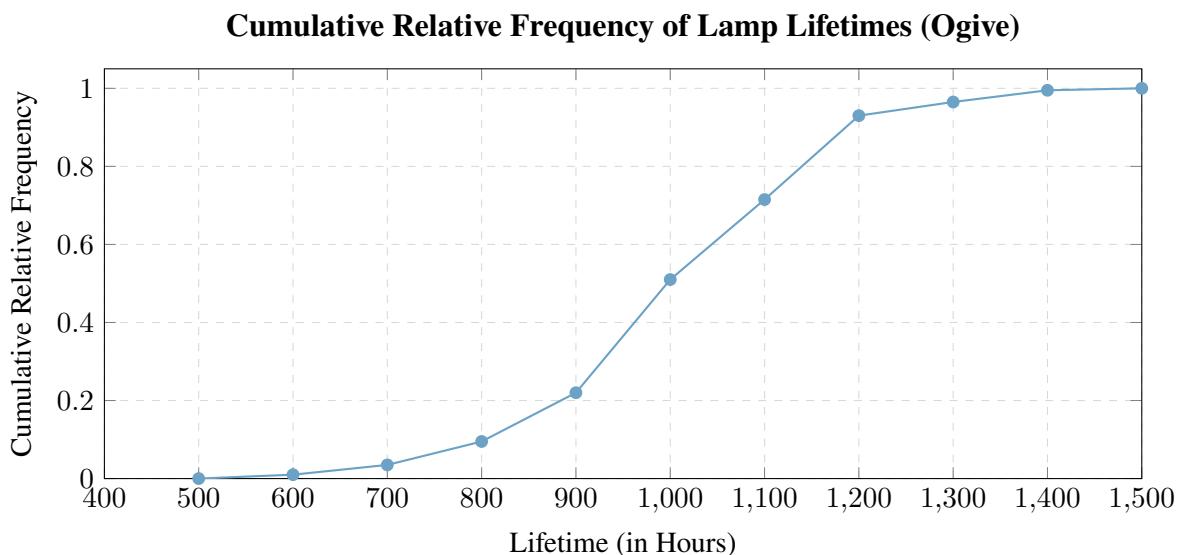


Figure 6.5: An ogive showing the cumulative relative frequency for the lamp lifetime data.

On an ogive, each point on the horizontal axis represents a data value, while the corresponding point on the vertical axis shows the proportion of the data that is less than or equal to that value. For example, Figure 6.5 presents the cumulative relative frequency plot for the lamp lifetime data. From this graph, we can quickly determine that approximately 50% of the lamps had a lifetime of less than 1000 hours, about 95% had a lifetime of less than 1200 hours (the actual value is 93%), and 100% of the lamps failed by 1500 hours.

6.3 Summarizing Data

Having seen how to visually and numerically summarize data distributions, we now focus on two fundamental aspects of describing a dataset: its *center* (or typical value) and its *dispersion* (or variability).

Measures of Central Tendency

The **mean**, often called the average, is the most commonly used measure of central tendency. **Definition 6.6** shows how to calculate the mean of a dataset and **Example 6.4** exemplifies .

Definition 6.6 (Arithmetic Mean)

For a dataset with n values x_1, x_2, \dots, x_n , the **arithmetic mean** is defined as:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i = \frac{x_1 + x_2 + \dots + x_n}{n}$$



Example 6.4 Response Time Analysis

A web application recorded the following response times (in milliseconds) for 10 requests:

$$\{120, 150, 180, 200, 220, 250, 280, 300, 350, 400\}$$

The mean response time is:

$$\bar{x} = \frac{120 + 150 + 180 + 200 + 220 + 250 + 280 + 300 + 350 + 400}{10} = \frac{2450}{10} = 245 \text{ milliseconds}$$

The mean incorporates every value in a dataset, which makes it a powerful measure of central tendency. It is most appropriate when the data values are of similar magnitude and when the distribution is reasonably symmetric without extreme outliers. Under these conditions, the mean provides a reliable summary of the dataset.

For example, in the analysis of web performance, the mean response time offers an overall indication of how quickly a system responds on average. However, if a few response times are much slower than the rest, the mean may give a misleading impression of the typical user experience.

Beyond its role as a summary statistic, the mean also serves as a foundation for many statistical methods. In particular, it is central to the definitions of variance and standard deviation, which measure the dispersion of data around the mean.

Remark: As mentioned earlier, and as we will revisit, the mean is highly sensitive to extreme values or outliers. Even a single unusually large or small observation can shift the mean considerably, reducing its ability to reflect the typical value of the dataset. It is therefore important to consider the distribution of the data before relying solely on the mean to describe the center of a dataset.

The **median** is the middle value when data is arranged in ascending order. It is less sensitive to outliers than the mean.

Definition 6.7 (Median)

For a dataset with n values arranged in ascending order, the median is defined as:

$$\text{median} = \begin{cases} x_{\frac{n+1}{2}} & \text{if } n \text{ is odd} \\ \frac{x_{\frac{n}{2}} + x_{\frac{n}{2}+1}}{2} & \text{if } n \text{ is even} \end{cases}$$

**Example 6.5** Median Response Time

Using the same response time data from the previous example:

$$\{120, 150, 180, 200, 220, 250, 280, 300, 350, 400\}$$

Since $n = 10$ (even), the median is the average of the 5th and 6th values:

$$\text{Median} = \frac{220 + 250}{2} = 235 \text{ milliseconds}$$

Example 6.6

Consider the same data with an extreme value:

$$\{120, 150, 180, 200, 220, 250, 280, 300, 350, 2000\}$$

The mean becomes: $\bar{x} = \frac{4050}{10} = 405$ milliseconds (significantly affected by the outlier)

The median remains: $\frac{220+250}{2} = 235$ milliseconds (unaffected by the outlier)

Because the median divides a dataset into two equal halves, it serves as a robust measure of central tendency, particularly when the data contain outliers or are skewed. Unlike the mean, which can be drawn toward extreme values, the median remains largely unaffected, offering a more reliable indication of the “typical” value in such situations.

The **mode** is the most frequently occurring value in a dataset.

Definition 6.8 (Mode)

The **mode** is the value that appears most frequently in a dataset. A dataset can have:

- **No mode:** If all values appear with equal frequency
- **One mode (unimodal):** If one value appears most frequently
- **Multiple modes (multimodal):** If two or more values tie for the highest frequency



Sometimes, when a data set has two modes, it is called a **bimodal** dataset.

Example 6.7

A software project has the following bug severity levels:

{High, Medium, Low, High, Critical, Medium, High, Low, High, Medium}

Counting frequencies:

- High: 4 occurrences
- Medium: 3 occurrences
- Low: 2 occurrences
- Critical: 1 occurrence

The mode is "High" since it appears most frequently.

The mode is especially useful when analyzing categorical or discrete data, where calculating a mean or median may not make sense. For example, in software engineering, the mode can help identify the most common error code returned by an API, the most frequently used programming language in a codebase, or the most reported type of bug in an issue tracker.

When to Use Each Measure

Measure	Best For	Limitations
Mean	Continuous data, symmetric distributions	Sensitive to outliers
Median	Data with outliers, skewed distributions	Less informative for symmetric data
Mode	Categorical data, discrete data	May not exist or be unique

Table 6.5: Comparison of Central Tendency Measures

Measures of Dispersion

While measures of central tendency tell us about the typical value, **measures of dispersion** describe how spread out the data are. Understanding variability is crucial in software engineering for assessing consistency and reliability.

The **range** is the simplest measure of dispersion.

Definition 6.9 (Range)

The **range** of a dataset is the difference between the maximum and minimum values:

$$\text{Range} = x_{\max} - x_{\min}$$



Example 6.8

For the response time data: {120, 150, 180, 200, 220, 250, 280, 300, 350, 400}

$$\text{Range} = 400 - 120 = 280 \text{ milliseconds}$$

Remark: The range is sensitive to outliers and doesn't provide information about the distribution of values between the extremes.

The **variance** and **standard deviation** are more sophisticated measures that consider all data points.

Definition 6.10 (Sample Variance and Standard Deviation)

For a sample with n values, the **sample variance**, denoted by s^2 , is defined as:

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

The **sample standard deviation**, denoted by s , is the square root of the variance:

$$s = \sqrt{s^2} = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}$$



While the mean identifies the center of a dataset, the **variance** and **standard deviation** measure its variability or spread. A small standard deviation indicates that data points are clustered tightly around the mean, whereas a large standard deviation signifies that they are more spread out.

The variance is the average of the squared deviations from the mean. Squaring each deviation serves two purposes: it ensures all values are positive, and it gives greater weight to larger, more significant deviations. However, this calculation leaves the variance in squared units (e.g., milliseconds squared), which are not intuitive to interpret.

To solve this, we use the standard deviation, which is simply the square root of the variance. This crucial step returns the measure to the original units of the data (e.g., milliseconds), making it directly interpretable. In short, the standard deviation represents the typical distance of a data point from the mean.

Example 6.9 Calculating Standard Deviation

For the response time data: $\{120, 150, 180, 200, 220, 250, 280, 300, 350, 400\}$ with $\bar{x} = 245$:

$$\begin{aligned} s^2 &= \frac{1}{9}[(120 - 245)^2 + (150 - 245)^2 + \cdots + (400 - 245)^2] \\ &= \frac{1}{9}[(-125)^2 + (-95)^2 + \cdots + (155)^2] \\ &= \frac{1}{9}[15625 + 9025 + \cdots + 24025] \\ &= \frac{1}{9} \times 82250 \approx 9138.89 \end{aligned}$$

$$s = \sqrt{9138.89} \approx 95.6 \text{ milliseconds}$$

The following algebraic identity is often helpful when computing variance or standard deviation by hand:

We start from the definition:

$$s^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

Using the identity

$$\sum_{i=1}^n (x_i - \bar{x})^2 = \sum_{i=1}^n x_i^2 - \frac{(\sum_{i=1}^n x_i)^2}{n}$$

we can rewrite sample variance as

$$s^2 = \frac{1}{n-1} \left(\sum_{i=1}^n x_i^2 - \frac{(\sum_{i=1}^n x_i)^2}{n} \right) = \frac{1}{n-1} \left(\sum_{i=1}^n x_i^2 - n\bar{x}^2 \right)$$

This is known as the **computational formula for variance**, as it allows you to calculate s^2 directly from $\sum x_i$ and $\sum x_i^2$ (or equivalently, from $\sum x_i^2$ and \bar{x}), without needing to compute each individual deviation from the mean.

The **coefficient of variation** allows comparison of variability across datasets with different scales.

Definition 6.11 (Coefficient of Variation)

The **coefficient of variation** is the ratio of standard deviation to the mean:

$$CV = \frac{s}{\bar{x}} \times 100\%$$



Example 6.10 Comparing Variability

Consider two systems:

- System A: Mean response time = 100ms, Standard deviation = 20ms
- System B: Mean response time = 500ms, Standard deviation = 100ms

$$CV_A = \frac{20}{100} \times 100\% = 20\%$$

$$CV_B = \frac{100}{500} \times 100\% = 20\%$$

Both systems have the same relative variability (20%), despite different absolute standard deviations.

Measures of Position

To introduce a measure of spread that is less sensitive to outliers, we first need to understand measures of position, which describe the relative standing of a data point within a dataset.

Definition 6.12 (Percentile)

The p -th **percentile** is the value below which $p\%$ of the data falls.



For example, if a response time is at the 95th percentile, it means that 95% of all response times are faster than this one.

Example 6.11 Percentile Calculation

For the response time data, to find the 90th percentile:

1. Sort data: {120, 150, 180, 200, 220, 250, 280, 300, 350, 400}
2. Position: $0.9 \times 10 = 9$ th position
3. 90th percentile = 350ms

Quartiles are specific, widely used percentiles that divide the data into four equal parts.

Definition 6.13 (Quartiles)

Quartiles divide an ordered dataset into four equal parts.

- **Q1 (First Quartile):** The 25th percentile. 25% of the data is less than this value.
- **Q2 (Second Quartile):** The 50th percentile. This is the **median** of the dataset.
- **Q3 (Third Quartile):** The 75th percentile. 75% of the data is less than this value.
- **Q4 (Fourth Quartile):** The 100th percentile. 100% of the data is less than this value.

**Example 6.12 Calculating Quartiles**

For the sorted response time data: {120, 150, 180, 200, 220, |250, 280, 300, 350, 400}

- The lower half is {120, 150, 180, 200, 220}. The median of this half is $Q_1 = 180$ ms.
- The upper half is {250, 280, 300, 350, 400}. The median of this half is $Q_3 = 300$ ms.

Percentiles divide data into 100 equal parts, while quartiles divide data into four equal parts.

The **interquartile range** is a robust measure of dispersion that is not affected by outliers.

Definition 6.14 (Interquartile Range)

The **interquartile range (IQR)** measures the spread of the middle 50% of the data and is defined as

$$IQR = Q_3 - Q_1,$$

where Q_1 and Q_3 are the first and third quartiles. A larger IQR indicates that the central portion of the data is more widely dispersed, while a smaller IQR indicates that it is more tightly clustered.

**Example 6.13 IQR Calculation**

For the response time data: {120, 150, 180, 200, 220, 250, 280, 300, 350, 400}

- Lower half: {120, 150, 180, 200, 220} $\rightarrow Q_1 = 180$
- Upper half: {250, 280, 300, 350, 400} $\rightarrow Q_3 = 300$
- $IQR = 300 - 180 = 120$ milliseconds

Summary Statistics

Before we conclude this chapter, we will review the summary statistics that we have covered in this section. Let us consider the response time data again, but now in the form of a table of 1000 values. **Table 6.6** shows the complete statistical summary of the data.

Example 6.14

For a dataset of 1000 web response times:

This table provides a complete picture of the data and is often a first step in the analysis of a dataset. Note that the skewness and kurtosis have not been discussed but will be explained in the next section. We have chosen to include them here for completeness as they are often used in the analysis of a dataset - although how to calculate them is beyond the scope of this book.

Measure	Value
Count	1000
Mean	245.3 ms
Median	238.0 ms
Mode	220.0 ms
Standard Deviation	95.6 ms
Variance	9138.9 ms ²
Minimum	45.0 ms
Maximum	1200.0 ms
Range	1155.0 ms
Q1	180.0 ms
Q3	300.0 ms
IQR	120.0 ms
Skewness	1.2 (positive skew)
Kurtosis	2.8 (leptokurtic)

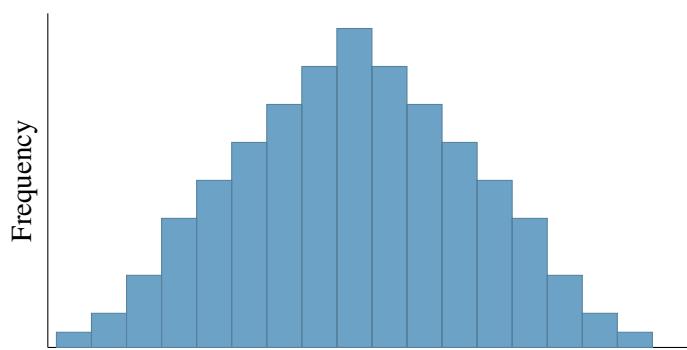
Table 6.6: Complete Statistical Summary

6.4 Understanding Data Distributions

Once we have visualized our data with histograms and box plots, we can begin to analyze the *shape* of the distribution. The shape tells us about the underlying patterns in our data and helps us choose the right statistical tools. In software engineering, understanding the distribution of metrics like response times or error rates is critical for setting performance baselines and detecting anomalies. First we review the concept of outliers.

Normal Data Sets

Many large data sets yield histograms with a distinctive overall shape: the frequencies peak near the sample median and decrease symmetrically on both sides, forming a characteristic **bell-shaped curve**. Data sets exhibiting this property are called **normal**, and their histograms are referred to as **normal histograms** (see [Figure 6.6](#)).

**Figure 6.6:** A histogram of a normal data set, showing the characteristic bell shape.

In software systems, perfectly normal distributions are rare, but some metrics can be **approximately normal**. For example, the latency of a highly optimized, internal microservice that performs a consistent task might be

approximately normal. A histogram of an approximately normal data set looks like [Figure 6.7](#).

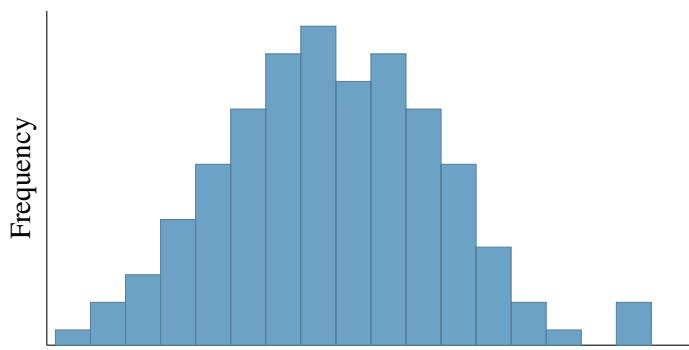


Figure 6.7: A histogram of a data set that is approximately normal. While not perfectly symmetric, it follows the general bell shape.

Many datasets observed in practice, from human height to measurement errors, follow this specific shape. Such data is said to be **normally distributed**.

Definition 6.15 (Normal Distribution)

A **normal distribution** is a symmetric, bell-shaped distribution where the mean, median, and mode are all equal and located at the center. Its shape is determined entirely by its mean (μ) and standard deviation (σ). 

When the bins of a histogram become smaller, the histogram becomes more and more similar to a normal distribution. This is because the normal distribution is a continuous distribution, while the histogram is a discrete distribution. This is a consequence of the **Law of Large Numbers**.

Theorem 6.1 (Law of Large Numbers)

The sample mean of a dataset will converge to the population mean as the sample size increases. 

When the sample size is large enough, the sample mean will be very close to the population mean, and the histogram will be very similar to a normal distribution and will get the following shape:

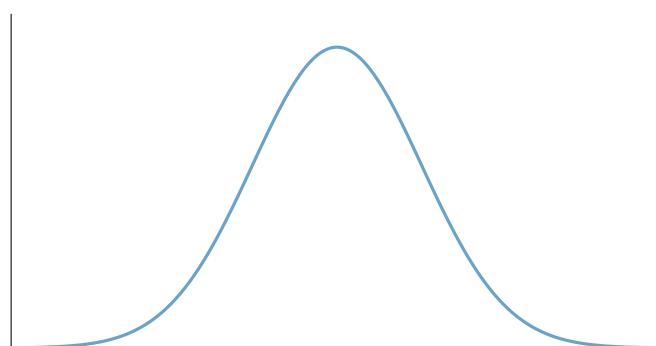


Figure 6.8: The normal distribution curve.

For data that is approximately normal, we can use a powerful rule of thumb to understand its spread without looking at every single data point. This is known as the **Empirical Rule**.

Theorem 6.2 (The Empirical Rule)

If a dataset is approximately normal with a sample mean \bar{x} and sample standard deviation s , then:

- Approximately **68.3%** of the observations lie within 1 standard deviation of the mean ($\bar{x} \pm s$).
- Approximately **95.4%** of the observations lie within 2 standard deviations of the mean ($\bar{x} \pm 2s$).
- Approximately **99.7%** of the observations lie within 3 standard deviations of the mean ($\bar{x} \pm 3s$). 

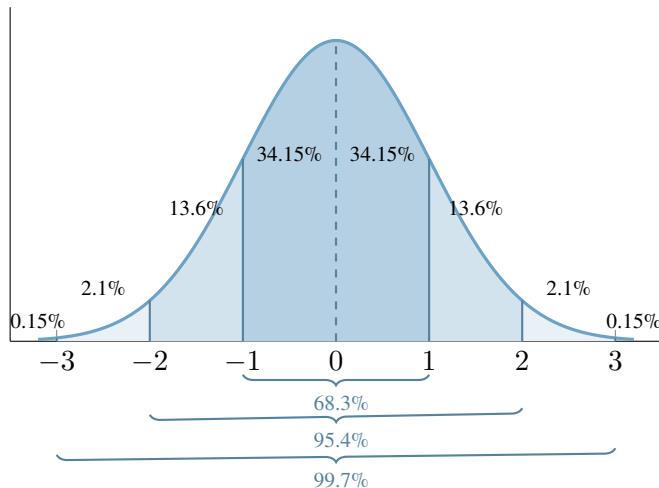


Figure 6.9: Empirical Rule on a normal curve. The inner bands mark $\pm 1\sigma$ (68.3%), $\pm 2\sigma$ (95.4%), and $\pm 3\sigma$ (99.7%).

Skewed Data Sets

When a data set is not symmetric about its sample median, it is described as **skewed**. A long tail extending to the right indicates the distribution is **skewed to the right**; a long tail to the left means it is **skewed to the left**. Thus, the histogram in Figure 6.10 is skewed left, while the one in Figure 6.11 is skewed right.

Definition 6.16 (Skewness)

Skewness measures the asymmetry of data distribution:

- **Positive skew (skewed to the right):** Tail extends to the right (mean > median)
- **Negative skew (skewed to the left):** Tail extends to the left (mean < median)
- **Symmetric:** Mean \approx median

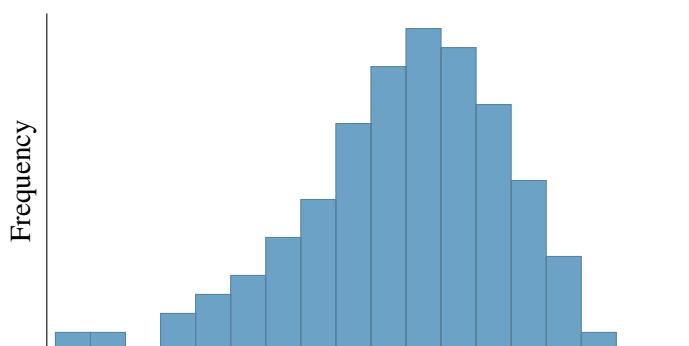


Figure 6.10: A histogram of a data set that is skewed to the left (negatively skewed). The tail of the distribution is on the left.

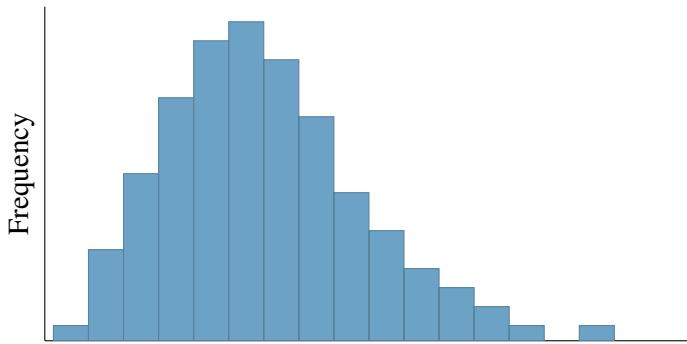


Figure 6.11: A histogram of a data set that is skewed to the right (positively skewed). The tail of the distribution is on the right.

Figure 6.12 shows the relationship between mean, median, and mode for right skewed, symmetric, and left skewed distributions.

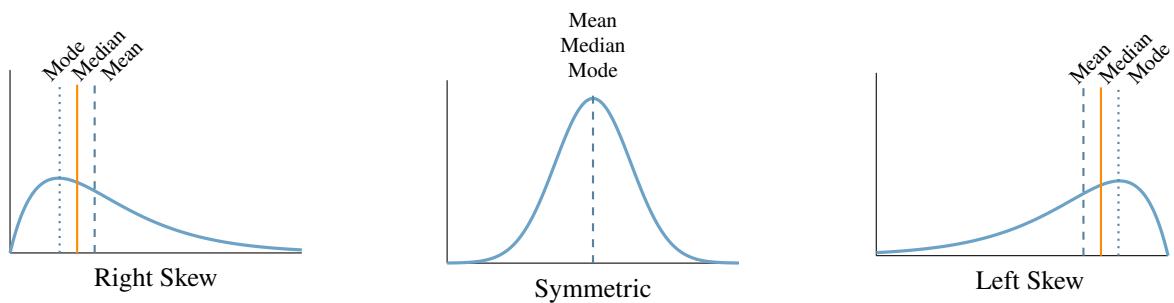


Figure 6.12: Relationship between mean, median, and mode for positively skewed, symmetric, and negatively skewed distributions.

Outlier Detection

As mentioned in section 6.1, outliers are data points that are significantly different from the rest of the data. They can arise due to measurement errors, data entry errors, or genuine variability in the data. Detecting outliers is important because they can distort statistical analyses.

A common method for identifying outliers is to use the interquartile range (IQR):

- **Mild outliers:** Any data point less than $Q_1 - 1.5 \times IQR$ or greater than $Q_3 + 1.5 \times IQR$.
- **Extreme outliers:** Any data point less than $Q_1 - 3 \times IQR$ or greater than $Q_3 + 3 \times IQR$.

Example 6.15 Detecting Outliers

Given the response time data: $\{120, 150, 180, 200, 220, 250, 280, 300, 350, 400\}$

- $Q_1 = 180, Q_3 = 300, IQR = 120$
- Lower fence: $Q_1 - 1.5 \times IQR = 180 - 180 = 0$
- Upper fence: $Q_3 + 1.5 \times IQR = 300 + 180 = 480$

All data points are between 0 and 480, so there are no mild outliers in this dataset.

Example 6.16

Suppose we have the following set of response times (in milliseconds) for a web application:

$$\{120, 150, 180, 200, 220, 250, 280, 300, 350, 400, 1200\}.$$

Let's identify if there is an upper outlier:

- Sorted data: 120, 150, 180, 200, 220, 250, 280, 300, 350, 400, 1200
 - $n = 11$
 - Median (Q_2): 250 (6th value)
 - Lower half: 120, 150, 180, 200, 220 $\rightarrow Q_1 = 180$
 - Upper half: 280, 300, 350, 400, 1200 $\rightarrow Q_3 = 350$
 - $IQR = Q_3 - Q_1 = 350 - 180 = 170$
 - Upper fence: $Q_3 + 1.5 \times IQR = 350 + 1.5 \times 170 = 350 + 255 = 605$

The value 1200 is greater than 605, so it is an upper outlier.

Box Plots

Box plots provide a compact summary of data distribution, highlighting the median, quartiles, and outliers. Figure 6.13 shows an example of a box plot.

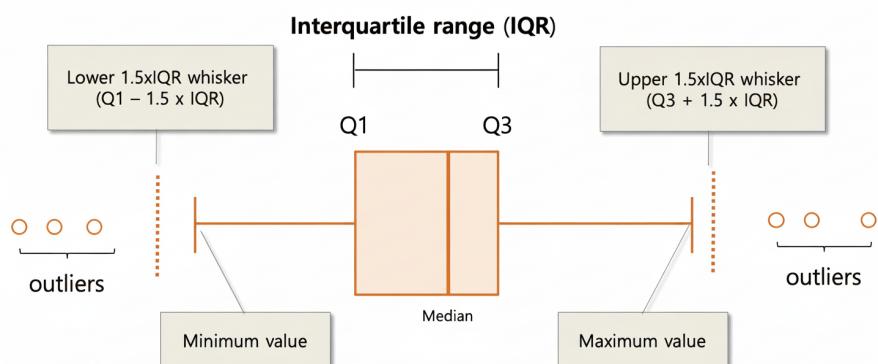


Figure 6.13: An example of a box plot, illustrating the median, quartiles, whiskers, and potential outliers.

Definition 6.17 (Box Plot Components)

A box plot displays five key statistics:

- **Minimum:** The smallest value (excluding outliers)
 - **Q1 (First Quartile):** 25% of data below this value
 - **Median (Q2):** 50% of data below this value
 - **Q3 (Third Quartile):** 75% of data below this value
 - **Maximum:** The largest value (excluding outliers)
 - **Outliers:** Points beyond $1.5 \times \text{IQR}$ from the box

A box plot provides a visual summary of the distribution of a dataset, making it easy to compare groups and spot outliers. Here's how to read a box plot.

- The **length of the box** shows the spread of the central half of the data. A longer box means more variability in the middle 50%.
- The **position of the median line** within the box indicates skewness. If the median is closer to Q_1 (the left side), the data are *skewed right* (longer tail on the right). If the median is closer to Q_3 (the right side), the data are *skewed left* (longer tail on the left).
- The **mean** typically lies in the direction of the skew. Therefore, if the median line is closer to Q_1 , the mean is to the *right* of the median (right-skewed). If the median line is closer to Q_3 , the mean is to the *left* of the median (left-skewed).
- The **mode**, median, and mean tend to align in this order: For left-skewed data: mean < median < mode. For right-skewed data: mode < median < mean.
- The **whiskers** show the range of the bulk of the data, excluding outliers.
- **Outliers** highlight unusually high or low values that may warrant further investigation.

Box plots are especially useful for comparing distributions across several groups or datasets, as they quickly reveal differences in medians, spreads, and the presence of outliers.

This concludes our discussion of data visualization and summary statistics, and in general our discussion of probability and statistics.

Chapter 7 Linear Equations in Linear Algebra

In 1949, Harvard Professor Wassily Leontief used one of the earliest computers, the Mark II, to solve a system of linear equations that modelled the U.S. economy. He divided the economy into 500 sectors, like coal, automotive, and communications, and described how each sector interacted using linear equations. Due to hardware limitations, he reduced the problem to 42 equations, which took the Mark II 56 hours to solve. This effort marked a milestone in computational applications, earning Leontief the 1973 Nobel Prize.

Leontief's work showcased how linear algebra, combined with computing, could solve large-scale problems — an idea that resonates even more today in the realm of software engineering. Linear algebra is fundamental to many modern technologies, particularly in Machine Learning (ML) and Artificial Intelligence (AI). From powering recommendation systems to optimising neural networks, linear algebra is at the core of algorithms that drive today's intelligent systems.

In ML and AI, large datasets are transformed into matrices, where linear algebra techniques like matrix factorisation, eigenvalue decomposition, and singular value decomposition (SVD) are used to extract insights and make predictions. Neural networks, the backbone of AI, rely on linear algebra to propagate data through layers and adjust weights during training. As a software engineer, mastering these concepts in linear algebra equips you to build scalable, intelligent systems capable of tackling today's most complex computational problems.

With the exponential growth in data and computing, the significance of linear algebra in software development, especially in AI and ML, continues to rise. It is a vital tool that bridges theoretical mathematics with real-world applications in tech.

7.1 Systems of Linear Equations

Linear equations and their systems form the foundation of linear algebra, a critical area in mathematics with extensive applications in software engineering. From computer graphics to machine learning algorithms, understanding how to model and solve linear systems is essential for developing efficient and effective software solutions. We begin with a definition:

Definition 7.1 (Linear Equations)

A linear equation in the variables x_1, x_2, \dots, x_n is an equation that can be written in the form

$$a_1x_1 + a_2x_2 + \dots + a_nx_n = b$$

where b and the coefficients a_1, \dots, a_n are constants.



Linear equations represent straight lines, planes, and hyperplanes in various dimensions, making them useful for modeling relationships in data and algorithms.

Example 7.1 A line in two-dimensional space given by $y = mx + b$ is a linear equation. It can be rewritten to fit our standard form:

$$-mx + y = b$$

Here, $a_1 = -m$, $a_2 = 1$, and b is the constant term.

Example 7.2 The general equation of a plane in three-dimensional space is

$$ax + by + cz = d$$

where a , b , c , and d are constants. This equation is linear in the variables x , y , and z .

Definition 7.2 (Systems of Linear Equations)

A **system of linear equations** (also called a **linear system**) is a collection of one or more linear equations involving the same set of variables.



Example 7.3 Consider the following system of linear equations in the variables x_1, x_2, x_3 :

$$2x_1 + 3x_2 + x_3 = 3$$

$$7x_2 - 4x_3 = 10$$

$$x_3 = 1$$

This system contains three equations with three unknowns.

To solve the system, we first note from the third equation that $x_3 = 1$. Substituting this into the second equation, we solve for x_2 :

$$7x_2 - 4(1) = 10$$

$$7x_2 = 14$$

$$x_2 = 2$$

Now, substituting $x_2 = 2$ and $x_3 = 1$ into the first equation, we solve for x_1 :

$$2x_1 + 3(2) + 1 = 3$$

$$2x_1 = -4$$

$$x_1 = -2$$

Thus, the solution to the system is $(x_1, x_2, x_3) = (-2, 2, 1)$.

Definition 7.3 (Solutions to a System of Linear Equations)

A **solution** of a linear system in the variables x_1, x_2, \dots, x_n is a list of numbers (s_1, s_2, \dots, s_n) that satisfies all equations of the system when substituted for the variables x_1, x_2, \dots, x_n , respectively. The set of all possible solutions is called its **solution set**. Two linear systems are called **equivalent** if they have the same solution set.



A system of linear equations has one of the following outcomes:

- (i) **No solutions**, when the equations are inconsistent.
- (ii) **Exactly one solution**, when there is a unique set of values satisfying all equations.
- (iii) **Infinitely many solutions**, when there are multiple sets of values that satisfy the equations.

We say a system is **consistent** if it has either one or infinitely many solutions. We say a system is **inconsistent** if it has no solution. We formalise this later in this chapter.

Remark: Determining whether a system is consistent addresses the *existence* of solutions. If solutions exist, we may further explore the *uniqueness* of these solutions.

Representing Systems with Matrices

Matrices provide a compact and efficient way to represent and manipulate systems of linear equations, which is particularly beneficial in software applications involving large datasets.

Definition 7.4

A **matrix** is a rectangular array of numbers arranged in rows and columns. The **coefficient matrix** of a linear system contains only the coefficients of the variables, while the **augmented matrix** includes an additional column for the constants from the right-hand side of the equations.



Example 7.4 For the linear system:

$$2x_1 + 3x_2 + x_3 = 3$$

$$7x_2 - 4x_3 = 10$$

$$x_3 = 1$$

the coefficient matrix is:

$$\begin{bmatrix} 2 & 3 & 1 \\ 0 & 7 & -4 \\ 0 & 0 & 1 \end{bmatrix}$$

and the augmented matrix is:

$$\left[\begin{array}{ccc|c} 2 & 3 & 1 & 3 \\ 0 & 7 & -4 & 10 \\ 0 & 0 & 1 & 1 \end{array} \right]$$

The vertical line between the coefficient part and the augmented part is optional and is used to visually separate the coefficients from the constants.

Definition 7.5

The **size** of a matrix is defined by the number of its rows and columns, expressed as *rows* \times *columns*. For instance, the coefficient matrix above is of size 3×3 , and the augmented matrix is of size 3×4 .



Solving Systems using Augmented Matrices

We now illustrate how to solve a system of linear equations using augmented matrix notation. This process will be formalised in the next section but here we offer an example. This method streamlines the process of solving systems by focusing on matrix manipulations.

Example 7.5 Let r_1 denote the first row, r_2 the second row, and r_3 the third row.

$$\left[\begin{array}{ccc|c} 2 & 3 & 1 & 3 \\ 0 & 7 & -4 & 10 \\ 0 & 0 & 1 & 1 \end{array} \right]$$

To simplify the system, we perform *elementary row operations*. Specifically, we adjust r_1 and r_2 as follows:

$$r_1 \mapsto r_1 - r_3,$$

$$r_2 \mapsto r_2 + 4r_3,$$

which gives:

$$\left[\begin{array}{cccc} 2 & 3 & 0 & 2 \\ 0 & 7 & 0 & 14 \\ 0 & 0 & 1 & 1 \end{array} \right]$$

Next, we scale r_2 by $\frac{1}{7}$:

$$r_2 \mapsto \frac{1}{7}r_2, \quad \left[\begin{array}{cccc} 2 & 3 & 0 & 2 \\ 0 & 1 & 0 & 2 \\ 0 & 0 & 1 & 1 \end{array} \right]$$

Now, we eliminate the 3 in the first row by performing:

$$r_1 \mapsto r_1 - 3r_2, \quad \left[\begin{array}{cccc} 2 & 0 & 0 & -4 \\ 0 & 1 & 0 & 2 \\ 0 & 0 & 1 & 1 \end{array} \right]$$

Finally, we scale r_1 by $\frac{1}{2}$ to get:

$$r_1 \mapsto \frac{1}{2}r_1, \quad \left[\begin{array}{cccc} 1 & 0 & 0 & -2 \\ 0 & 1 & 0 & 2 \\ 0 & 0 & 1 & 1 \end{array} \right]$$

We can now reinterpret the matrix as a linear system. From the augmented matrix, we find:

$$x_1 = -2,$$

$$x_2 = 2,$$

$$x_3 = 1.$$

This can be expressed compactly as:

$$(x_1, x_2, x_3) = (-2, 2, 1).$$

Thus, the solution is identical to the one we obtained through substitution.

Example 7.6 Determine if the following system is consistent:

$$x_2 + 4x_3 = 2$$

$$x_1 - 3x_2 + 2x_3 = 6$$

$$x_1 - 2x_2 + 6x_3 = 9$$

Solution: First, we write the augmented matrix of the system:

$$\left[\begin{array}{cccc} 0 & 1 & 4 & 2 \\ 1 & -3 & 2 & 6 \\ 1 & -2 & 6 & 9 \end{array} \right]$$

To simplify the matrix, we interchange r_1 and r_2 :

$$r_1 \leftrightarrow r_2, \quad \begin{bmatrix} 1 & -3 & 2 & 6 \\ 0 & 1 & 4 & 2 \\ 1 & -2 & 6 & 9 \end{bmatrix}$$

Next, we eliminate the 1 in the first column of r_3 by performing:

$$r_3 \mapsto r_3 - r_1, \quad \begin{bmatrix} 1 & -3 & 2 & 6 \\ 0 & 1 & 4 & 2 \\ 0 & 1 & 4 & 3 \end{bmatrix}$$

Then, we eliminate the 1 in the second column of r_3 :

$$r_3 \mapsto r_3 - r_2, \quad \begin{bmatrix} 1 & -3 & 2 & 6 \\ 0 & 1 & 4 & 2 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

The third row now corresponds to the equation:

$$0x_1 + 0x_2 + 0x_3 = 1$$

which simplifies to:

$$0 = 1$$

This is a contradiction, so the system is **inconsistent** and has no solution. ◀

Example 7.7 Give a solution of the following system (if one exists). Is it unique?

$$x_1 + 2x_2 + 3x_3 = 4$$

$$3x_1 + 6x_2 + 9x_3 = 12$$

Solution:

$$\begin{bmatrix} 1 & 2 & 3 & 4 \\ 3 & 6 & 9 & 12 \end{bmatrix}$$

We observe that r_2 is a multiple of r_1 :

$$r_2 = 3r_1$$

To simplify the system, we perform the following elementary row operation to eliminate redundancy:

$$r_2 \mapsto r_2 - 3r_1,$$

which yields:

$$\begin{bmatrix} 1 & 2 & 3 & 4 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

The second row now consists entirely of zeros, indicating that it does not provide any new information. This means the system is underdetermined and has infinitely many solutions.

We can express the first equation in terms of x_1 :

$$x_1 = 4 - 2x_2 - 3x_3$$

Letting x_2 and x_3 be free variables (parameters), we set:

$$x_2 = s, \quad x_3 = t, \quad \text{where } s, t \in \mathbb{R}$$

Substituting back, we find:

$$x_1 = 4 - 2s - 3t,$$

$$x_2 = s,$$

$$x_3 = t.$$

This can be expressed compactly as:

$$(x_1, x_2, x_3) = (4 - 2s - 3t, s, t), \quad \text{for all } s, t \in \mathbb{R}.$$

The system has infinitely many solutions parameterized by s and t . Therefore, the solution is not unique. ◀

Example 7.8 Choose h and k such that the following system

$$x_1 - 3x_2 = 1$$

$$2x_1 + hx_2 = k$$

has

- (i) a unique solution,
- (ii) many solutions, and
- (iii) no solution.

Solution:

$$\left[\begin{array}{ccc} 1 & -3 & 1 \\ 2 & h & k \end{array} \right]$$

To simplify the system, we perform an elementary row operation to eliminate x_1 from the second equation. Specifically, we adjust r_2 as follows:

$$r_2 \mapsto r_2 - 2r_1,$$

which yields:

$$\left[\begin{array}{ccc} 1 & -3 & 1 \\ 0 & h+6 & k-2 \end{array} \right]$$

Now, we analyze the resulting system based on the value of $h + 6$.

Case 1: $h + 6 \neq 0$

When $h + 6 \neq 0$, we can solve for x_2 from the second equation:

$$(h+6)x_2 = k-2 \implies x_2 = \frac{k-2}{h+6}$$

Substituting x_2 back into the first equation:

$$x_1 - 3x_2 = 1 \implies x_1 = 1 + 3x_2 = 1 + 3\left(\frac{k-2}{h+6}\right)$$

Thus, we obtain a unique solution for x_1 and x_2 . **The system has a unique solution when $h + 6 \neq 0$.**

Case 2: $h + 6 = 0$

When $h + 6 = 0$, i.e., $h = -6$, the second equation becomes:

$$0x_2 = k - 2$$

This simplifies to:

$$0 = k - 2$$

We have two subcases:

Subcase 2a: $k - 2 = 0$

If $k = 2$, the equation becomes $0 = 0$, which is always true. The second equation provides no new information, so the system reduces to:

$$x_1 - 3x_2 = 1$$

Here, x_2 is a free variable. Solving for x_1 :

$$x_1 = 1 + 3x_2$$

The system has infinitely many solutions when $h = -6$ and $k = 2$.

Subcase 2b: $k - 2 \neq 0$

If $k \neq 2$, the equation becomes $0 = k - 2$, which is a contradiction since $0 \neq k - 2$. **Thus, the system has no solution when $h = -6$ and $k \neq 2$.**

Conclusion:

- (i) **No solution** when $h = -6$ and $k \neq 2$.
- (ii) **A unique solution** when $h \neq -6$.
- (iii) **Infinitely many solutions** when $h = -6$ and $k = 2$.



7.2 Elementary Row Operations and Echelon Forms

In Example 7.5-Example 7.8, we performed a series of operations known as **elementary row operations**. These are fundamental transformations that simplify systems without changing their solution sets.

The three types of elementary row operations are:

- **Replacement:** Replace one row by the sum of itself and a multiple of another row.
- **Interchange:** Swap two rows.
- **Scaling:** Multiply all entries in a row by a nonzero constant.

Two matrices are called **row equivalent** if one can be transformed into the other through a sequence of elementary row operations.

Remark: Two linear systems are equivalent (i.e., they have the same solution set) if their augmented matrices are row equivalent. Performing elementary row operations on an augmented matrix does not change the solution

set of the system.

As you may have noticed in the examples, the matrices resulting from these operations often have a special pattern — lots of zeros below certain entries. This isn't just a coincidence but a goal in the process. We call this the **echelon form** of a matrix. By organizing a matrix into echelon form, we:

- Make solving systems of linear equations easier.
- Systematically simplify the matrix, reducing the complexity of calculations.
- Reveal key properties about the system, like existence and number of solutions or whether certain equations are dependent.

Definition 7.6 (Echelon Forms)

A matrix is in **echelon form** if it satisfies the following conditions:

1. All zero rows are at the bottom.
2. Each leading entry of a row is in a column to the right of the leading entry of the row above it.
3. All entries in a column below a leading entry are zeros.

In addition to echelon form, a matrix may also be in **reduced row echelon form** (RREF), which has the following properties:

4. The leading entry in each nonzero row is 1.
5. Each leading 1 is the only nonzero entry in its column.



We say that an echelon matrix U is an echelon form of the matrix A if U is row equivalent to A . Similarly, we say that a reduced echelon matrix U is the reduced echelon form of the matrix A if U is row equivalent to A .

The significance of putting the augmented matrix of a linear system in echelon form is explained by the following theorem.

Theorem 7.1 (Existence Theorem)

A linear system is consistent if and only if an echelon form of the augmented matrix has no row of the form

$$\left[\begin{array}{cccc} 0 & \dots & 0 & b \end{array} \right]$$

where b is nonzero.



We saw an example of an inconsistent system in [Example 7.6](#) because the echelon form of the augmented matrix has a row of the form $\left[\begin{array}{ccc|c} 0 & 0 & 0 & 1 \end{array} \right]$. This row indicates that the system has no solution.

Example 7.9 The augmented matrix of the linear system used as the main example in the preceding section,

$$\left[\begin{array}{cccc} 2 & 3 & 1 & 3 \\ 0 & 7 & -4 & 10 \\ 0 & 0 & 1 & 1 \end{array} \right],$$

is already in echelon form. Since it has no row of the form mentioned in the theorem, we know immediately that this system is consistent. The leading entries are 2, 7, and 1, and all entries below them are zeros. This matrix is not in reduced row echelon form however because the leading entries are not all 1.

Recall that we performed a sequence of row operations on the preceding matrix to get

$$\begin{bmatrix} 1 & 0 & 0 & -2 \\ 0 & 1 & 0 & 2 \\ 0 & 0 & 1 & 1 \end{bmatrix}$$

which is in reduced echelon form. This allowed us easily to see the solutions of this system, which is the main advantage of putting the matrix in this form. This brings us to the following important theorem.

Theorem 7.2 (Uniqueness of Reduced Echelon Form)

Each matrix is row equivalent to one and only one reduced echelon matrix



A pivot position in a matrix A is a location in A that corresponds to a leading 1 in the reduced echelon form of A . A pivot column is a column of A that contains a pivot position. Notice that you can obtain a pivot by scaling the leading entry of a row to be 1. Therefore, the pivot column is the column of the leading entry.

We are now ready to present the algorithm for solving a system of linear equations using matrices.

The Row Reduction Algorithm

Here we describe an algorithm for turning any matrix into an equivalent (reduced) echelon matrix. This algorithm is the foundation of solving systems of linear equations using matrices.

1. Begin with the leftmost nonzero column. This is a pivot column, with the pivot position at the top.
2. Select a nonzero entry in the pivot column as a pivot. If necessary, interchange rows to move this entry into the pivot position.
3. Use row replacement operations to create zeros in all positions below the pivot.
4. Apply steps 1-3 to the submatrix of all entries below and to the right of the pivot position. Repeat this process until there are no more nonzero rows to modify. (At this point we have reached an echelon form of the matrix.)
5. Beginning with the rightmost pivot and working upward and to the left, create zeros above each pivot using row operations. If a pivot is not 1, make it 1 by a scaling operation. (This step produces the reduced echelon form of the matrix.)

Solutions of Linear Systems

Let A be the coefficient matrix of a linear system. The pivot columns in the matrix correspond to what we call **basic variables**. The nonpivot columns correspond to what we call **free variables**.

Example 7.10 Suppose the augmented matrix of a linear system has been reduced to the following form:

$$\left[\begin{array}{ccccc} \blacksquare & * & * & * & * \\ 0 & \blacksquare & * & * & * \\ 0 & 0 & 0 & \blacksquare & * \\ 0 & 0 & 0 & 0 & 0 \end{array} \right]$$

where \blacksquare represents any nonzero number, and $*$ represents any number (including 0). The basic variables of this system are x_1, x_2 , and x_4 . The only free variable is x_3 .

Theorem 7.3 (Uniqueness Theorem)

If a linear system is consistent, then the solution set contains either

- (i) a unique solution, when there are no free variables, or
- (ii) infinitely many solutions, when there is at least one free variable.



Example 7.11 Suppose the following matrix is the augmented matrix of a linear system in the variables x_1, x_2 , and x_3 . Row reduce the matrix to echelon form to determine if it is consistent.

$$\left[\begin{array}{cccc} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 9 & 10 & 11 & 12 \end{array} \right]$$

If it is consistent, find the reduced echelon form and write the solution set using free variables as parameters.

Solution: To simplify the system, we perform *elementary row operations*.

Step 1: Eliminate the 5 in r_2

We adjust r_2 as follows:

$$r_2 \mapsto r_2 - 5r_1,$$

which gives:

$$\left[\begin{array}{cccc} 1 & 2 & 3 & 4 \\ 0 & -4 & -8 & -12 \\ 9 & 10 & 11 & 12 \end{array} \right]$$

Step 2: Eliminate the 9 in r_3

We adjust r_3 as follows:

$$r_3 \mapsto r_3 - 9r_1,$$

which gives:

$$\left[\begin{array}{ccc|c} 1 & 2 & 3 & 4 \\ 0 & -4 & -8 & -12 \\ 0 & -8 & -16 & -24 \end{array} \right]$$

Step 3: Eliminate the -8 in r_3

We adjust r_3 again to eliminate the entry in the second column:

$$r_3 \mapsto r_3 - 2r_2,$$

which gives:

$$\left[\begin{array}{cccc} 1 & 2 & 3 & 4 \\ 0 & -4 & -8 & -12 \\ 0 & 0 & 0 & 0 \end{array} \right]$$

Step 4: Scale r_2 to obtain a leading 1

We scale r_2 by $-\frac{1}{4}$:

$$r_2 \mapsto -\frac{1}{4}r_2,$$

which gives:

$$\left[\begin{array}{ccc|c} 1 & 2 & 3 & 4 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 0 & 0 \end{array} \right]$$

Step 5: Eliminate the 2 in r_1

We adjust r_1 as follows:

$$r_1 \mapsto r_1 - 2r_2,$$

which gives:

$$\left[\begin{array}{ccc|c} 1 & 0 & -1 & -2 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 0 & 0 \end{array} \right]$$

At this point, the matrix is in *reduced row-echelon form*.

Step 6: Interpret the matrix as a linear system

From the augmented matrix, we have:

$$r_1 : x_1 - x_3 = -2,$$

$$r_2 : x_2 + 2x_3 = 3,$$

$$r_3 : 0 = 0.$$

Since the third row corresponds to the equation $0 = 0$, which is always true, the system is **consistent**.

Step 7: Express the solution using a free variable

$$x_1 = -2 + x_3,$$

$$x_2 = 3 - 2x_3,$$

$$x_3 = x_3.$$



Chapter 8 Vectors and Matrices

A **vector** is a mathematical object that has both **magnitude** (length) and **direction**. In two or three dimensions, a vector is often represented as an arrow pointing from one point to another. For example, the vector $\mathbf{v} = (x_1, x_2)$ in 2D space describes a movement from the origin $(0, 0)$ to the point (x_1, x_2) .

Vectors are used to represent quantities like velocity, force, and displacement, which require both magnitude and direction to fully describe them. Vectors can be added together, scaled by a number, and decomposed into components.

Matrices, on the other hand, are rectangular arrays of numbers arranged in rows and columns. They are used to represent and solve systems of linear equations, perform transformations, and encode relationships between sets of vectors.

8.1 Vectors in \mathbb{R}^2 and \mathbb{R}^n

In linear algebra, vectors are often represented as **column matrices**. For example, the vector $\mathbf{v} = (x_1, x_2)$ can be written as a column matrix:

$$\mathbf{v} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

We say that two vectors are **equal** if and only if their corresponding entries are equal.

Example 8.1 The following are vectors in \mathbb{R}^2 (a.k.a. the plane consisting of ordered pairs of real numbers):

$$\mathbf{u} = \begin{bmatrix} 3 \\ 5 \end{bmatrix}, \quad \mathbf{v} = \begin{bmatrix} 5 \\ 3 \end{bmatrix}$$

They are not equal because their corresponding entries do not match.

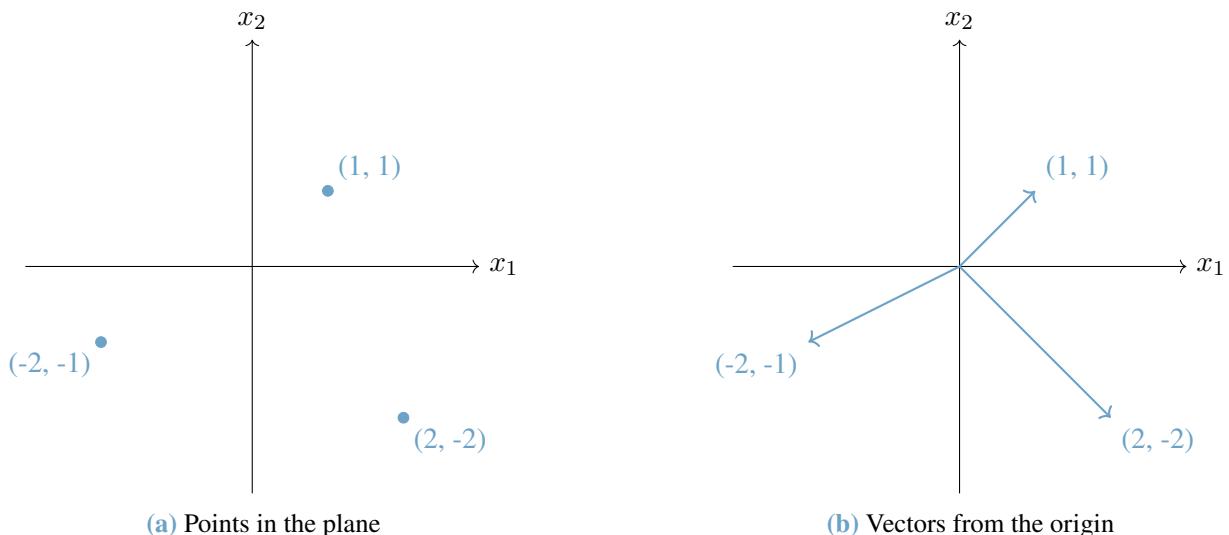


Figure 8.1: Points and vectors in the plane

The sum of two vectors \mathbf{u} and \mathbf{v} in \mathbb{R}^2 , denoted $\mathbf{u} + \mathbf{v}$, is obtained by adding the corresponding entries of \mathbf{u} and \mathbf{v} . Given a real number c , the scalar multiple of \mathbf{u} by c , denoted $c\mathbf{u}$, is obtained by multiplying each entry in \mathbf{u} by c .

Example 8.2 If \mathbf{u} and \mathbf{v} are as in the preceding example, then

$$\begin{aligned}\mathbf{u} + \mathbf{v} &= \begin{bmatrix} 3 \\ 5 \end{bmatrix} + \begin{bmatrix} 5 \\ 3 \end{bmatrix} = \begin{bmatrix} 3+5 \\ 5+3 \end{bmatrix} = \begin{bmatrix} 8 \\ 8 \end{bmatrix}, \text{ and} \\ 6\mathbf{u} &= 6 \begin{bmatrix} 3 \\ 5 \end{bmatrix} = \begin{bmatrix} 6 \cdot 3 \\ 6 \cdot 5 \end{bmatrix} = \begin{bmatrix} 18 \\ 30 \end{bmatrix}\end{aligned}$$

Remark: It is often helpful to identify a vector $\begin{bmatrix} a \\ b \end{bmatrix} \in \mathbb{R}^2$ with a geometric point (a, b) in the plane in order to get a picture of what we are working with. Please see [Figure 8.1](#) for an example.

If \mathbf{u} and \mathbf{v} in \mathbb{R}^2 are thought of as points in the plane, then $\mathbf{u} + \mathbf{v}$ corresponds to the fourth vertex of the parallelogram whose other vertices are $\mathbf{0}$, \mathbf{u} , and \mathbf{v} . Note: By $\mathbf{0}$, we mean the zero vector, or the vector whose entries are all zero. In \mathbb{R}^2 , we have $\mathbf{0} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$.

We call this the *parallelogram law of addition* and it can be seen in [Figure 8.2](#).

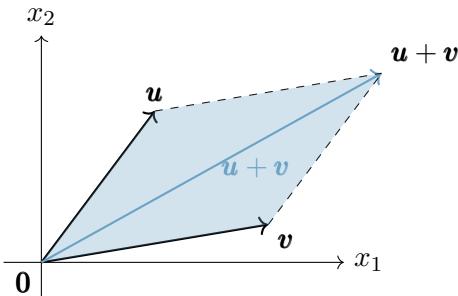


Figure 8.2: The parallelogram law of vector addition in \mathbb{R}^2 .

These ideas generalise to higher-dimensional spaces. More specifically, we can define \mathbb{R}^n as follows.

Definition 8.1

For each positive integer n , we let \mathbb{R}^n denote the collection of ordered n -tuples with each entry in \mathbb{R} . We often write these elements as $n \times 1$ matrices. We define addition and scalar multiplication of vectors in \mathbb{R}^n in the same way as we do for \mathbb{R}^2 . That is, we go coordinate-by-coordinate.



Example 8.3 If $u_1, u_2, \dots, u_n \in \mathbb{R}$, then

$$\mathbf{u} = \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{bmatrix} \in \mathbb{R}^n.$$

Example 8.4 If \mathbf{u} and \mathbf{v} are in \mathbb{R}^n (with entries denoted u_1, \dots, u_n and v_1, \dots, v_n , respectively), and $c \in \mathbb{R}$, then

$$\mathbf{u} + \mathbf{v} = \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{bmatrix} + \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix} = \begin{bmatrix} u_1 + v_1 \\ u_2 + v_2 \\ \vdots \\ u_n + v_n \end{bmatrix}, \text{ and } c\mathbf{u} = c \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_n \end{bmatrix} = \begin{bmatrix} cu_1 \\ cu_2 \\ \vdots \\ cu_n \end{bmatrix}$$

Algebraic Properties of \mathbb{R}^n

Let $\mathbf{u}, \mathbf{v}, \mathbf{w} \in \mathbb{R}^n$ and $c, d \in \mathbb{R}$. Then the following properties hold:

1. **Commutative Property of Addition:** $\mathbf{u} + \mathbf{v} = \mathbf{v} + \mathbf{u}$.
2. **Associative Property of Addition:** $(\mathbf{u} + \mathbf{v}) + \mathbf{w} = \mathbf{u} + (\mathbf{v} + \mathbf{w})$.
3. **Additive Identity:** There exists a vector $\mathbf{0} \in \mathbb{R}^n$ such that $\mathbf{u} + \mathbf{0} = \mathbf{u}$ for all $\mathbf{u} \in \mathbb{R}^n$.
4. **Additive Inverse:** For each $\mathbf{u} \in \mathbb{R}^n$, there exists a vector $-\mathbf{u} \in \mathbb{R}^n$ such that $\mathbf{u} + (-\mathbf{u}) = \mathbf{0}$.
5. **Distributive Property:** $c(\mathbf{u} + \mathbf{v}) = cu + cv$ and $(c + d)\mathbf{u} = cu + du$.
6. **Associative Property of Scalar Multiplication:** $c(d\mathbf{u}) = (cd)\mathbf{u}$.
7. **Multiplicative Identity:** $1\mathbf{u} = \mathbf{u}$.

A **linear combination** is a way to combine vectors using scalar multiplication and addition. Given a set of vectors, we multiply each by a scalar and then sum the results. Linear combinations help us understand how vectors relate to each other and whether one vector can be expressed in terms of others.

Definition 8.2 (Linear Combinations)

Given a set of vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_p \in \mathbb{R}^n$ and scalars $c_1, c_2, \dots, c_p \in \mathbb{R}$, the vector \mathbf{y} given by

$$\mathbf{y} = c_1\mathbf{v}_1 + \dots + c_p\mathbf{v}_p$$

is called a linear combination of $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_p$ with weights c_1, c_2, \dots, c_p .



Example 8.5 Let $\mathbf{v}_1 = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$ and $\mathbf{v}_2 = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$. Some linear combinations of \mathbf{v}_1 and \mathbf{v}_2 include

$$\mathbf{0} = 0\mathbf{v}_1 + 0\mathbf{v}_2$$

$$\begin{bmatrix} 3 \\ 0 \end{bmatrix} = \mathbf{v}_1 + 2\mathbf{v}_2$$

$$\begin{bmatrix} -5 \\ -1 \end{bmatrix} = -2\mathbf{v}_1 - 3\mathbf{v}_2, \text{ and}$$

$$\begin{bmatrix} 5 \\ 1 \end{bmatrix} = 2\mathbf{v}_1 + 3\mathbf{v}_2$$

Vector Equations and Linear Systems

It is often the case that we wish to know if some vector \mathbf{b} can be formed as a linear combination of some other set of vectors $\mathbf{a}_1, \dots, \mathbf{a}_n$. The process for figuring this out is given by the following.

Using Matrices to Determine Linear Combinations

A vector equation

$$x_1\mathbf{a}_1 + x_2\mathbf{a}_2 + \cdots + x_n\mathbf{a}_n = \mathbf{b}$$

has the same solution set as the linear system whose augmented matrix is

$$\left[\begin{array}{ccccc} \mathbf{a}_1 & \mathbf{a}_2 & \dots & \mathbf{a}_n & \mathbf{b} \end{array} \right].$$

More specifically, if the \mathbf{a}_i 's are in \mathbb{R}^m with

$$\mathbf{a}_1 = \begin{bmatrix} a_{11} \\ a_{21} \\ \vdots \\ a_{m1} \end{bmatrix}, \quad \mathbf{a}_2 = \begin{bmatrix} a_{12} \\ a_{22} \\ \vdots \\ a_{m2} \end{bmatrix}, \dots, \quad \mathbf{a}_n = \begin{bmatrix} a_{1n} \\ a_{2n} \\ \vdots \\ a_{mn} \end{bmatrix}, \quad \text{and} \quad \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_m \end{bmatrix}$$

then you would row reduce the matrix

$$\left[\begin{array}{ccccc} a_{11} & a_{12} & \dots & a_{1n} & b_1 \\ a_{21} & a_{22} & \dots & a_{2n} & b_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ a_{m1} & a_{m2} & \dots & a_{mn} & b_m \end{array} \right]$$

$\uparrow \quad \uparrow \quad \dots \quad \uparrow \quad \uparrow$
 $\mathbf{a}_1 \quad \mathbf{a}_2 \quad \dots \quad \mathbf{a}_n \quad \mathbf{b}$

to determine if there is some set of weights x_1, \dots, x_n that work.

Example 8.6 Let

$$\mathbf{a}_1 = \begin{bmatrix} 2 \\ -1 \\ 1 \end{bmatrix}, \quad \mathbf{a}_2 = \begin{bmatrix} 0 \\ 8 \\ -2 \end{bmatrix}, \quad \mathbf{a}_3 = \begin{bmatrix} 6 \\ 5 \\ 1 \end{bmatrix}, \quad \text{and} \quad \mathbf{b} = \begin{bmatrix} 10 \\ 3 \\ 7 \end{bmatrix}$$

We determine if \mathbf{b} is a linear combination of $\mathbf{a}_1, \mathbf{a}_2, \mathbf{a}_3$, i.e. if there is some set of weights x_1, x_2, x_3 such that $x_1\mathbf{a}_1 + x_2\mathbf{a}_2 + x_3\mathbf{a}_3 = \mathbf{b}$. By the above, we translate this question to the matrix setting.

$$\left[\begin{array}{cccc} 2 & 0 & 6 & 10 \\ -1 & 8 & 5 & 3 \\ 1 & -2 & 1 & 7 \end{array} \right] \quad r_1 \leftrightarrow r_3 \quad \left[\begin{array}{cccc} 1 & -2 & 1 & 7 \\ -1 & 8 & 5 & 3 \\ 2 & 0 & 6 & 10 \end{array} \right]$$

$$r_2 \mapsto r_2 + r_1 \quad \left[\begin{array}{cccc} 1 & -2 & 1 & 7 \\ 0 & 6 & 6 & 10 \\ 2 & 0 & 6 & 10 \end{array} \right]$$

$$r_3 \mapsto r_3 - 2r_1 \quad \left[\begin{array}{cccc} 1 & -2 & 1 & 7 \\ 0 & 6 & 6 & 10 \\ 0 & 4 & 4 & -4 \end{array} \right]$$

$$r_3 \mapsto r_3 - \frac{4}{6}r_2 \quad \left[\begin{array}{cccc} 1 & -2 & 1 & 7 \\ 0 & 6 & 6 & 10 \\ 0 & 0 & 0 & \blacksquare \end{array} \right].$$

Since the symbol \blacksquare denotes something nonzero, we see that there's no solution, i.e. that \mathbf{b} is not a linear combination of $\mathbf{a}_1, \mathbf{a}_2, \mathbf{a}_3$.

Example 8.7 Let

$$\mathbf{a}_1 = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}, \quad \mathbf{a}_2 = \begin{bmatrix} -4 \\ 6 \\ -4 \end{bmatrix}, \quad \mathbf{a}_3 = \begin{bmatrix} -6 \\ 7 \\ 5 \end{bmatrix}, \quad \text{and } \mathbf{b} = \begin{bmatrix} 11 \\ -5 \\ 9 \end{bmatrix}$$

We determine if \mathbf{b} is a linear combination of $\mathbf{a}_1, \mathbf{a}_2, \mathbf{a}_3$. We reduce the corresponding matrix to echelon form:

$$\left[\begin{array}{cccc} 1 & -4 & -6 & 11 \\ 0 & 6 & 7 & -5 \\ 1 & -4 & 5 & 9 \end{array} \right] \quad r_3 \mapsto r_3 - r_1 \quad \left[\begin{array}{cccc} 1 & -4 & -6 & 11 \\ 0 & 6 & 7 & -5 \\ 0 & 0 & 11 & -2 \end{array} \right]$$

and we see that there is a solution. Now we find what weights x_1, x_2, x_3 work by finding the reduced echelon form.

$$\begin{aligned} r_3 \mapsto \frac{1}{11}r_3 &\quad \left[\begin{array}{cccc} 1 & -4 & -6 & 11 \\ 0 & 6 & 7 & -5 \\ 0 & 0 & 1 & -\frac{2}{11} \end{array} \right] \\ r_2 \mapsto r_2 - 7r_3 &\quad \left[\begin{array}{cccc} 1 & -4 & 0 & \frac{109}{11} \\ 0 & 6 & 0 & -\frac{41}{11} \\ 0 & 0 & 1 & -\frac{2}{11} \end{array} \right] \\ r_1 \mapsto r_1 + 6r_3 &\quad \left[\begin{array}{cccc} 1 & -4 & 0 & \frac{109}{11} \\ 0 & 1 & 0 & -\frac{41}{66} \\ 0 & 0 & 1 & -\frac{2}{11} \end{array} \right] \\ r_2 \mapsto \frac{1}{6}r_2 &\quad \left[\begin{array}{cccc} 1 & -4 & 0 & \frac{109}{11} \\ 0 & 1 & 0 & -\frac{41}{66} \\ 0 & 0 & 1 & -\frac{2}{11} \end{array} \right] \\ r_1 \mapsto r_1 + 4r_2 &\quad \left[\begin{array}{cccc} 1 & 0 & 0 & \frac{245}{38} \\ 0 & 1 & 0 & -\frac{41}{66} \\ 0 & 0 & 1 & -\frac{2}{11} \end{array} \right], \end{aligned}$$

so $(x_1, x_2, x_3) = \left(\frac{245}{38}, -\frac{41}{66}, -\frac{2}{11} \right)$. Since there are no free variables, this is the unique solution.

The **span** of a set of vectors is the collection of all possible linear combinations of those vectors. In other words, it's the set of all vectors you can reach by scaling and adding the given vectors. The span gives us insight into the "space" those vectors cover:

Definition 8.3 (Span of a Set of Vectors)

If $\mathbf{v}_1, \dots, \mathbf{v}_p$ are in \mathbb{R}^n , then the set of all linear combinations of $\mathbf{v}_1, \dots, \mathbf{v}_p$ is denoted by $\text{Span}\{\mathbf{v}_1, \dots, \mathbf{v}_p\}$ and is called the subset of \mathbb{R}^n spanned by $\mathbf{v}_1, \dots, \mathbf{v}_p$. In other words, the span of $\mathbf{v}_1, \dots, \mathbf{v}_p$ is all vectors that can be written in the form

$$c_1\mathbf{v}_1 + \cdots + c_p\mathbf{v}_p$$

with c_1, \dots, c_p scalars.



8.2 Matrix Equations

In the previous sections, we explored different ways of representing linear relationships. For instance, we looked at individual linear equations, such as

$$2x_1 + 3x_2 = 5,$$

and systems of linear equations, like

$$x_1 + 2x_2 = 4$$

$$3x_1 - x_2 = 2.$$

We also expressed these systems as vector equations, such as

$$x_1 \begin{bmatrix} 1 \\ 3 \end{bmatrix} + x_2 \begin{bmatrix} 2 \\ -1 \end{bmatrix} = \begin{bmatrix} 4 \\ 2 \end{bmatrix}.$$

Now, it's time to introduce a powerful new form: the **matrix equation**. This compact representation allows us to handle systems of linear equations efficiently using matrix notation. In this form, our system becomes

$$A\mathbf{x} = \mathbf{b},$$

where A is a matrix, \mathbf{x} is a vector of variables, and \mathbf{b} is the result vector. Matrix equations give us a structured, algebraic approach to solving systems.

Definition 8.4 (Matrix Equation)

If A is an $m \times n$ matrix with columns $\mathbf{a}_1, \dots, \mathbf{a}_n$, and if $\mathbf{x} \in \mathbb{R}^n$, then the product of A and \mathbf{x} , denoted by $A\mathbf{x}$, is

$$A\mathbf{x} = \begin{bmatrix} \mathbf{a}_1 & \mathbf{a}_2 & \dots & \mathbf{a}_n \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = x_1\mathbf{a}_1 + x_2\mathbf{a}_2 + \dots + x_n\mathbf{a}_n$$



Example 8.8

$$\begin{aligned} \begin{bmatrix} 4 & 1 & 2 \\ 8 & 0 & 3 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \\ 4 \end{bmatrix} &= 2 \begin{bmatrix} 4 \\ 8 \end{bmatrix} + 1 \begin{bmatrix} 1 \\ 0 \end{bmatrix} + 4 \begin{bmatrix} 2 \\ 3 \end{bmatrix} \\ &= \begin{bmatrix} 8 \\ 16 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix} + \begin{bmatrix} 8 \\ 12 \end{bmatrix} \\ &= \begin{bmatrix} 17 \\ 28 \end{bmatrix} \end{aligned}$$

Building a bit on the main result from the preceding section, we have the following theorem.

Theorem 8.1

If A is an $m \times n$ matrix, with columns $\mathbf{a}_1, \dots, \mathbf{a}_n \in \mathbb{R}^m$ and if $\mathbf{b} \in \mathbb{R}^m$, the matrix equation

$$A\mathbf{x} = \mathbf{b}$$

has the same solution set as the vector equation

$$x_1\mathbf{a}_1 + x_2\mathbf{a}_2 + \cdots + x_n\mathbf{a}_n = \mathbf{b}$$

which, in turn, has the same solution set as the system of linear equations with augmented matrix

$$\left[\begin{array}{cccc} \mathbf{a}_1 & \mathbf{a}_2 & \dots & \mathbf{a}_n & \mathbf{b} \end{array} \right].$$



This theorem is essential because it ties together the three forms of representing and solving systems of linear equations: matrix equations, vector equations, and systems of equations with augmented matrices. It shows that no matter which form we use, they all lead to the same solution set.

By stating that the matrix equation $A\mathbf{x} = \mathbf{b}$ is equivalent to the vector equation $x_1\mathbf{a}_1 + x_2\mathbf{a}_2 + \cdots + x_n\mathbf{a}_n = \mathbf{b}$, we can interpret solving a matrix equation as finding the right combination of the columns of A that yields \mathbf{b} . This connection allows us to visualize the problem in terms of vector spaces and linear combinations.

Moreover, the theorem shows that this same process is reflected in the augmented matrix, where row reduction reveals the solution through elementary row operations. So whether we approach the problem algebraically, geometrically, or algorithmically, we are dealing with the same underlying structure. This unifying perspective simplifies our approach to solving systems and provides flexibility in how we choose to represent and manipulate the problem.

Remark: The equation $A\mathbf{x} = \mathbf{b}$ has a solution if and only if \mathbf{b} is a linear combination of the columns of A .

The following theorem tells us when the column vectors of an $m \times n$ matrix (the columns are vectors in \mathbb{R}^m , since there are m rows) can be used to generate all of \mathbb{R}^m . This basically summarises several things we have already seen in different contexts.

Theorem 8.2

Let A be an $m \times n$ matrix. Then the following statements are equivalent:

- (a) For each $\mathbf{b} \in \mathbb{R}^m$, the equation $A\mathbf{x} = \mathbf{b}$ has a solution.
- (b) Each $\mathbf{b} \in \mathbb{R}^m$ is a linear combination of the columns of A .
- (c) The columns of A span \mathbb{R}^m .
- (d) The matrix A has a pivot position in every row.



This theorem shows that several seemingly different ideas are in fact equivalent. The existence of a solution to the matrix equation $A\mathbf{x} = \mathbf{b}$ (statement (a)) depends on whether the columns of A can form any vector in \mathbb{R}^m (statement (b)). Geometrically, this means the columns span the entire space \mathbb{R}^m (statement (c)).

Finally, the presence of a pivot position in every row (statement (d)) gives an algebraic condition that guarantees the span of the columns covers \mathbb{R}^m , ensuring that there is always a solution to $A\mathbf{x} = \mathbf{b}$.

This equivalence is a powerful tool, as it connects solutions, linear combinations, and geometric interpretations, while also providing a concrete method (checking for pivot positions) to verify these properties.

Row-Vector Rule for Computing $A\mathbf{x}$

Assuming the product $A\mathbf{x}$ is defined, the i -th entry in $A\mathbf{x}$ is the sum of the products of corresponding entries from row i of A and the vector \mathbf{x} . In other words, the i -th entry of $A\mathbf{x}$ is the dot product of the vector forming the i -th row of A and the vector \mathbf{x} .

This rule simplifies the matrix multiplication process by breaking it down into smaller, familiar operations — dot products — making it easier to compute and understand. It also highlights the connection between matrix multiplication and linear combinations, as the product $A\mathbf{x}$ is a linear combination of the rows of A with weights given by the entries of \mathbf{x} .

Example 8.9 Let

$$A = \begin{bmatrix} 1 & -1 & 4 \\ 5 & 0 & 2 \end{bmatrix}, \text{ and } \mathbf{x} = \begin{bmatrix} 2 \\ 1 \\ 5 \end{bmatrix}$$

Then

$$\begin{bmatrix} 1 & -1 & 4 \\ 5 & 0 & 2 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \\ 5 \end{bmatrix} = \begin{bmatrix} 1 \cdot 2 + (-1) \cdot 1 + 4 \cdot 5 \\ 5 \cdot 2 + 0 \cdot 1 + 2 \cdot 5 \end{bmatrix} = \begin{bmatrix} 21 \\ 20 \end{bmatrix}$$

Notice that the number of columns in A must match the number of rows in \mathbf{x} , for otherwise the dot product would not make sense!

The next theorem captures two important properties of matrix-vector multiplication: **distributivity** and **scalar multiplication**. These properties mirror the familiar rules of algebra but now apply in the context of matrices and vectors.

Theorem 8.3

If A is an $m \times n$ matrix, \mathbf{u} and \mathbf{v} are vectors in \mathbb{R}^n , and c is a scalar, then:

- (a) $A(\mathbf{u} + \mathbf{v}) = A\mathbf{u} + A\mathbf{v}$
- (b) $A(c\mathbf{u}) = c(A\mathbf{u})$



Property (a) shows that matrix multiplication distributes over vector addition. It means that multiplying A by the sum of two vectors is the same as multiplying A by each vector separately and then adding the results. Property (b) illustrates that scalar multiplication commutes with matrix multiplication. Multiplying a vector by a scalar first, then applying the matrix, gives the same result as applying the matrix first and then multiplying the resulting vector by the scalar.

Example 8.10 Let

$$A = \begin{bmatrix} 1 & 5 & -2 & 0 \\ -3 & 1 & 9 & -5 \\ 4 & -8 & -1 & 7 \end{bmatrix}, \quad \mathbf{p} = \begin{bmatrix} 3 \\ -2 \\ 0 \\ -4 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} -7 \\ 9 \\ 0 \end{bmatrix}$$

It can be shown that \mathbf{p} is a solution of $A\mathbf{x} = \mathbf{b}$. Use this fact to write \mathbf{b} as a linear combination of the columns of A .

Solution: Since \mathbf{p} is a solution to $A\mathbf{x} = \mathbf{b}$, we have:

$$A\mathbf{p} = \mathbf{b}.$$

We can express \mathbf{b} as a linear combination of the columns of A using the entries of \mathbf{p} . Let $\mathbf{a}_1, \mathbf{a}_2, \mathbf{a}_3$, and \mathbf{a}_4 denote the columns of A , so:

$$A = [\mathbf{a}_1 \mid \mathbf{a}_2 \mid \mathbf{a}_3 \mid \mathbf{a}_4].$$

Thus,

$$\mathbf{b} = p_1\mathbf{a}_1 + p_2\mathbf{a}_2 + p_3\mathbf{a}_3 + p_4\mathbf{a}_4.$$

Now, we compute each term one at a time.

Compute $p_1\mathbf{a}_1$:

$$p_1\mathbf{a}_1 = 3 \begin{bmatrix} 1 \\ -3 \\ 4 \end{bmatrix} = \begin{bmatrix} 3 \\ -9 \\ 12 \end{bmatrix}.$$

Compute $p_2\mathbf{a}_2$:

$$p_2\mathbf{a}_2 = (-2) \begin{bmatrix} 5 \\ 1 \\ -8 \end{bmatrix} = \begin{bmatrix} -10 \\ -2 \\ 16 \end{bmatrix}.$$

Compute $p_3\mathbf{a}_3$:

$$p_3\mathbf{a}_3 = 0 \begin{bmatrix} -2 \\ 9 \\ -1 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}.$$

Compute $p_4\mathbf{a}_4$:

$$p_4\mathbf{a}_4 = (-4) \begin{bmatrix} 0 \\ -5 \\ 7 \end{bmatrix} = \begin{bmatrix} 0 \\ 20 \\ -28 \end{bmatrix}.$$

Add the computed vectors to find \mathbf{b} :

$$\mathbf{b} = p_1\mathbf{a}_1 + p_2\mathbf{a}_2 + p_3\mathbf{a}_3 + p_4\mathbf{a}_4 = \begin{bmatrix} 3 \\ -9 \\ 12 \end{bmatrix} + \begin{bmatrix} -10 \\ -2 \\ 16 \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} + \begin{bmatrix} 0 \\ 20 \\ -28 \end{bmatrix}.$$

Now, compute the sum step by step.

Step 1: Add $p_1\mathbf{a}_1$ and $p_2\mathbf{a}_2$:

$$\begin{bmatrix} 3 \\ -9 \\ 12 \end{bmatrix} + \begin{bmatrix} -10 \\ -2 \\ 16 \end{bmatrix} = \begin{bmatrix} -7 \\ -11 \\ 28 \end{bmatrix}.$$

Step 2: Add $p_3\mathbf{a}_3$ (which is zero) to the result:

$$\begin{bmatrix} -7 \\ -11 \\ 28 \end{bmatrix} + \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} = \begin{bmatrix} -7 \\ -11 \\ 28 \end{bmatrix}.$$

Step 3: Add $p_4\mathbf{a}_4$ to the result:

$$\begin{bmatrix} -7 \\ -11 \\ 28 \end{bmatrix} + \begin{bmatrix} 0 \\ 20 \\ -28 \end{bmatrix} = \begin{bmatrix} -7 \\ 9 \\ 0 \end{bmatrix}.$$

Final Result:

$$\mathbf{b} = \begin{bmatrix} -7 \\ 9 \\ 0 \end{bmatrix}.$$

Conclusion:

We have expressed \mathbf{b} as a linear combination of the columns of A using the entries of \mathbf{p} :

$$\mathbf{b} = 3\mathbf{a}_1 - 2\mathbf{a}_2 + 0\mathbf{a}_3 - 4\mathbf{a}_4.$$



8.3 Solution Sets of Linear Systems

Solution sets of linear systems play a crucial role in the study of linear algebra and will reappear in various contexts throughout the subject. In this section, we use vector notation to provide clear, explicit, and geometric descriptions of these solution sets.

A linear system is said to be *homogeneous* if it can be written in the form

$$A\mathbf{x} = \mathbf{0}$$

where A is an $m \times n$ matrix and $\mathbf{0}$ is the zero vector in \mathbb{R}^m . Such a system always has the *trivial solution*, namely $\mathbf{x} = \mathbf{0}$ (the zero vector in \mathbb{R}^n , not \mathbb{R}^m).

Since homogeneous systems always have the trivial solution, the interesting question is whether or not they have *nontrivial solutions*.

Remark: The homogeneous equation $A\mathbf{x} = \mathbf{0}$ has a nontrivial solution if and only if the equation has at least one free variable.

Example 8.11 Determine if the following homogeneous system has a nontrivial solution. If it does, describe

the solution set using the free variable(s).

$$2x_1 + x_2 - 3x_3 = 0$$

$$x_1 - x_2 + x_3 = 0$$

$$-2x_1 + 5x_2 - 7x_3 = 0$$

Solution: We let A be the coefficient matrix of the system, and reduce the augmented matrix $\left[\begin{array}{cc|c} A & 0 \end{array} \right]$ to echelon form (notice how we put the second equation first in order to simplify the computation a bit):

$$\left[\begin{array}{cccc} 1 & -1 & 1 & 0 \\ 2 & 1 & -3 & 0 \\ -2 & 5 & -7 & 0 \end{array} \right] \sim \left[\begin{array}{cccc} 1 & -1 & 1 & 0 \\ 0 & 3 & -5 & 0 \\ 0 & 3 & -5 & 0 \end{array} \right] \sim \left[\begin{array}{cccc} 1 & -1 & 1 & 0 \\ 0 & 3 & -5 & 0 \\ 0 & 0 & 0 & 0 \end{array} \right]$$

so we see that x_3 is a free variable, and as a result we have nontrivial solutions. Now we get the reduced echelon form:

$$\left[\begin{array}{cccc} 1 & -1 & 1 & 0 \\ 0 & 3 & -5 & 0 \\ 0 & 0 & 0 & 0 \end{array} \right] \sim \left[\begin{array}{cccc} 1 & -1 & 1 & 0 \\ 0 & 1 & -\frac{5}{3} & 0 \\ 0 & 0 & 0 & 0 \end{array} \right] \sim \left[\begin{array}{cccc} 1 & 0 & -\frac{2}{3} & 0 \\ 0 & 1 & -\frac{5}{3} & 0 \\ 0 & 0 & 0 & 0 \end{array} \right]$$

This gives us the equations

$$x_1 - \frac{2}{3}x_3 = 0 \implies x_1 = \frac{2}{3}x_3$$

$$x_2 - \frac{5}{3}x_3 = 0 \implies x_2 = \frac{5}{3}x_3$$

In vector form, the solution \mathbf{x} of the equation $A\mathbf{x} = 0$ is written

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} \frac{2}{3}x_3 \\ \frac{5}{3}x_3 \\ x_3 \end{bmatrix} = x_3 \begin{bmatrix} \frac{2}{3} \\ \frac{5}{3} \\ 1 \end{bmatrix}$$

Since x_3 can be anything, in geometric terms this solution set describes the line in \mathbb{R}^3 extending from the origin through the point $\left(\frac{2}{3}, \frac{5}{3}, 1\right)$. ◀

Parametric Vector Form

Whenever a solution set is described explicitly with vectors (as in the preceding example), we say that the solution is in **parametric vector form**.

Example 8.12 Describe all solutions of the homogeneous equation

$$x_1 - 2x_2 - 5x_3 = 0$$

Solution:

Writing this system in a matrix, we would see that x_2 and x_3 are free variables, and x_1 is a basic variable with $x_1 = 2x_2 + 5x_3$. Hence the general solution is

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 2x_2 + 5x_3 \\ x_2 \\ x_3 \end{bmatrix} = x_2 \begin{bmatrix} 2 \\ 1 \\ 0 \end{bmatrix} + x_3 \begin{bmatrix} 5 \\ 0 \\ 1 \end{bmatrix}$$

where the right-most part of this equation is the **parametric vector** form of the solution set. In this case, the entire solution set is a plane in \mathbb{R}^3 that contains the two lines extending from $(0, 0, 0)$ to $(2, 1, 0)$, and from $(0, 0, 0)$ to $(5, 0, 1)$. ◀

When a *nonhomogeneous* linear system has many solutions, the general solution can be written in parametric vector form as one vector plus arbitrary linear combinations of vectors that satisfy the corresponding homogeneous system.

Example 8.13 We reconsider the homogeneous system at the beginning of this section, except this time it will be nonhomogeneous (i.e., the right side will not be all zeroes):

$$2x_1 + x_2 - 3x_3 = 2$$

$$x_1 - x_2 + x_3 = 1$$

$$-2x_1 + 5x_2 - 7x_3 = -2$$

As before, we perform row operations on the augmented matrix $[A \ b]$ where A is the same coefficient matrix and

$$\mathbf{b} = \begin{bmatrix} 2 \\ 1 \\ -2 \end{bmatrix}$$

and we find that

$$\left[\begin{array}{cccc|c} 1 & -1 & 1 & 1 & 2 \\ 2 & 1 & -3 & 2 & 1 \\ -2 & 5 & -7 & -2 & -2 \end{array} \right] \sim \left[\begin{array}{cccc|c} 1 & 0 & -\frac{2}{3} & 1 & 2 \\ 0 & 1 & -\frac{5}{3} & 0 & 1 \\ 0 & 0 & 0 & 0 & -2 \end{array} \right]$$

Hence,

$$x_1 = 1 + \frac{2}{3}x_3, \text{ and}$$

$$x_2 = \frac{5}{3}x_3$$

This can be written as

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 1 + \frac{2}{3}x_3 \\ \frac{5}{3}x_3 \\ x_3 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} + x_3 \begin{bmatrix} \frac{2}{3} \\ \frac{5}{3} \\ 1 \end{bmatrix}$$

which has the form we claimed. This is the equation of the line through the point $(1, 0, 0)$ that is parallel to the line extending from $(0, 0, 0)$ through $(\frac{2}{3}, \frac{5}{3}, 1)$.

We summarise the general situation with the following theorem. It might be best to think of the conclusion as shifting all homogeneous solutions by the vector \mathbf{p} .

Theorem 8.4

Suppose $A\mathbf{x} = \mathbf{b}$ is consistent for some \mathbf{b} , and let \mathbf{p} be a solution. Then the solution set of $A\mathbf{x} = \mathbf{b}$ is the set of all vectors of the form

$$\mathbf{w} = \mathbf{p} + \mathbf{v}_h,$$

where \mathbf{v}_h is any solution of the homogeneous equation $A\mathbf{x} = \mathbf{0}$. 

Example 8.14 Write the general solution of

$$x_1 - 2x_2 - 5x_3 = 3$$

in parametric vector form. In geometric terms, what does this solution set look like in comparison to the solution set of the equation $x_1 - 2x_2 - 5x_3 = 0$ that we saw earlier?

Solution: Let us solve the equation:

$$x_1 - 2x_2 - 5x_3 = 3$$

To express the general solution, we solve for x_1 in terms of x_2 and x_3 :

$$x_1 = 3 + 2x_2 + 5x_3,$$

$$x_2 = x_2,$$

$$x_3 = x_3.$$

We can express the solution set in *parametric vector form*:

$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 3 \\ 0 \\ 0 \end{bmatrix} + x_2 \begin{bmatrix} 2 \\ 1 \\ 0 \end{bmatrix} + x_3 \begin{bmatrix} 5 \\ 0 \\ 1 \end{bmatrix}, \quad \text{for all } s, t \in \mathbb{R}.$$

Geometric Interpretation: The solution set represents a plane in \mathbb{R}^3 . This plane is parallel to the plane defined by the homogeneous equation:

$$x_1 - 2x_2 - 5x_3 = 0$$

The homogeneous solution set can be expressed as:

$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = x_2 \begin{bmatrix} 2 \\ 1 \\ 0 \end{bmatrix} + x_3 \begin{bmatrix} 5 \\ 0 \\ 1 \end{bmatrix}, \quad \text{for all } s, t \in \mathbb{R}.$$

Comparing both solution sets, we observe that the original solution is the homogeneous solution shifted by the vector $\begin{bmatrix} 3 \\ 0 \\ 0 \end{bmatrix}$. Geometrically, this means the plane defined by $x_1 - 2x_2 - 5x_3 = 3$ is parallel to, but not the same as, the plane defined by $x_1 - 2x_2 - 5x_3 = 0$. The two planes are offset along the x_1 -axis by 3 units. The general solution of the equation $x_1 - 2x_2 - 5x_3 = 3$ is a plane in \mathbb{R}^3 that is parallel to the solution set of the homogeneous equation $x_1 - 2x_2 - 5x_3 = 0$, but shifted away from the origin. 

Example 8.15 Find the parametric equation of the line through $\mathbf{a} = \begin{bmatrix} 3 \\ -2 \end{bmatrix}$ that is also parallel to $\mathbf{b} = \begin{bmatrix} -7 \\ 6 \end{bmatrix}$.

Solution: The line through \mathbf{a} parallel to \mathbf{b} is given by the equation $\mathbf{r} = \mathbf{a} + t\mathbf{b}$, where t is a parameter. Substituting the given vectors, we have

$$\begin{aligned}\mathbf{r} &= \begin{bmatrix} 3 \\ -2 \end{bmatrix} + t \begin{bmatrix} -7 \\ 6 \end{bmatrix} \\ &= \begin{bmatrix} 3 - 7t \\ -2 + 6t \end{bmatrix}\end{aligned}$$

This is the parametric equation of the line through \mathbf{a} that is parallel to \mathbf{b} . ◀

8.4 Linear Independence

We conclude this section by introducing one of the most important concepts in linear algebra: *linear independence*. This concept is fundamental to understanding the structure of vectors and their relationships in vector spaces.

Definition 8.5

An indexed set of vectors $\{\mathbf{v}_1, \dots, \mathbf{v}_p\} \subset \mathbb{R}^n$ is said to be **linearly independent** if the vector equation

$$x_1\mathbf{v}_1 + \cdots + x_p\mathbf{v}_p = \mathbf{0}$$

has only the trivial solution, i.e., if the only solution is $(x_1, \dots, x_p) = (0, \dots, 0)$. Likewise, the set $\{\mathbf{v}_1, \dots, \mathbf{v}_p\}$ is said to be **linearly dependent** if there exist weights c_1, \dots, c_p , not all zero, such that

$$c_1\mathbf{v}_1 + \cdots + c_p\mathbf{v}_p = \mathbf{0}$$

We call such an equation a *linear dependence relation* when the weights are not all zero. ♣

Remark: A set of vectors cannot be both linearly independent and linearly dependent, but it must be one of them!

Determining if a set of vectors is linearly independent is tantamount to solving the matrix equation

$$A\mathbf{x} = \mathbf{0},$$

where the columns of A are given by the vectors. The set of vectors is linearly independent if and only if the only solution is $\mathbf{x} = \mathbf{0}$. Otherwise, there is some linear dependence relation.

Example 8.16 Determine if the following vectors in \mathbb{R}^3 are linearly independent:

$$\mathbf{v}_1 = \begin{bmatrix} 2 \\ 1 \\ 2 \end{bmatrix}, \quad \mathbf{v}_2 = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \quad \mathbf{v}_3 = \begin{bmatrix} 0 \\ 1 \\ 1 \end{bmatrix}$$

Solution: We begin by reducing the corresponding augmented matrix to echelon form:

$$\begin{bmatrix} 2 & 1 & 0 & 0 \\ 1 & 1 & 1 & 0 \\ 2 & 1 & 1 & 0 \end{bmatrix} \sim \begin{bmatrix} 1 & 1 & 1 & 0 \\ 2 & 1 & 0 & 0 \\ 2 & 1 & 1 & 0 \end{bmatrix} \sim \begin{bmatrix} 1 & 1 & 1 & 0 \\ 0 & -1 & -2 & 0 \\ 0 & -1 & -1 & 0 \end{bmatrix} \sim \begin{bmatrix} 1 & 1 & 1 & 0 \\ 0 & -1 & -2 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

from which we see that there is a solution. Now we want to know if the solution is the trivial solution $(0, 0, 0)$. Normally, we would continue row operations until we reach reduced echelon form, but we can be smarter about this. Notice first of all that there are no bad rows (so there is at least one solution), which we expect since a homogeneous system always has at least the trivial solution. Since there are no free variables, we see there is exactly one solution. This tells us that the solution must be

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

Hence the set of vectors $\{\mathbf{v}_1, \mathbf{v}_2, \mathbf{v}_3\}$ is linearly independent. ◀

In general, it is very useful to determine if some set of vectors is linearly independent, so it is good to have some theorems to handle this problem quickly in certain special cases.

Theorem 8.5

- (a) If a set of vectors contains the zero vector, then the set is linearly dependent.
- (b) If a set of vectors contains a scalar multiple of another vector, then the set is linearly dependent.
- (c) If a set of vectors contains more vectors than there are entries in each vector, then the set is linearly dependent.



The last part of the theorem is particularly useful, as it allows us to quickly determine if a set of vectors is linearly dependent by checking if one of the vectors is a linear combination of the others. We can formalise this idea in the following corollary.

Corollary 8.1

An indexed set $S = \{\mathbf{v}_1, \dots, \mathbf{v}_p\}$ of two or more vectors is linearly dependent if and only if at least one of the vectors in S is a linear combination of the others. If S is linearly dependent and $\mathbf{v}_1 \neq \mathbf{0}$, then some \mathbf{v}_j (with $1 < j \leq p$) is a linear combination of the preceding vectors $\mathbf{v}_1, \dots, \mathbf{v}_{j-1}$.



Remark: The corollary tells us that if a set of vectors is linearly dependent, then at least one of the vectors is redundant, as it can be expressed as a linear combination of the others. This redundancy is what causes the linear dependence. This does not mean that every vector in S is a linear combination of the others, but only that *at least one* vector in a linearly dependent set is a linear combination of others.

Example 8.17 Let

$$\mathbf{u} = \begin{bmatrix} 2 \\ 1 \\ 1 \end{bmatrix}, \quad \mathbf{v} = \begin{bmatrix} 4 \\ 3 \\ 5 \end{bmatrix}, \quad \mathbf{w} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \quad \mathbf{z} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}.$$

- (a) Is any pair of these vectors (e.g. $\{\mathbf{u}, \mathbf{v}\}$) linearly dependent? Explain.
- (b) Does the answer to part (a) tell us that $\{\mathbf{u}, \mathbf{v}, \mathbf{w}, \mathbf{z}\}$ is linearly independent?
- (c) Is $\{\mathbf{u}, \mathbf{v}, \mathbf{w}, \mathbf{z}\}$ linearly dependent? You should be able to answer this question without any computation.

Solution:

- (a) To determine if any pair of these vectors is linearly dependent, we check whether one vector is a scalar multiple of the other.

Checking the pair $\{\mathbf{u}, \mathbf{v}\}$: Assume there exists a scalar k such that:

$$\mathbf{v} = k\mathbf{u}.$$

Compute k using the first component:

$$k = \frac{v_1}{u_1} = \frac{4}{2} = 2.$$

Verify with the second component:

$$v_2 = ku_2 \implies 3 = 2 \times 1 \implies 3 = 2.$$

Since $3 \neq 2$, \mathbf{v} is not a scalar multiple of \mathbf{u} .

Checking the pair $\{\mathbf{u}, \mathbf{w}\}$: Assume there exists a scalar k such that:

$$\mathbf{u} = k\mathbf{w}.$$

Compute k using the first component:

$$k = \frac{u_1}{w_1} = \frac{2}{1} = 2.$$

Verify with the second component:

$$u_2 = kw_2 \implies 1 = 2 \times 1 \implies 1 = 2.$$

Since $1 \neq 2$, \mathbf{u} is not a scalar multiple of \mathbf{w} .

Checking the pair $\{\mathbf{w}, \mathbf{z}\}$: Since $z_1 = 0$ and $w_1 = 1$, assuming $\mathbf{w} = k\mathbf{z}$ leads to:

$$w_1 = kz_1 \implies 1 = k \times 0 \implies 1 = 0,$$

which is a contradiction.

- (b) No, the fact that no pair is linearly dependent does not imply that the entire set is linearly independent.

Linear independence requires that the only solution to:

$$c_1\mathbf{u} + c_2\mathbf{v} + c_3\mathbf{w} + c_4\mathbf{z} = \mathbf{0}$$

is $c_1 = c_2 = c_3 = c_4 = 0$.

- (c) Yes, the set is linearly dependent. In \mathbb{R}^3 , any set of more than three vectors must be linearly dependent because the maximum number of linearly independent vectors in \mathbb{R}^3 is three.

Conclusion:

- (a) No pair of these vectors is linearly dependent.
- (b) The answer to part (a) does not guarantee that the entire set is linearly independent.
- (c) The set $\{\mathbf{u}, \mathbf{v}, \mathbf{w}, \mathbf{z}\}$ is linearly dependent without the need for further computation.



Chapter 9 Matrix Algebra

In the previous chapters, we explored concepts like linear equations, echelon forms, row reductions, linear independence, vector equations, and matrix equations. Now, to take our problem-solving skills to the next level, we need to dive into matrix algebra. Matrix operations allow us to handle complex systems more efficiently, building on the foundations we've established. The tools and techniques in this chapter will streamline the way we work with multiple matrices and lead us toward deeper insights, including key ideas like the Invertible Matrix Theorem.

9.1 Matrix Arithmetic

Matrix arithmetic like adding, subtracting, scaling, and multiplying let you combine and manipulate matrices in simple ways. These operations are the building blocks for more complex matrix manipulations, like solving systems of equations and finding inverses.

Addition and Scalar Multiplication

We begin with the definition of matrix addition.

Definition 9.1 (Matrix Addition)

Let A and B be $m \times n$ matrices. The **sum** of A and B , denoted $A + B$, is the $m \times n$ matrix whose entries are obtained by adding the corresponding entries of A and B .

Given matrices

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \vdots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}, \quad B = \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1n} \\ b_{21} & b_{22} & \cdots & b_{2n} \\ \vdots & \vdots & \vdots & \vdots \\ b_{m1} & b_{m2} & \cdots & b_{mn} \end{bmatrix}$$

both of the same dimension $m \times n$, the sum $A + B$ is thus defined as

$$A + B = \begin{bmatrix} a_{11} + b_{11} & a_{12} + b_{12} & \cdots & a_{1n} + b_{1n} \\ a_{21} + b_{21} & a_{22} + b_{22} & \cdots & a_{2n} + b_{2n} \\ \vdots & \vdots & \vdots & \vdots \\ a_{m1} + b_{m1} & a_{m2} + b_{m2} & \cdots & a_{mn} + b_{mn} \end{bmatrix}$$



Remark: Similar considerations apply for subtraction of matrices, although not mentioned here explicitly.

Next is the definition of scalar-matrix multiplication.

Definition 9.2 (Scalar-Matrix Multiplication)

Let A be an $m \times n$ matrix and c be a scalar. The **product** of c and A , denoted cA , is the $m \times n$ matrix whose entries are obtained by multiplying each entry of A by c , and is thus defined as

$$cA = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \vdots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix} = \begin{bmatrix} ca_{11} & ca_{12} & \cdots & ca_{1n} \\ ca_{21} & ca_{22} & \cdots & ca_{2n} \\ \vdots & \vdots & \vdots & \vdots \\ ca_{m1} & ca_{m2} & \cdots & ca_{mn} \end{bmatrix}$$



Example 9.1 Given matrices

$$A = \begin{bmatrix} -5 & 2 & 0 \\ 7 & -3 & 4 \\ -1 & 3 & 2 \end{bmatrix}, \quad B = \begin{bmatrix} 0 & -1 & 8 \\ 6 & -14 & 2 \\ 9 & 5 & 1 \end{bmatrix}$$

find the sum $A + B$ and the product $3A$.

Solution: The sum $A + B$ is obtained by adding the corresponding entries of A and B :

$$\begin{aligned} A + B &= \begin{bmatrix} -5 & 2 & 0 \\ 7 & -3 & 4 \\ -1 & 3 & 2 \end{bmatrix} + \begin{bmatrix} 0 & -1 & 8 \\ 6 & -14 & 2 \\ 9 & 5 & 1 \end{bmatrix} \\ &= \begin{bmatrix} (-5) + (0) & (2) + (-1) & (0) + (8) \\ (7) + (6) & (-3) + (-14) & (4) + (2) \\ (-1) + (9) & (3) + (5) & (2) + (1) \end{bmatrix} \\ &= \begin{bmatrix} -5 & 1 & 8 \\ 13 & -17 & 6 \\ 8 & 8 & 3 \end{bmatrix} \end{aligned}$$

The product $3A$ is obtained by multiplying each entry of A by 3:

$$\begin{aligned} 3A &= 3 \begin{bmatrix} -5 & 2 & 0 \\ 7 & -3 & 4 \\ -1 & 3 & 2 \end{bmatrix} \\ &= \begin{bmatrix} 3(-5) & 3(2) & 3(0) \\ 3(7) & 3(-3) & 3(4) \\ 3(-1) & 3(3) & 3(2) \end{bmatrix} \\ &= \begin{bmatrix} -15 & 6 & 0 \\ 21 & -9 & 12 \\ -3 & 9 & 6 \end{bmatrix} \end{aligned}$$



Example 9.2 Given A and B below, find $3A - 2B$.

$$A = \begin{bmatrix} 1 & -2 & 5 \\ 0 & -3 & 9 \\ 4 & -6 & 7 \end{bmatrix}, B = \begin{bmatrix} 5 & 0 & -11 \\ 3 & -5 & 1 \\ -1 & -9 & 0 \end{bmatrix}$$

Solution: We compute:

$$\begin{aligned} 3A - 2B &= \begin{bmatrix} 3 & -6 & 15 \\ 0 & -9 & 27 \\ 12 & -18 & 21 \end{bmatrix} - \begin{bmatrix} 10 & 0 & -22 \\ 6 & -10 & 2 \\ -2 & -18 & 0 \end{bmatrix} \\ &= \begin{bmatrix} -7 & -6 & 37 \\ -6 & 1 & 25 \\ 14 & 0 & 21 \end{bmatrix} \end{aligned}$$



Before moving on to matrix multiplication, we need to state some basic algebraic properties of matrix addition and scalar multiplication.

Theorem 9.1

Let A, B, C be matrices of the same size and let α, β be scalars. Then

- (a) $A + B = B + A$
- (b) $(A + B) + C = A + (B + C)$
- (c) $A + 0 = A$
- (d) $\alpha(A + B) = \alpha A + \alpha B$
- (e) $(\alpha + \beta)A = \alpha A + \beta A$
- (f) $\alpha(\beta A) = (\alpha\beta)A$



Matrix Multiplication

Matrix multiplication is a bit more complex than addition and scalar multiplication. The product of two matrices is defined only when the number of columns in the first matrix is equal to the number of rows in the second matrix. The product of two matrices A and B is a new matrix C whose entries are determined by the dot product of the rows of A and the columns of B . Note that \mathbf{x} is a column vector.

Let B be an $n \times p$ matrix and A be an $m \times n$ matrix. If $\mathbf{x} \in \mathbb{R}^p$, then multiplying B by \mathbf{x} produces a new vector $B\mathbf{x}$ in \mathbb{R}^n . Once we have this result, we can further multiply it by A , giving us $A(B\mathbf{x})$, which is a vector in \mathbb{R}^m .

Thus, for any vector \mathbf{x} in \mathbb{R}^p , this process produces a corresponding vector in \mathbb{R}^m . This two-step operation—first multiplying by B and then by A —is referred to as the composition of A and B , and is usually written as AB . Therefore, we have:

$$(AB)\mathbf{x} = A(B\mathbf{x})$$

To compute the matrix resulting from this composition, we multiply A by B directly, following the rules of matrix multiplication. The result is a matrix $C = AB$, where each entry of C is determined by the interactions between the rows of A and the columns of B .

In essence, C represents the matrix that captures the combined effect of applying both B and A to any vector $\mathbf{x} \in \mathbb{R}^p$, without needing to break it down into intermediate steps.

Definition 9.3

For $A \in \mathbb{R}^{m \times n}$ and $B \in \mathbb{R}^{n \times p}$, with $B = [\mathbf{b}_1 \ \mathbf{b}_2 \ \dots \ \mathbf{b}_p]$, we define the product AB by the formula

$$AB = [Ab_1 \ Ab_2 \ \dots \ Ab_p]$$


The product AB is defined only when the number of columns of A equals the number of rows of B . The following diagram is useful for remembering this:

$$(m \times n) \cdot (n \times p) \rightarrow m \times p$$

The diagram shows that the number of columns in the first matrix must equal the number of rows in the second matrix (the **blue arrow**). The result is a new matrix with the number of rows from the first matrix and the number of columns from the second matrix (the **red arrow**).

Example 9.3 For A and B below compute AB and BA .

$$A = \begin{bmatrix} 1 & 2 & -2 \\ 1 & 1 & -3 \end{bmatrix}, \quad B = \begin{bmatrix} -4 & 2 & 4 & -4 \\ -1 & -5 & -3 & 3 \\ -4 & -4 & -3 & -1 \end{bmatrix}$$

Solution: First $AB = [Ab_1 \ Ab_2 \ Ab_3 \ Ab_4]$:

$$\begin{aligned} AB &= \begin{bmatrix} 1 & 2 & -2 \\ 1 & 1 & -3 \end{bmatrix} \begin{bmatrix} -4 & 2 & 4 & -4 \\ -1 & -5 & -3 & 3 \\ -4 & -4 & -3 & -1 \end{bmatrix} \\ &= \begin{bmatrix} 2 & \text{Row 1 of } A \text{ and Column 1 of } B \\ 7 & \text{Row 2 of } A \text{ and Column 1 of } B \end{bmatrix} \\ &= \begin{bmatrix} 2 & 0 & \text{Row 1 of } A \text{ and Column 2 of } B \\ 7 & 9 & \text{Row 2 of } A \text{ and Column 2 of } B \end{bmatrix} \\ &= \begin{bmatrix} 2 & 0 & 4 & \text{Row 1 of } A \text{ and Column 3 of } B \\ 7 & 9 & 10 & \text{Row 2 of } A \text{ and Column 3 of } B \end{bmatrix} \\ &= \begin{bmatrix} 2 & 0 & 4 & 4 & \text{Row 1 of } A \text{ and Column 4 of } B \\ 7 & 9 & 10 & 2 & \text{Row 2 of } A \text{ and Column 4 of } B \end{bmatrix} \\ &= \begin{bmatrix} 2 & 0 & 4 & 4 \\ 7 & 9 & 10 & 2 \end{bmatrix} \end{aligned}$$

On the other hand, BA is not defined! B has 4 columns and A has 2 rows. Thus, the number of columns in B is not equal to the number of rows in A . 

Example 9.4 Example Compute the matrix AB , where

$$A = \begin{bmatrix} 1 & 2 \\ 2 & 3 \\ 3 & 4 \end{bmatrix}, \quad \text{and } B = \begin{bmatrix} 2 & -1 \\ 3 & 1 \end{bmatrix}$$

Solution: By the definition of multiplication given above,

$$\begin{aligned} AB &= \left[A \begin{bmatrix} 2 \\ 3 \end{bmatrix} A \begin{bmatrix} -1 \\ 1 \end{bmatrix} \right] \\ &= \begin{bmatrix} 1 & 2 \\ 2 & 3 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} 2 \\ 3 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 2 & 3 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} -1 \\ 1 \end{bmatrix} \\ &= \begin{bmatrix} 1 \cdot 2 + 2 \cdot 3 \\ 2 \cdot 2 + 3 \cdot 3 \\ 3 \cdot 2 + 4 \cdot 3 \end{bmatrix} \begin{bmatrix} 1 \cdot (-1) + 2 \cdot 1 \\ 2 \cdot (-1) + 3 \cdot 1 \\ 3 \cdot (-1) + 4 \cdot 1 \end{bmatrix} \\ &= \begin{bmatrix} 8 & 1 \\ 13 & 1 \\ 18 & 1 \end{bmatrix} \end{aligned}$$



Example 9.5 If A is 3×5 and B is 5×2 , what are the sizes of AB and BA (assuming they are defined)?

Solution:

- The product AB is defined since the number of columns of A matches the number of rows of B (5). The resulting matrix AB is a 3×2 matrix.
- The product BA is not defined, since the number of columns of B (2) does not match the number of rows of A (3).



The next example illustrate that even if both AB and BA are defined, they are not necessarily equal.

Example 9.6 For $A = \begin{bmatrix} -4 & 4 & 3 \\ 3 & -3 & -1 \\ -2 & -1 & 1 \end{bmatrix}$ and $B = \begin{bmatrix} -1 & -1 & 0 \\ -3 & 0 & -2 \\ -2 & 1 & -2 \end{bmatrix}$ compute AB and BA .

Solution: First AB :

$$\begin{aligned} AB &= \begin{bmatrix} -4 & 4 & 3 \\ 3 & -3 & -1 \\ -2 & -1 & 1 \end{bmatrix} \begin{bmatrix} -1 & -1 & 0 \\ -3 & 0 & -2 \\ -2 & 1 & -2 \end{bmatrix} \\ &= \begin{bmatrix} -14 \\ 8 \\ 3 \end{bmatrix} \\ &= \begin{bmatrix} -14 & 7 \\ 8 & -4 \\ 3 & 3 \end{bmatrix} \\ &= \begin{bmatrix} -14 & 7 & -14 \\ 8 & -4 & 8 \\ 3 & 3 & 0 \end{bmatrix} \end{aligned}$$

Next BA :

$$\begin{aligned} BA &= \begin{bmatrix} -1 & -1 & 0 \\ -3 & 0 & -2 \\ -2 & 1 & -2 \end{bmatrix} \begin{bmatrix} -4 & 4 & 3 \\ 3 & -3 & -1 \\ -2 & -1 & 1 \end{bmatrix} \\ &= \begin{bmatrix} 1 \\ 16 \\ 15 \end{bmatrix} \\ &= \begin{bmatrix} 1 & -1 \\ 16 & -10 \\ 15 & -9 \end{bmatrix} \\ &= \begin{bmatrix} 1 & -1 & -2 \\ 16 & -10 & -11 \\ 15 & -9 & -9 \end{bmatrix} \end{aligned}$$

We see that $AB \neq BA$. ◀

In regular arithmetic the multiplicative identity is 1. In matrix algebra, the multiplicative identity is the identity matrix, denoted by I . The identity matrix is a *square* matrix with 1s on the diagonal and 0s elsewhere. The size of the identity matrix is determined by the context, and is usually clear from the context. For example, I_2 is a 2×2 identity matrix, and I_3 is a 3×3 identity matrix and in general $I_n \in \mathbb{R}^{n \times n}$ is an $n \times n$ identity matrix:

$$I_n = \begin{bmatrix} 1 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ \vdots & \vdots & \vdots & \cdots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{bmatrix}$$

We can now state the following theorem.

Theorem 9.2

Let A, B, C be matrices, of appropriate dimensions, and let α be a scalar. Then

- (a) $A(BC) = (AB)C$
- (b) $A(B + C) = AB + AC$
- (c) $(B + C)A = BA + CA$
- (d) $\alpha(AB) = (\alpha A)B = A(\alpha B)$
- (e) $I_n A = A I_n = A$



We conclude this section by looking at the k th power of a matrix.

Definition 9.4

Let A be a square matrix, i.e. $A \in \mathbb{R}^{n \times n}$. The k th power of A , denoted A^k , is defined as the product of A with itself k times. That is,

$$A^k = \underbrace{AAA \cdots A}_{k \text{ times}}$$

where A appears k times on the right-hand side.



Example 9.7 Compute A^3 if

$$A = \begin{bmatrix} -2 & 3 \\ 1 & 0 \end{bmatrix}$$

Solution:

Compute A^2 :

$$A^2 = \begin{bmatrix} -2 & 3 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} -2 & 3 \\ 1 & 0 \end{bmatrix} = \begin{bmatrix} 7 & -6 \\ -2 & 3 \end{bmatrix}$$

And then A^3 :

$$\begin{aligned} A^3 &= A^2 A = \begin{bmatrix} 7 & -6 \\ -2 & 3 \end{bmatrix} \begin{bmatrix} -2 & 3 \\ 1 & 0 \end{bmatrix} \\ &= \begin{bmatrix} -20 & 21 \\ 7 & -6 \end{bmatrix} \end{aligned}$$

We could also do:

$$A^3 = A A^2 = \begin{bmatrix} -2 & 3 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} 7 & -6 \\ -2 & 3 \end{bmatrix} = \begin{bmatrix} -20 & 21 \\ 7 & -6 \end{bmatrix}$$



9.2 Matrix Transpose

We begin with the definition of the transpose of a matrix.

Definition 9.5

Given a matrix $A \in \mathbb{R}^{m \times n}$, the transpose of A is the matrix A^T whose i th column is the i th row of A .



If A is $m \times n$ then A^T is $n \times m$. For example, if

$$A = \begin{bmatrix} 0 & -1 & 8 & -7 & -4 \\ -4 & 6 & -10 & -9 & 6 \\ 9 & 5 & -2 & -3 & 5 \\ -8 & 8 & 4 & 7 & 7 \end{bmatrix}$$

then

$$A^T = \begin{bmatrix} 0 & -4 & 9 & -8 \\ -1 & 6 & 5 & 8 \\ 8 & -10 & -2 & 4 \\ -7 & -9 & -3 & 7 \\ -4 & 6 & 5 & 7 \end{bmatrix}$$

Example 9.8 Compute $(AB)^T$ and $B^T A^T$ if

$$A = \begin{bmatrix} -2 & 1 & 0 \\ 3 & -1 & -3 \end{bmatrix}, \quad B = \begin{bmatrix} -2 & 1 & 2 \\ -1 & -2 & 0 \\ 0 & 0 & -1 \end{bmatrix}.$$

Solution: First, compute $(AB)^T$:

$$\begin{aligned} AB &= \begin{bmatrix} -2 & 1 & 0 \\ 3 & -1 & -3 \end{bmatrix} \begin{bmatrix} -2 & 1 & 2 \\ -1 & -2 & 0 \\ 0 & 0 & -1 \end{bmatrix} \\ &= \begin{bmatrix} 3 & -4 & -4 \\ -5 & 5 & 9 \end{bmatrix} \end{aligned}$$

and then $(AB)^T$:

$$\begin{aligned} (AB)^T &= \begin{bmatrix} 3 & -4 & -4 \\ -5 & 5 & 9 \end{bmatrix}^T \\ &= \begin{bmatrix} 3 & -5 \\ -4 & 5 \\ -4 & 9 \end{bmatrix} \end{aligned}$$

Next, compute $B^T A^T$:

$$\begin{aligned} B^T A^T &= \begin{bmatrix} -2 & -1 & 0 \\ 1 & -2 & 0 \\ 2 & 0 & -1 \end{bmatrix} \begin{bmatrix} -2 & 3 \\ 1 & -1 \\ 0 & -3 \end{bmatrix} \\ &= \begin{bmatrix} 3 & -5 \\ -4 & 5 \\ -4 & 9 \end{bmatrix} \end{aligned}$$

We see that $(AB)^T = B^T A^T$.

The following theorem summarises the properties of the transpose of a matrix.

Theorem 9.3

Let A and B be matrices of appropriate dimensions and let α be a scalar. Then

- (a) $(A^T)^T = A$
- (b) $(A + B)^T = A^T + B^T$
- (c) $(\alpha A)^T = \alpha A^T$
- (d) $(AB)^T = B^T A^T$



A consequence of property (4) is that

$$(A_1 A_2 \dots A_k)^T = A_k^T A_{k-1}^T \cdots A_2^T A_1^T$$

and as a special case

$$(A^k)^T = (A^T)^k$$

9.3 Invertible Matrices

The inverse of a square matrix $A \in \mathbb{R}^{n \times n}$ extends the concept of the reciprocal for a nonzero real number $a \in \mathbb{R}$. More precisely, the inverse of a non-zero number $a \in \mathbb{R}$ is the unique number $c \in \mathbb{R}$ such that $ac = ca = 1$. The inverse of $a \neq 0$, typically written as $a^{-1} = \frac{1}{a}$, enables solving the equation $ax = b$:

$$ax = b \Rightarrow a^{-1}ax = a^{-1}b \Rightarrow x = a^{-1}b.$$

This concept extends to square matrices, where the inverse of a matrix $A \in \mathbb{R}^{n \times n}$ is a matrix $C \in \mathbb{R}^{n \times n}$ such that $AC = CA = I_n$, where I_n is the identity matrix of size $n \times n$. The inverse of a matrix is denoted by A^{-1} . We can now define the invertible matrix.

Definition 9.6

A square matrix $A \in \mathbb{R}^{n \times n}$ is invertible (or **nonsingular**) if there exists a matrix $C \in \mathbb{R}^{n \times n}$ such that $AC = CA = I_n$. The matrix C is called the inverse of A and is denoted by A^{-1} . Thus, $A^{-1}A = AA^{-1} = I_n$.



Example 9.9 Given A and C below, show that C is the inverse of A .

$$A = \begin{bmatrix} 1 & -3 & 0 \\ -1 & 2 & -2 \\ -2 & 6 & 1 \end{bmatrix}, \quad C = \begin{bmatrix} -14 & -3 & -6 \\ -5 & -1 & -2 \\ 2 & 0 & 1 \end{bmatrix}$$

Solution: We need to show that $AC = CA = I_3$. First, compute AC :

$$\begin{aligned} AC &= \begin{bmatrix} 1 & -3 & 0 \\ -1 & 2 & -2 \\ -2 & 6 & 1 \end{bmatrix} \begin{bmatrix} -14 & -3 & -6 \\ -5 & -1 & -2 \\ 2 & 0 & 1 \end{bmatrix} \\ &= \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} = I_3 \end{aligned}$$

Next, compute CA :

$$\begin{aligned} CA &= \begin{bmatrix} -14 & -3 & -6 \\ -5 & -1 & -2 \\ 2 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & -3 & 0 \\ -1 & 2 & -2 \\ -2 & 6 & 1 \end{bmatrix} \\ &= \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} = I_3 \end{aligned}$$

We see that $C = A^{-1}$ is the inverse of A . ◀

The following theorem summarizes the relationship between the matrix inverse and matrix multiplication and matrix transpose.

Theorem 9.4

Let A and B be invertible $n \times n$ matrices. Then:

(a) A^{-1} is invertible, with

$$(A^{-1})^{-1} = A$$

(b) The product AB is invertible, with

$$(AB)^{-1} = B^{-1}A^{-1}$$

(c) The transpose of A is also invertible, i.e. A^T is invertible, with

$$(A^T)^{-1} = (A^{-1})^T$$



We will now consider how we can find the inverse of a matrix.

Finding the Determinant and Inverse of a 2×2 Matrix

The inverse of a matrix is not always easy to find. However, for a 2×2 matrix, we can use a formula to find the inverse.

Theorem 9.5

Let $A = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$. If $ad - bc \neq 0$, then the inverse of A is given by

$$A^{-1} = \frac{1}{ad - bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$$



Remark: The quantity $ad - bc$ is called the determinant of A , and we write

$$\det(A) = |A| = ad - bc$$

The determinant of a 2×2 matrix is a scalar quantity that provides information about the matrix. If the determinant is zero, then the matrix is not invertible. If the determinant is non-zero, then the matrix is invertible. These formulae are only valid 2×2 matrices.

Example 9.10 We let

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

and compute A^{-1} . By the theorem, we have

$$\begin{aligned} A^{-1} &= \frac{1}{1 \cdot 4 - 2 \cdot 3} \begin{bmatrix} 4 & -2 \\ -3 & 1 \end{bmatrix} \\ &= -\frac{1}{2} \begin{bmatrix} 4 & -2 \\ -3 & 1 \end{bmatrix} \\ &= \begin{bmatrix} -2 & 1 \\ \frac{3}{2} & -\frac{1}{2} \end{bmatrix} \end{aligned}$$

Now we confirm that $AA^{-1} = I_2$:

$$\begin{aligned} AA^{-1} &= \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} -2 & 1 \\ \frac{3}{2} & -\frac{1}{2} \end{bmatrix} \\ &= \begin{bmatrix} 1(-2) + 2\left(\frac{3}{2}\right) & 1(1) + 2\left(-\frac{1}{2}\right) \\ 3(-2) + 4\left(\frac{3}{2}\right) & 3(1) + 4\left(-\frac{1}{2}\right) \end{bmatrix} \\ &= \begin{bmatrix} -2 + 3 & 1 - 1 \\ -6 + 6 & 3 - 2 \end{bmatrix} \\ &= \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \end{aligned}$$

Example 9.11 Find the inverse of $A = \begin{bmatrix} 1 & 3 \\ -1 & -2 \end{bmatrix}$ if it exists.

Solution: We first compute the determinant of A :

$$\det(A) = 1(-2) - 3(-1) = (-2) + 3 = 1$$

Since the determinant is non-zero, the inverse of A exists. By the formula, we have

$$A^{-1} = \frac{1}{1} \begin{bmatrix} -2 & -3 \\ 1 & 1 \end{bmatrix}$$

We can confirm that $AA^{-1} = I_2$:

$$AA^{-1} = \begin{bmatrix} 1 & 3 \\ -1 & -2 \end{bmatrix} \begin{bmatrix} -2 & -3 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = I_2$$



For larger matrices, we need to use other methods to find the inverse.

Finding the Inverse of a $n \times n$ Matrix

An **elementary matrix** is a matrix that is obtained by performing a single elementary row operation (replacement, swap, or scaling) on an identity matrix. For example, the matrix

$$E_1 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ -2 & 0 & 1 \end{bmatrix}$$

is an elementary matrix since it is obtained from I_3 via the single elementary row operation $r_3 \mapsto r_3 - 2r_1$.

Example 9.12

Let A be a general 3×3 matrix

$$A = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}$$

and we let E_1, E_2 , and E_3 be the elementary matrices

$$E_1 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ -2 & 0 & 1 \end{bmatrix}, \quad E_2 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix}, \quad \text{and} \quad E_3 = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 3 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Notice, then, that E_1 corresponds to $r_3 \mapsto r_3 - 2r_1$, E_2 corresponds to $r_2 \longleftrightarrow r_3$, and E_3 corresponds to $r_2 \mapsto 3r_2$. We also have the following products:

$$E_1A = \begin{bmatrix} a & b & c \\ d & e & f \\ g - 2a & h - 2b & i - 2c \end{bmatrix}, \quad E_2A = \begin{bmatrix} a & b & c \\ g & h & i \\ d & e & f \end{bmatrix}, \quad \text{and} \quad E_3A = \begin{bmatrix} a & b & c \\ 3d & 3e & 3f \\ g & h & i \end{bmatrix}.$$

Notice that E_1A is the matrix obtained by performing the row operation $r_3 \mapsto r_3 - 2r_1$ on A . In general, multiplying A on the left by an elementary matrix is the same as performing the corresponding row operation on A . We can also represent a sequence of row operations by multiplication of several elementary matrices. For example,

$$E_2E_1A = \begin{bmatrix} a & b & c \\ g - 2a & h - 2b & i - 2c \\ d & e & f \end{bmatrix}$$

corresponds to performing $r_3 \mapsto r_3 - 2r_1$ followed by $r_2 \longleftrightarrow r_3$ on the matrix A . As we have observed before, row operations are reversible. It follows that elementary matrices are also invertible. This leads to the following theorem.

Theorem 9.6

An $n \times n$ matrix A is invertible if and only if A is row equivalent to I_n . In this case, any sequence of elementary row operations that reduces A to I_n also transforms I_n into A^{-1} .



This theorem leads to a nice algorithm for finding the inverse of an $n \times n$ matrix, assuming such an inverse exists.

Algorithm for Finding the Inverse of an $n \times n$ Matrix

Let A be an $n \times n$ matrix. To find the inverse of A , follow these steps:

1. Form the augmented matrix $[A|I_n]$.
2. Perform row operations on $[A|I_n]$ to reduce A to I_n .
3. The matrix on the right side of the augmented matrix is A^{-1} .

If A is not invertible, then the algorithm will not be able to reduce A to I_n .

Example 9.13 Find the inverse of $A = \begin{bmatrix} 1 & 0 & 3 \\ 1 & 1 & 0 \\ -2 & 0 & -7 \end{bmatrix}$ if it exists.

Solution: Solution. Form the augmented matrix $[A \ I_3]$ and row reduce:

$$\begin{array}{c} \left[\begin{array}{ccc|ccc} 1 & 0 & 3 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 & 1 & 0 \\ -2 & 0 & -7 & 0 & 0 & 1 \end{array} \right] \quad r_2 \leftrightarrow -r_1 + r_2 \quad \left[\begin{array}{ccc|ccc} 1 & 0 & 3 & 1 & 0 & 0 \\ 0 & 1 & -3 & -1 & 1 & 0 \\ -2 & 0 & -7 & 0 & 0 & 1 \end{array} \right] \\ \quad r_3 \leftrightarrow 2r_1 + r_3 \quad \left[\begin{array}{ccc|ccc} 1 & 0 & 3 & 1 & 0 & 0 \\ 0 & 1 & -3 & -1 & 1 & 0 \\ 0 & 0 & -1 & 2 & 0 & 1 \end{array} \right] \\ \\ \quad r_3 \leftrightarrow -r_3 \quad \left[\begin{array}{ccc|ccc} 1 & 0 & 3 & 1 & 0 & 0 \\ 0 & 1 & -3 & -1 & 1 & 0 \\ 0 & 0 & 1 & -2 & 0 & -1 \end{array} \right] \\ \\ \quad r_2 \leftrightarrow 3r_3 + r_2 \quad \left[\begin{array}{ccc|ccc} 1 & 0 & 0 & 7 & 0 & 3 \\ 0 & 1 & 0 & -7 & 1 & -3 \\ 0 & 0 & 1 & -2 & 0 & -1 \end{array} \right] \\ \quad r_1 \leftrightarrow -3r_3 + r_1 \quad \left[\begin{array}{ccc|ccc} 1 & 0 & 0 & 7 & 0 & 3 \\ 0 & 1 & 0 & -7 & 1 & -3 \\ 0 & 0 & 1 & -2 & 0 & -1 \end{array} \right] \end{array}$$

Therefore, $\text{rref } A = I_3$, confirming that A is invertible. The inverse is

$$A^{-1} = \begin{bmatrix} 7 & 0 & 3 \\ -7 & 1 & -3 \\ -2 & 0 & -1 \end{bmatrix}$$

Verify:

$$AA^{-1} = \begin{bmatrix} 1 & 0 & 3 \\ 1 & 1 & 0 \\ -2 & 0 & -7 \end{bmatrix} \begin{bmatrix} 7 & 0 & 3 \\ -7 & 1 & -3 \\ -2 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



Example 9.14 Find the inverse of $A = \begin{bmatrix} 1 & 0 & 1 \\ 1 & 1 & -2 \\ -2 & 0 & -2 \end{bmatrix}$ if it exists.

Solution: Form the augmented matrix $[A \ I_3]$ and row reduce:

$$\begin{array}{c} \left[\begin{array}{ccc|ccc} 1 & 0 & 1 & 1 & 0 & 0 \\ 1 & 1 & -2 & 0 & 1 & 0 \\ -2 & 0 & -2 & 0 & 0 & 1 \end{array} \right] \quad r_2 \leftrightarrow -r_1 + r_2 \quad \left[\begin{array}{ccc|ccc} 1 & 0 & 1 & 1 & 0 & 0 \\ 0 & 1 & -3 & -1 & 1 & 0 \\ -2 & 0 & -2 & 0 & 0 & 1 \end{array} \right] \\ \quad r_3 \leftrightarrow 2r_1 + r_3 \quad \left[\begin{array}{ccc|ccc} 1 & 0 & 1 & 1 & 0 & 0 \\ 0 & 1 & -3 & -1 & 1 & 0 \\ 0 & 0 & 0 & 2 & 0 & 1 \end{array} \right] \end{array}$$

We need not go further since the $\text{rref}(A)$ is not I_3 . Therefore, A is not invertible.



Example 9.15 Find the inverse of the matrix

$$A = \begin{bmatrix} 1 & 2 & 1 \\ 4 & 5 & 3 \\ 0 & 0 & 2 \end{bmatrix}$$

if it exists.

Solution: (We omit the row reduction steps for brevity.)

$$\begin{aligned}
 \left[\begin{array}{cccccc} 1 & 2 & 1 & 1 & 0 & 0 \\ 4 & 5 & 3 & 0 & 1 & 0 \\ 0 & 0 & 2 & 0 & 0 & 1 \end{array} \right] &\sim \left[\begin{array}{cccccc} 1 & 2 & 1 & 1 & 0 & 0 \\ 0 & -3 & -1 & -4 & 1 & 0 \\ 0 & 0 & 2 & 0 & 0 & 1 \end{array} \right] \\
 &\sim \left[\begin{array}{cccccc} 1 & 2 & 1 & 1 & 0 & 0 \\ 0 & -3 & -1 & -4 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 & \frac{1}{2} \end{array} \right] \\
 &\sim \left[\begin{array}{cccccc} 1 & 2 & 0 & 1 & 0 & -\frac{1}{2} \\ 0 & -3 & 0 & -4 & 1 & \frac{1}{2} \\ 0 & 0 & 1 & 0 & 0 & \frac{1}{2} \end{array} \right] \\
 &\sim \left[\begin{array}{cccccc} 1 & 2 & 0 & 1 & 0 & -\frac{1}{2} \\ 0 & 1 & 0 & \frac{4}{3} & -\frac{1}{3} & -\frac{1}{6} \\ 0 & 0 & 1 & 0 & 0 & \frac{1}{2} \end{array} \right] \\
 &\sim \left[\begin{array}{cccccc} 1 & 0 & 0 & -\frac{5}{3} & \frac{2}{3} & -\frac{1}{6} \\ 0 & 1 & 0 & \frac{4}{3} & -\frac{1}{3} & -\frac{1}{6} \\ 0 & 0 & 1 & 0 & 0 & \frac{1}{2} \end{array} \right].
 \end{aligned}$$

Hence

$$A^{-1} = \left[\begin{array}{ccc} -\frac{5}{3} & \frac{2}{3} & -\frac{1}{6} \\ \frac{4}{3} & -\frac{1}{3} & -\frac{1}{6} \\ 0 & 0 & \frac{1}{2} \end{array} \right].$$



With the method for finding matrix inverses established, we now have a powerful tool for approaching linear systems. We will see how matrix inverses can be used to solve systems of equations efficiently, offering a straightforward way to find solutions when a matrix is invertible.

Solving Matrix Equations with Inverses

Just like we can solve an equation by using the reciprocal of a non-zero number as we saw previously, we can do something similar with matrices. When a matrix A is invertible, we can use its inverse to solve matrix equations involving A . This brings us to the theorem below, which shows how we can find solutions for equations like $A\mathbf{x} = \mathbf{b}$ when A has an inverse.

Theorem 9.7

Let A be an invertible matrix. Then the equation $A\mathbf{x} = \mathbf{b}$ has a unique solution given by $\mathbf{x} = A^{-1}\mathbf{b}$.



Example 9.16 We can use the inverse from [Example 9.10](#) to solve the linear system

$$x_1 + 2x_2 = 5$$

$$3x_1 + 4x_2 = 6.$$

We think of this in matrix terms as $A\mathbf{x} = \mathbf{b}$, where

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \quad \text{and} \quad \mathbf{b} = \begin{bmatrix} 5 \\ 6 \end{bmatrix}$$

The theorem tells us that $\mathbf{x} = A^{-1}\mathbf{b}$ is a solution, i.e. that

$$\begin{aligned} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} &= \begin{bmatrix} -2 & 1 \\ \frac{3}{2} & -\frac{1}{2} \end{bmatrix} \begin{bmatrix} 5 \\ 6 \end{bmatrix} \\ &= \begin{bmatrix} -2(5) + 1(6) \\ \frac{3}{2}(5) - \frac{1}{2}(6) \end{bmatrix} \\ &= \begin{bmatrix} -4 \\ \frac{9}{2} \end{bmatrix} \end{aligned}$$

Example 9.17 Use the result from [Example 9.9](#) to solve the linear system $A\mathbf{x} = \mathbf{b}$ where

$$A = \begin{bmatrix} 1 & -3 & 0 \\ -1 & 2 & -2 \\ -2 & 6 & 1 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} 1 \\ -3 \\ -1 \end{bmatrix}$$

Solution: We showed in [Example 9.9](#) that

$$A^{-1} = \begin{bmatrix} -14 & -3 & -6 \\ -5 & -1 & -2 \\ 2 & 0 & 1 \end{bmatrix}$$

Therefore, the unique solution to the linear system $A\mathbf{x} = \mathbf{b}$ is

$$A^{-1}\mathbf{b} = \begin{bmatrix} -14 & -3 & -6 \\ -5 & -1 & -2 \\ 2 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ -3 \\ -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix}$$

Verify:

$$\begin{bmatrix} 1 & -3 & 0 \\ -1 & 2 & -2 \\ -2 & 6 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ -3 \\ -1 \end{bmatrix}$$



We conclude this chapter with one of the most important theorems in linear algebra, The Invertible Matrix Theorem.

9.4 The Invertible Matrix Theorem

The Invertible Matrix Theorem provides a comprehensive list of equivalent statements for a square matrix to be invertible and it enables us to determine the truth value of one statement by checking the truth value of another.

Theorem 9.8

Let A be an $n \times n$ matrix. The following statements are equivalent:

- (a) A is invertible.
- (b) A is row equivalent to I_n .
- (c) A has n pivot positions (i.e. one for each row and column).
- (d) The equation $A\mathbf{x} = \mathbf{0}$ has only the trivial solution.
- (e) The columns of A are linearly independent.
- (f) The equation $A\mathbf{x} = \mathbf{b}$ has a unique solution for each $\mathbf{b} \in \mathbb{R}^n$.
- (g) The columns of A span \mathbb{R}^n .
- (h) $\det(A) \neq 0$.
- (i) There is an $n \times n$ matrix C such that $CA = I$.
- (j) There is an $n \times n$ matrix D such that $AD = I$.
- (k) A^T is invertible.



Note how the theorem connects the properties of a matrix to its invertibility, thus summarising the main elements of our discussion on linear systems, independence, vectors, and matrices.

Example 9.18 Decide if the following matrix is invertible:

$$A = \begin{bmatrix} 2 & 2 & 2 \\ 1 & 3 & 1 \\ 4 & 4 & 6 \end{bmatrix}$$

Solution: Performing row operations, we see

$$\begin{aligned} A &\sim \begin{bmatrix} 1 & 3 & 1 \\ 2 & 2 & 2 \\ 4 & 4 & 6 \end{bmatrix} \\ &\sim \begin{bmatrix} 1 & 3 & 1 \\ 0 & -4 & 0 \\ 0 & -8 & 2 \end{bmatrix} \\ &\sim \begin{bmatrix} 1 & 3 & 1 \\ 0 & -4 & 0 \\ 0 & 0 & 2 \end{bmatrix} \end{aligned}$$

which has 3 pivots, so by (c) we have that A is invertible. ◀

This concludes our discussion on matrices.

Chapter 10 Differential Calculus

In both practical experience and modern engineering, we frequently encounter situations where one quantity changes in response to another. A car's fuel consumption reflects how distance depends on the amount of fuel used. A sprinter's velocity describes how position varies with time. In agriculture, yield changes with the amount of fertiliser, and in economics, demand responds to price. The study of such relationships lies at the heart of *differential calculus*, the mathematical discipline devoted to understanding how quantities vary together.

The concept that captures this variation is the *rate of change*.

- The change of distance with respect to time is *velocity*.
- The change of profit with respect to price is *marginal profit*.
- The change of loss with respect to model parameters is the *gradient*.

When a rate of change remains constant, the relationship between the quantities is linear, and the graph forms a straight line. When the rate of change varies, the graph becomes curved, and the slope depends on the specific point of measurement. This naturally leads to the idea of a *tangent line*, which locally approximates a curve at a single point.

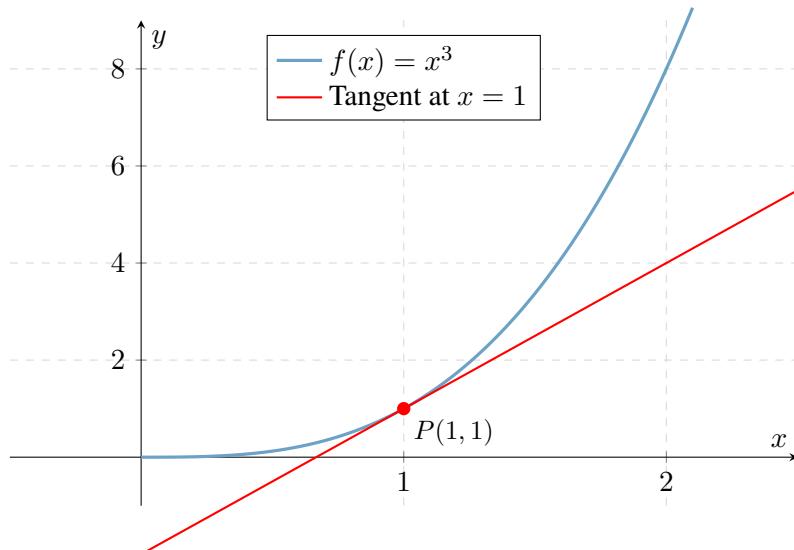


Figure 10.1: The tangent line to the graph of $f(x) = x^3$ at the point $P(1, 1)$.

Remark: For a straight line, the slope is constant. For a curved graph, the slope varies from point to point, so we speak of the *instantaneous rate of change*, the slope of the tangent at a specific point.

The same principle extends to computational and data-driven systems. A program's runtime changes with input size, a model's accuracy varies with the number of training epochs, and network throughput depends on available bandwidth. In machine learning, the process of training a model resembles a hiker searching for the lowest point in a vast, foggy landscape. The “altitude” represents the model’s error, and progress requires moving downhill. Determining the correct direction among millions of possible ones depends on knowing the local slope — the *derivative* — which expresses the instantaneous rate of change.

Differential calculus provides the tools to describe, compute, and interpret such changes with precision. It

transforms intuitive ideas such as being faster, steeper, or more efficient into exact mathematical expressions. This chapter introduces the fundamental ideas of differential calculus, beginning with the concepts of *limits* and *continuity*, which formalise the notion of approaching a point infinitely closely. Building on these, we develop the definition and interpretation of the *derivative* as a measure of instantaneous change and explore its applications in contexts ranging from physical motion to the optimisation of modern algorithms.

10.1 Limits and Continuity: The Foundation of Calculus

Before we can measure an instantaneous rate of change, we must first build a formal understanding of what it means to get "infinitesimally close" to a point without necessarily being at that point. This is the concept of a **limit**. Building on limits, **continuity** gives us the language to describe functions that are predictable and well-behaved, without sudden jumps or breaks.

Limits

A limit describes the value that a function "approaches" as its input gets closer and closer to a specific point. This idea is fundamental not just in calculus, but in software engineering, where we often analyze the behavior of systems under approaching conditions.

For example, when analyzing an algorithm's runtime, we ask what happens to the execution time as the input size 'n' approaches infinity. In networking, we might model the theoretical maximum throughput as a limit when latency approaches zero.

Definition 10.1 (Limit)

Let $f(x)$ be a function defined near a point c . We say that the **limit** of $f(x)$ as x approaches c is L , written as

$$\lim_{x \rightarrow c} f(x) = L$$

if we can make the value of $f(x)$ arbitrarily close to L by choosing an x that is sufficiently close to c , but not equal to c .



In many simple cases, the limit is just the value of the function at that point. For example, for $f(x) = x^2$, the limit as x approaches 3 is simply $3^2 = 9$. The real power of limits, however, is in handling cases where the function is undefined at the point of interest.

Example 10.1 A Limit at an Undefined Point

Consider the function $f(x) = \frac{\sin(x)}{x}$. This function is fundamental in signal processing and is known as the sinc function. Notice that $f(0)$ is undefined because it results in division by zero. However, we can ask what value the function approaches as x gets very close to 0.

By plugging in values very close to 0, we can observe a trend:

$$\begin{aligned}f(0.1) &= \frac{\sin(0.1)}{0.1} \approx 0.9983 \\f(0.01) &= \frac{\sin(0.01)}{0.01} \approx 0.99998 \\f(-0.01) &= \frac{\sin(-0.01)}{-0.01} \approx 0.99998\end{aligned}$$

The function appears to approach 1. In calculus, it can be formally proven that:

$$\lim_{x \rightarrow 0} \frac{\sin(x)}{x} = 1$$

This ability to analyze behavior at a problematic point is essential for creating robust numerical algorithms.

Example 10.2 Asymptotic Analysis with Limits

Suppose you want to analyze whether the function $f(n) = 5n^2 + 4n + 7$ is $\mathcal{O}(n^2)$ as n grows large, as is common in the study of algorithm complexity. One way to formalize this is to examine the limit:

$$\lim_{n \rightarrow \infty} \frac{f(n)}{n^2} = \lim_{n \rightarrow \infty} \frac{5n^2 + 4n + 7}{n^2} = \lim_{n \rightarrow \infty} \left(5 + \frac{4}{n} + \frac{7}{n^2} \right)$$

As n becomes very large, the terms $\frac{4}{n}$ and $\frac{7}{n^2}$ both approach zero, so the whole expression approaches 5. Because this limit is a finite constant, it confirms that $f(n)$ grows at the same rate as n^2 . This type of limit calculation is at the heart of asymptotic analysis, formally justifying Big- \mathcal{O} claims about algorithm runtime.

Continuity

Continuity formalizes the idea of a function being "well-behaved" or predictable. Intuitively, a function is continuous if you can draw its graph without lifting your pen from the paper. In software, we often assume our systems are continuous: a small change in input should lead to a small, predictable change in output, not a sudden, drastic jump.

Definition 10.2 (Continuity)

A function f is **continuous** at a point c if the following three conditions are met:

1. $f(c)$ is defined.
2. $\lim_{x \rightarrow c} f(x)$ exists.
3. $\lim_{x \rightarrow c} f(x) = f(c)$.

A function is continuous on an interval if it is continuous at every point in that interval.



Discontinuities often represent important events in software systems, such as state transitions or threshold triggers.

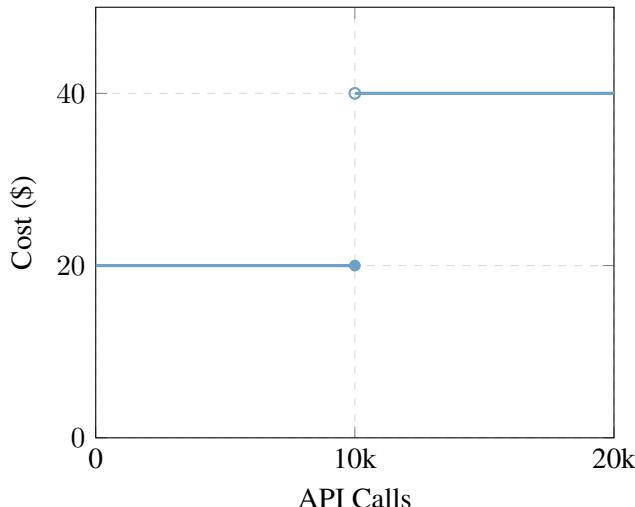
Example 10.3 Jump Discontinuity in a SaaS Pricing Model

Consider a cloud service that charges based on the number of API calls per month. The pricing model is tiered: \$20 for up to 10,000 calls, and \$40 for any usage above 10,000 calls. The cost function $C(x)$, where x is the

number of calls, has a **jump discontinuity** at $x = 10,000$.

$$C(x) = \begin{cases} 20 & \text{if } x \leq 10000 \\ 40 & \text{if } x > 10000 \end{cases}$$

As shown in [Figure 10.2](#), the cost suddenly jumps at the threshold. Understanding discontinuities is crucial for modeling any system with threshold-based logic, from auto-scaling policies to billing systems.



[Figure 10.2](#): The cost function exhibits a jump discontinuity at 10,000 API calls as the price tier changes.

10.2 The Derivative: Measuring the Rate of Change

We often talk about average rates of change. For example, if a data transfer takes 10 seconds and moves 500 MB, the average transfer rate is 50 MB/s. But this average tells us nothing about fluctuations during the transfer. The **derivative** is the tool that lets us move from this average rate to the **instantaneous rate of change** at any specific moment.

We may thus say that differential calculus is about determining how one quantity changes in relation to another. Geometrically, the derivative of a function at a point is the slope of the line tangent to the function's graph at that point. To find the slope of a tangent, we approximate it using secant lines between two nearby points on the curve. As the points move closer together, the secant slope approaches the tangent slope. As shown in [Figure 10.3](#), this tangent line is the limit of secant lines passing through two points on the curve as the points get infinitesimally close.

In [Figure 10.3](#) as the point Q moves along the curve toward P (from Q_1 to Q_2 to P), the slope of the secant line PQ approaches the slope of the tangent line to $f(x) = x^3$ at P . Note how the tangent only touches the curve at P , while the secant passes through two points.

This leads to the formal definition of the derivative.

Definition 10.3 (The Derivative)

The **derivative** of a function f with respect to x , denoted $f'(x)$, is the function

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

provided the limit exists. If $f'(x)$ exists, we say that f is **differentiable** at x .



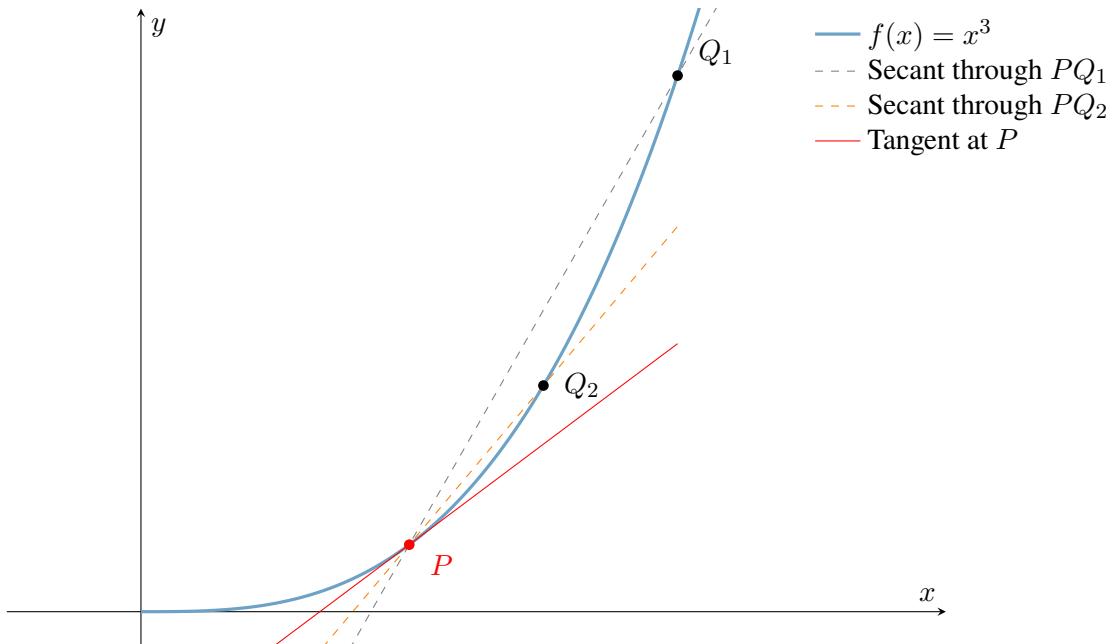


Figure 10.3: Illustration of the tangent line as a limit of secant lines.

The expression $\frac{f(x+h)-f(x)}{h}$ is the **difference quotient**, representing the average rate of change over the interval from x to $x + h$. The derivative is the limit of this average rate as the interval size h shrinks to zero.

Example 10.4 Derivative from First Principles

Let's find the derivative of $f(x) = x^2$ using the limit definition.

Solution: We start with the difference quotient and simplify:

$$\frac{f(x+h) - f(x)}{h} = \frac{(x+h)^2 - x^2}{h} = \frac{x^2 + 2xh + h^2 - x^2}{h} = \frac{2xh + h^2}{h} = 2x + h$$

Finally, we take the limit as $h \rightarrow 0$:

$$f'(x) = \lim_{h \rightarrow 0} (2x + h) = 2x$$

Thus, the derivative of $f(x) = x^2$ is $f'(x) = 2x$. This tells us the slope of the parabola at any point x . ◀

Notation

The derivative of f can be written in several equivalent forms:

$$f'(x), \quad \frac{df}{dx}, \quad \frac{dy}{dx}, \quad Df(x), \quad \frac{d}{dx}f(x).$$

10.3 Rules of Differentiation

Differential calculus provides a way to find the exact derivative of a function directly from its formula, without relying on graphs or numerical estimation. In practice, this is done using a set of simple rules that allow us to compute the derivative of nearly any function we are likely to encounter. In this section, we will introduce these rules, explain their meaning, and show how to apply them in practice.

Constant and Power Rules

Perhaps the simplest functions in mathematics are the constant functions and the functions of the form x^n .

Theorem 10.1 (Constant and Power Rules)

1. **Constant Rule:** If c is constant, then $\frac{d}{dx}(c) = 0$,
2. **Power Rule:** For any real n , $\frac{d}{dx}x^n = nx^{n-1}$.



Example 10.5

- If $f(x) = x^7$, then $f'(x) = 7x^6$,
- If $y = x^{-0.5}$, then $\frac{dy}{dx} = -0.5x^{-1.5}$,
- $\frac{d}{dx}x^{-3} = -3x^{-4}$,
- If $g(x) = 3.2$, then $g'(x) = 0$,
- If $f(t) = \sqrt{t} = t^{\frac{1}{2}}$, then $f'(t) = \frac{1}{2}t^{-\frac{1}{2}} = \frac{1}{2\sqrt{t}}$,
- If $h(u) = -13.29$, then $h'(u) = 0$,
- $\frac{d}{dx}x^3 = 3x^2$, $\frac{d}{dx}x^{-2} = -2x^{-3}$.

In the examples above, we have used the constant and power rules to find the derivatives of several simple functions. However, it is important to remember what these calculations actually represent. The derivative tells us the slope of the tangent line to the graph of a function at any given point.

For instance, if $f(x) = x^2$, then $f'(x) = 2x$. To find the slope of the tangent to the graph of x^2 at $x = 1$, we substitute $x = 1$ into the derivative, giving $f'(1) = 2 \times 1 = 2$. Similarly, the slope of the tangent at $x = -0.5$ is $f'(-0.5) = 2 \times (-0.5) = -1$.

These results are illustrated in [Figure 10.4](#), where the tangent lines at $x = 1$ and $x = -0.5$ have slopes of 2 and -1, respectively.

Example 10.6

Let $f(x) = (x^2 + 1)^3$. Then

$$f'(x) = 3(x^2 + 1)^2 \cdot 2x = 6x(x^2 + 1)^2.$$

Example 10.7

Find the slope of the tangent to the graph of the function $g(t) = t^4$ at the point on the graph where $t = -2$.

Solution:

The derivative is $g'(t) = 4t^3$, and so the slope of the tangent line at $t = -2$ is $g'(-2) = 4 \times (-2)^3 = -32$. ◀

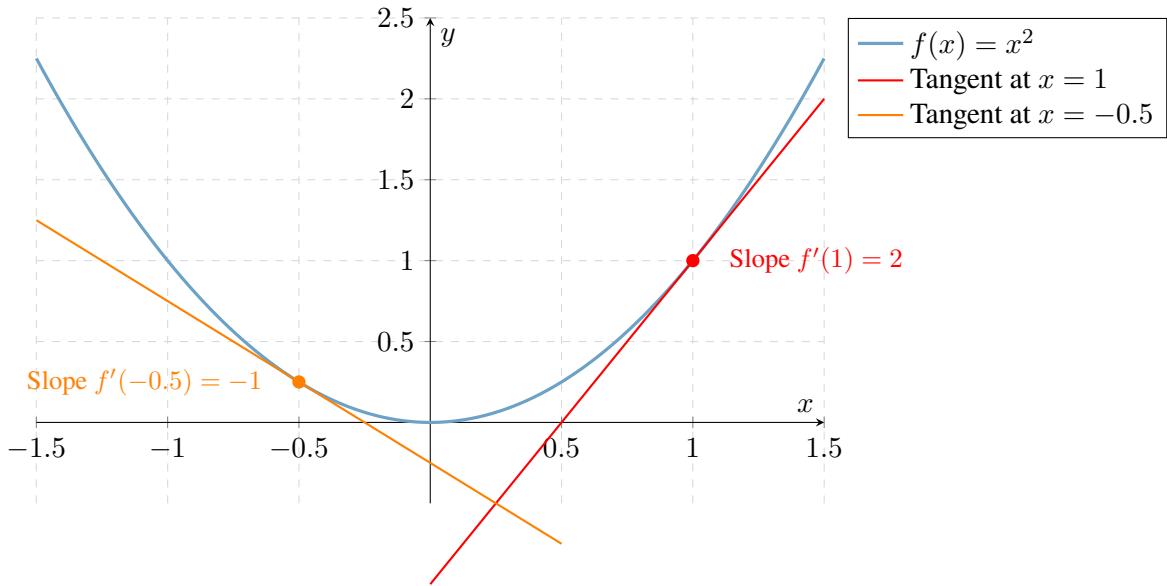


Figure 10.4: The tangent lines to the graph of $f(x) = x^2$ at different points have different slopes.

Example 10.8

Find the equation of the line tangent to the graph of $y = f(x) = x^{\frac{1}{2}}$ at the point $x = 4$.

Solution:

$f(4) = 4^{\frac{1}{2}} = \sqrt{4} = 2$, so the coordinates of the point on the graph are $(4, 2)$. The derivative is

Any non vertical line has equation of the form $y = mx + b$ where m is the slope and b the vertical intercept.

In this case the slope is $\frac{1}{4}$, so $m = \frac{1}{4}$, and the equation is $y = \frac{x}{4} + b$. Because the line passes through the point $(4, 2)$ we know that $y = 2$ when $x = 4$.

Substituting we get $2 = \frac{4}{4} + b$, so that $b = 1$. The equation is therefore $y = \frac{x}{4} + 1$. ◀

We now know how to differentiate any function that is a power of the variable. Examples are functions like x^3 and $t^{-1.3}$. You will come across functions that do not at first appear to be a power of the variable, but can be rewritten in this form. One of the simplest examples is the function

$$f(t) = \sqrt{t}$$

which can also be written in the form

$$f(t) = t^{\frac{1}{2}}.$$

The derivative is then

$$f'(t) = \frac{t^{-\frac{1}{2}}}{2} = \frac{1}{2\sqrt{t}}$$

Similarly, if

$$h(s) = \frac{1}{s} = s^{-1}$$

then

$$h'(s) = -s^{-2} = -\frac{1}{s^2}$$

Example 10.9

If $f(x) = \frac{1}{\sqrt[3]{x}} = x^{-\frac{1}{3}}$ then $f'(x) = -\frac{1}{3}x^{-\frac{4}{3}}$.

Example 10.10

If $y = \frac{1}{x\sqrt{x}} = x^{-\frac{3}{2}}$ then $\frac{dy}{dx} = -\frac{3}{2}x^{-\frac{5}{2}}$.

Scalar Rule and Linearity

So far we know how to differentiate powers of the independent variable. Many of the functions that arise in applications are built from such powers in simple ways. For instance, the function $3x^2$ is merely a scalar multiple of x^2 , yet neither the constant and power rules tell us how to differentiate $3x^2$. Likewise, these rules do not explain how to differentiate expressions such as $x^2 + x^3$ or $x^2 - x^3$.

Theorem 10.2 (Scalar and Linearity Rules)

3. **Scalar Rule:** $\frac{d}{dx}[c \cdot f(x)] = c \cdot f'(x)$.
4. **Linearity Rule:** $\frac{d}{dx}[f(x) \pm g(x)] = f'(x) \pm g'(x)$.



The scalar rule and linearity rules address this gap. They describe how to differentiate functions that are formed by multiplying a function by a constant, and by adding or subtracting functions. These rules allow us to extend our differentiation techniques from simple powers to a broad class of functions constructed from them.

Example 10.11

- If $f(x) = 3x^2$ then $f'(x) = 3 \times \frac{d}{dx}x^2 = 6x$.
- If $g(t) = 3t^2 + 2t^{-2}$ then $g'(t) = \frac{d}{dt}3t^2 + \frac{d}{dt}2t^{-2} = 6t - 4t^{-3}$.
- If $y = \frac{3}{\sqrt[3]{x}} - 2x\sqrt[3]{x} = 3x^{-\frac{1}{3}} - 2x^{\frac{4}{3}}$ then $\frac{dy}{dx} = -\frac{3}{2}x^{-\frac{5}{3}} - \frac{8}{3}x^{\frac{1}{3}}$.
- If $y = -0.3x^{-0.4}$ then $\frac{dy}{dx} = 0.12x^{-1.4}$.
- $\frac{d}{dx}2x^{0.3} = 0.6x^{-0.7}$.

Caution. Although the linearity rule states that $\frac{d}{dx}(f(x) \pm g(x)) = f'(x) \pm g'(x)$, this property does *not* extend to products or quotients. In general,

$$\frac{d}{dx}(f(x)g(x)) \neq f'(x)g'(x), \quad \frac{d}{dx}\left(\frac{f(x)}{g(x)}\right) \neq \frac{f'(x)}{g'(x)}.$$

To differentiate $f(x)g(x)$ or $\frac{f(x)}{g(x)}$ we cannot simply differentiate the factors and then multiply or divide the results. The correct techniques are the *product rule* and the *quotient rule*, which are developed in the next section.

The Product and Quotient Rules

Another common way of combining functions is to multiply them, thereby forming a *product*. The product rule provides the method for differentiating functions constructed in this way.

On the other hand the quotient rule allows us to differentiate functions which are formed by dividing one function by another, i.e. by forming quotients of functions. The quotient rule provides the method for differentiating functions constructed in this way.

Theorem 10.3 (Product and Quotient Rules)

- 5. Product Rule:** If $f(x)$ and $g(x)$ are differentiable, then

$$\frac{d}{dx}(f(x)g(x)) = f'(x)g(x) + f(x)g'(x).$$

- 6. Quotient Rule:** If $f(x)$ and $g(x)$ are differentiable and $g(x) \neq 0$, then

$$\frac{d}{dx}\left(\frac{f(x)}{g(x)}\right) = \frac{f'(x)g(x) - f(x)g'(x)}{[g(x)]^2}.$$



Example 10.12 Product Rule

- If $y = (x+2)(x^2+3)$ then $y' = (x+2)2x + 1(x^2+3)$.
- If $f(x) = \sqrt{x}(x^3 - 3x^2 + 7)$ then $f'(x) = \sqrt{x}(3x^2 - 6x) + \frac{1}{2}x^{-\frac{1}{2}}(x^3 - 3x^2 + 7)$.
- If $z = (t^2 + 3)(\sqrt{t} + t^3)$ then $\frac{dz}{dt} = (t^2 + 3)\left(\frac{1}{2}t^{-\frac{1}{2}} + 3t^2\right) + 2t(\sqrt{t} + t^3)$.

Example 10.13 Quotient Rule

- If $y = \frac{2x^2+3x}{x^3+1}$, then $\frac{dy}{dx} = \frac{(x^3+1)(4x+3)-(2x^2+3x)3x^2}{(x^3+1)^2}$.
- If $g(t) = \frac{t^2+3t+1}{\sqrt{t+1}}$ then $g'(t) = \frac{(\sqrt{t+1})(2t+3)-(t^2+3t+1)\left(\frac{1}{2}t^{-\frac{1}{2}}\right)}{(\sqrt{t+1})^2}$.

In the quotient rule, because of the minus sign in the numerator (i.e. in the top line) it is important to get the terms in the numerator in the correct order. This is often a source of mistakes, so be careful. Decide on your own way of remembering the correct order of the terms.

The Chain Rule

The chain rule is arguably the most important differentiation rule, especially in machine learning. It tells us how to find the derivative of a composite function — a function nested inside another. In software, we constantly compose functions, and the chain rule is the mathematical tool for analyzing how changes propagate through these compositions.

Theorem 10.4 (Chain Rule)

1. If $h(x) = f(g(x))$ is a composite function, then its derivative is the derivative of the outer function (evaluated at the inner function) multiplied by the derivative of the inner function.

$$h'(x) = f'(g(x)) \cdot g'(x).$$



Example 10.14

Find the derivative of $h(x) = (x^3 + 2x)^5$.

Solution: This is a composition where the outer function is $f(u) = u^5$ and the inner function is $g(x) = x^3 + 2x$. Their derivatives are $f'(u) = 5u^4$ and $g'(x) = 3x^2 + 2$. Applying the chain rule:

$$\begin{aligned} h'(x) &= f'(g(x)) \cdot g'(x) \\ &= 5(x^3 + 2x)^4 \cdot (3x^2 + 2) \end{aligned}$$

Example 10.15

Differentiate $(3x^2 - 5)^3$.

Solution: The first step is always to identify that the expression represents a composite function and then to separate it into its outer and inner components. In this example, the outer function is $(\cdot)^3$, whose derivative is $3(\cdot)^2$, and the inner function is $3x^2 - 5$, whose derivative is $6x$. Applying the composite function rule therefore gives

$$\frac{d}{dx}(3x^2 - 5)^3 = 3(3x^2 - 5)^2 \times 6x = 18x(3x^2 - 5)^2.$$

Alternatively, we may introduce the substitution $u = 3x^2 - 5$ and write $y = u^3$. Then

$$\frac{dy}{dx} = \frac{dy}{du} \times \frac{du}{dx} = 3u^2 \times 6x = 18x(3x^2 - 5)^2.$$

Example 10.16

Find $\frac{dy}{dx}$ if $y = \sqrt{x^2 + 1}$.

Solution: The outer function is $\sqrt{\cdot}$, whose derivative is $\frac{1}{2\sqrt{\cdot}}$, and the inner function is $x^2 + 1$, whose derivative is $2x$. Applying the chain rule therefore gives

$$\frac{dy}{dx} = \frac{1}{2\sqrt{x^2 + 1}} \times 2x = \frac{x}{\sqrt{x^2 + 1}}.$$

The chain rule is the core mechanism behind the **backpropagation** algorithm used to train neural networks. An error at the output of a network is a composition of many nested functions (the layers). Backpropagation uses the chain rule repeatedly to calculate the derivative of the error with respect to each weight in the network, telling us how to adjust each weight to improve the model.

Example 10.17 A Multi-Rule Problem

Find the derivative of $h(x) = \frac{x^2}{(3x+1)^4}$.

Solution: This problem requires the quotient rule, and the denominator requires the chain rule. Let $f(x) = x^2$ and $g(x) = (3x + 1)^4$. Then $f'(x) = 2x$. To find $g'(x)$, we use the chain rule: $g'(x) = 4(3x + 1)^3 \cdot 3 =$

$12(3x + 1)^3$. Now, apply the quotient rule:

$$\begin{aligned} h'(x) &= \frac{f'(x)g(x) - f(x)g'(x)}{[g(x)]^2} \\ &= \frac{(2x)(3x+1)^4 - (x^2)(12(3x+1)^3)}{((3x+1)^4)^2} \\ &= \frac{(2x)(3x+1) - 12x^2}{(3x+1)^5} \quad (\text{after factoring out and canceling } (3x+1)^3) \\ &= \frac{6x^2 + 2x - 12x^2}{(3x+1)^5} = \frac{2x - 6x^2}{(3x+1)^5} \end{aligned}$$



Summary of Rules of Differentiation

The rules of differentiation are summarized in the following theorem.

Theorem 10.5 (Rules of Differentiation)

1. **Constant Rule:** If c is constant, then $\frac{d}{dx}c = 0$.
2. **Power Rule:** For any real n , $\frac{d}{dx}x^n = nx^{n-1}$.

$$\frac{d}{dx}x^n = nx^{n-1}.$$
3. **Scalar Rule:** If c is constant, then $\frac{d}{dx}[c \cdot f(x)] = c \cdot f'(x)$.
4. **Linearity:** For differentiable f, g and constants a, b , $\frac{d}{dx}(af + bg) = af' + bg'$.
5. **Product Rule:** If f, g are differentiable, then $\frac{d}{dx}(fg) = f'g + fg'$.
6. **Quotient Rule:** If f, g are differentiable and $g(x) \neq 0$, then $\frac{d}{dx}\left(\frac{f}{g}\right) = \frac{f'g - fg'}{g^2}$.
7. **Chain Rule:** If f, g are differentiable, then $\frac{d}{dx}(f(g(x))) = f'(g(x))g'(x)$.



These rules do not apply to every function. For instance, the derivative of e^x is e^x , while the derivative of $\ln x$ is $\frac{1}{x}$, and neither of these fits the patterns described so far. The most common differentiable functions that fall outside these earlier rules are listed below.

- $\frac{d}{dx}e^x = e^x$
- $\frac{d}{dx}\ln x = \frac{1}{x}$
- $\frac{d}{dx}\sin x = \cos x$
- $\frac{d}{dx}\cos x = -\sin x$

Example 10.18

Differentiate $\ln(x^2 + 3x + 1)$.

Solution: We solve this by using the chain rule and our knowledge of the derivative of $\ln x$.

$$\begin{aligned}\frac{d}{dx} \ln(x^2 + 3x + 1) &= \frac{d}{dx} (\ln u) \quad (\text{where } u = x^2 + 3x + 1) \\ &= \frac{d}{du} (\ln u) \times \frac{du}{dx} \quad (\text{by the chain rule}) \\ &= \frac{1}{u} \times \frac{du}{dx} \\ &= \frac{1}{x^2 + 3x + 1} \times \frac{d}{dx} (x^2 + 3x + 1) \\ &= \frac{1}{x^2 + 3x + 1} \times (2x + 3) \\ &= \frac{2x + 3}{x^2 + 3x + 1}\end{aligned}$$



Example 10.19

Find $\frac{d}{dx} (e^{3x^2})$.

Solution: This is an application of the chain rule together with our knowledge of the derivative of e^x .

$$\begin{aligned}\frac{d}{dx} (e^{3x^2}) &= \frac{de^u}{dx} \quad \text{where } u = 3x^2 \\ &= \frac{de^u}{du} \times \frac{du}{dx} \quad \text{by the chain rule} \\ &= e^u \times \frac{du}{dx} \\ &= e^{3x^2} \times \frac{d}{dx} (3x^2) \\ &= 6xe^{3x^2}\end{aligned}$$



Example 10.20

Find $\frac{d}{dx} (e^{x^2+2x})$.

Solution: Again, we use our knowledge of the derivative of e^x together with the chain rule.

$$\begin{aligned}\frac{d}{dx} (e^{x^2+2x}) &= \frac{de^u}{dx} \quad (\text{where } u = x^3 + 2x) \\ &= e^u \times \frac{du}{dx} \quad (\text{by the chain rule}) \\ &= e^{x^3+2x} \times \frac{d}{dx} (x^3 + 2x) \\ &= (3x^2 + 2) \times e^{x^3+2x}\end{aligned}$$



Example 10.21

Differentiate $\ln(2x^3 + 5x^2 - 3)$.

Solution: We solve this by using the chain rule and our knowledge of the derivative of $\ln x$.

$$\begin{aligned}\frac{d}{dx} \ln(2x^3 + 5x^2 - 3) &= \frac{d \ln u}{dx} \quad (\text{where } u = 2x^3 + 5x^2 - 3) \\ &= \frac{d \ln u}{du} \times \frac{du}{dx} \quad (\text{by the chain rule}) \\ &= \frac{1}{u} \times \frac{du}{dx} \\ &= \frac{1}{2x^3 + 5x^2 - 3} \times \frac{d}{dx}(2x^3 + 5x^2 - 3) \\ &= \frac{1}{2x^3 + 5x^2 - 3} \times (6x^2 + 10x) \\ &= \frac{6x^2 + 10x}{2x^3 + 5x^2 - 3}\end{aligned}$$



There are two shortcuts to differentiating functions involving exponents and logarithms. The four examples above gave

$$\begin{aligned}\frac{d}{dx}(\ln(x^2 + 3x + 1)) &= \frac{2x + 3}{x^2 + 3x + 1} \\ \frac{d}{dx}(e^{3x^2}) &= 6xe^{3x^2} \\ \frac{d}{dx}(e^{x^2+2x}) &= (3x^2 + 2)e^{x^2+2x} \\ \frac{d}{dx}(\ln(2x^3 + 5x^2 - 3)) &= \frac{6x^2 + 10x}{2x^3 + 5x^2 - 3}\end{aligned}$$

These examples suggest the general rules

$$\begin{aligned}\frac{d}{dx}(e^{f(x)}) &= f'(x)e^{f(x)} \\ \frac{d}{dx}(\ln f(x)) &= \frac{f'(x)}{f(x)}\end{aligned}$$

These rules arise from the chain rule and the fact that $\frac{dx^x}{dx} = e^x$ and $\frac{d \ln x}{dx} = \frac{1}{x}$. They can speed up the process of differentiation but it is not necessary that you remember them. If you forget, just use the chain rule as in the examples above.

10.4 Applications of Derivatives

The development of mathematics ranks among the greatest achievements of human thought, and the emergence of calculus — both differential and integral — marks a turning point in that history. Differential calculus in particular has countless practical applications across science and engineering, far too many to enumerate. Its importance is reflected in the fact that virtually every quantitative discipline relies on its methods.

Within elementary mathematics, differential calculus is used primarily in two areas: analysing and sketching curves, and solving optimisation problems. In this section, we offer a brief introduction to how differential calculus is applied in optimisation.

Stationary points - the idea behind optimisation

Let us return to our metaphor of the mountain range. Imagine standing somewhere among peaks and valleys, perhaps even with your eyes closed or shrouded in fog. As you wander along, you notice that whenever you're hiking uphill or downhill, you know you are not at the highest summit or the lowest valley—you are either still climbing or still descending.

But at the top of a peak, there is a brief, level stretch where the slope under your feet is zero. At a bottom of a valley, there is also a level spot. If you find yourself on perfectly flat ground—at least for a moment—you have a clue: you might be at a summit, a valley, or perhaps at a gentle inflection point along a ridge. In the vast terrain of a mountain range, it is only at these level patches, where the slope vanishes, that you could possibly be at an extreme point: a maximum or minimum.

This is the key insight behind optimisation in calculus: to find the highest or lowest points (the peaks or valleys) of a function—in other words, to solve optimisation problems—we need only to examine the locations where the “slope” is zero. These locations are known as stationary points.

Definition 10.4

For a function $y = f(x)$ the points on the graph where the graph has zero slope are called stationary points. In other words stationary points are where $f'(x) = 0$.



To find the stationary points of a function we differentiate, set the derivative equal to zero and solve the equation.

Example 10.22

Find the stationary points of the function $f(x) = 2x^3 + 3x^2 - 12x + 17$.

Solution:

$f'(x) = 6x^2 + 6x - 12$. Setting $f'(x) = 0$ and solving we obtain

$$6x^2 + 6x - 12 = 0$$

$$x^2 + x - 2 = 0$$

$$(x - 1)(x + 2) = 0$$

$$x = 1, -2$$

This gives us the values of x for which the function f is stationary. The corresponding values of the function are found by substituting 1 and -2 into the function.

They are $f(1) = 2 \times 1^3 + 3 \times 1^2 - 12 \times 1 + 17 = 10$ and $f(-2) = 2 \times (-2)^3 + 3 \times (-2)^2 - 12 \times (-2) + 17 = 37$. The stationary points are therefore $(1, 10)$ and $(-2, 37)$.



Example 10.23

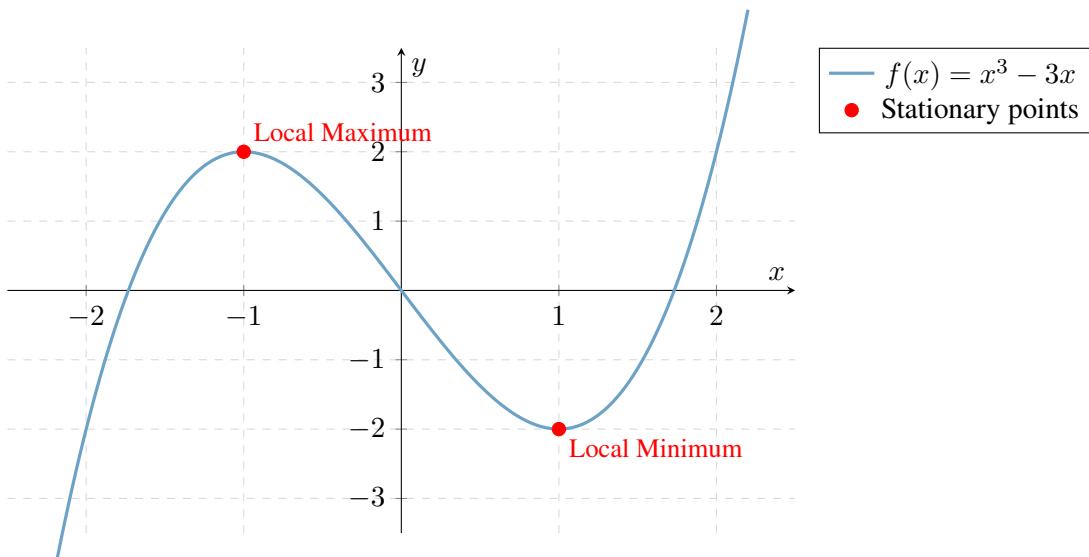
Find the stationary points of the function $g(t) = e^{t^2}$.

Solution: Differentiating and setting the derivative equal to zero we obtain the equation $g'(t) = 2te^{t^2} = 0$. Since e^{t^2} is never zero, the only solution to this equation is where $2t = 0$, ie $t = 0$. Substituting into the formula for g we obtain the function value $g(0) = e^{0^2} = 1$. Thus the stationary point is $(0, 1)$.



Types of stationary points

Returning to our mountain range metaphor, stationary points can be thought of as the special spots where, just for a moment, the ground feels perfectly level beneath your feet. Sometimes, this happens at the very top of a hill—where, if you reach out in any direction, you only sense downhill slopes ahead. This is called a local maximum: the summit of your immediate surroundings, though not necessarily the tallest peak in the whole range. Other times, the level spot is at the bottom of a valley—everything around you slopes upward. This is called a local minimum: the lowest place in your local vicinity, even if deeper valleys exist elsewhere. The “local” label emphasizes that we’re talking about the highest or lowest point nearby, not globally. You might also hear the terms “relative maximum” and “relative minimum” used for these. [Figure 10.5](#) illustrates a function as a landscape with both a hilltop (local maximum) and a valley floor (local minimum). Notice that, at each of these points—just like on a level patch in the mountains—the slope is zero.



[Figure 10.5](#): A curve with a local maximum at $x = -1$ and a local minimum at $x = 1$.

Local maxima and local minima are not the only types of stationary points. There is a third kind. [Figure 10.6](#) shows a stationary point that is neither a local maximum nor a local minimum. This type of stationary point is called a stationary point of inflection or simply *inflection point*.

The first derivative test

Let us return to [Example 10.22](#). For the function $f(x) = 2x^3 + 3x^2 - 12x + 17$ we identified stationary points at $(1, 10)$ and $(-2, 37)$. The natural question is: what type of stationary points are these? Without a sketch of the graph it is not immediately clear whether they correspond to local maxima, local minima, or stationary points of inflection. Drawing an accurate graph would resolve the question, but doing so can require substantial effort. Instead, we seek a method that allows us to classify a stationary point without needing to plot the entire function.

Several approaches exist, but in this section we focus on one technique, known as the *first derivative test*. The idea is to examine the behaviour of the function immediately to the left and immediately to the right of the stationary point.

To understand the principle, imagine a person standing at a point on a landscape where the ground is level.

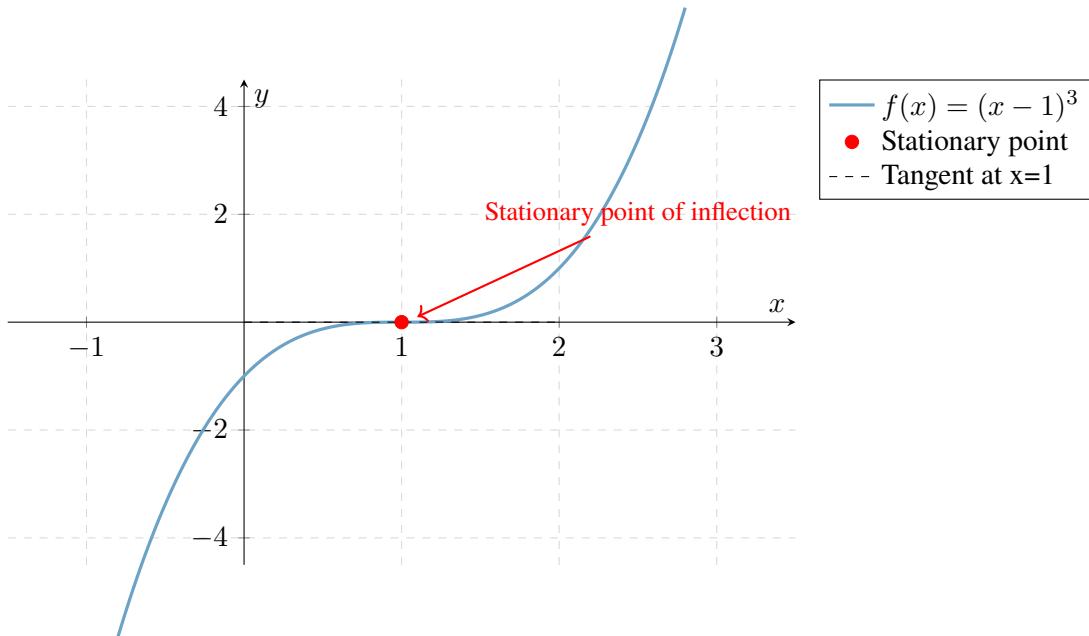


Figure 10.6: A function with a stationary point of inflection at $x = 1$. The tangent line at this point is horizontal, but it is neither a local maximum nor a local minimum.

Without seeing the surroundings, the person wishes to determine whether this point is the top of a hill, the bottom of a valley, or neither. One way to decide is to take a small step backward and observe the slope, and then take a small step forward and observe the slope again.

If the ground slopes upward behind and downward ahead, the level point must be the top of a hill, corresponding to a local maximum. If the ground slopes downward behind and upward ahead, the point must be the bottom of a valley, corresponding to a local minimum. The remaining possibility—where the slopes have the same sign on both sides—indicates a stationary point of inflection.

This reasoning captures the essence of the first derivative test: by analysing the sign of $f'(x)$ immediately on either side of a stationary point, we can determine its character without relying on a graph.

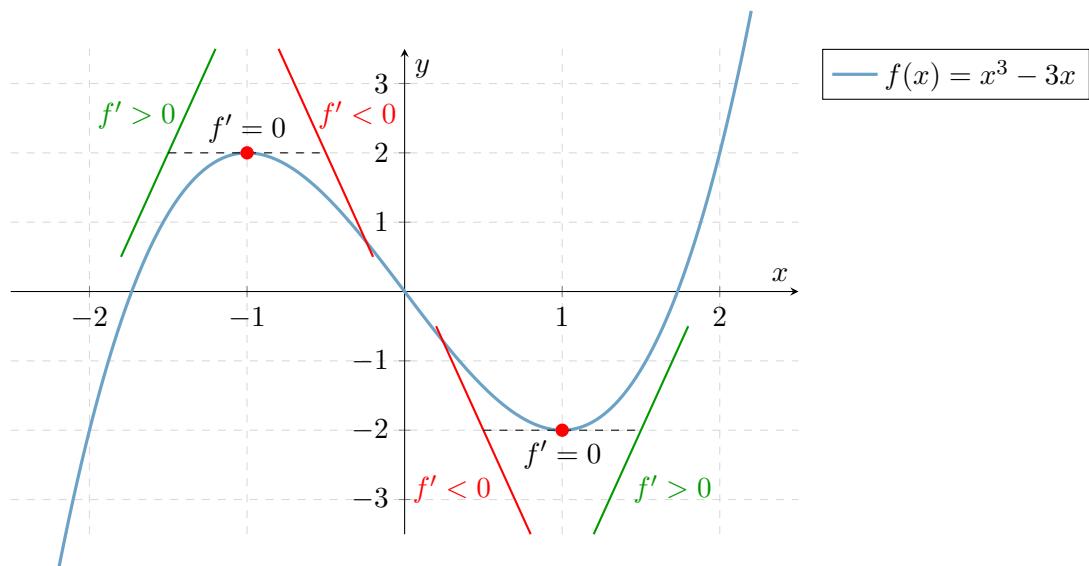


Figure 10.7: Illustration of the first derivative test. At a local maximum, the derivative changes from positive to negative. At a local minimum, it changes from negative to positive.

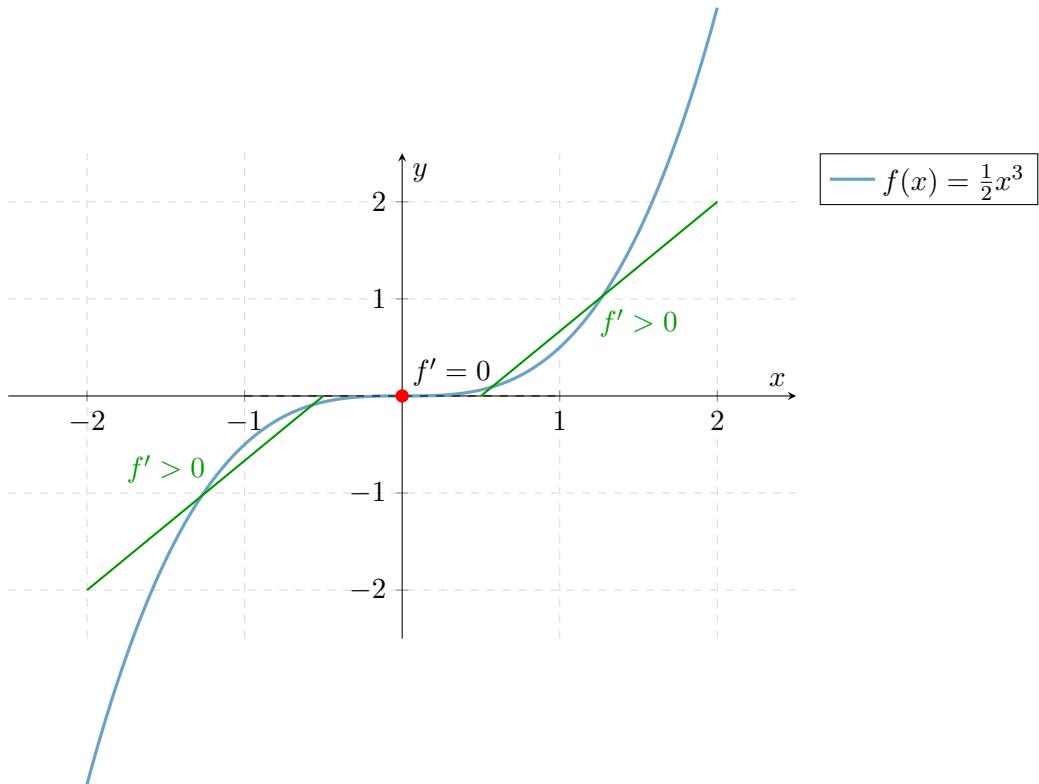


Figure 10.8: Illustration of a stationary point of inflection. The derivative is zero at $x = 0$, but it is positive on both sides, so the point is neither a local maximum nor a minimum.

We can summarize these considerations in the following theorem:

Theorem 10.6 (The First Derivative Test)

Let c be a critical point of a continuous function f .

- If $f'(x)$ changes from negative to positive at c , then f has a **local minimum** at c .
- If $f'(x)$ changes from positive to negative at c , then f has a **local maximum** at c .
- If $f'(x)$ does not change sign at c , then f has no local extremum at c ; it is a stationary point of inflection.



This test is visualized in [Figure 10.7](#) and [Figure 10.8](#).

Optimisation

Now that we've laid the groundwork, we can approach some optimisation problems. When we want to maximise a function $f(x)$ within a specified interval for x , our goal is to find the largest value that $f(x)$ achieves on that interval. Importantly, this maximum value does not always coincide with a stationary point. As shown in [Figure 10.9](#), consider searching for the greatest and least values of a function on the interval $2 \leq x \leq 7$. Within this interval, the function has two stationary points—one corresponding to a local maximum, the other to a local minimum. However, the highest value of the function in this interval actually occurs at the endpoint $x = 7$, which is not a stationary point. It takes this value simply because, outside of the interval, larger x values are excluded from consideration. In contrast, the minimum value in this example is located at a stationary point within the interval. With this understanding, we can now explain the systematic steps for finding the maximum or minimum of a function in a given region.

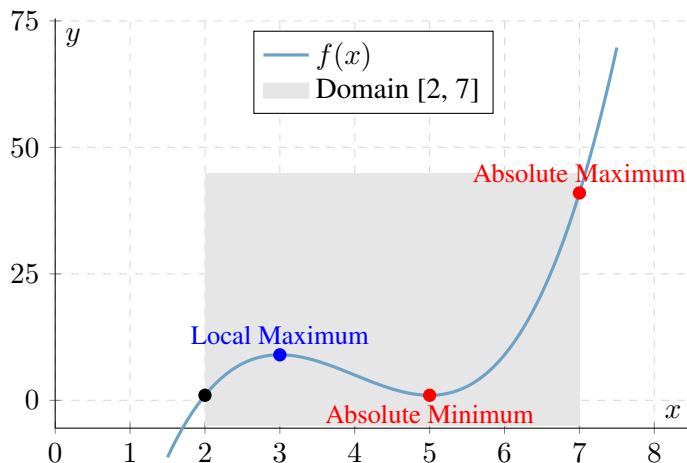


Figure 10.9: On the restricted domain $[2, 7]$, the absolute maximum occurs at an endpoint ($x = 7$), not at the local maximum. The absolute minimum occurs at an interior stationary point ($x = 5$).

The location of maxima and minima

A function $f(x)$ may or may not have a maximum or minimum value in a particular region of x values. However, if they do exist the maximum and the minimum values must occur at one of three places:

1. At the endpoints (if they exist) of the region under consideration.
2. Inside the region at a stationary point.
3. Inside the region at a point where the derivative does not exist.

Remark:

1. It is straightforward to find cases where a function does not attain a maximum or minimum within a certain range. For instance, the function $f(x) = x$ lacks both a maximum and a minimum over the interval $-\infty < x < \infty$; its graph increases without bound as x increases. According to Point 1 above, note that in the region $-\infty < x < \infty$, there are actually no endpoints. As another illustration, for the region $x \geq 1$, there is only a single endpoint at $x = 1$.
2. Regarding Point 3, this text does not cover situations where the derivative fails to exist at certain points. Nevertheless, keep in mind that such points can occur and could be where a maximum or minimum is located. For more details, refer to advanced calculus resources.

Now that we know exactly where the maxima or minima can occur, we can give a procedure for finding them.

Procedure for finding the maximum or minimum values of a function.

1. Find the endpoints of the region under consideration (if there are any).
2. Find all the stationary points in the region.
3. Find all points in the region where the derivative does not exist.
4. Substitute each of these into the function and see which gives the greatest (or smallest) function value.

Example 10.24

Find the minimum value and the maximum value of the function $f(x) = x^2 e^x$ for $-4 \leq x \leq 1$.

Solution: We will follow the procedure outlined above. The endpoints are -4 and 1. Differentiating we obtain $f'(x) = x^2 e^x + 2xe^x = x(x+2)e^x$. Setting $f'(x) = 0$ and solving we get stationary points at $x = 0$ and $x = -2$. There are no points where the derivative does not exist. Therefore the maximum and minimum values will be found at one of the points $x = -4, -2, 0, 1$. Substituting we obtain $f(-4) \approx 0.29$, $f(-2) \approx 0.54$, $f(0) = 0$ and $f(1) = e \approx 2.7$. therefore the maximum value occurs at $x = 1$ and is equal to e , and the minimum value occurs at $x = 0$ and is 0.

Example 10.25

Find the maximum and minimum values of the function $g(t) = \frac{1}{3}t^3 - t + 2$ for $0 \leq t \leq 3$.

Solution: The endpoints are $t = 0$ and $t = 3$. Differentiating and equating to zero we get $g'(t) = t^2 - 1 = (t-1)(t+1) = 0$ so the stationary points are at $t = -1, 1$. Since -1 is not in the region, the possible locations of the maximum and the minimum are $t = 0, 1, 3$. Substituting into g we obtain $g(0) = 2$, $g(1) = \frac{4}{3}$ and $g(3) = 8$. The maximum is therefore $g(3) = 8$ and the minimum is $g(1) = \frac{4}{3}$.

Example 10.26

Suppose you are tuning a machine learning model, and you want to find the learning rate α (between 0 and 1) that minimizes the loss function $L(\alpha) = (\alpha - 0.3)^2 + 1$. What value of α gives the minimum loss?

Solution: We will apply the procedure for finding absolute extrema on the closed interval $[0, 1]$. The function to minimize is the loss function $L(\alpha) = (\alpha - 0.3)^2 + 1$.

1. **Identify Endpoints:** The region under consideration is $0 \leq \alpha \leq 1$, so the endpoints are $\alpha = 0$ and $\alpha = 1$.

2. **Find Stationary Points:** We first find the derivative of the loss function with respect to α :

$$\frac{dL}{d\alpha} = \frac{d}{d\alpha} [(\alpha - 0.3)^2 + 1] = 2(\alpha - 0.3)$$

Next, we set the derivative to zero to find the stationary points:

$$2(\alpha - 0.3) = 0 \implies \alpha = 0.3$$

The only stationary point is $\alpha = 0.3$, which lies within our interval $[0, 1]$.

3. **Identify Points Where Derivative is Undefined:** The derivative $2(\alpha - 0.3)$ is a simple linear function and is defined for all values of α . There are no such points.

4. **Compare Values at Candidate Points:** We now evaluate the loss function $L(\alpha)$ at all the candidate points we have found: the endpoints and the stationary point.

- At the left endpoint, $\alpha = 0$:

$$L(0) = (0 - 0.3)^2 + 1 = 0.09 + 1 = 1.09$$

- At the stationary point, $\alpha = 0.3$:

$$L(0.3) = (0.3 - 0.3)^2 + 1 = 0 + 1 = 1.00$$

- At the right endpoint, $\alpha = 1$:

$$L(1) = (1 - 0.3)^2 + 1 = (0.7)^2 + 1 = 0.49 + 1 = 1.49$$

By comparing these values, we see that the smallest loss is 1.00, which occurs at the stationary point.

Conclusion: The minimum loss occurs when the learning rate is $\alpha = 0.3$.

Bibliography

- Cormen, T., Leiserson, C., Rivest, R., & Stein, C. (2022). *Introduction to algorithms* (4th). MIT Press.
- Lay, D. (2003). *Linear algebra and its applications*. Pearson Education.
- Montgomery, D. (2013). *Applied statistics and probability for engineers, 6th edition*. John Wiley; Sons, Incorporated.
- Pishro-Nik, H. (2014). *Introduction to probability, statistics, and random processes*. Kappa Research, LLC.
- Rosen, K. H. (2012). *Discrete mathematics and its applications* (7th). McGraw-Hill Education.
- Ross, S. M. (2020). *Introduction to probability and statistics for engineers and scientists* (6th). Academic Press.
- Thomas, C. *Introduction to differential calculus* [Accessed online]. Mathematics Learning Centre, University of Sydney. Accessed online. Sydney, NSW, 1997.

Appendix A: Important Concepts

This appendix is a collection of important mathematical concepts that are frequently used in software engineering. The content of this appendix is based on the concepts in this book.

Proposition A.1 (Order of Operations)

To evaluate mathematical expressions, operations are performed in the following order:

1. **Brackets (Parentheses):** First, perform all operations inside brackets or parentheses.
2. **Exponents and Radicals:** Next, evaluate exponents (powers) and radicals (roots).
3. **Multiplication and Division:** Then, perform multiplication and division from left to right.
4. **Addition and Subtraction:** Finally, execute addition and subtraction from left to right.



Proposition A.2 (Rules for Calculations with Fractions)

For $a, b, c, m \in \mathbb{R}$, with $a, b, c, m \neq 0$ where required, the following identities hold:

$$(1) \quad \frac{a}{b} \times m = \frac{am}{b}$$

$$(2) \quad \frac{a}{b} \div m = \frac{a}{bm}$$

$$(3) \quad m \div \frac{a}{b} = \frac{mb}{a}$$

$$(4) \quad \frac{a}{b} \times \frac{c}{a} = \frac{c}{b}$$

$$(5) \quad \frac{a}{b} \div \frac{c}{a} = \frac{a^2}{bc}$$

$$(6) \quad \frac{a}{b} = \frac{ac}{bc}$$

$$(7) \quad \frac{a}{b} + \frac{c}{a} = \frac{a^2 + bc}{ab}$$



Proposition A.3 (Properties of Integer Exponents)

Let $n, m \in \mathbb{Z}$. Then the following hold (with $x, y \in \mathbb{R}$ and nonzero where stated):

$$(1) \quad x^n \cdot x^m = x^{n+m},$$

$$(2) \quad \frac{x^n}{x^m} = x^{n-m} \quad \text{with } x \neq 0,$$

$$(3) \quad x^n \cdot y^n = (xy)^n,$$

$$(4) \quad \frac{x^n}{y^n} = \left(\frac{x}{y}\right)^n \quad \text{with } y \neq 0,$$

$$(5) \quad (x^n)^m = x^{nm},$$

$$(6) \quad x^1 = x.$$



Proposition A.4 (More Properties of Integer Exponents)

Let $n, m \in \mathbb{Z}$. Then the following hold (with $x, y \in \mathbb{R}$ and nonzero where stated):

$$(7) \quad x^0 = 1 \quad x \neq 0$$

$$(8) \quad \frac{1}{x^m} = x^{-m} \quad x \neq 0$$



Rules for rearranging formulae

The following operations can be performed on both sides of the formula:

- Add the same quantity to both sides
- Subtract the same quantity from both sides
- Multiply both sides by the same quantity - remember to multiply all terms
- Divide both sides by the same quantity - remember to divide all terms
- Apply a function to both sides, such as squaring or finding the reciprocal

Definition A.1 (Injective and Surjective Functions)

A function $f : A \rightarrow B$ is called **one-to-one** (or **injective**) if different elements in A map to different elements in B . A function $f : A \rightarrow B$ is called **onto** (or **surjective**) if every element in B is the image of at least one element in A .



Definition A.2 (Inverse Functions)

Let f be a one-to-one correspondence from the set A to the set B . The inverse function of f is the function that assigns to an element b belonging to B the unique element a in A such that $f(a) = b$. The inverse function of f is denoted by f^{-1} . Hence, $f^{-1}(b) = a$ when $f(a) = b$.

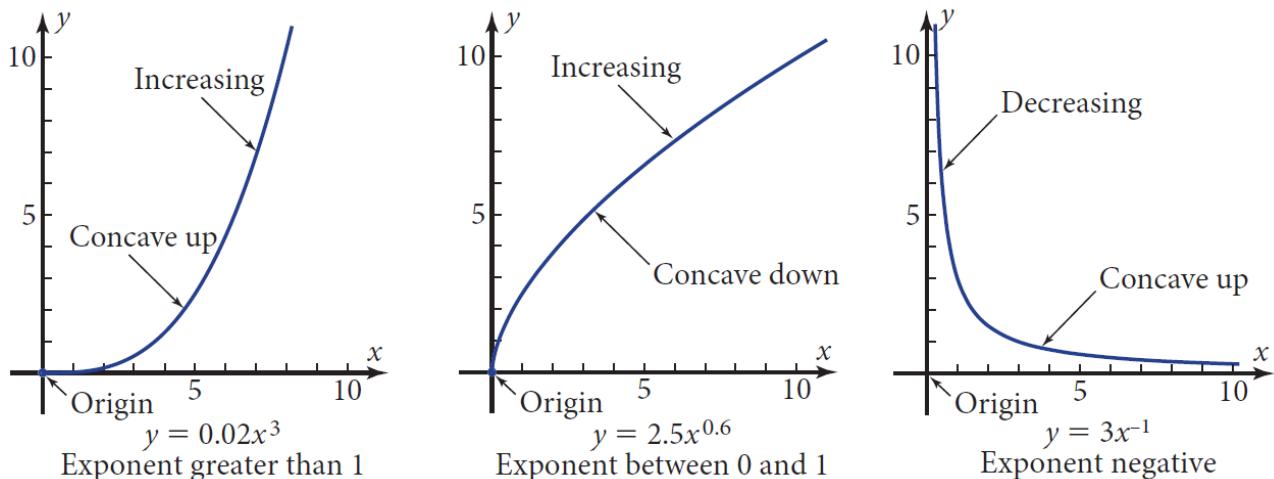


Figure A.1: Power functions

Definition A.3 (Base-10 Logarithms)

$$\log x = y \iff 10^y = x$$

Verbally: $\log x$ is the exponent in the power of 10 that gives x



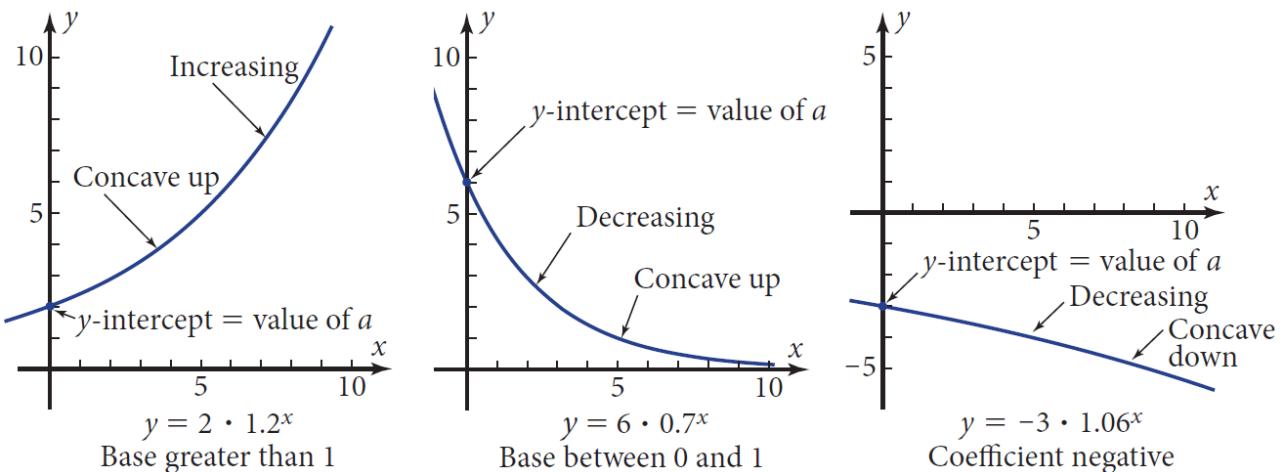


Figure A.2: Exponential functions

Properties of base-10 logarithms

- Log of a Product:

$$\log xy = \log x + \log y$$

Verbally: The log of a product equals the sum of the logs of the factors.

- Log of a Quotient:

$$\log \frac{x}{y} = \log x - \log y$$

Verbally: The log of a quotient equals the log of the numerator minus the log of the denominator.

- Log of a Power:

$$\log x^y = y \log x$$

Verbally: The log of a power equals the exponent times the log of the base.

Definition A.4 (Common Logarithm and Natural Logarithm)

Common: The symbol $\log x$ means $\log_{10} x$.

Natural: The symbol $\ln x$ means $\log_e x$, where e is a constant equal to $2.71828182845\dots$



The Change-of-Base Property of Logarithms

$$\log_a x = \frac{\log_b x}{\log_b a} \quad \text{or} \quad \log_a x = \frac{1}{\log_b a} (\log_b x)$$

Properties of Logarithms

The Logarithm of a Power:

$$\log_b x^y = y \log_b x$$

The Logarithm of a Product:

$$\log_b(xy) = \log_b x + \log_b y$$

The Logarithm of a Quotient:

$$\log_b \frac{x}{y} = \log_b x - \log_b y$$

Numeral system	Symbols	Base	Additional information
Decimal	0-9	10	-
Binary	0, 1	2	-
Hexadecimal	0-9, A-F	16	A ≡ 10, B ≡ 11, C ≡ 12, D ≡ 13, E ≡ 14, F ≡ 15
Octal	0-7	8	-

Table A.1: Summary of Common Numeral Systems

Decimal number	In powers of 2	Power of 2				Binary number
		3	2	1	0	
8	= 2^3	1	0	0	0	1000
7	= $2^2 + 2^1 + 2^0$	0	1	1	1	111
6	= $2^2 + 2^1$	0	1	1	0	110
5	= $2^2 + 2^0$	0	1	0	1	101
4	= 2^2	0	1	0	0	100
3	= $2^1 + 2^0$	0	0	1	1	11
2	= 2^1	0	0	1	0	10
1	= 2^0	0	0	0	1	1

Table A.2: Decimal Numbers in Binary Representation

Proposition A.5 (Binary Addition Rules)

$$0 + 0 = 0, \quad 0 + 1 = 1, \quad 1 + 0 = 1, \quad 1 + 1 = 10$$



Proposition A.6 (Binary Multiplication Rules)

$$0 \times 0 = 0$$

$$0 \times 1 = 0$$

$$1 \times 0 = 0$$

$$1 \times 1 = 1$$

$$1 \times 10_2 = 10_2 \quad (\text{multiplying by base } 10_2 \text{ adds a 0 to the end})$$



Proposition A.7 (XOR Operation)

XOR produces a 1 if the two bits being compared are different and a 0 if they are the same:

$$0 \oplus 0 = 0, \quad 0 \oplus 1 = 1, \quad 1 \oplus 0 = 1, \quad 1 \oplus 1 = 0$$

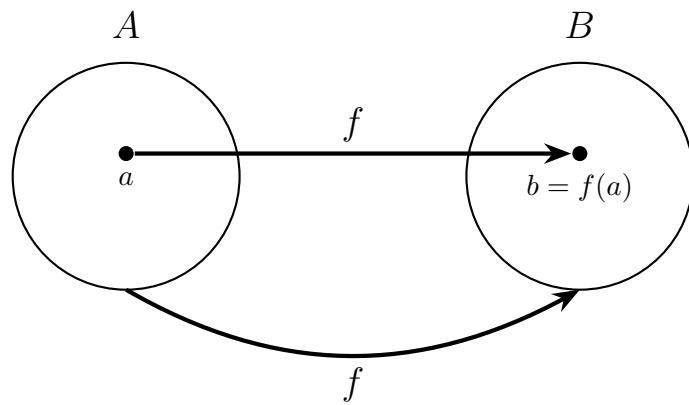


Figure A.3: A function f mapping an element a from set A to an element $b = f(a)$ in set B .

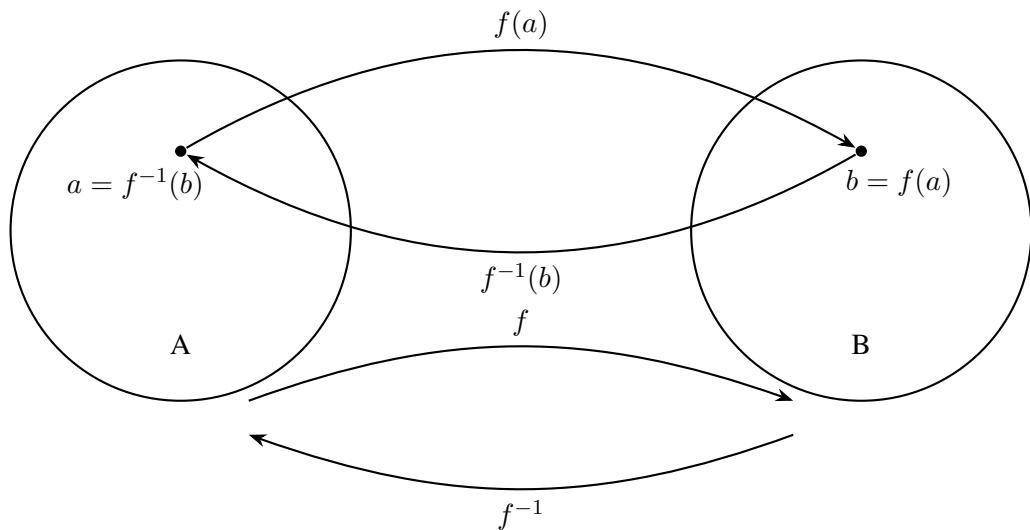


Figure A.4: The function f^{-1} is the inverse of function f .

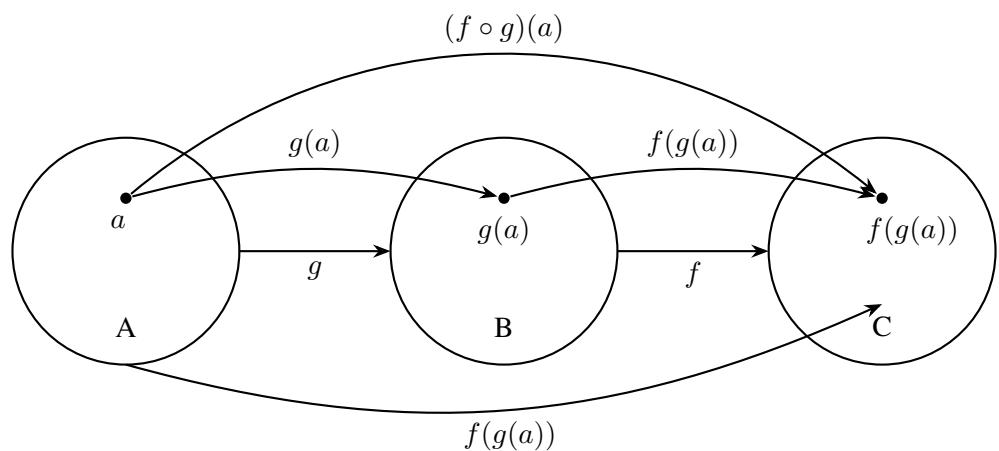


Figure A.5: The composition of functions f and g , denoted $f \circ g$, is the function that results from applying g and then f .